



# NoSQL Application Development with JSON and MapR-DB

Dale Kim, Sr. Director, Industry Solutions

Tug Grall, Technical Evangelist

August 24, 2016




# Our Speakers



Dale Kim  
Sr. Director, Industry Solutions



Tugdual Grall  
Technical Evangelist




# Quick NoSQL Overview




# Relational Databases Were Not Designed for Big Data

- RDBMSs are the default choice for applications
  - But large, rapidly changing, and/or diverse data sets add cost/time pressures



- This forces trade-offs with your data



Throwing away data to preserve performance?

- Or significant costs




Throwing extra money at the problem?

# The Common NoSQL Data Models

Key-Value




Wide Column



Document (JSON)


```
{  
  "fname": "John",  
  "lname": "Doe"  
}
```

Graph



JSON documents are popular because they easily model:

- Hierarchical/nested data
- Evolving data
- Varying data



# Let's Dive into JSON



# JSON is a Powerful Construct

```
function Thing()  
{  
  // public  
  this.get_name = function() {  
    return name;  
  }  
  
  this.set_name = function(tname) {  
    name = tname;  
  }  
  
  this.increment = (function () {  
    var count = 0;  
    return function () {return ++count;}  
  }) ();  
  
  // private  
  var name = null;  
}
```

```
var thing = new Thing();  
  
thing.set_name("John");  
  
// alert(thing.name); //error, private  
variable  
  
var counter = thing.counter;
```



# JSON Allows Easy Variation across Records

```
{
  "_id" : "rp-prod132546",
  "name" : "Marvel T2 Athena",
  "brand" : "Pinarello",
  "category" : "bike",
  "type" : "Road Bike",
  "price" : 2949.99,

  "size" : "55cm",
  "wheel_size" : "700c",
  "frameset" : {
    "frame" : "Carbon Toryaca",
    "fork" : "Onda 2V C"
  },
  "groupset" : {
    "chainset" : "Camp. Athena 50/34",
    "brake" : "Camp."
  },
  "wheelset" : {
    "wheels" : "Camp. Zonda",
    "tyres" : "Vittoria Pro"
  }
}
```



bike

```
{
  "_id" : "rp-prod106702",
  "name" : " Ultegra SPD-SL 6800",
  "brand" : "Shimano",
  "category" : "pedals",
  "type" : "Components",
  "price" : 112.99,

  "features" : [
    "Low profile design increases ...",
    "Supplied with floating SH11 cleats",
    "Weight: 260g (pair)"
  ]
}
```



pedal

```
{
  "_id" : "rp-prod113104",
  "name" : "Bianchi Pride Jersey SS15",
  "brand" : "Nalini",
  "category" : "Jersey",
  "type" : "Clothing",
  "price" : 76.99,


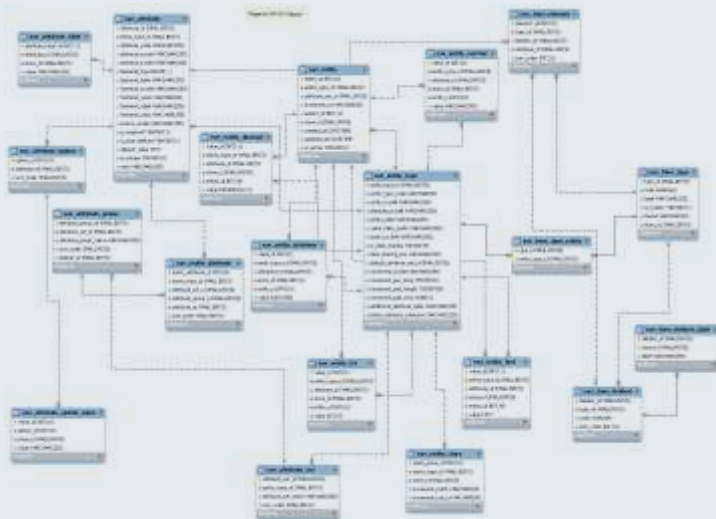
  "features" : [
    "100% Polyester",
    "3/4 hidden zip",
    "3 rear pocket"
  ],
  "color" : "black"
}
```



jersey



# How Variability Is Handled in an RDBMS




```
SELECT * FROM (
  SELECT
    ce.sku,
    ea.attribute_id,
    ea.attribute_code,
    CASE ea.backend_type
      WHEN 'varchar' THEN ce_varchar.value
      WHEN 'int' THEN ce_int.value
      WHEN 'text' THEN ce_text.value
      WHEN 'decimal' THEN ce_decimal.value
      WHEN 'datetime' THEN ce_datetime.value
      ELSE ea.backend_type
    END AS value,
    ea.is_required AS required
  FROM catalog_product_entity AS ce
  LEFT JOIN eav_attribute AS ea
    ON ce.entity_type_id = ea.entity_type_id
  LEFT JOIN catalog_product_entity_varchar AS ce_varchar
    ON ce.entity_id = ce_varchar.entity_id
    AND ea.attribute_id = ce_varchar.attribute_id
    AND ea.backend_type = 'varchar'
  LEFT JOIN catalog_product_entity_text AS ce_text
    ON ce.entity_id = ce_text.entity_id
    AND ea.attribute_id = ce_text.attribute_id
    AND ea.backend_type = 'text'
  LEFT JOIN catalog_product_entity_decimal AS ce_decimal
    ON ce.entity_id = ce_decimal.entity_id
    AND ea.attribute_id = ce_decimal.attribute_id
    AND ea.backend_type = 'decimal'
  LEFT JOIN catalog_product_entity_datetime AS ce_datetime
    ON ce.entity_id = ce_datetime.entity_id
    AND ea.attribute_id = ce_datetime.attribute_id
    AND ea.backend_type = 'datetime'
  WHERE ce.sku = 'rp-prod132546'
) AS tab
WHERE tab.value != '';
```

“Entity Value Attribute” pattern

To get a single product

# Product Catalog - NoSQL/Document

```
{  
  "_id" : "rp-prod132546",  
  "name" : "Marvel T2 Athena",  
  "brand" : "Pinarello",  
  "category" : "bike",  
  "type" : "Road Bike",  
  "price" : 2949.99,
```




```
products  
  .findById("rp-prod132546")
```

```
  "size" : "55cm",  
  "wheel_size" : "700c",  
  "frameset" : {  
    "frame" : "Carbon Toryaca",  
    "fork" : "Onda 2V C"  
  },  
  "groupset" : {  
    "chainset" : "Camp. Athena 50/34",  
    "brake" : "Camp."  
  },  
  "wheelset" : {  
    "wheels" : "Camp. Zonda",  
    "tyres" : "Vittoria Pro"  
  }  
}
```

Store the product “as a business object”


To get a single product



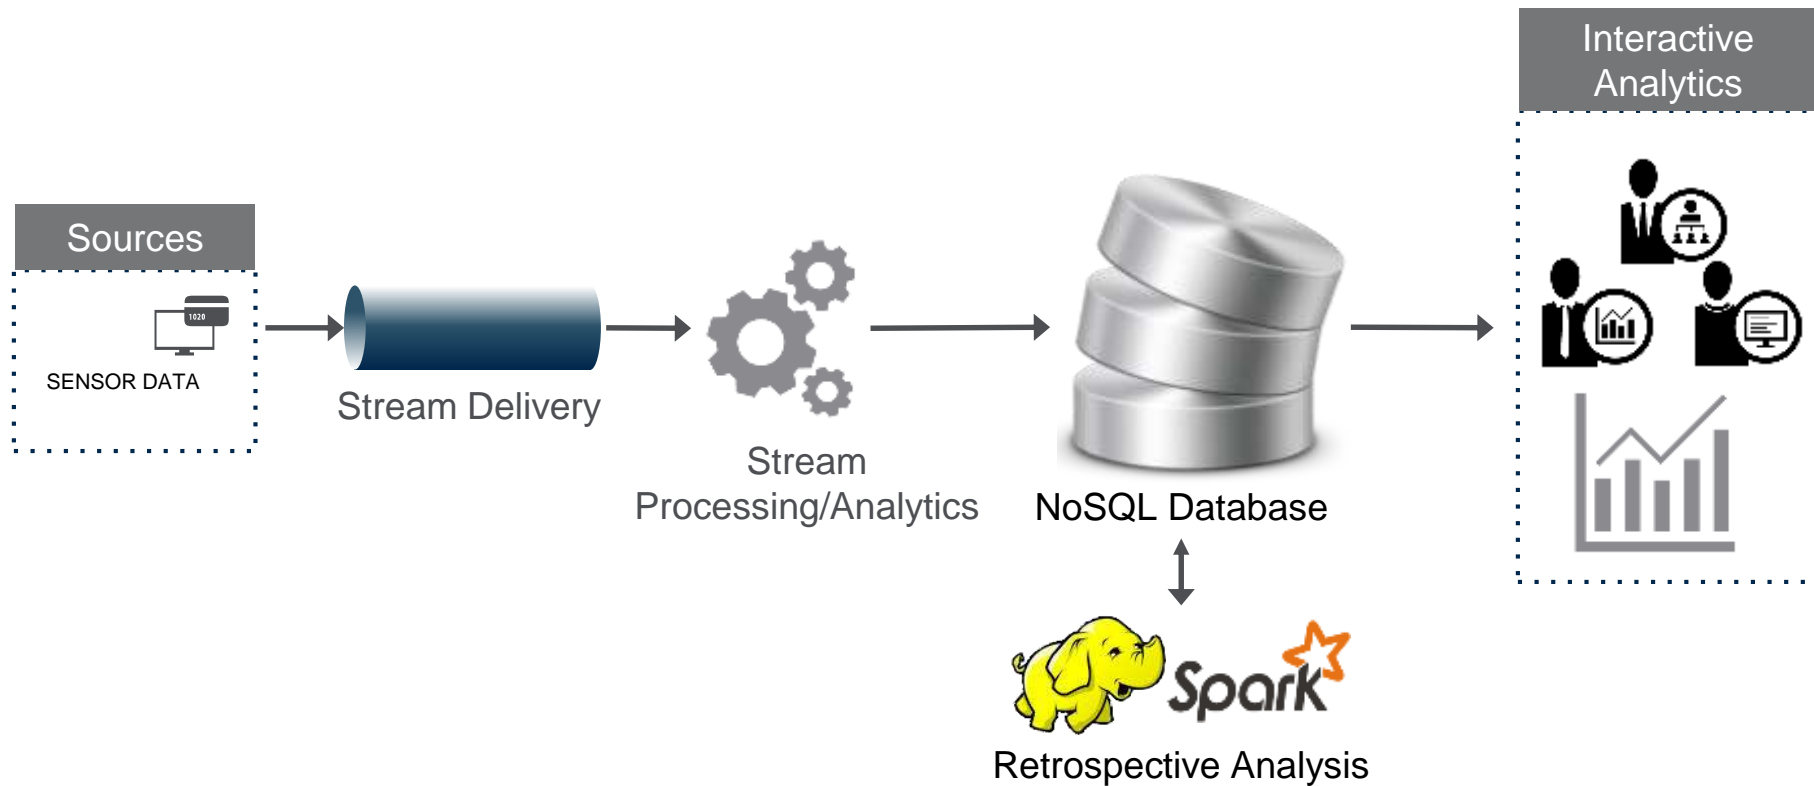
# Example Use Cases for JSON/NoSQL




# Example Use Case 1: Data Lake



# Example Use Case 2: IoT and Predictive Analytics





# What Should Native JSON Support Include?



# Recognize JSON Elements at Server-Side

```
{  
  order_num: 5555,  
  products: [  
    {  
      product_id: 348752,  
      quantity: 1,  
      unit_price: 149.99,  
      total_price: 149.99  
    },  
    {  
      product_id: 439322,  
      quantity: 1,  
      unit_price: 99.99,  
      total_price: 99.99  
    },  
    {  
      product_id: 953923,  
      quantity: 1,  
      unit_price: 49.99,  
      total_price: 49.99  
    },  
  ],  
}
```



Reads/writes at element level

- Granular disk reads/writes
- Less network traffic
- Higher concurrency




Any new elements added on demand

- No predefined schemas
- Easy to store evolving data



versus



# Organize Elements for Different Policies

```
/
{a:
  {a1:
    {b1: "v1",
      b2: [
        {c1: "v1",
          c2: "v2"}
      ]
    },
    a2:
    {
      e1: "v1",
      e2: <inline jpg>
    }
  }
}
```

- Control layout for faster data access
- Different TTL requirements
- Separate replication settings
- Efficient access controls

Column Family 1

Column Family 2





# Fine Grained Security within JSON Document

```
{  
  "fname": "John",  
  "lname": "Doe",  
  "address": "111 Main St.",  
  "city": "San Jose",  
  "state": "CA",  
  "zip": "95134",  
  "credit_cards": [  
    {"issuer": "Visa",  
     "number": "4444555566667777"},  
    {"issuer": "MasterCard",  
     "number": "5555666677778888"}  
  ]  
}
```

← Entire document

← Element: "fname"

← Array: "credit\_cards"

← Sub-element in array element:  
"credit\_cards[\*].number"

Specify different permissions levels within the document.




# Extensions for Comprehensive Data Type Support

- NULL
- Boolean
- String
- Map
- Array
- Float, Double
- Binary
- Byte, Short, Int, Long
- Date
- Decimal
- Interval
- Time
- Timestamp

## Examples:

```
{  
  "sample_int": {"$numberLong": 2147483647},  
  "sample_date": {"$dateDay": "2016-02-22"},  
  "sample_decimal": {"$decimal": "1234567890.23456789"},  
  "sample_time": {"$time": "10:26:12.487"},  
  "sample_timestamp": {"$date": "2016-02-22T10:26:12.487+Z"}  
}
```

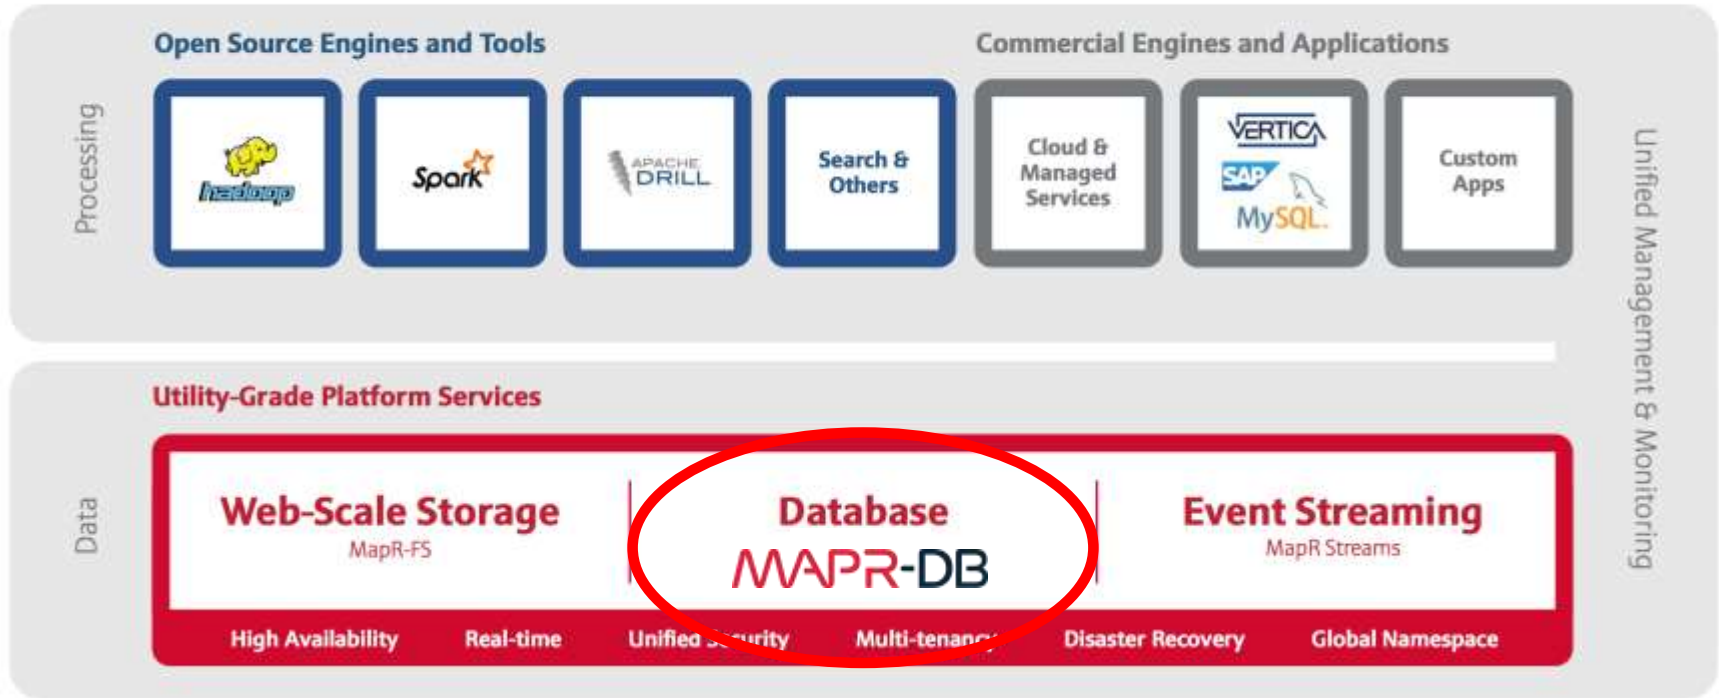




# How Does This Relate to MapR?

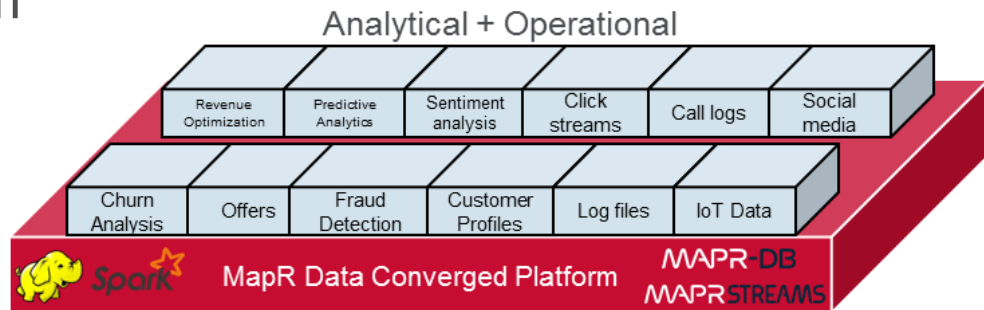


# MapR Converged Data Platform



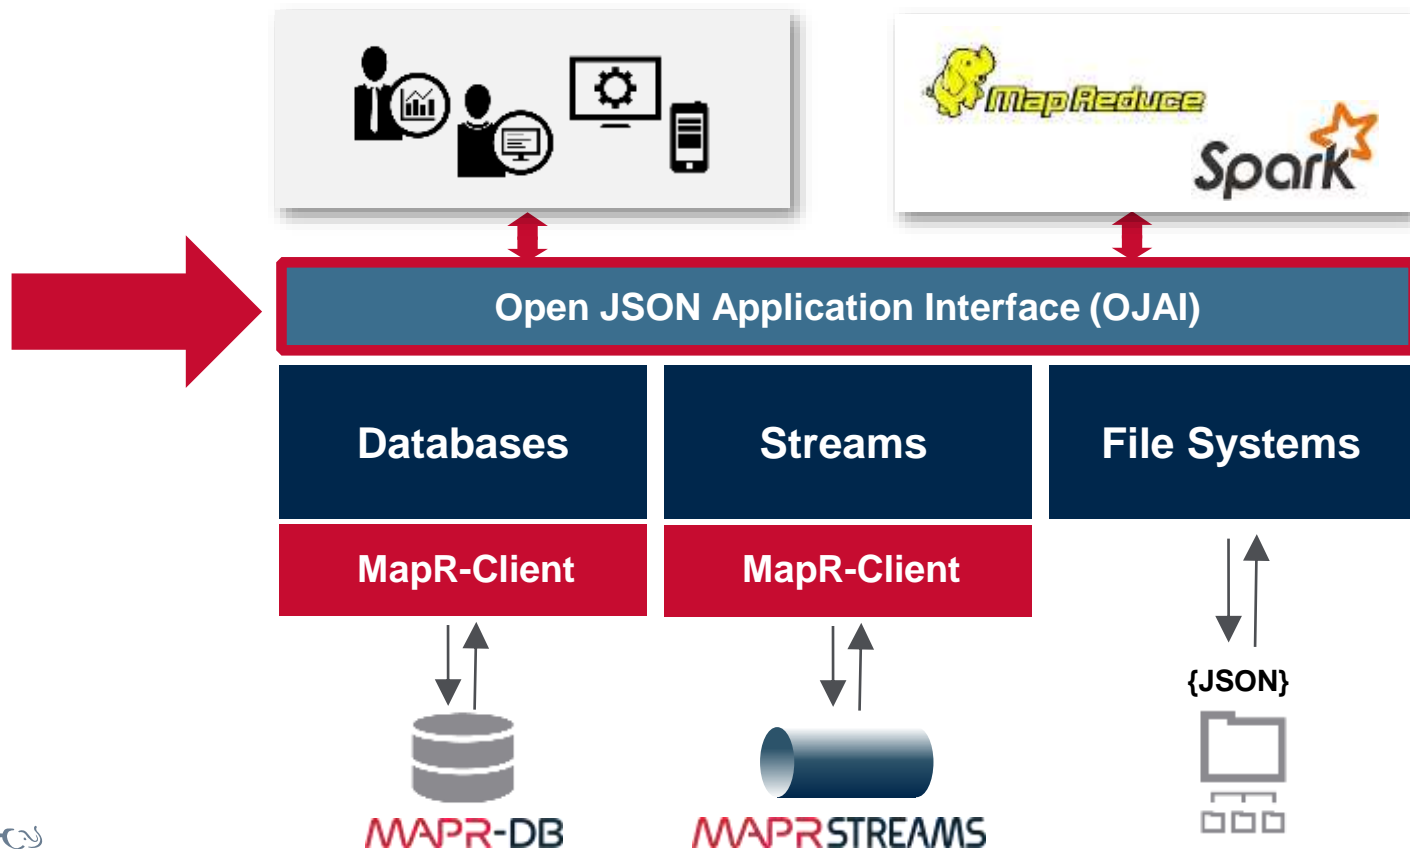
# A Single Application Development Platform

- Business applications and analytics on the same platform
  - No data movement
  - Real-time access to data
- Architectural simplicity
- Optimized stack for greater efficiency
- “Hadoop-scale” for multi-petabyte deployments



Analytics as it happens, no cross-cluster copying

# Open Source OJAI API for JSON-Based Applications



# Familiar JSON Paradigm – Similar API Constructs

## MapR-DB

## Other Document Database

```
Document record = Json.newDocument()  
    .set("firstName", "John")  
    .set("lastName", "Doe")  
    .set("age", 50);
```

```
table.insert("jdoe", record);
```

```
BasicDBObject doc = new BasicDBObject  
    ("firstName", "John")  
    .append("lastName", "Doe")  
    .append("age", 50);
```

```
coll.insert(doc);
```



# Multi-Master Global Deployments

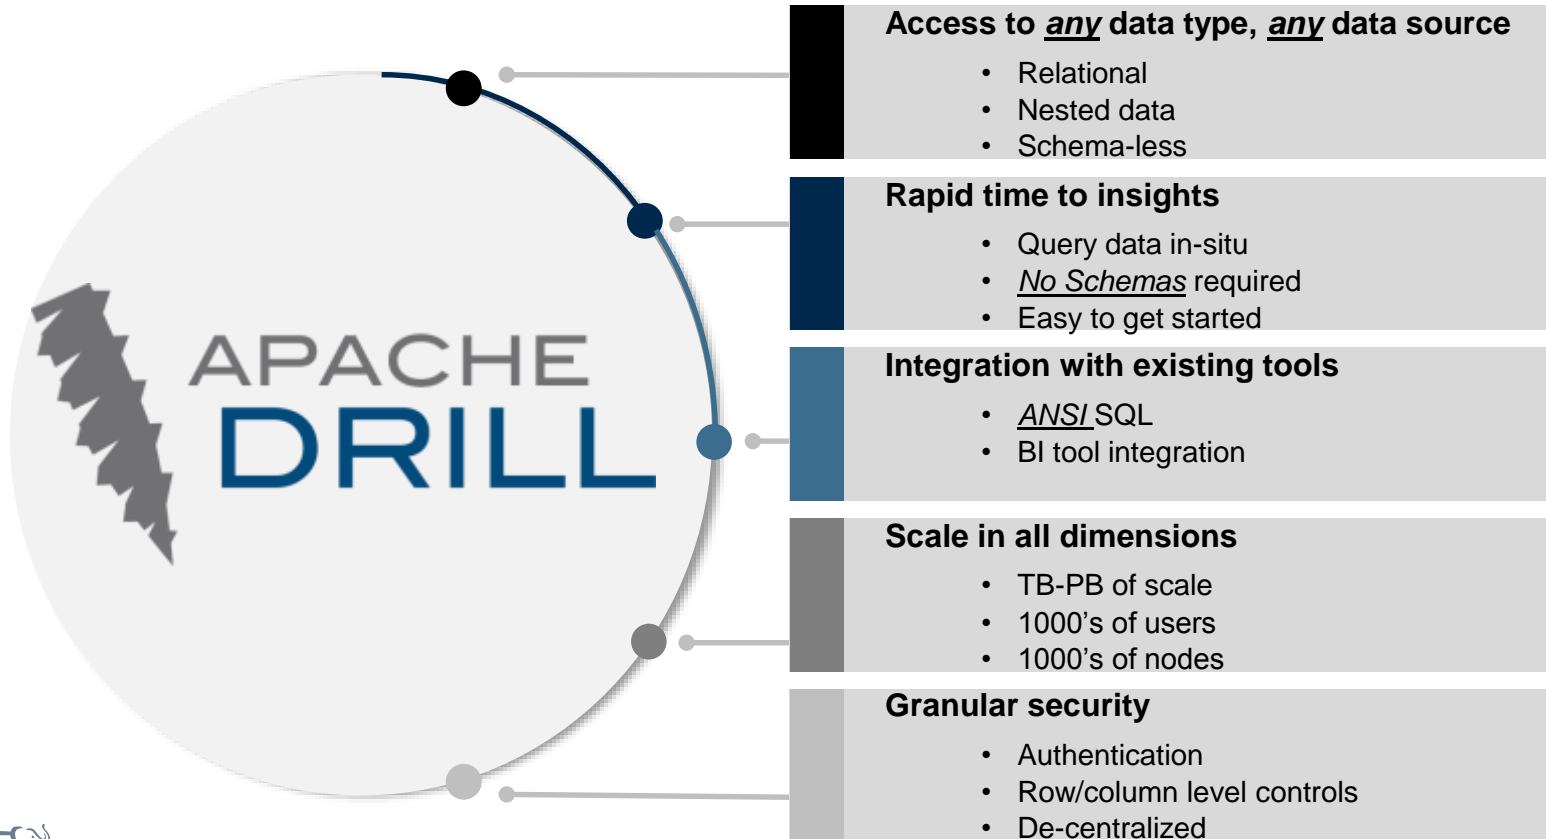


## Active-active replication


- **Faster data access** – minimize network latency on global data with local clusters
- **Reduced risk of data loss** – real-time, bi-directional replication for synchronized data across active clusters
- **Application failover** – upon any cluster failure, applications continue via redirection to another cluster



# Apache Drill for SQL Querying on MapR-DB

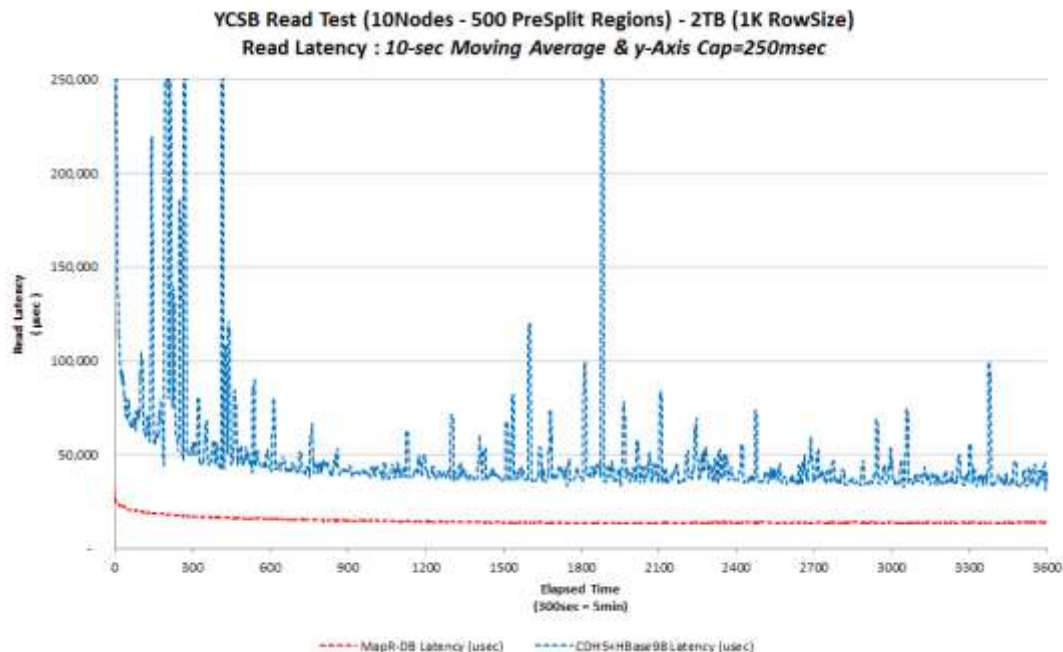


# Flash/SSD/NVMe Optimization



# Fast, Automatic Optimizations

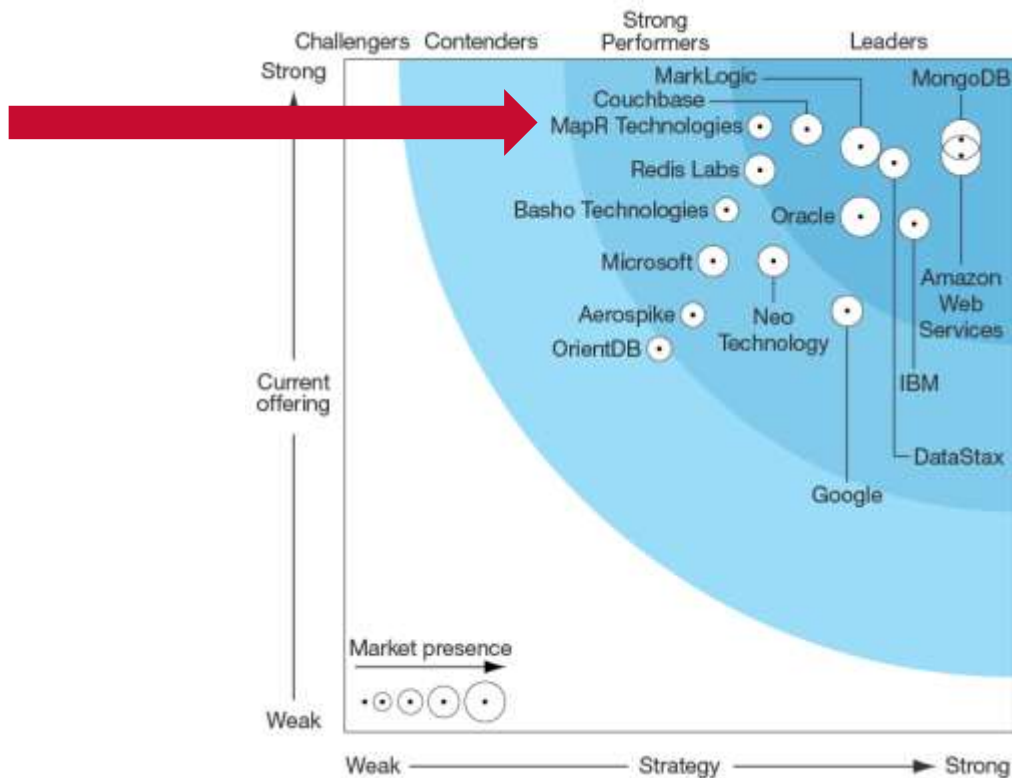
- No manual cleanup tasks
- Automatic sharding
- No compaction delays
- No anti-entropy task needed to sync replicas (strong consistency)



Red plots show MapR-DB response times.  
Blue plots show other NoSQL response times.




# Independent Third-Party Evaluation Results



The Forrester Wave™ is copyrighted by Forrester Research, Inc. Forrester and Forrester Wave™ are trademarks of Forrester Research, Inc. The Forrester Wave™ is a graphical representation of Forrester's call on a market and is plotted using a detailed spreadsheet with exposed scores, weightings, and comments. Forrester does not endorse any vendor, product, or service depicted in the Forrester Wave. Information is based on best available resources. Opinions reflect judgment at the time and are subject to change.





# Demo



# Q&A

Engage with us!

## 1. MapR-DB

Learn more and get started using MapR-DB

<https://www.mapr.com/products/mapr-db-in-hadoop-nosql>

Free on-demand training:

<https://www.mapr.com/ODT>

## 2. Ask Questions:

- Ask Us Anything about MapR-DB in the MapR Community now through Fri (Aug 26)
- <https://community.mapr.com/>

**CONVERGE** *community*  
Powered by MapR

