



Scale Your SQL Database with SQL, Hadoop, Drill, JSON, NoSQL and HBase

Carol McDonald cmcDonald@mapr.com

<http://answers.mapr.com/>





Free Hadoop On Demand Training

- <https://www.mapr.com/services/mapr-academy/big-data-hadoop-online-training>

DEV 320 - HBase Data Model and Architecture

[Register](#)

This course is intended for data analysts, data architects and application developers. DEV 320 provides you with a thorough understanding of the HBase data model and architecture, which is required before going on to designing HBase schemas and developing HBase applications.

[Learn More](#)

DEV 325 - HBase Schema Design

[Register](#)

NEW!
Targeted towards data analysts, data architects and application developers, the goal of this course is to enable you to design HBase schemas based on design guidelines. You will learn about the various elements of schema design and how to design for data access patterns. The course offers an in-depth look at designing row keys, avoiding hot-spotting and designing column families. It discusses how to transition from a relational model to an HBase model. You will learn the differences between tall tables and wide tables. Concepts are conveyed through lectures, hands-on labs and analysis of scenarios.

[Learn More](#)

DA 410 - Drill Essentials

[Register](#)

NEW!
This introductory Drill course, targeted at Data Analysts, Scientists and SQL programmers, covers how to use Drill to explore known or unknown data without writing code. You will write SQL queries on a variety of data types including structured data in a Hive table, semi-structured data in HBase or MapR-DB, and complex data file types, such as Parquet and JSON.



Objectives of this session

- What is HBase?
 - Why do we need NoSQL / HBase?
 - Overview of HBase & HBase data model
 - HBase Architecture and data flow
- Design considerations when migrating from RDBMS to HBase
- MapReduce and HBase
- Hive and HBase
- Drill and HBase

Prerequisite for Hands-On-Labs (before Lecture)

- Install a one-node MapR Sandbox on your laptop
 - Download install the MapR Sandbox:
 - <https://cwiki.apache.org/confluence/display/DRILL/Installing+the+Apache+Drill+Sandbox>



[Drill Wiki](#) / [Apache Drill Wiki](#) / [Apache Drill Tutorial](#)

Installing the Apache Drill Sandbox

Added by Bridget Bevens, last edited by Bridget Bevens on Sep 25, 2014

Prerequisites

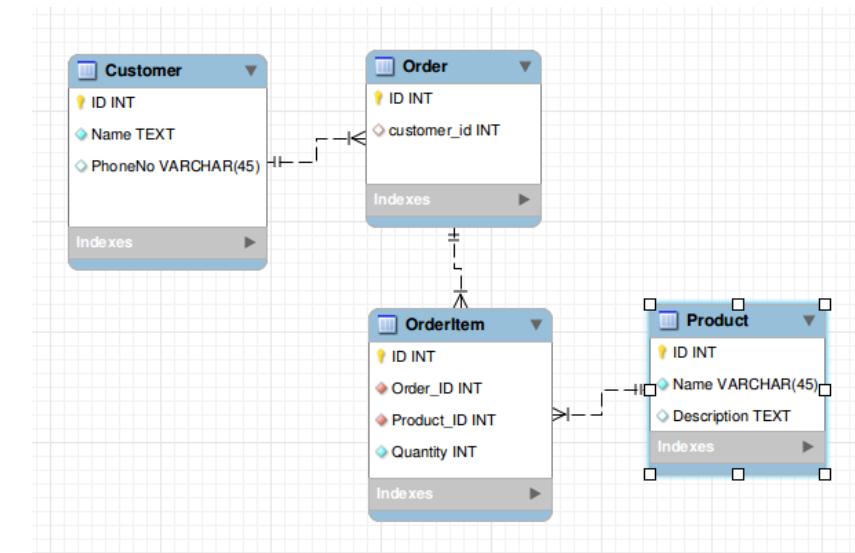
The MapR Sandbox with Apache Drill runs on VMware Player and VirtualBox, free desktop applications Mac, or Linux PC. Before you install the MapR Sandbox with Apache Drill, verify that the host system m

- VMware Player or VirtualBox is installed.
- At least 20 GB free hard disk space, at least 4 physical cores, and 8 GB of RAM is available. Perf
- Uses one of the following 64-bit x86 architectures:
 - A 1.3 GHz or faster AMD CPU with segment-limit support in long mode
 - A 1.3 GHz or faster Intel CPU with VT-x support
- If you have an Intel CPU with VT-x support, verify that VT-x support is enabled in the host system support vary depending on the system vendor. See the VMware knowledge base article at <http://> determine if VT-x support is enabled.

Why do we need NoSQL / HBase?

Relational Model

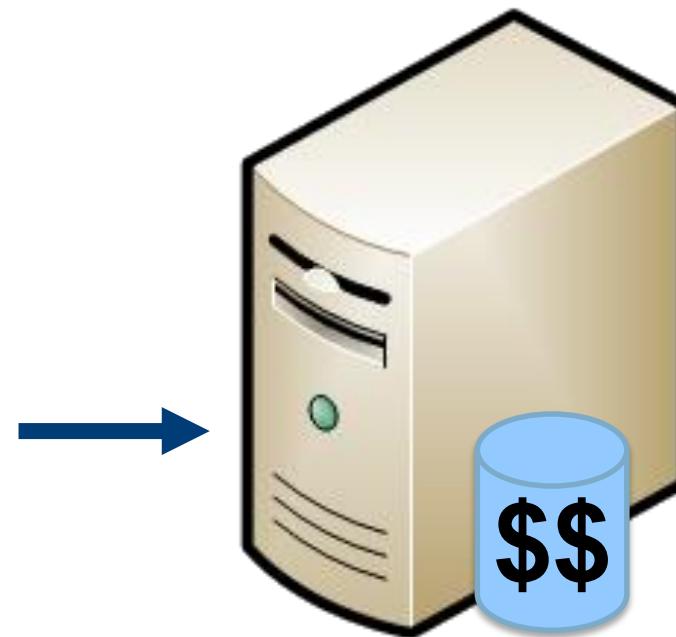
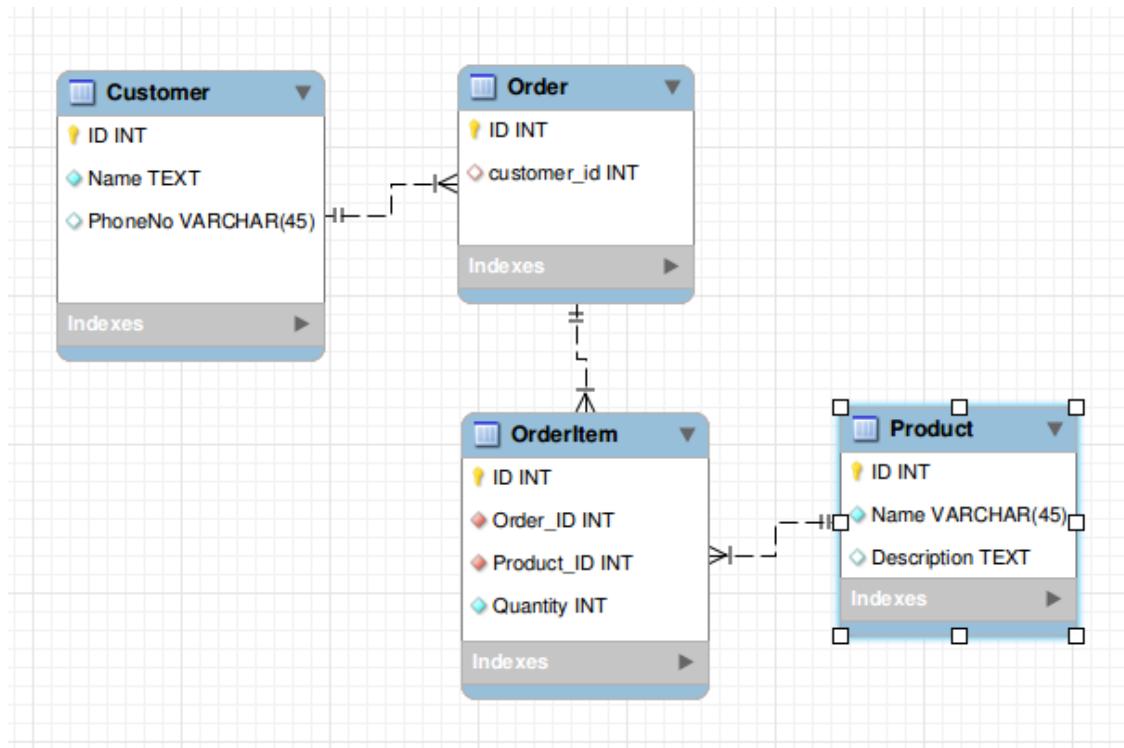
- Pros
 - Standard persistence model
 - Transactions handle
 - concurrency , consistency
 - efficient and robust structure for storing data





Relational Databases vs HBase - Scaling

RDBMS - Scale UP approach

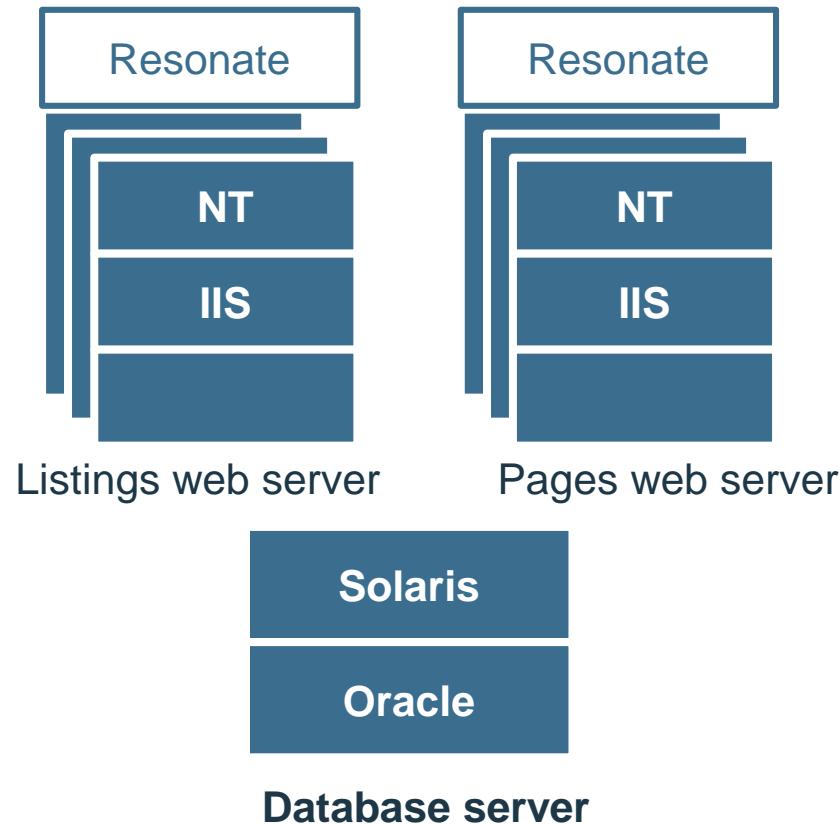


Vertical scale = **big box**



RDBMS Scaling Up Example - eBay

- Back End Oracle Database server **scaled vertically to a larger machine** (Sun E10000)



<http://www.addsimplicity.com/downloads/eBaySDForum2006-11-29.pdf>



What changed to bring on NoSQL?

Lots of data, the need to scale horizontally

Horizontal scale: split table by rows into partitions across a cluster

Key	colB	colC
val	val	val
xxx	val	val

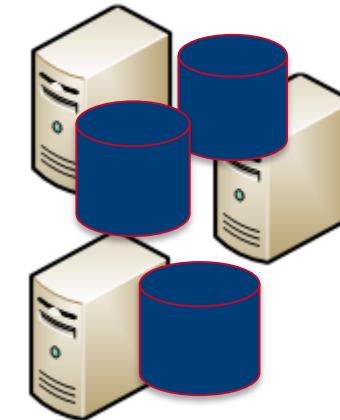
Key	colB	colC
val	val	val
xxx	val	val

Key	colB	colC
val	val	val
xxx	val	val

id 1-1000

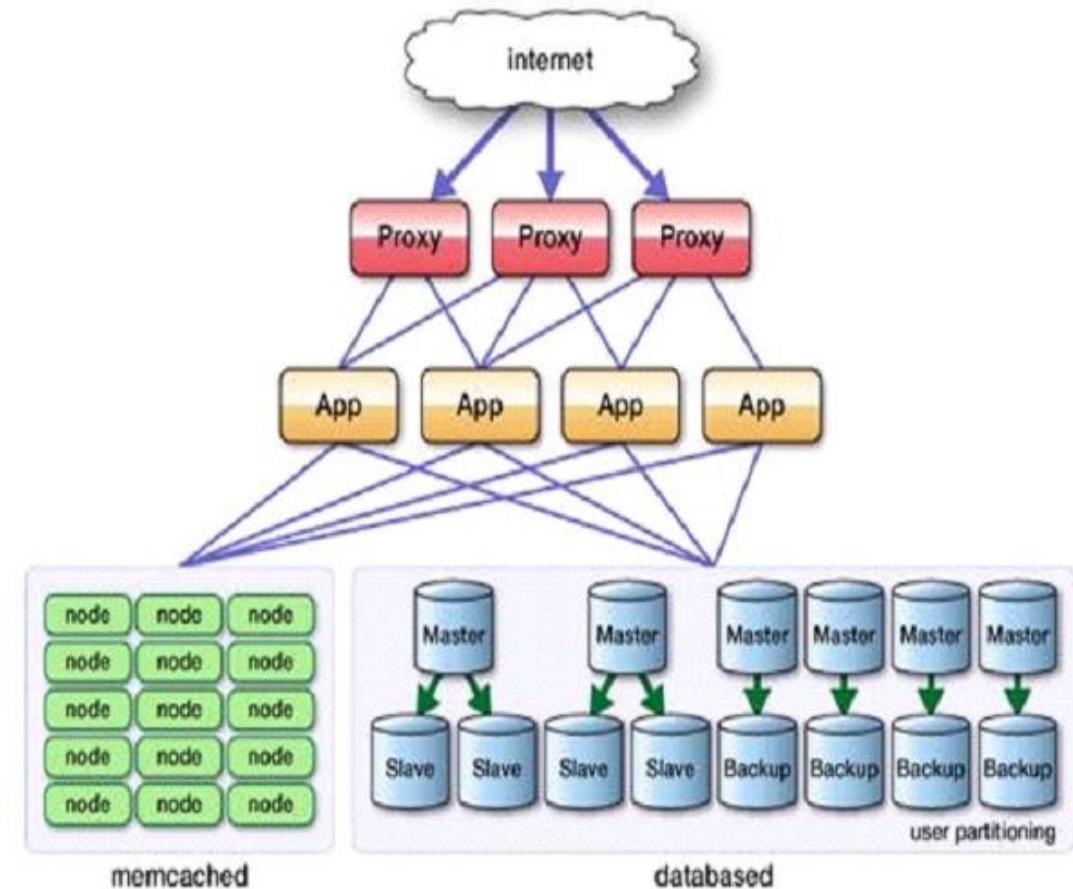
id 1000-2000

id 2000=3000



- **Horizontal scaling**
 - Cheaper than vertical
 - parallel execution
- Relational databases were **not designed to do this** automatically

- 9000 memcache instances
- **4000 Shards mysql**

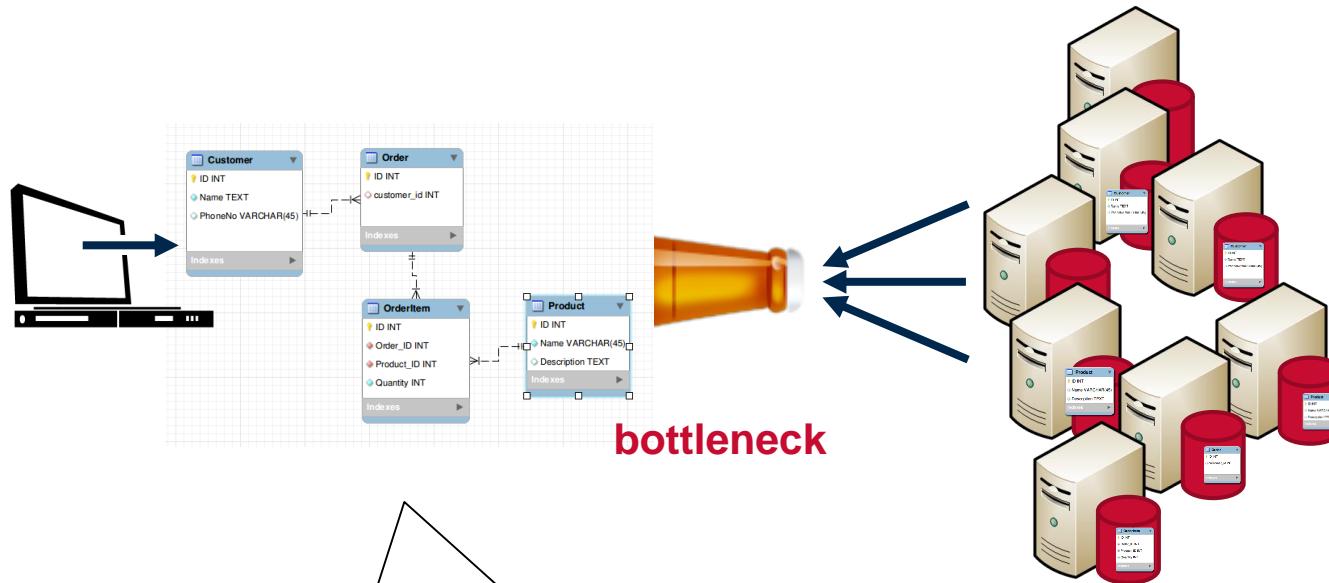


<http://gigaom.com/2011/07/07/facebook-trapped-in-mysql-fate-worse-than-death/>



Relational Databases vs. HBase – Data Storage Model

RDBMS



Distributed Joins, Transactions
do not scale

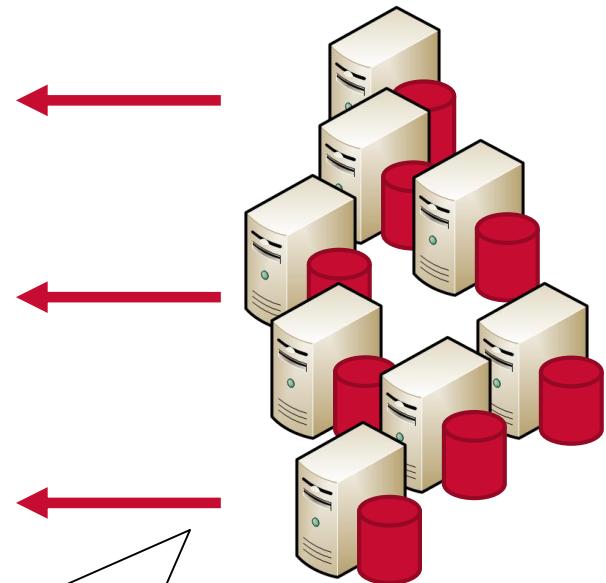
Storage Model

Key	colB	colC
val	val	val
xxx	val	val

Key	colB	colC
val	val	val
xxx	val	val

Key	colB	colC
val	val	val
xxx	val	val

HBase



Data that is accessed together is stored together



Hbase designed for Distribution, Scale, Speed



Google

Big Table

- Distributed Storage System
- Paper published in 2006.

Google File System

MapReduce

Runs on commodity hardware
Designed to Scale



HBase is a ColumnFamily oriented Database

Customer id

Customer Address data

Customer order data

RowKey	CF1			CF2		
	colA	colB	colC	colA	colB	colC
axxx	Val		val		val	
gxxx		val			val	

Data is **accessed and stored together**:

- RowKey is the primary index
- Column Families group similar data by **row key**



HBase is a Distributed Database



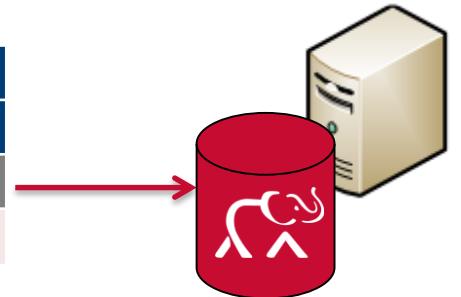
Put, Get by Key



Key Range	
XXXX	
XXXX	

CF1		
colA	colB	colC
val		val
	val	

CF2		
colA	colB	colC
val		val
	val	



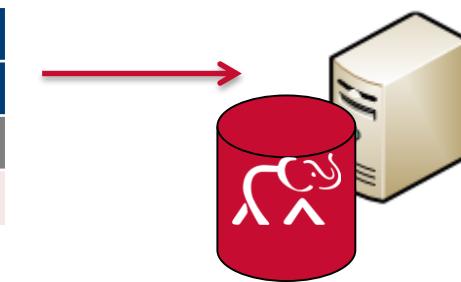
Data is **automatically distributed** across the cluster

- **Key range** is used for horizontal partitioning

Key Range	
XXXX	
XXXX	

CF1		
colA	colB	colC
val		val
	val	

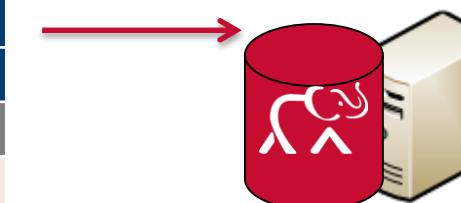
CF2		
colA	colB	colC
val		val
	val	



Key Range	
XXXX	
XXXX	

CF1		
colA	colB	colC
val		val
	val	

CF2		
colA	colB	colC
val		val
	val	

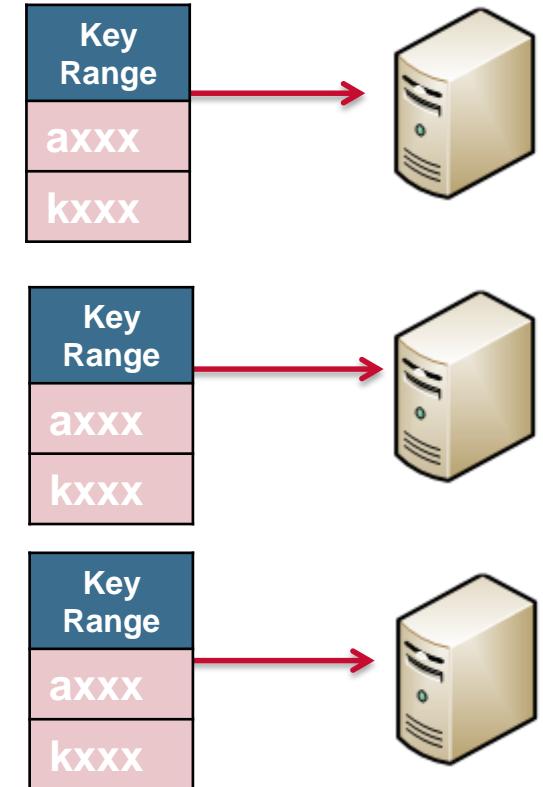




Column Family Databases

distributed data stored and accessed together:

- Pros
 - **scalable**
 - **Fast Writes and Reads**
- Cons
 - **No** joins
 - No dynamic queries
 - Need to **know how** data will be **queried in advance**



Hbase Data Model





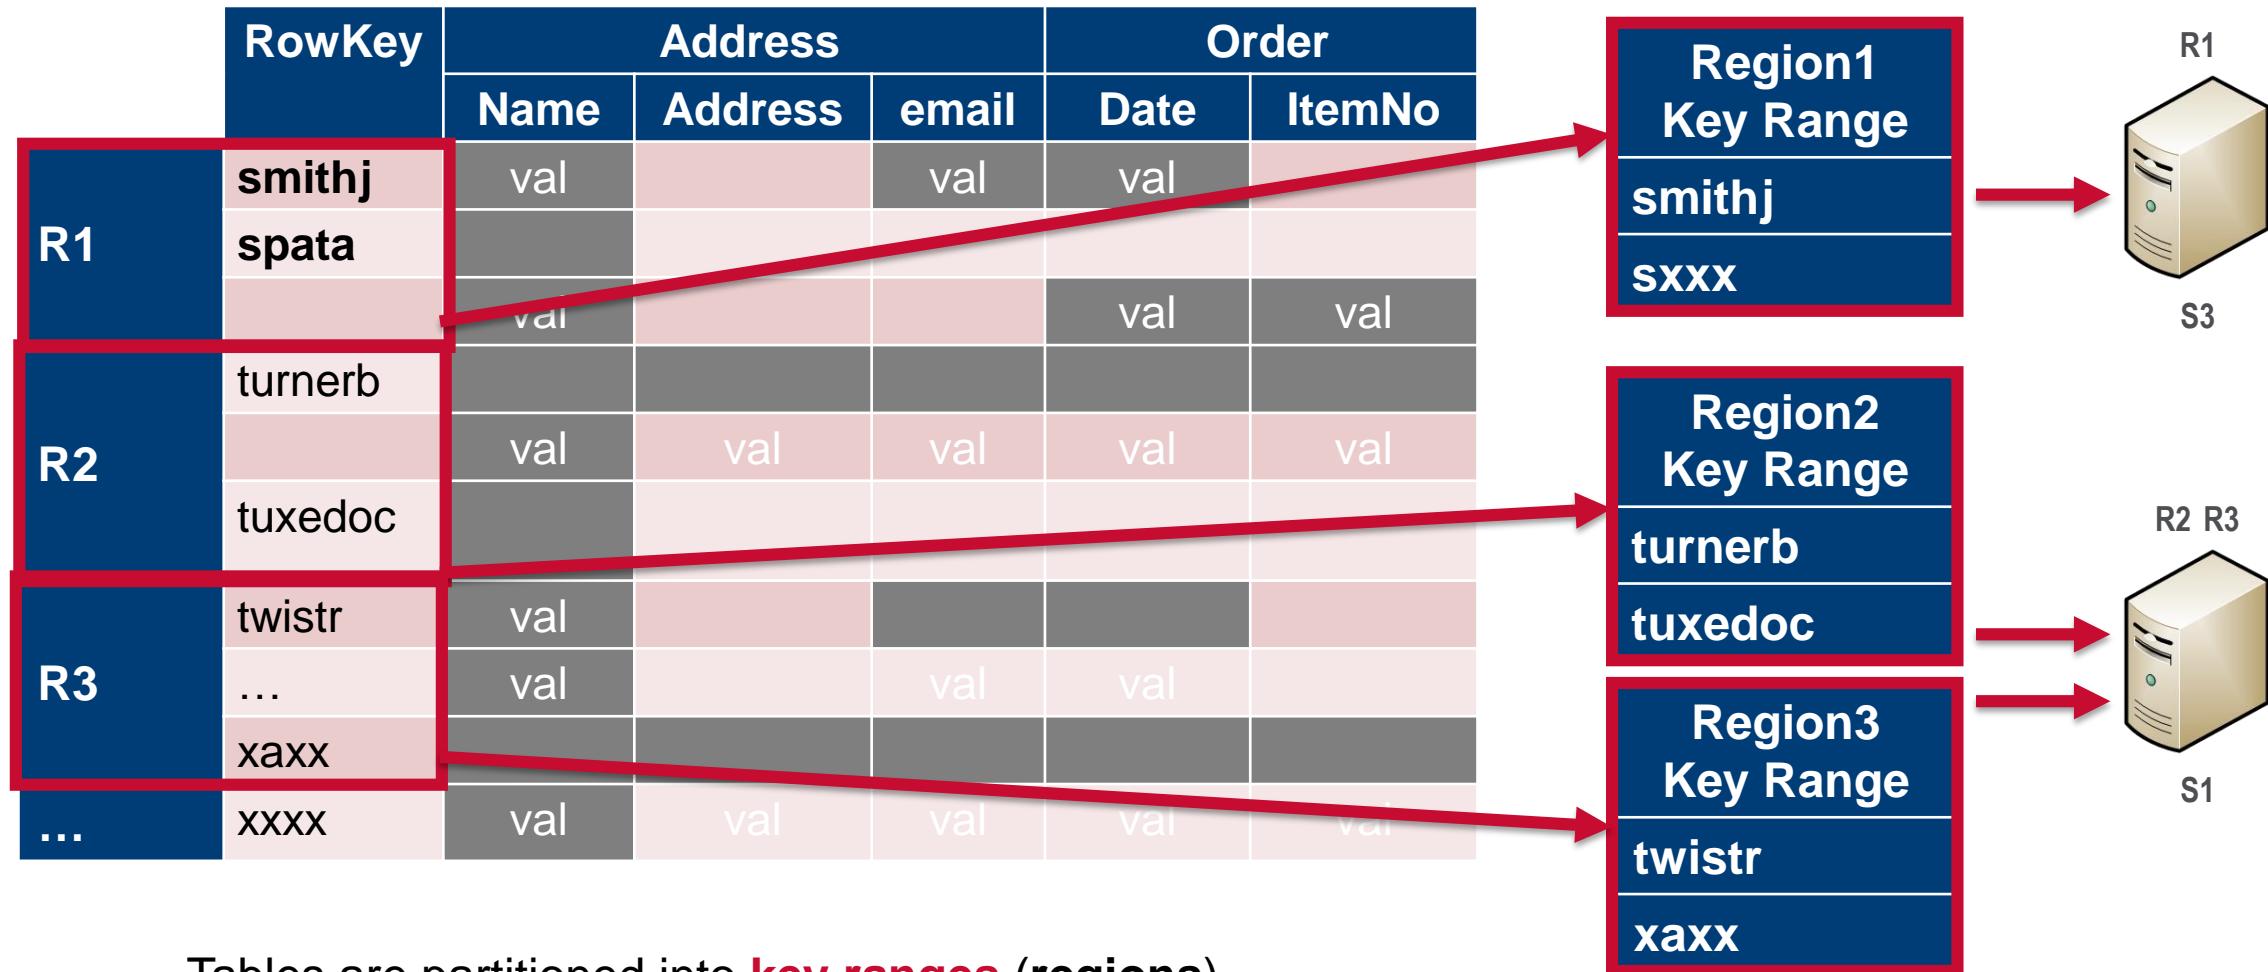
HBase Data Model- Row Keys

RowKey	Address			Order				...
	street	city	state	Date	ItemNo	Ship Address	Cost	
smithj	val		val	val			val	
spata								
sxxxx	val			val	val	val		
...								
turnerb	val	val	val	val	val	val	val	
...								
	val							
twistr	val		val	val			val	
...								
zaxx	val	val	val	val	val	val		
zxxx	val						val	

Row Keys: identify the rows in an HBase table.



Tables are split into Regions = contiguous keys

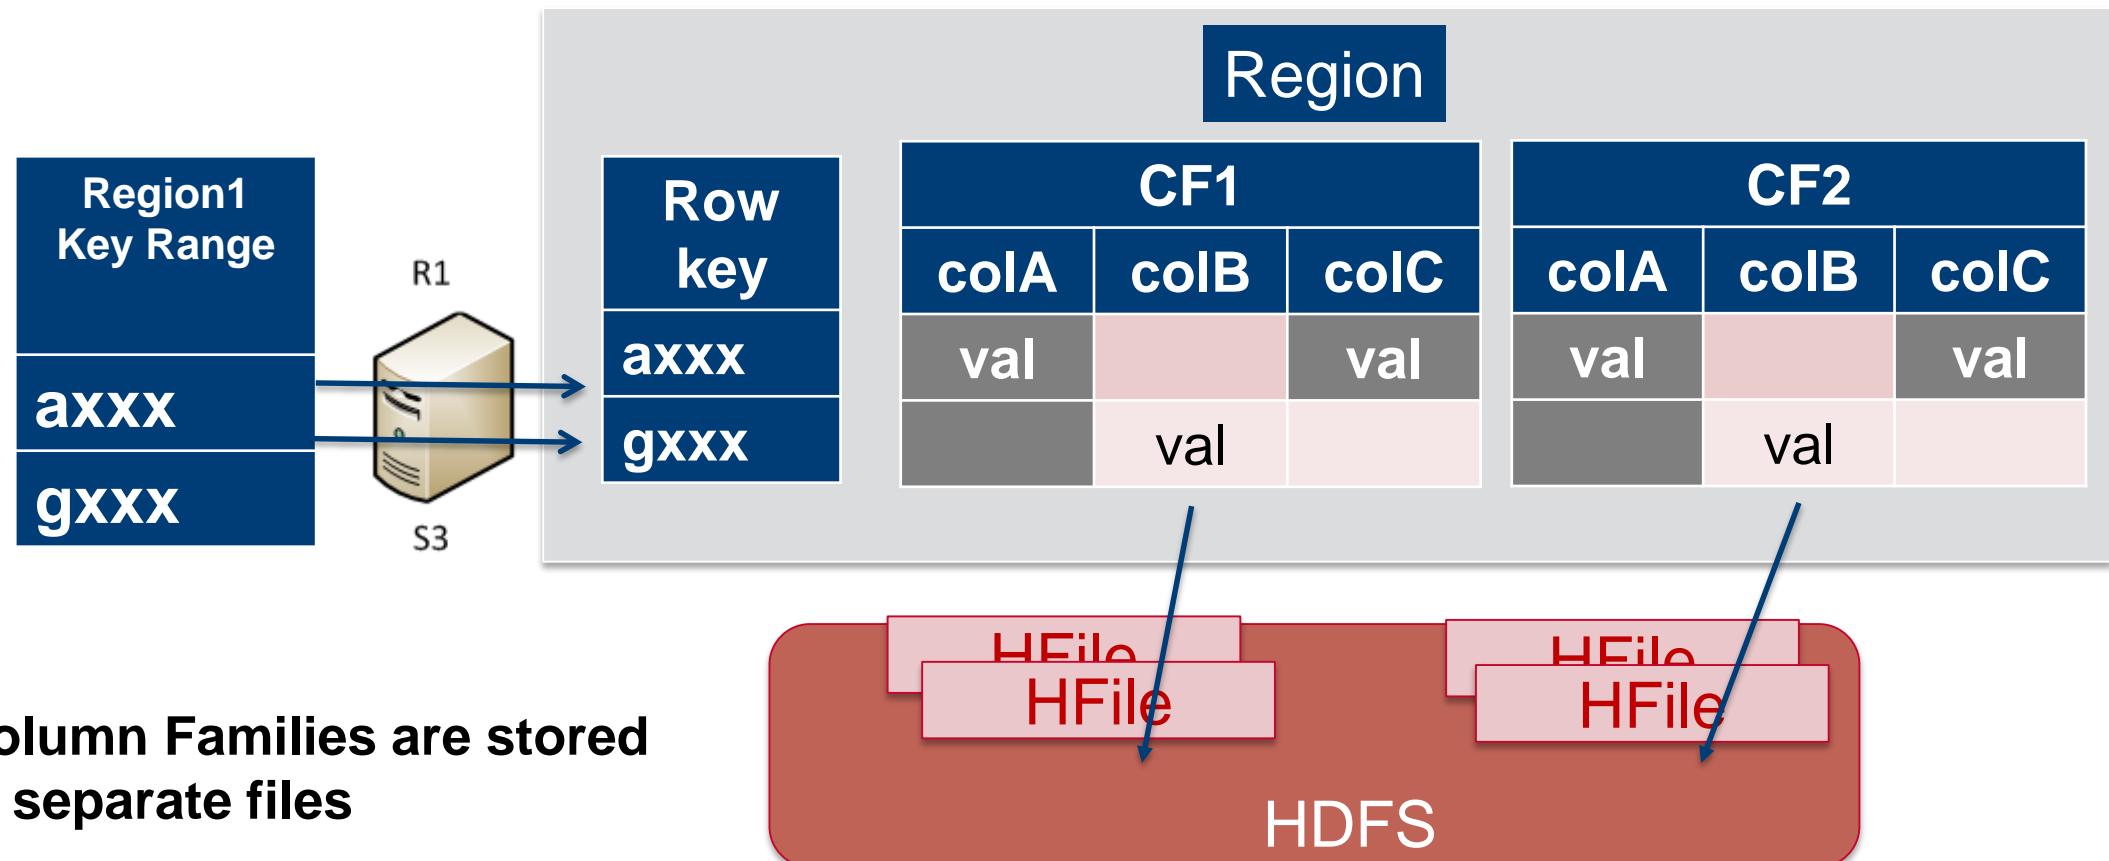


Tables are partitioned into **key ranges (regions)**
Region= served by nodes (Region Servers)
Regions are spread across cluster



Column Families are stored Separately

- column families are stored and can be accessed separately





HBase Data Model - Cells

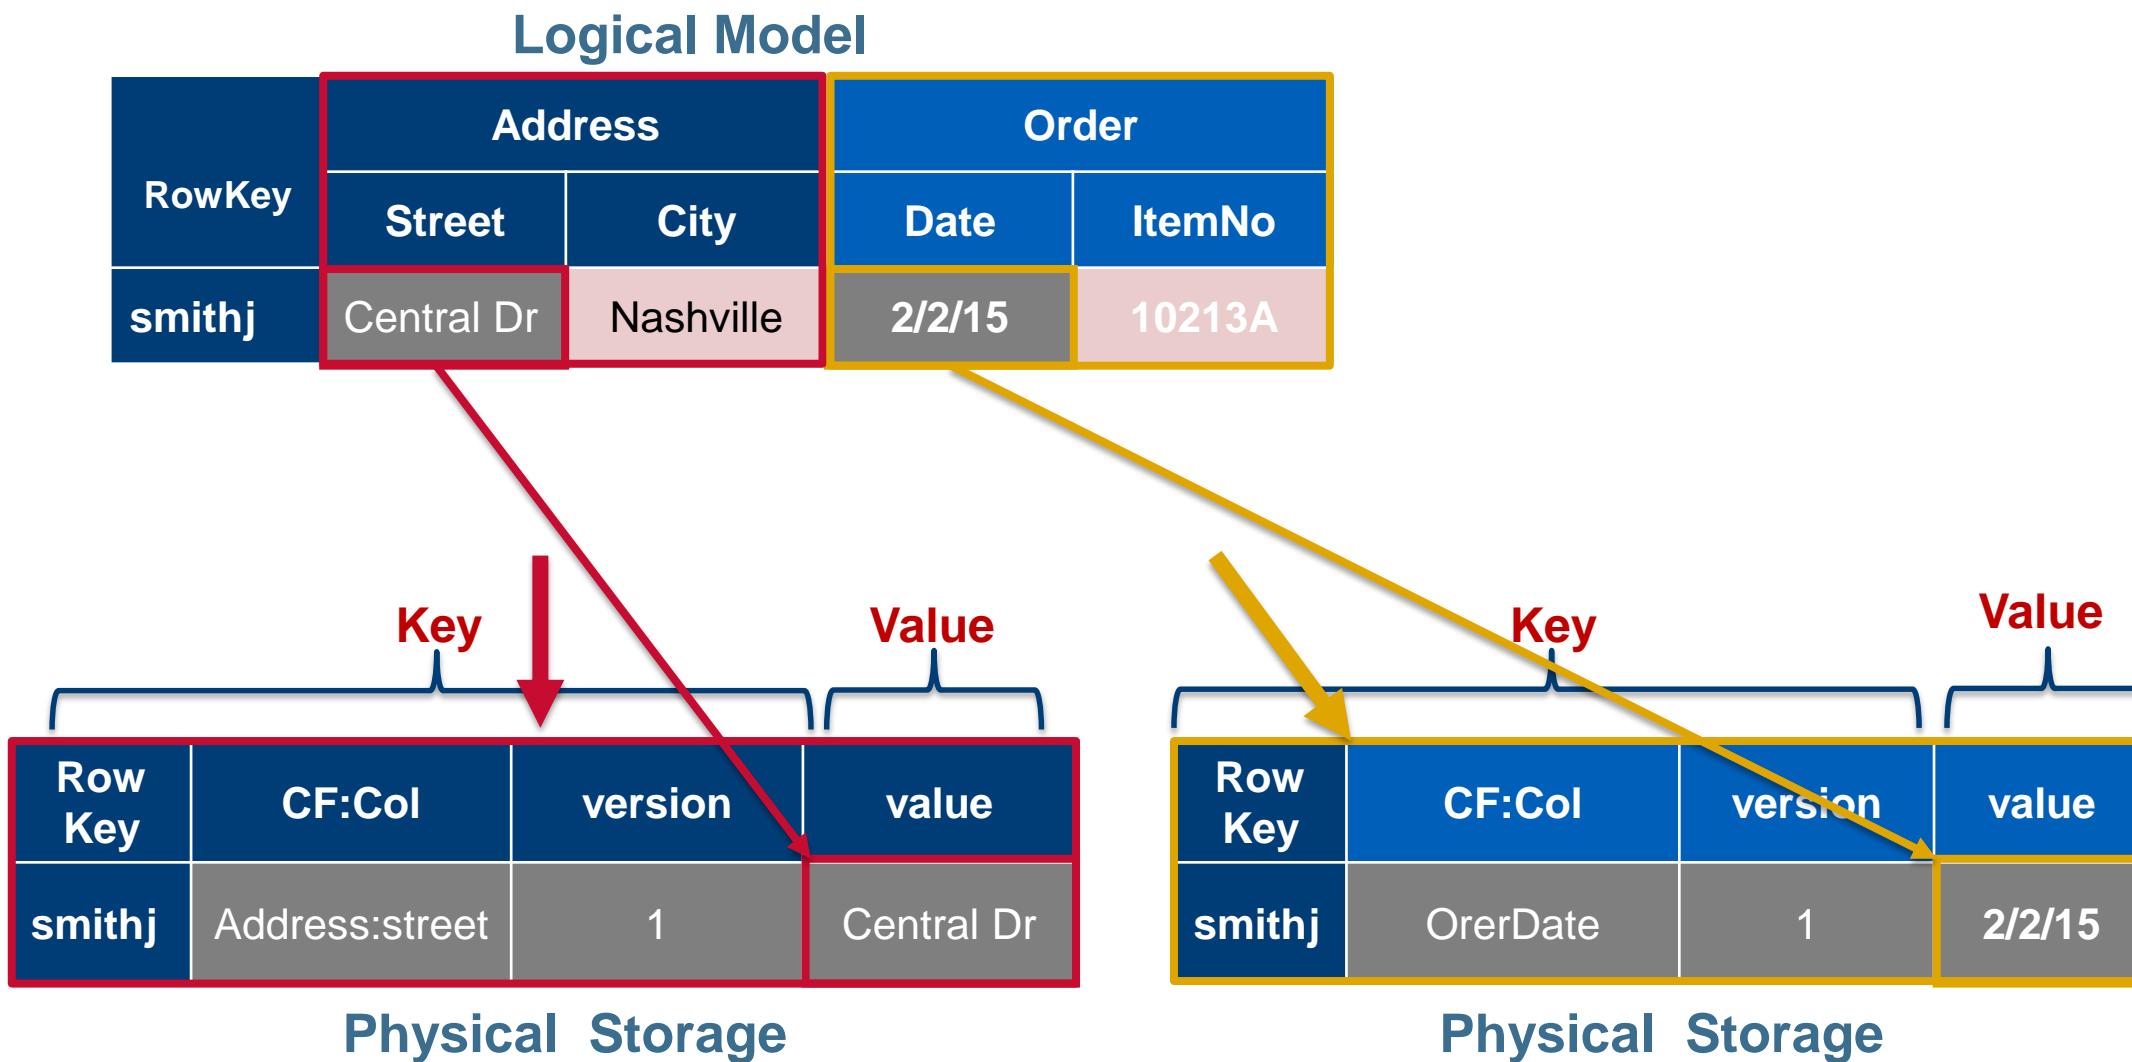
- Data is stored in **Key Value** format
- Value for each **cell** is specified by complete coordinates:
 - (Row key, ColumnFamily, Column Qualifier, timestamp) => Value

Cell Coordinates= **Key**

Row key	Column Family	Column Qualifier	Timestamp	Value
Smithj	Address	city	1391813876369	Nashville



Logical Data Model vs Physical Data Storage





Sparse Data with Cell Versions

	CF1:colA	CF1:colB	CF1:colC
Row1	@time7: value3 @time6: value1 @time5: value1		
Row10	@time2: value1	@time2: value1	
Row11	@time6: value2 @time5: value1		
Row2	@time4: value1		@time4: value1



Versioned Data

Number of versions can be configured. Default number equal to 1

Key	CF:Col	version	value
smithj	Address:street	v3	19 th Ave
smithj	Address:street	v2	Main St
smithj	Address:street	v1	Central Dr

Logical Data Model vs Physical Data Storage

Row Key	CF1		CF2	
	ca	cb	ca	cd
ra	1		2	
rb		3,4		
rc	5		6,7	8

Logical Model

Column families are stored separately
Row keys, Qualifiers are sorted lexicographically

Physical Storage

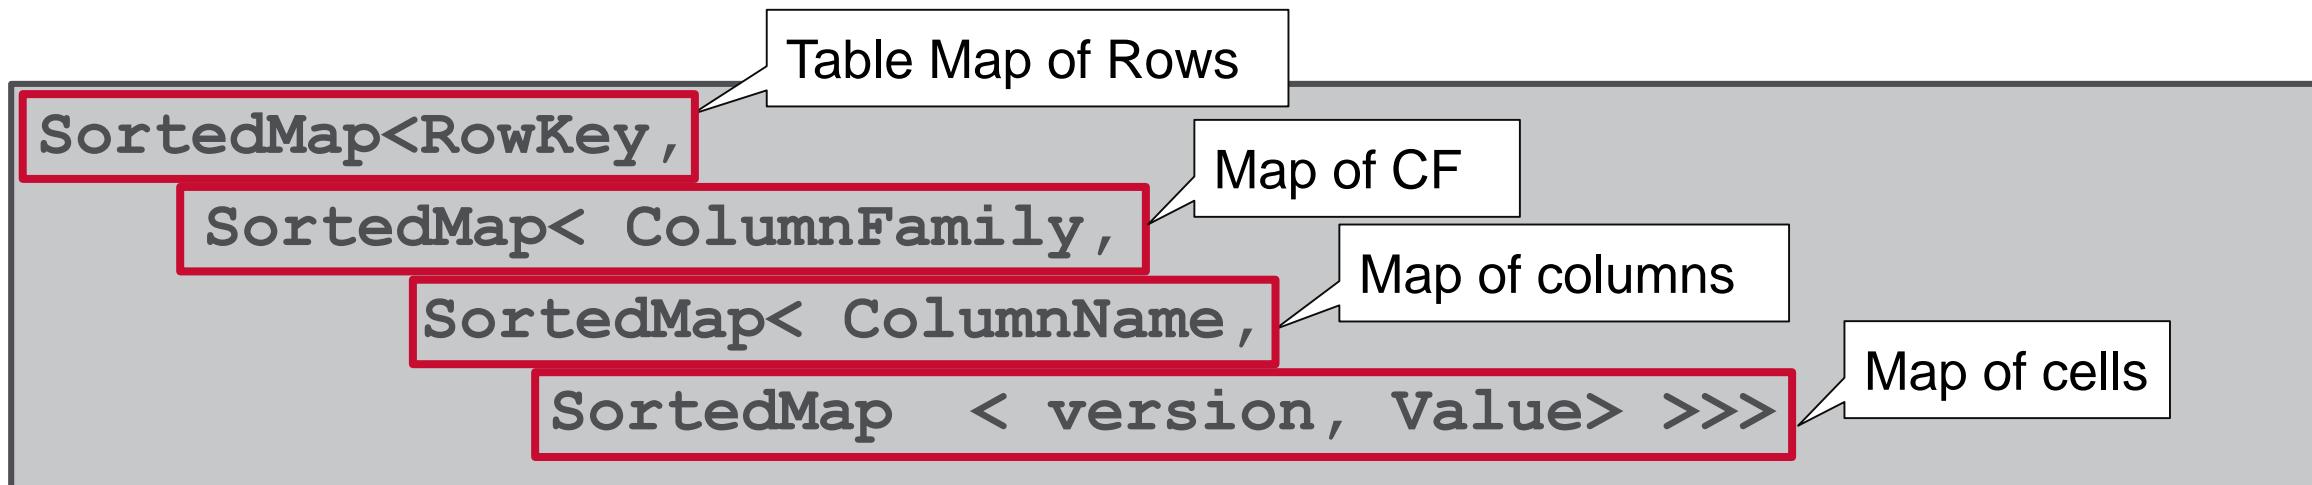
Key	Value	
	CF1:Col	version
ra	cf1:ca	1
rb	cf1:cb	2
rb	cf1:cb	1
rc	cf1:ca	1

Key	Value	
	CF2:Col	version
ra	cf2:ca	1
rc	cf2:ca	2
rc	cf2:ca	1
rc	cf2:cd	1



HBase Table is a Sorted Map of Maps

In Java this is the table:



Key	CF:Col	version	value
smithj	Address:street	v2	Main St
smithj	Address:street	v1	Central Dr
spata	Address:street	v1	High Ave
turnerb	Address:street	v1	Cedar St

Key	CF:Col	version	value
smithj	Order:Date	v1	2/2/15
spata	Order:Date	v1	1/31/14
turnerb	OrderDate	v1	7/8/14



HBase Table - SortedMap

```
<smithj,<Address, <street, <v1, Central Dr>>
    <street, <v2, Main St>>
    <Order <Date, <v1, 2/2/15>>>
<spata,<Address, <street, <v1,High Ave>>
    <Order <Date, <v1, 1/31/14>>>
<turnerb,<Address, <street, <v1,Cedar St>>
    <Order <Date, <v1, 7/8/14>>>
```

Key	CF:Col	version	value
smithj	Address:street	v2	Main St
smithj	Address:street	v1	Central Dr
spata	Address:street	v1	High Ave
turnerb	Address:street	v1	Cedar St

Key	CF:Col	version	value
smithj	Order:Date	v1	2/2/15
shawa	Order:Date	v1	1/31/14
turnerb	OrderDate	v1	7/8/14



Basic Table Operations

- **Create Table, define Column Families before data is imported**
 - but not the rows keys or number/names of columns
- Low level API, technically more demanding
- **Basic data access operations (CRUD):**

put	Inserts data into rows (both create and update)
get	Accesses data from one row
scan	Accesses data from a range of rows
delete	Delete a row or a range of rows or columns



Create HBase Table – Using HBase Shell

```
hbase> create '/user/user01/Customer', {NAME =>'Address' } , {NAME =>'Order' }

hbase> put '/user/user01/Customer', 'smithj', 'Address:street', 'Central Dr'

hbase> put '/user/user01/Customer', 'smithj', 'Order:Date', '2/2/15'

hbase> put '/user/user01/Customer', 'spata', 'Address:city', 'Columbus'

hbase> put '/user/user01/Customer', 'spata', 'Order:Date', '1/31/14'
```

Row Key	Address			Order	
	street	city	state	Date	ItemNo
smithj	Central Dr	Nashville	TN	2/2/15	10213A
spata	High Ave	Columbus	OH	1/31/14	23401V
turnerb	Cedar St	Seattle	WA	7/8/14	10938A



Create HBase Table – Using HBase Shell

```
hbase> get '/user/user01/Customer', 'smithj'
```

```
hbase> scan '/user/user01/Customer'
```

```
hbase> describe '/user/user01/Customer'
```



HBase Architecture

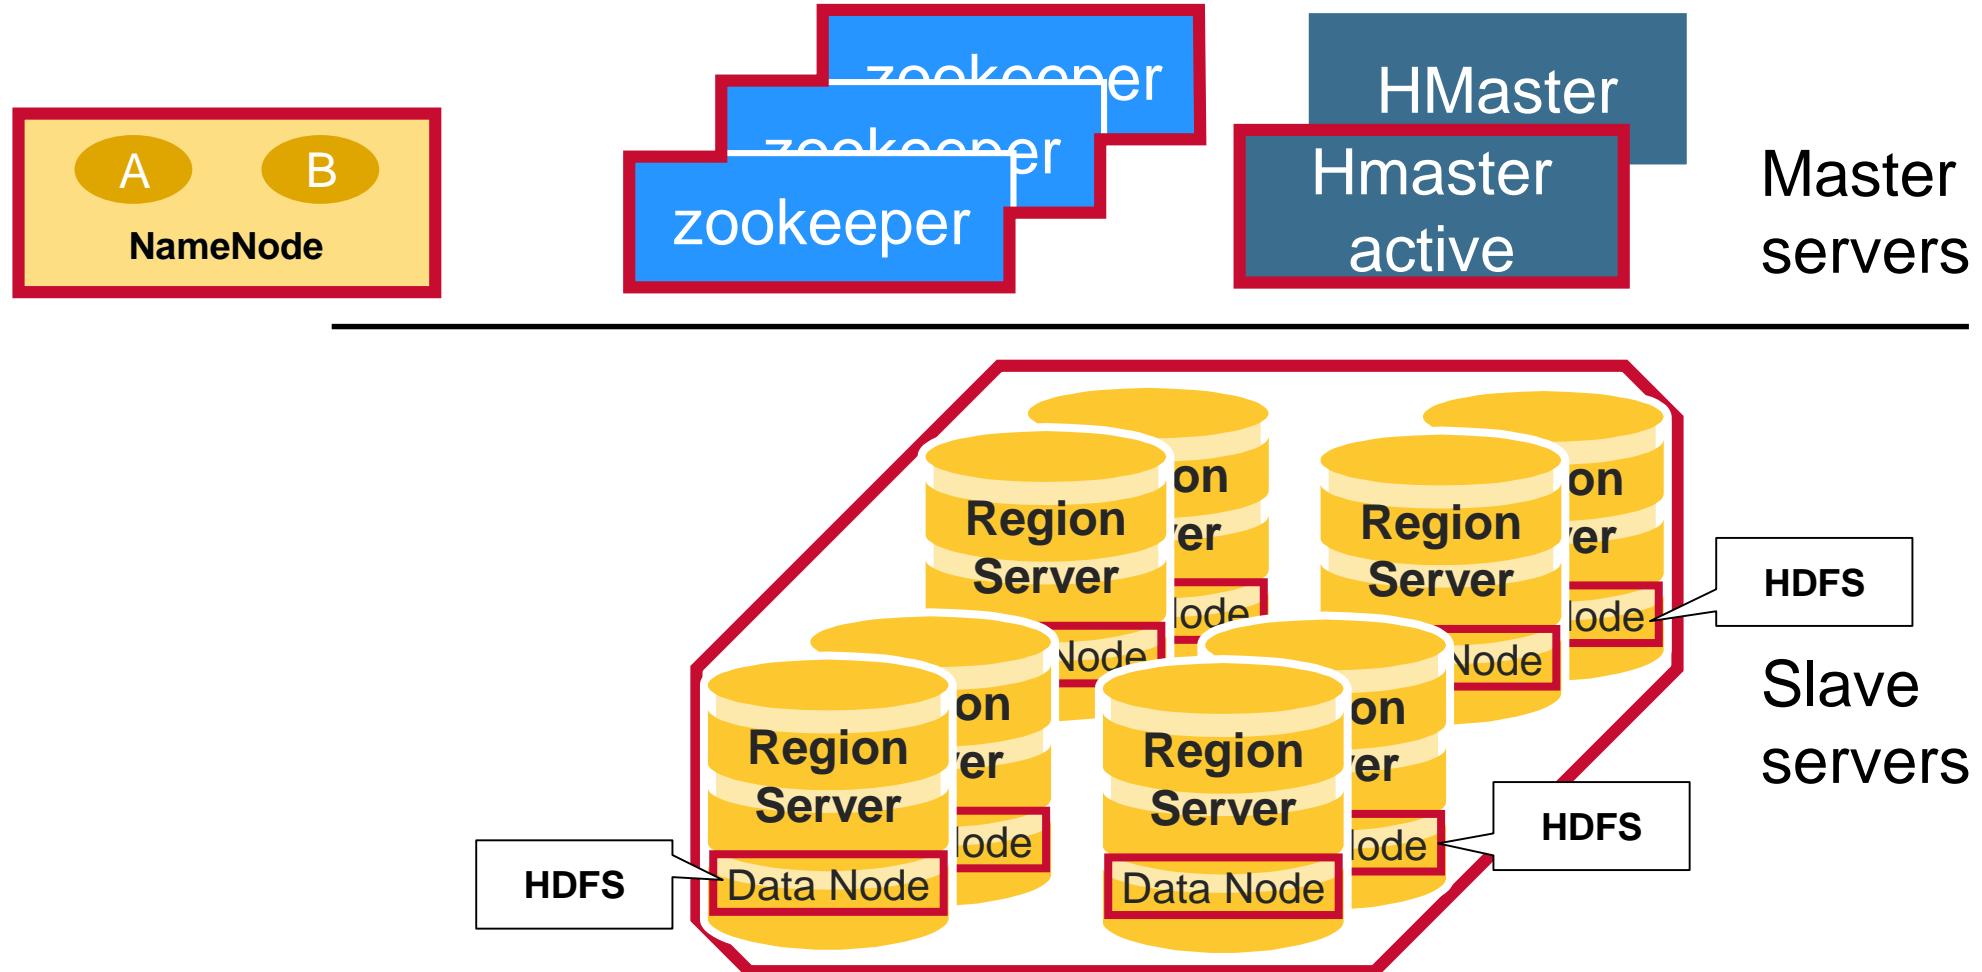
Data flow for Writes, Reads

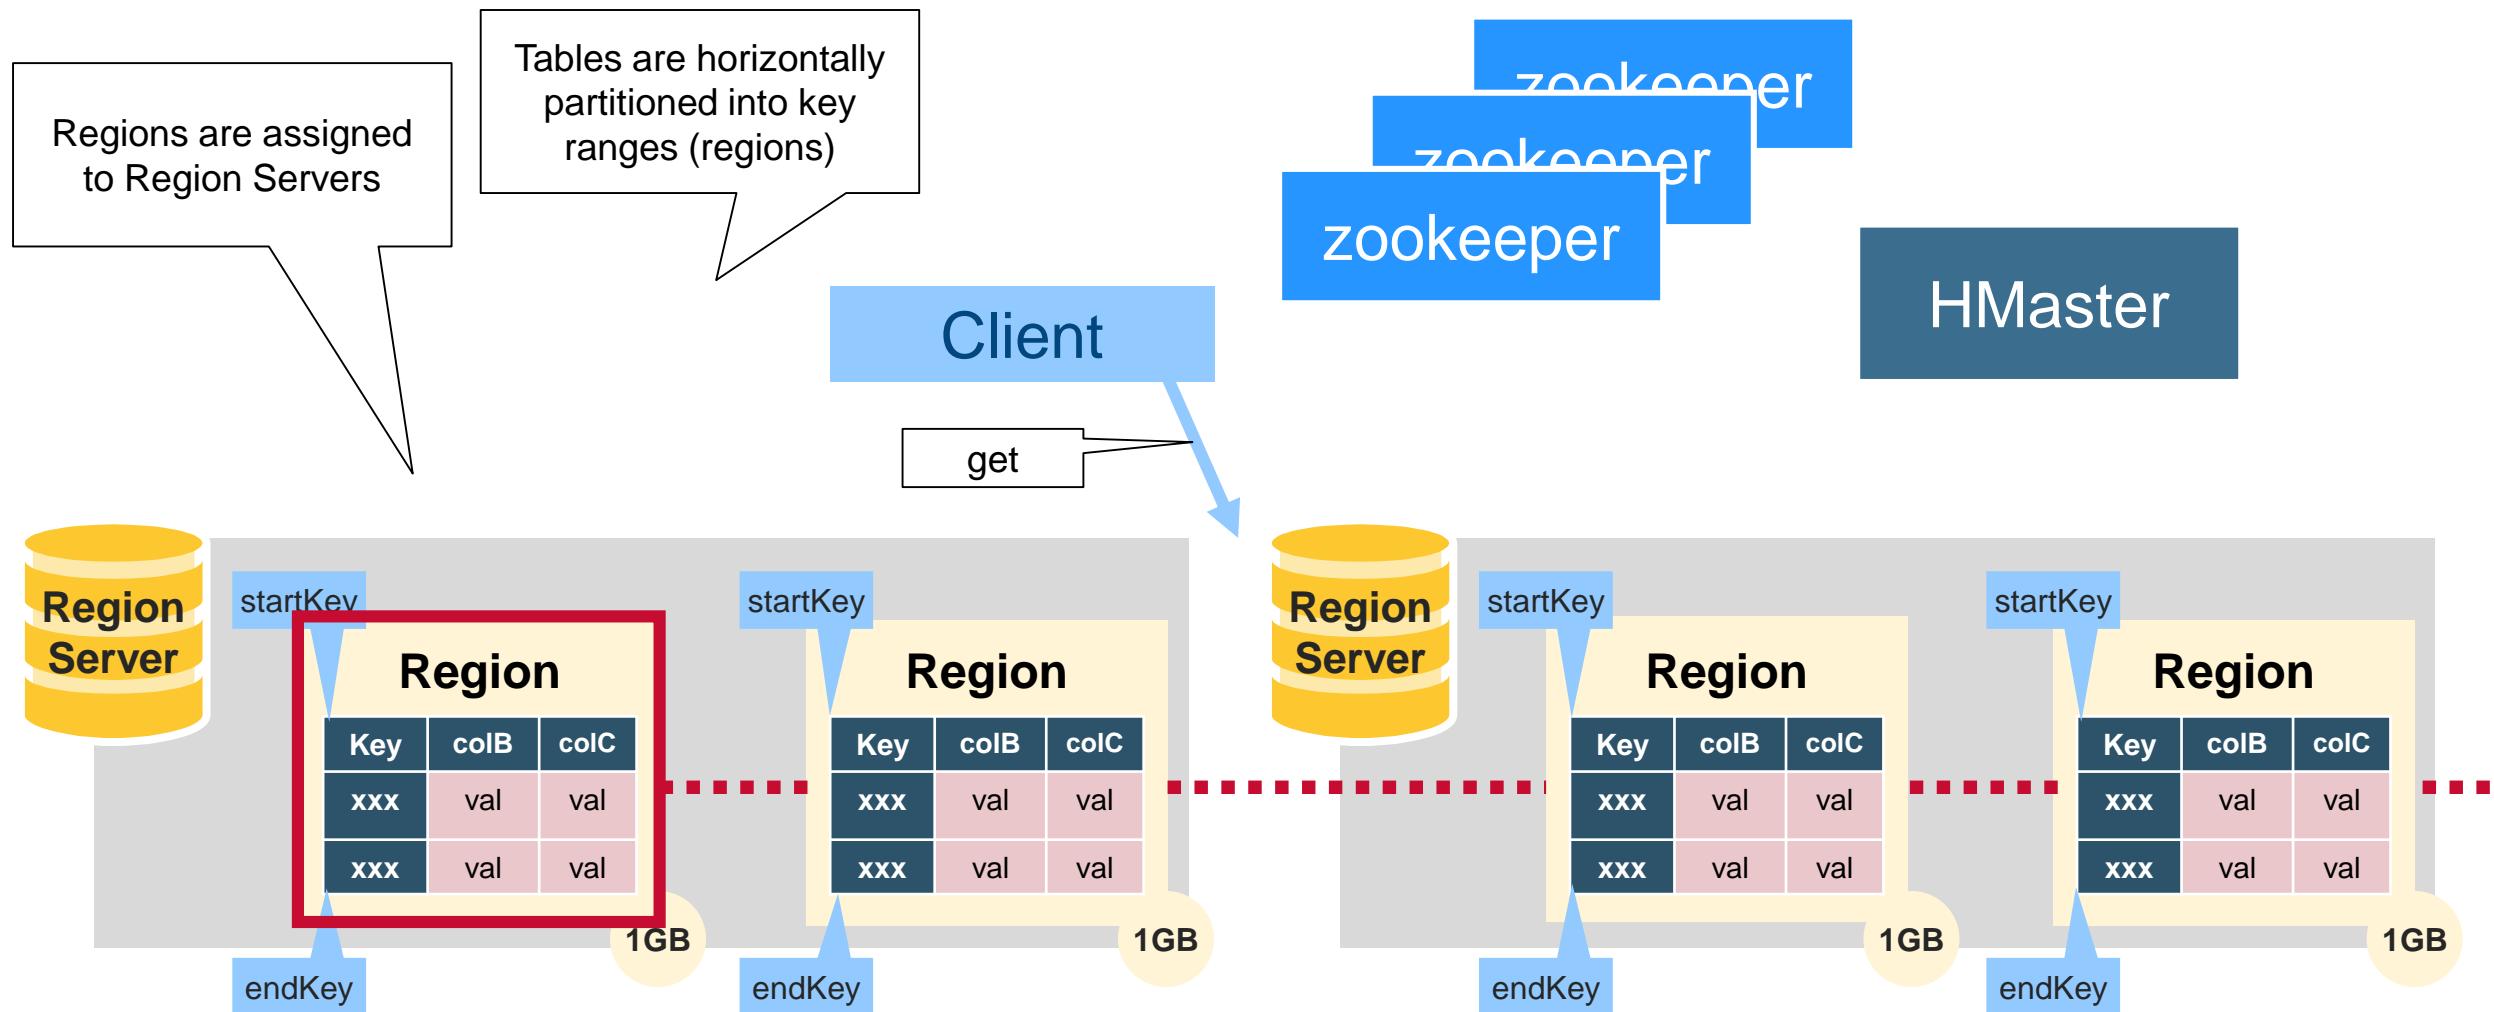
Designed to Scale





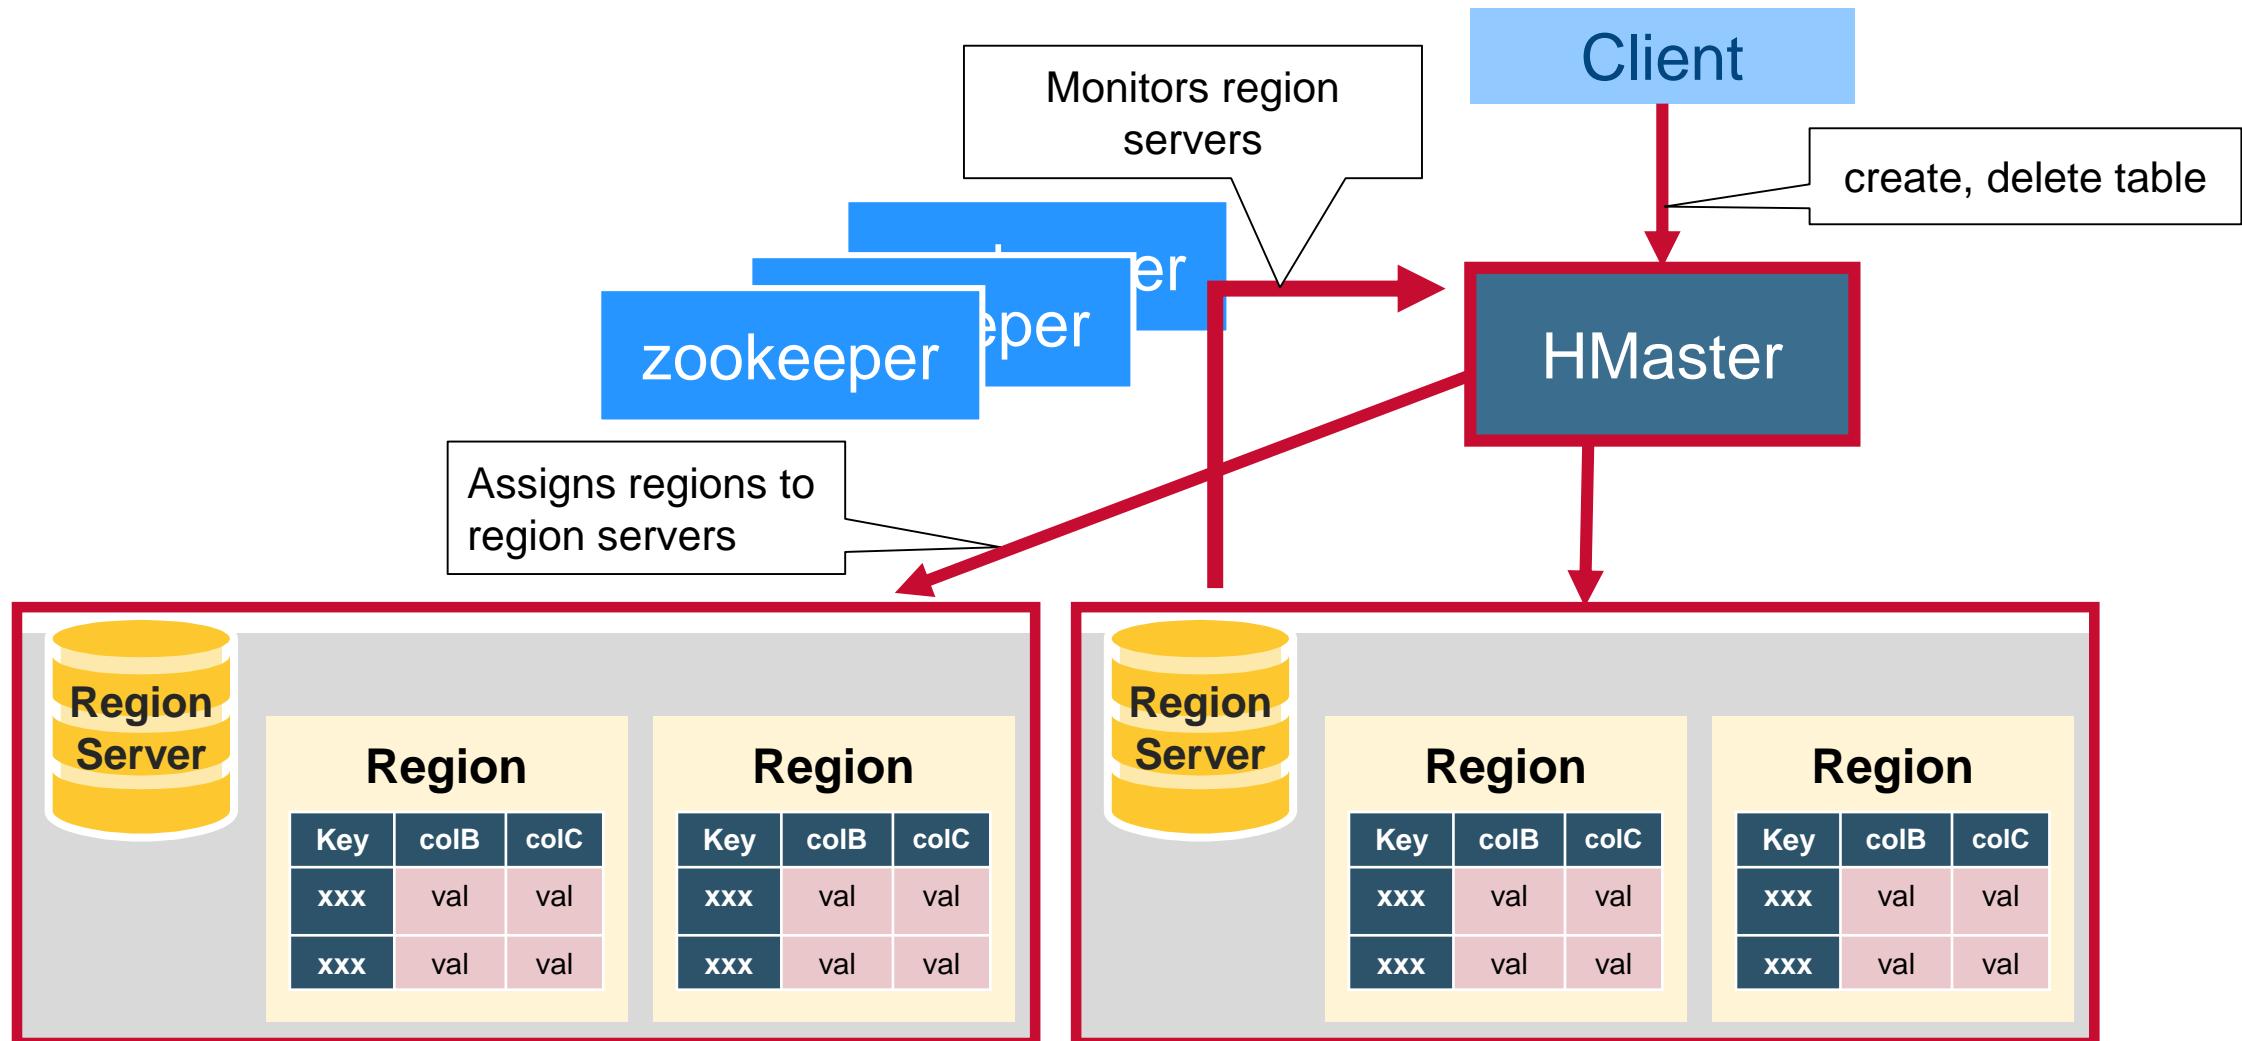
HBase Architectural Components



 Regions

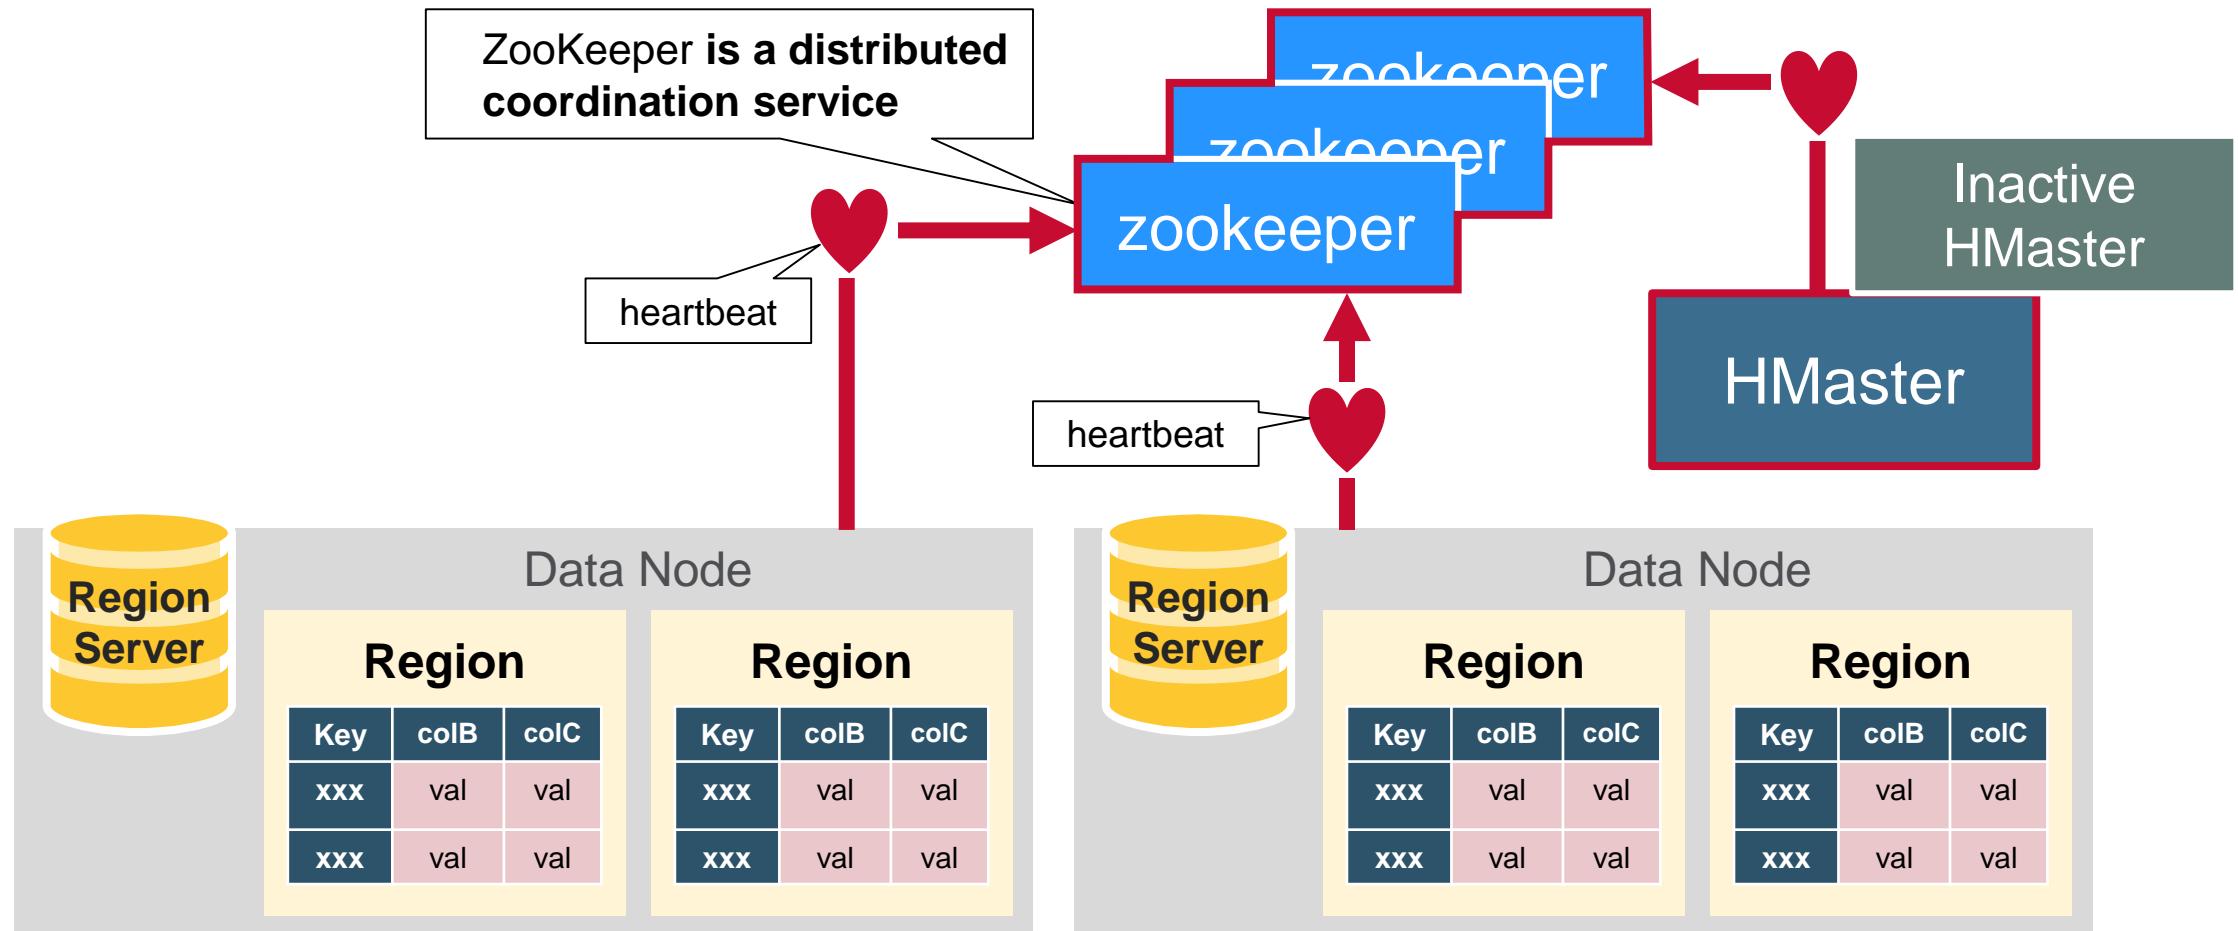


HBase HMaster



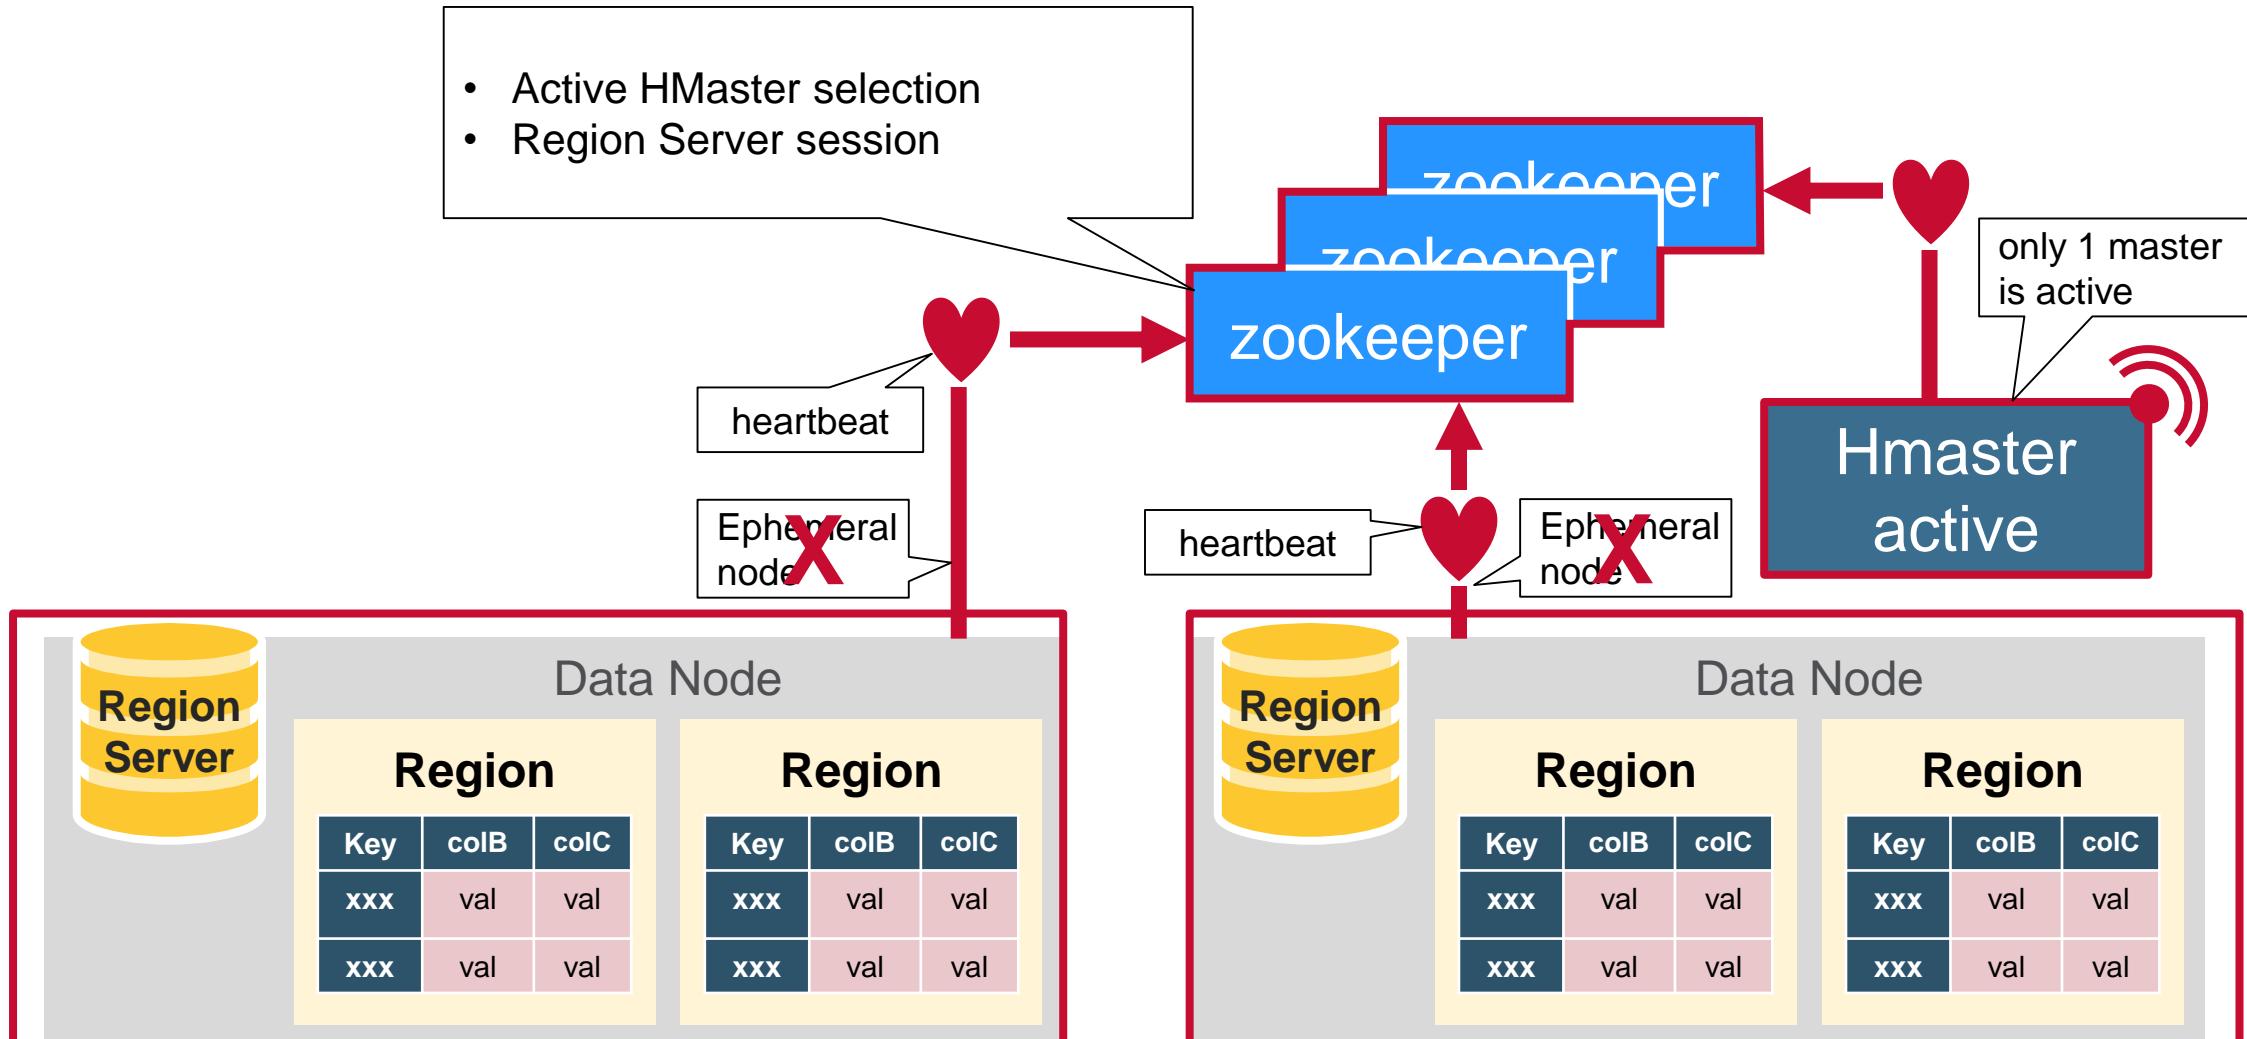


Zookeeper The Coordinator



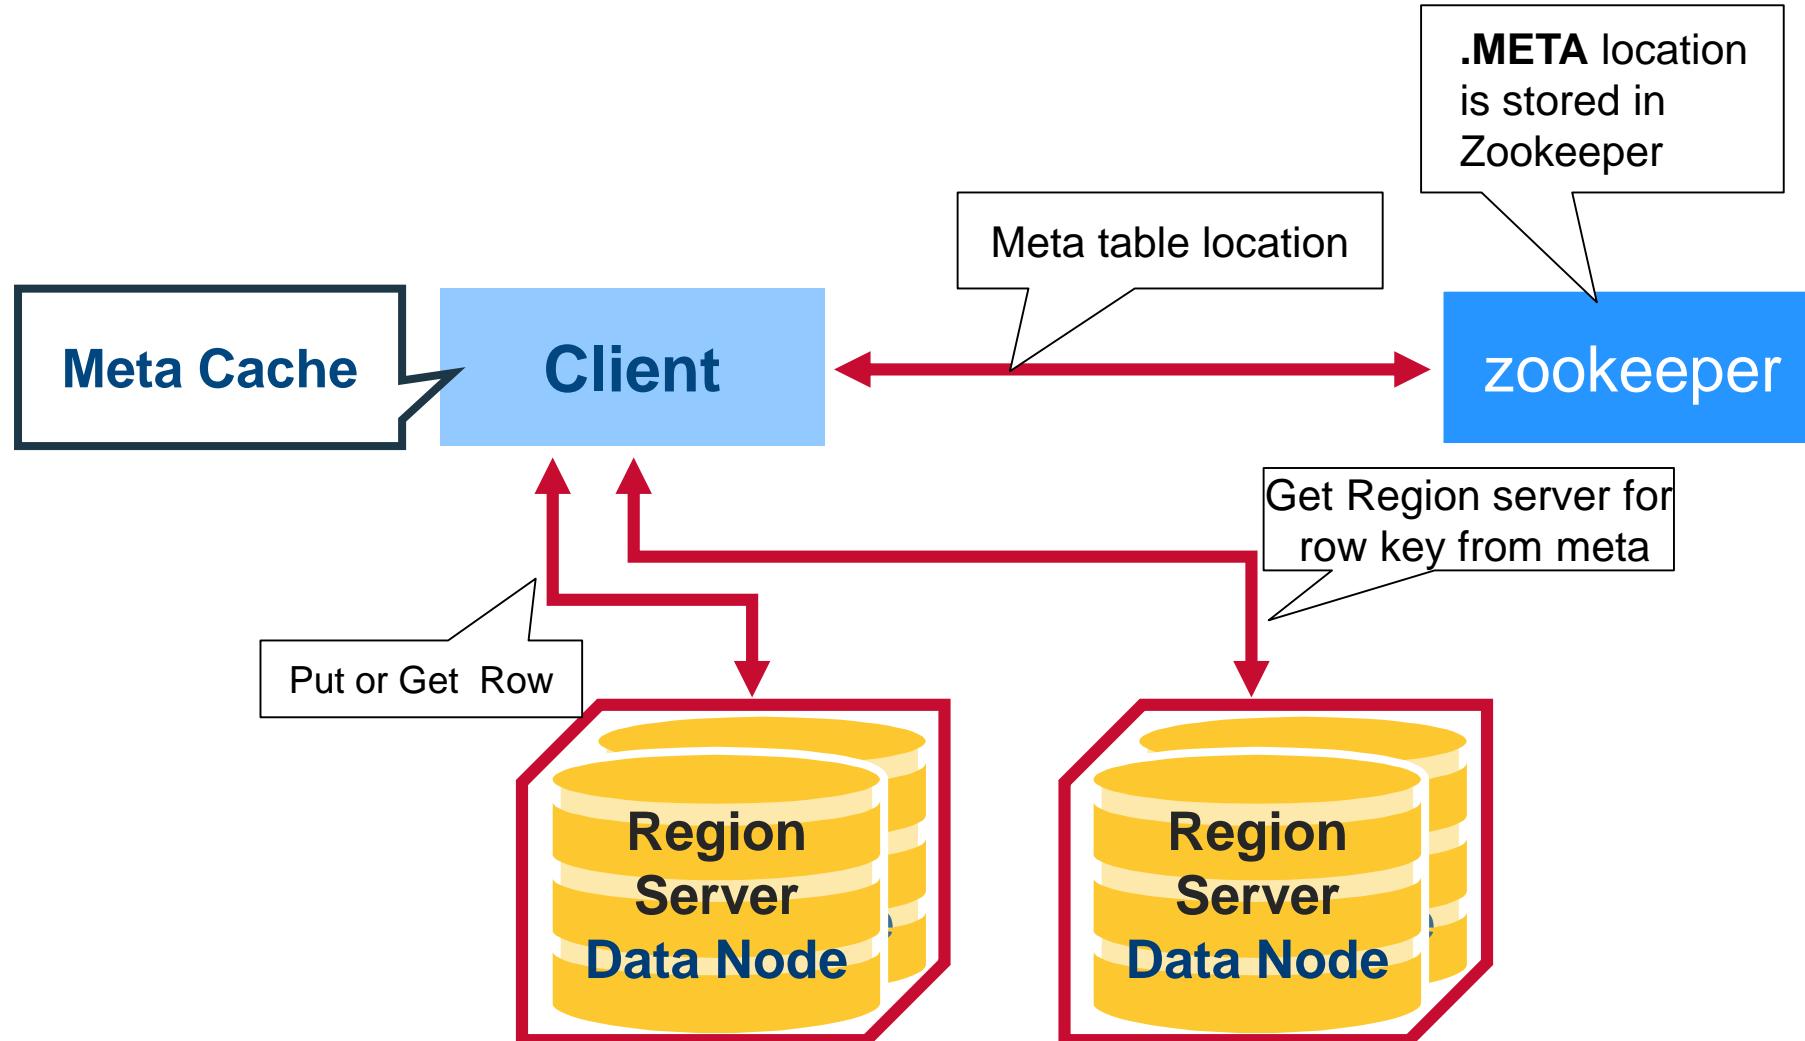


How the Components Work Together





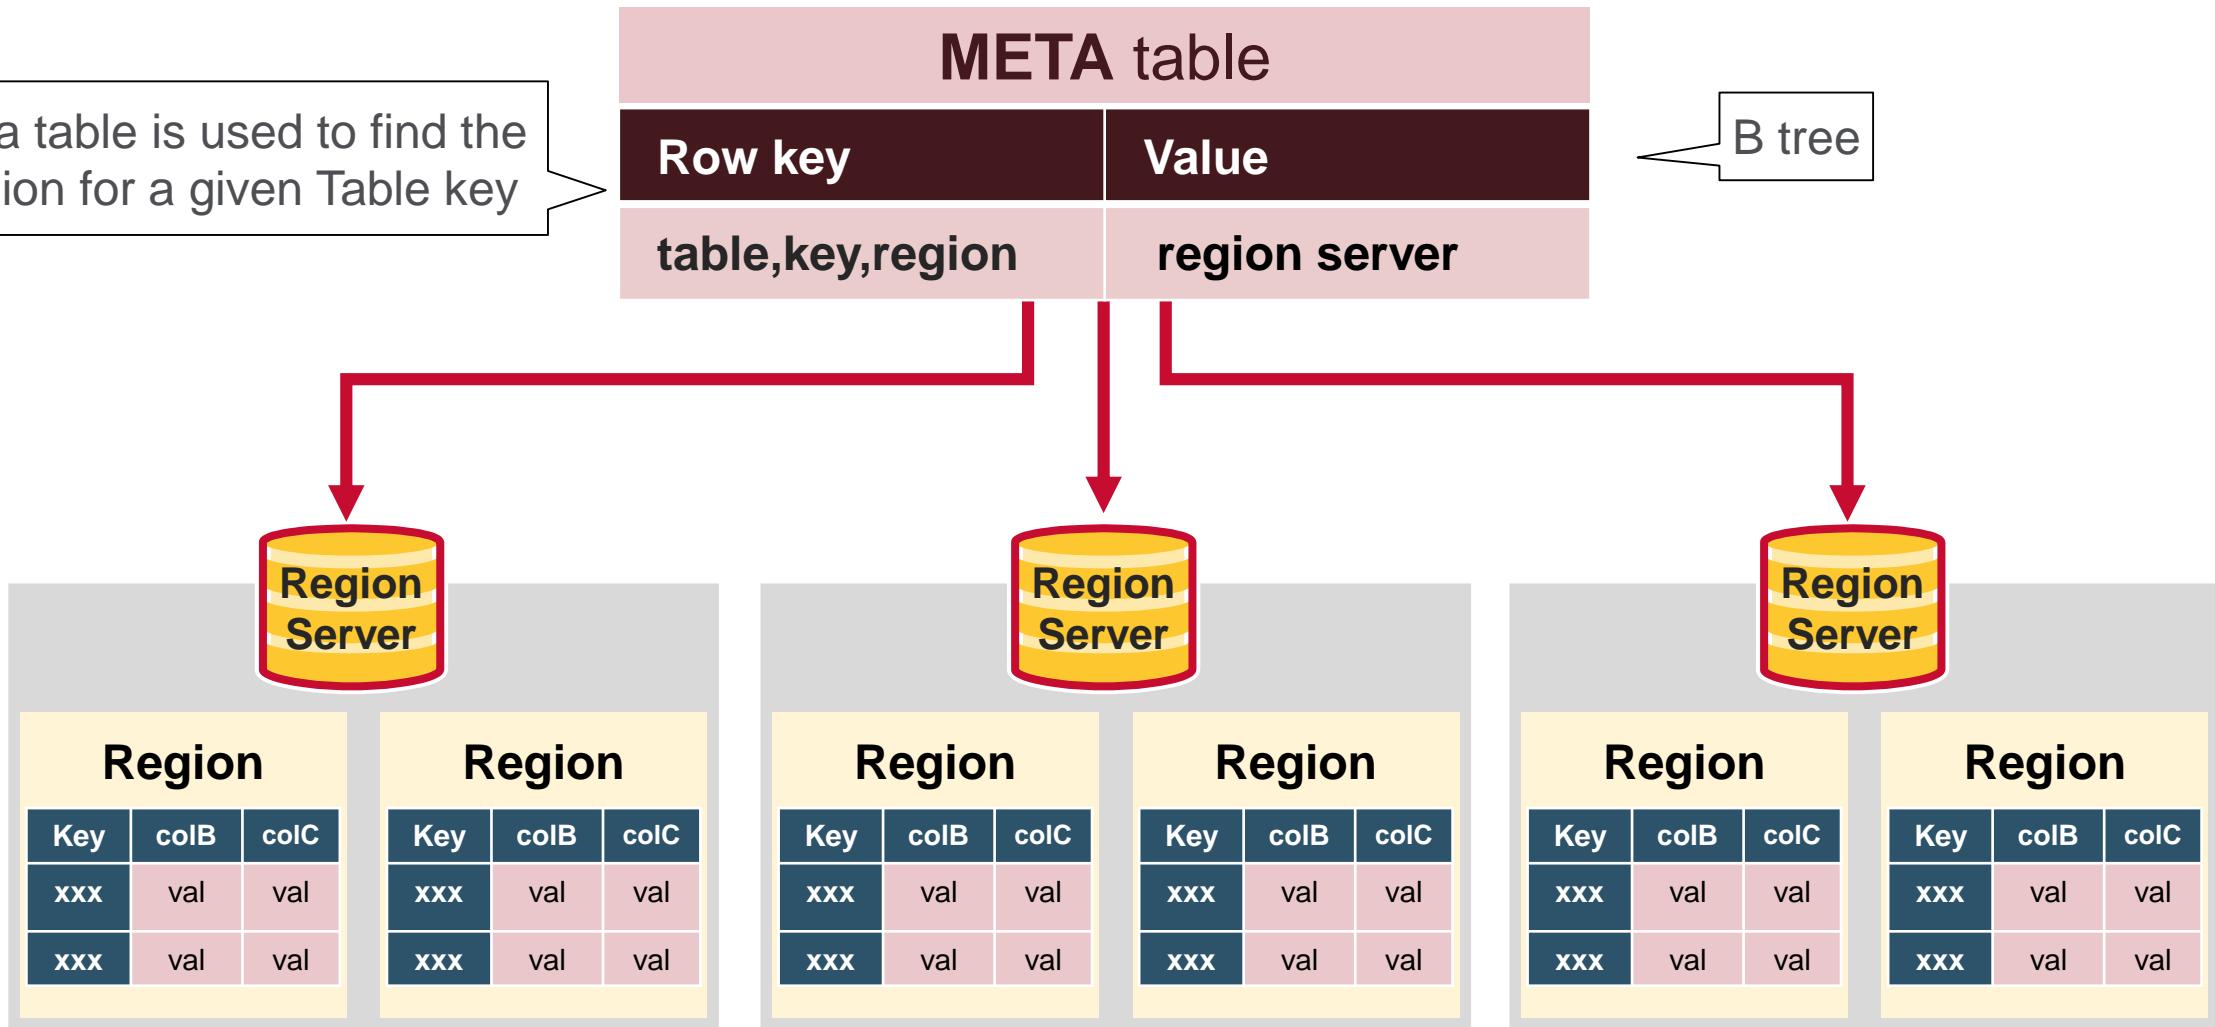
HBase First Read or Write





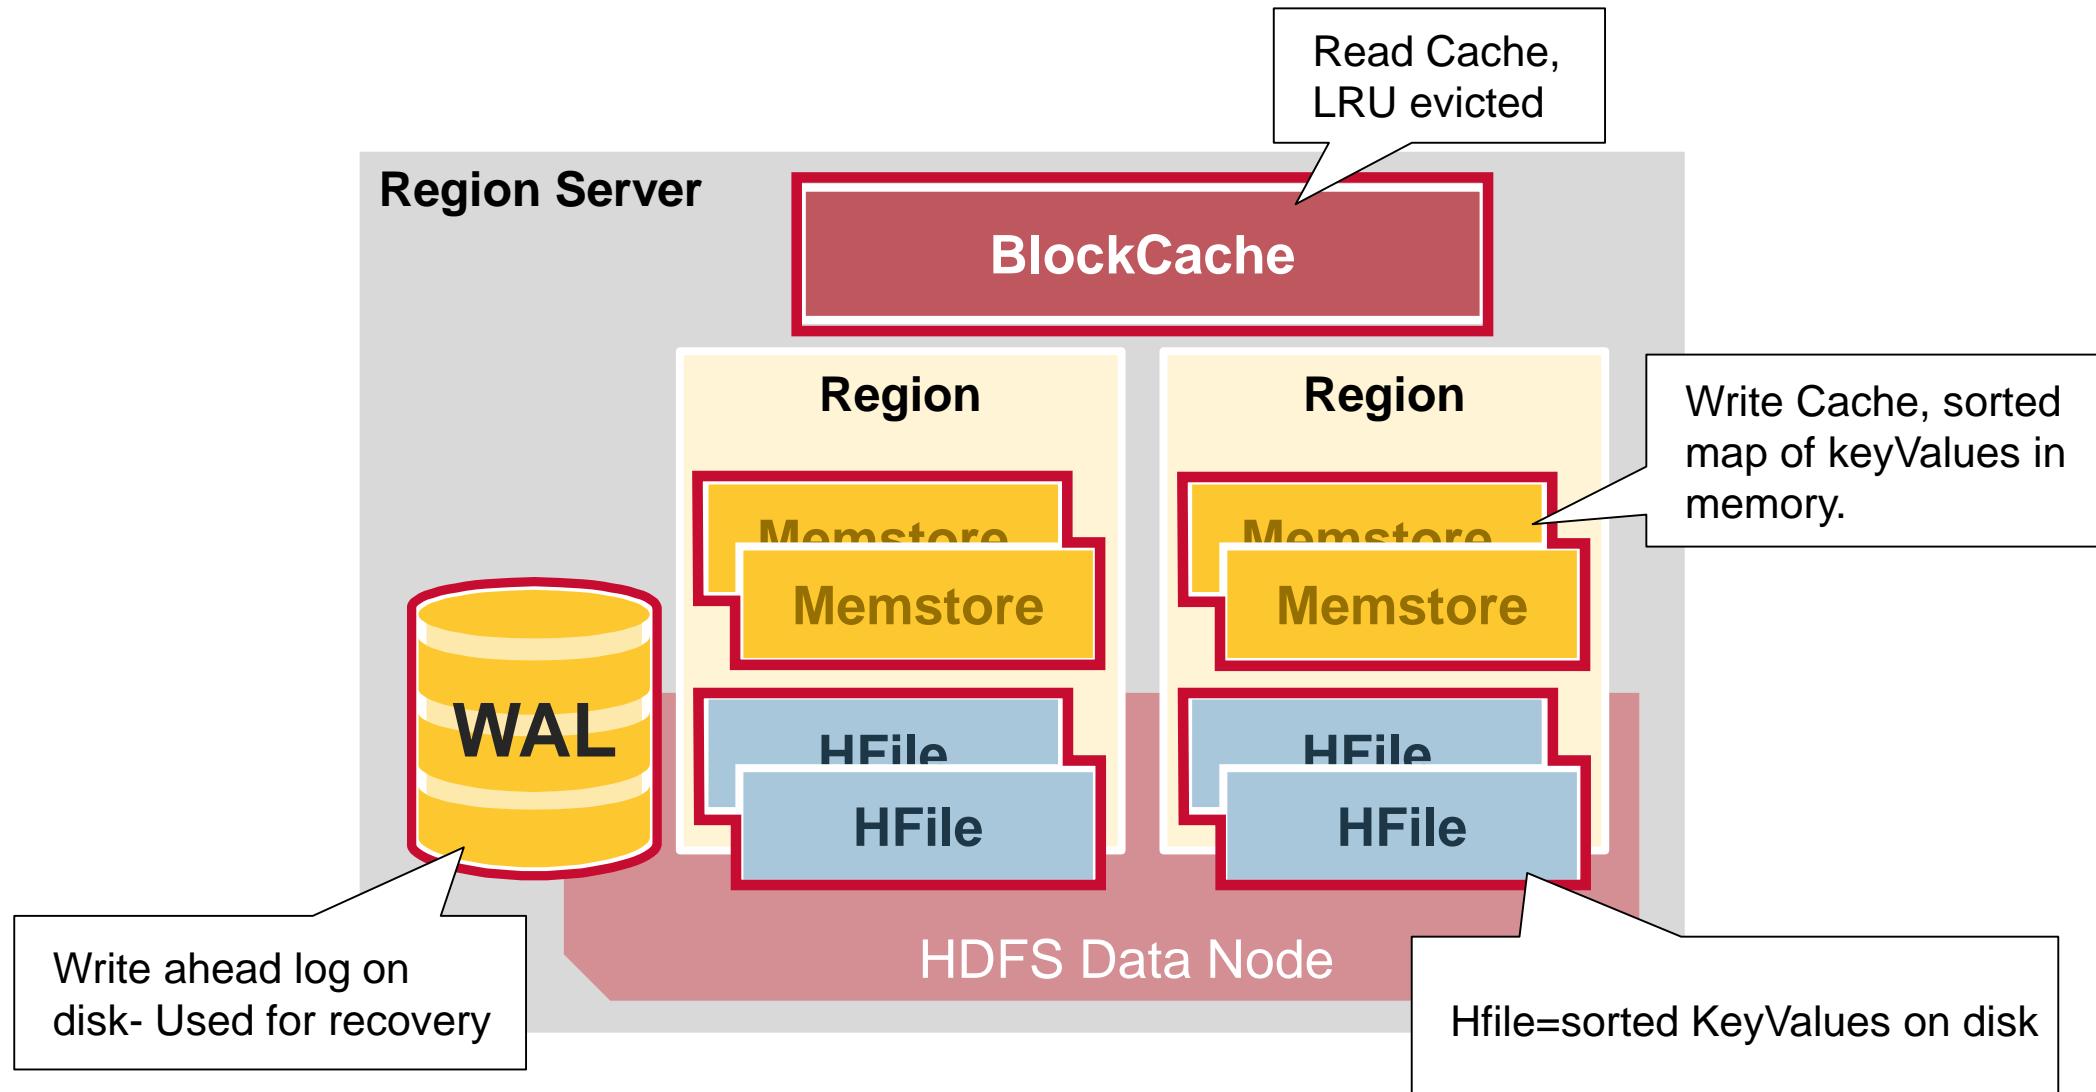
HBase Meta Table

Meta table is used to find the Region for a given Table key



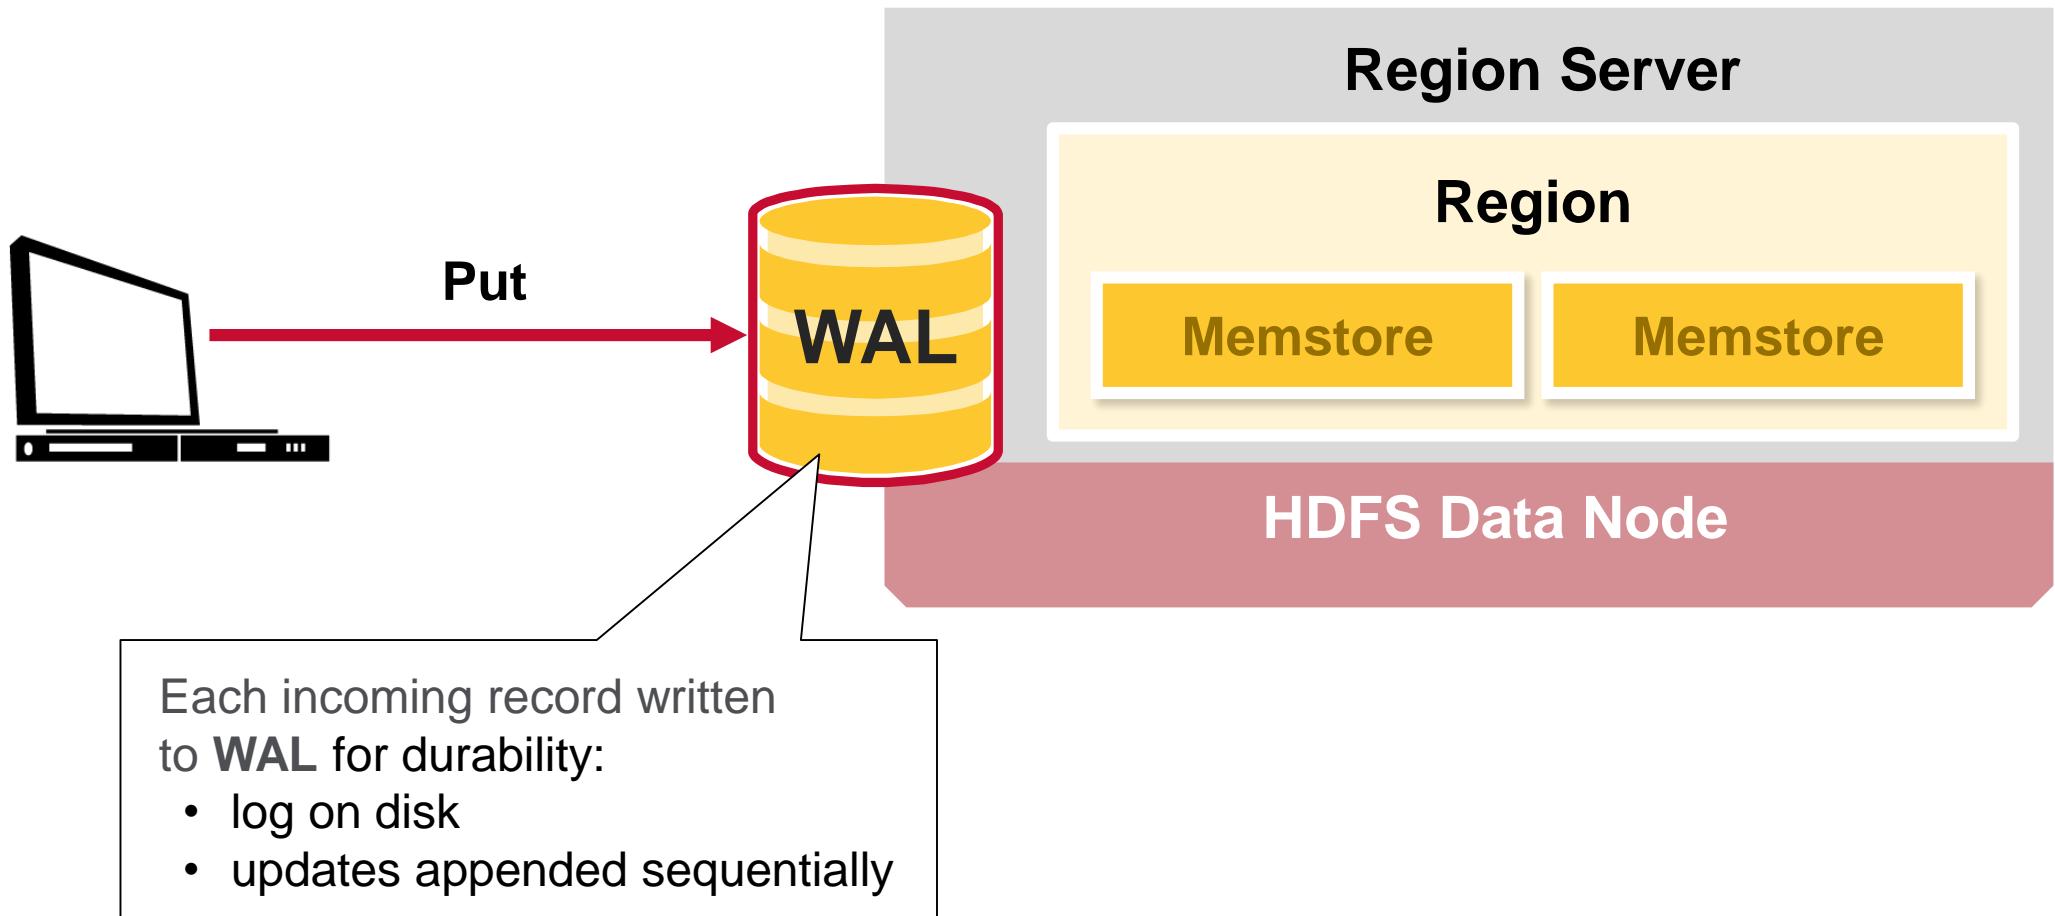


Region Server Components



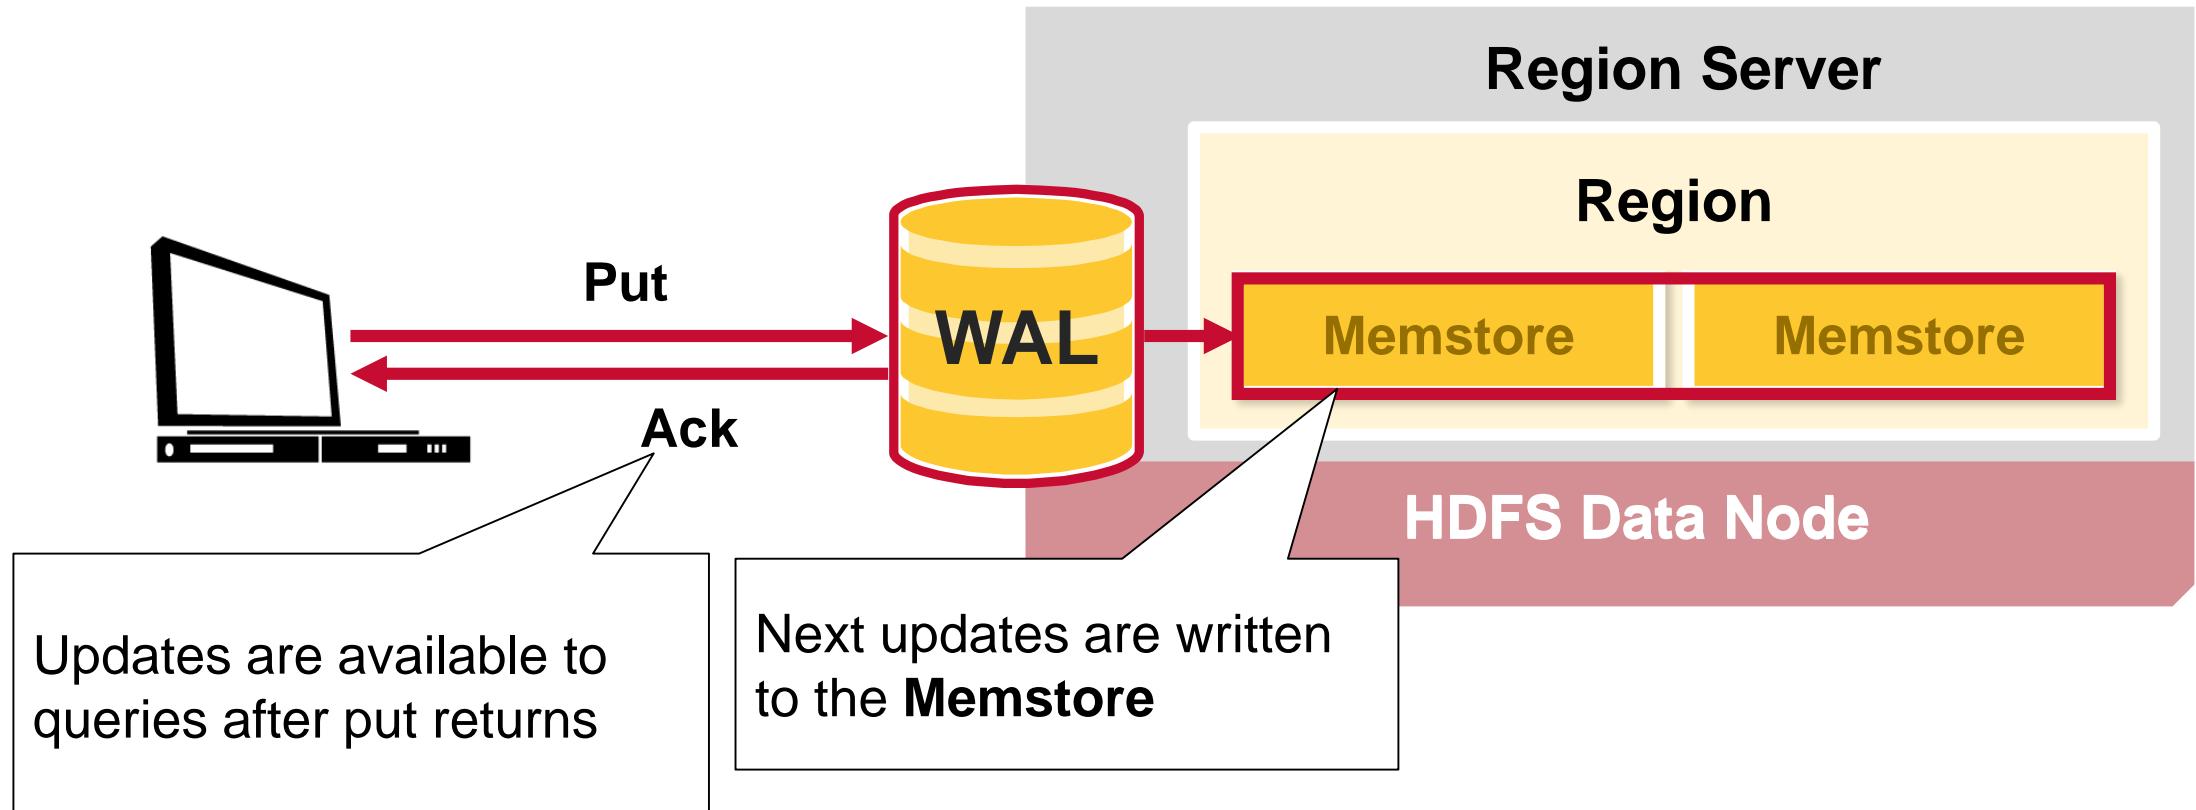


HBase Write Steps



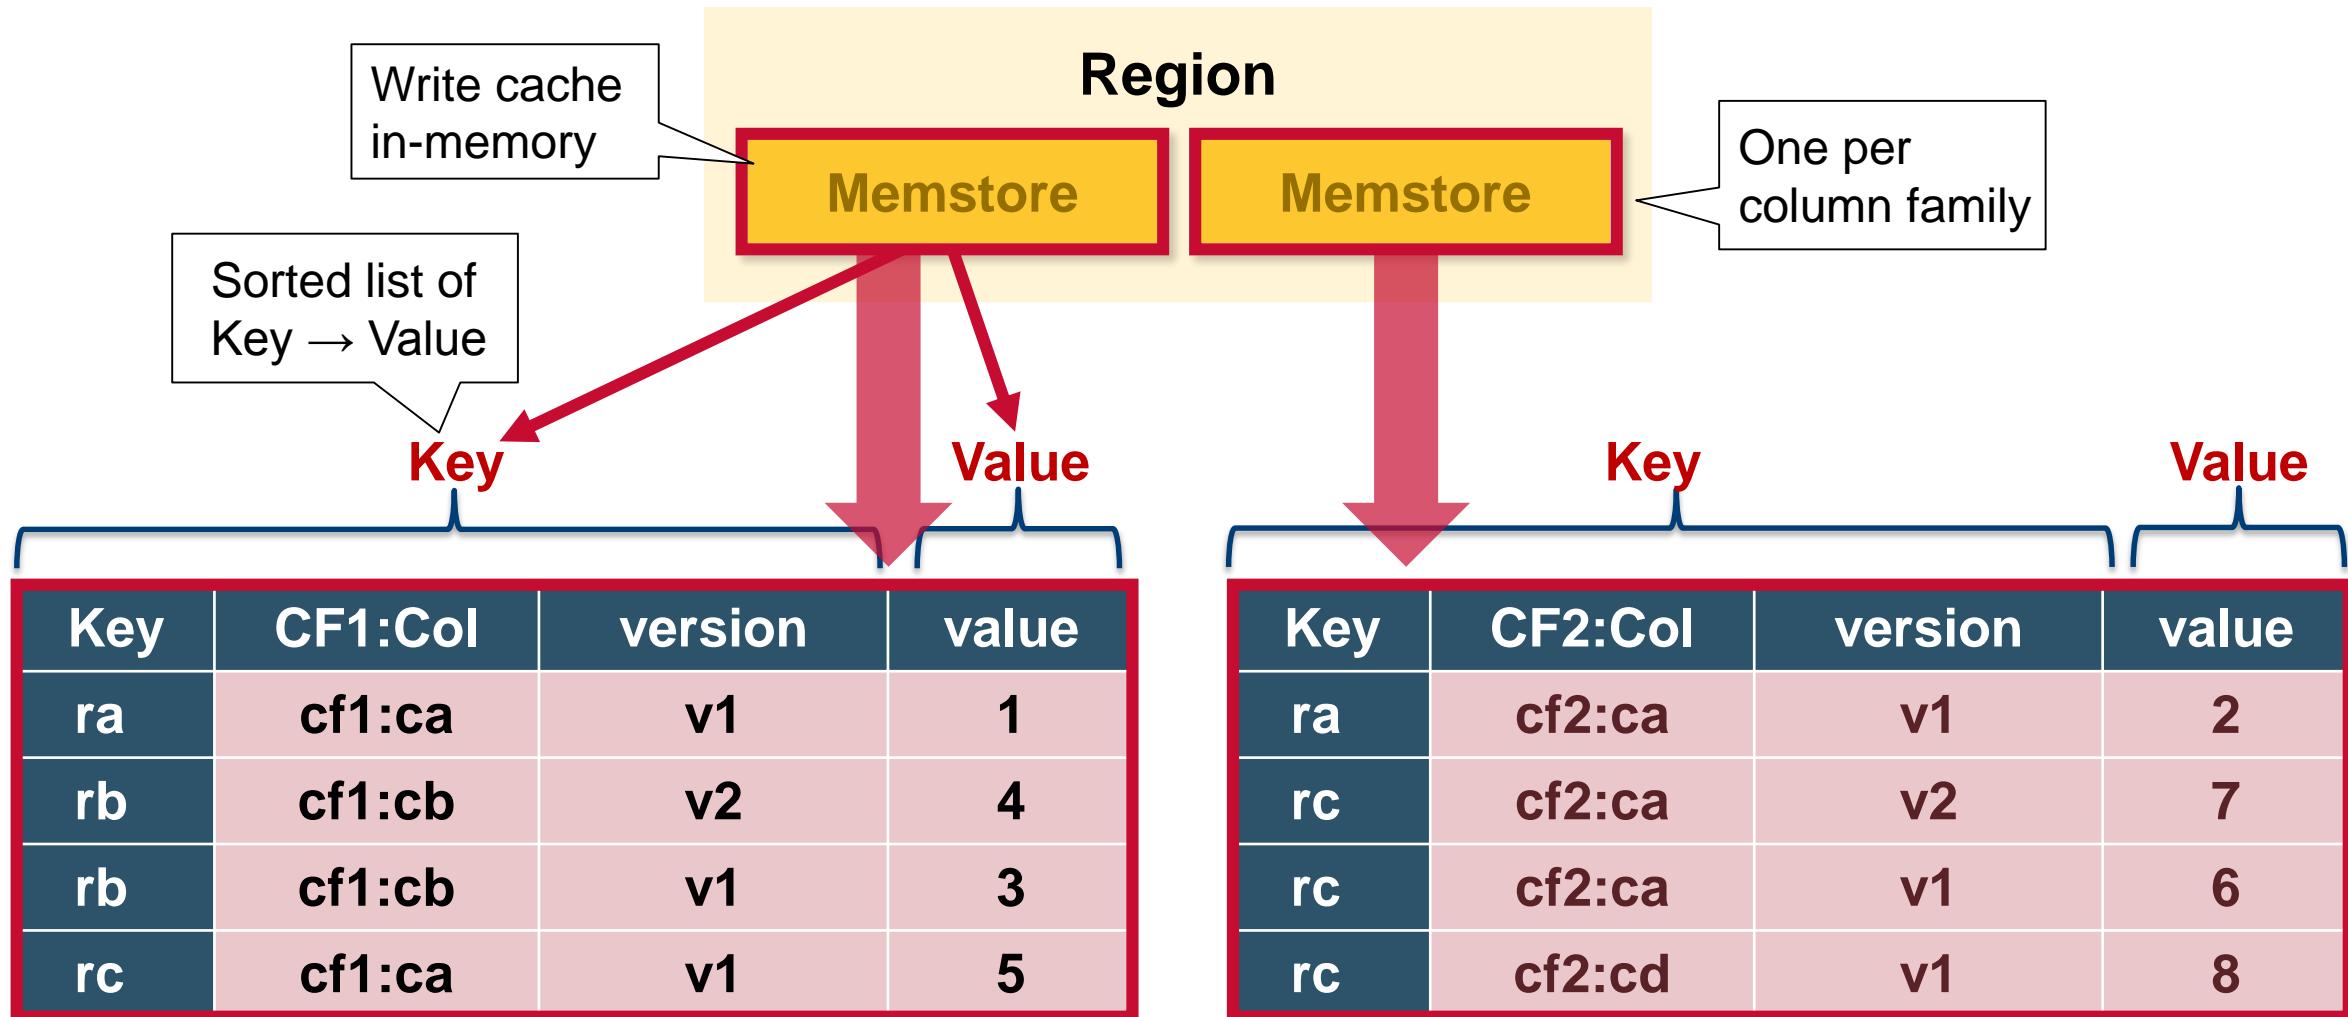


HBase Write Steps – (2)



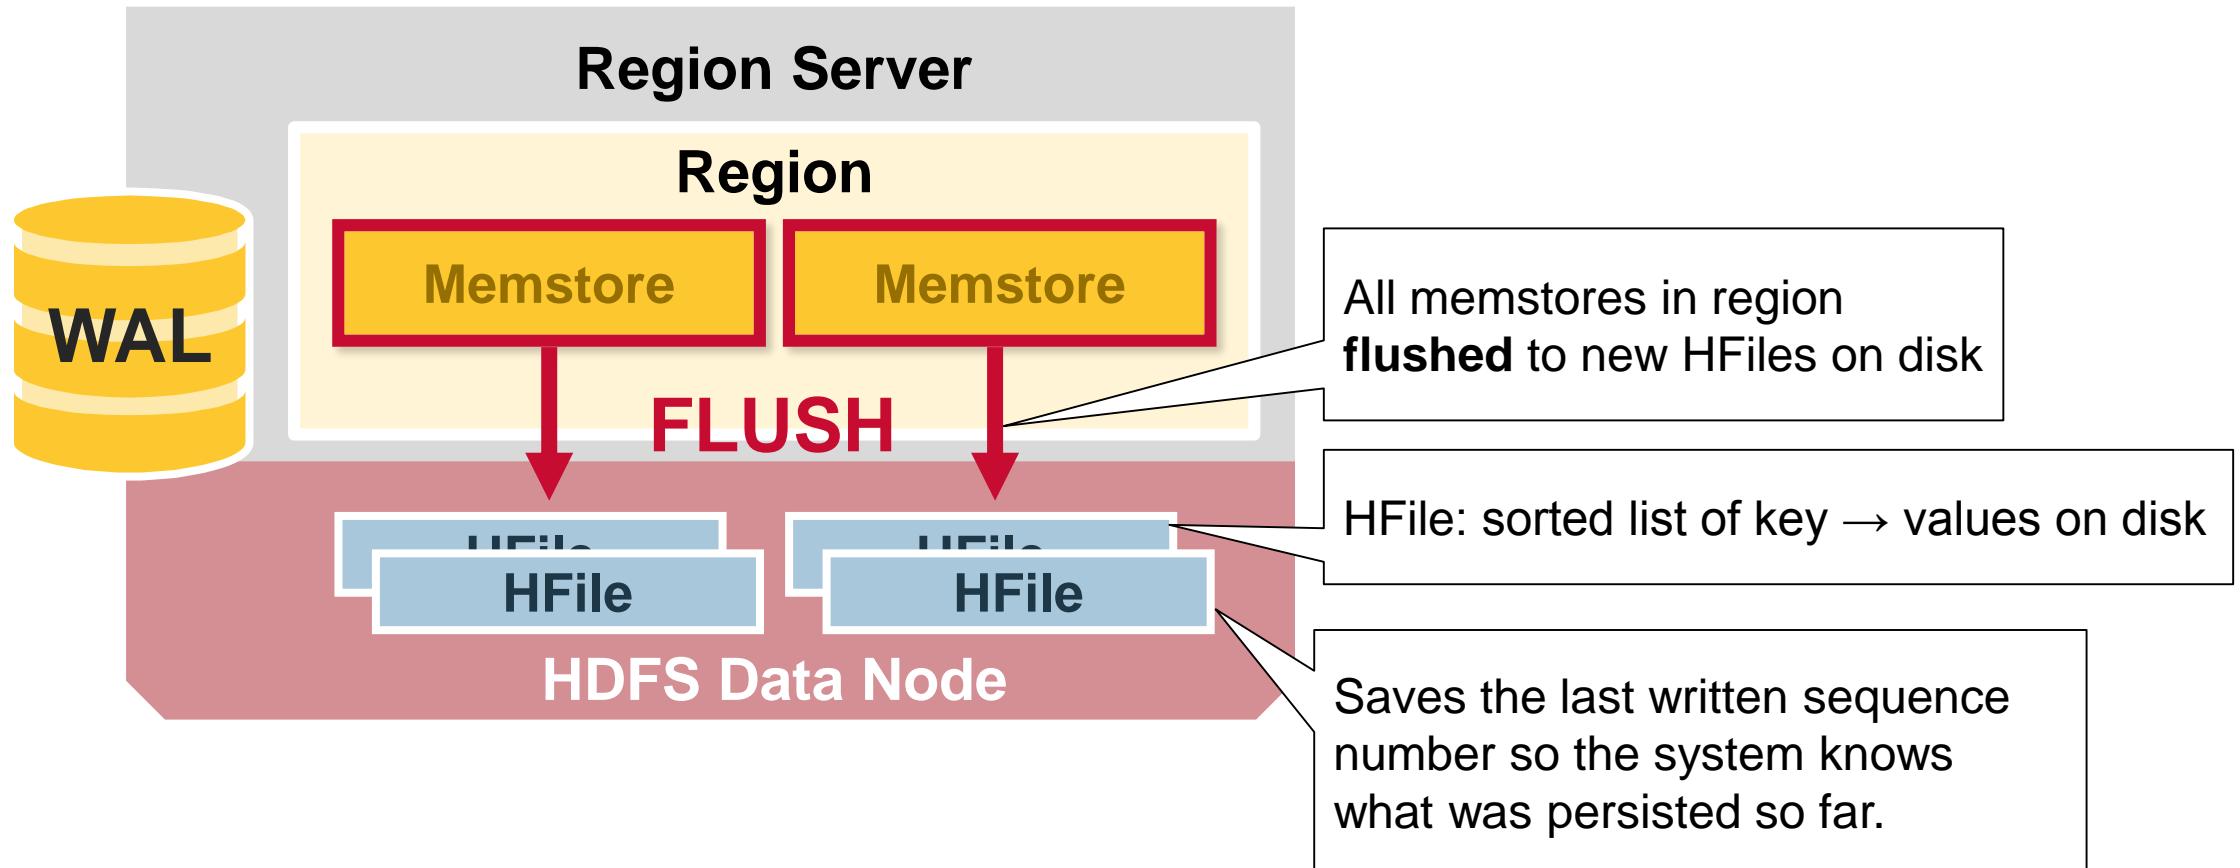


HBase Memstore



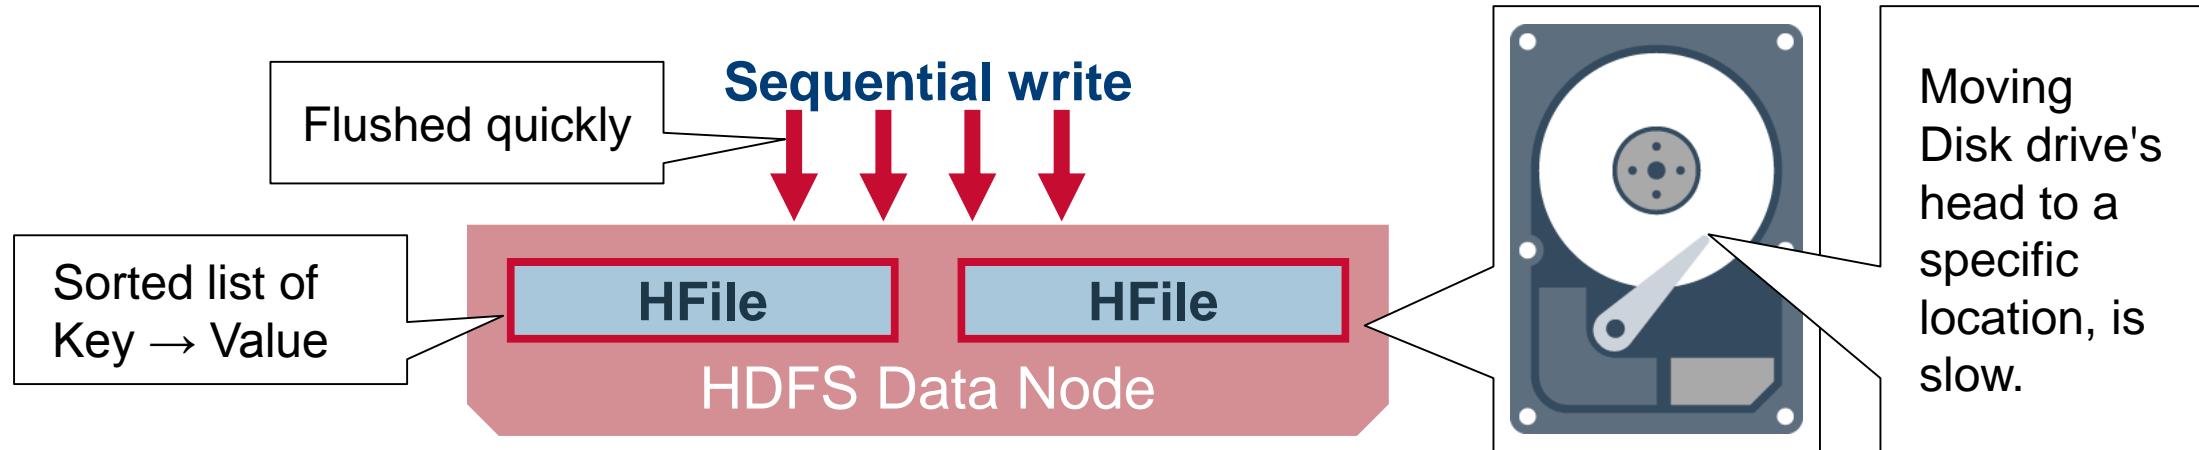


HBase Region Flush





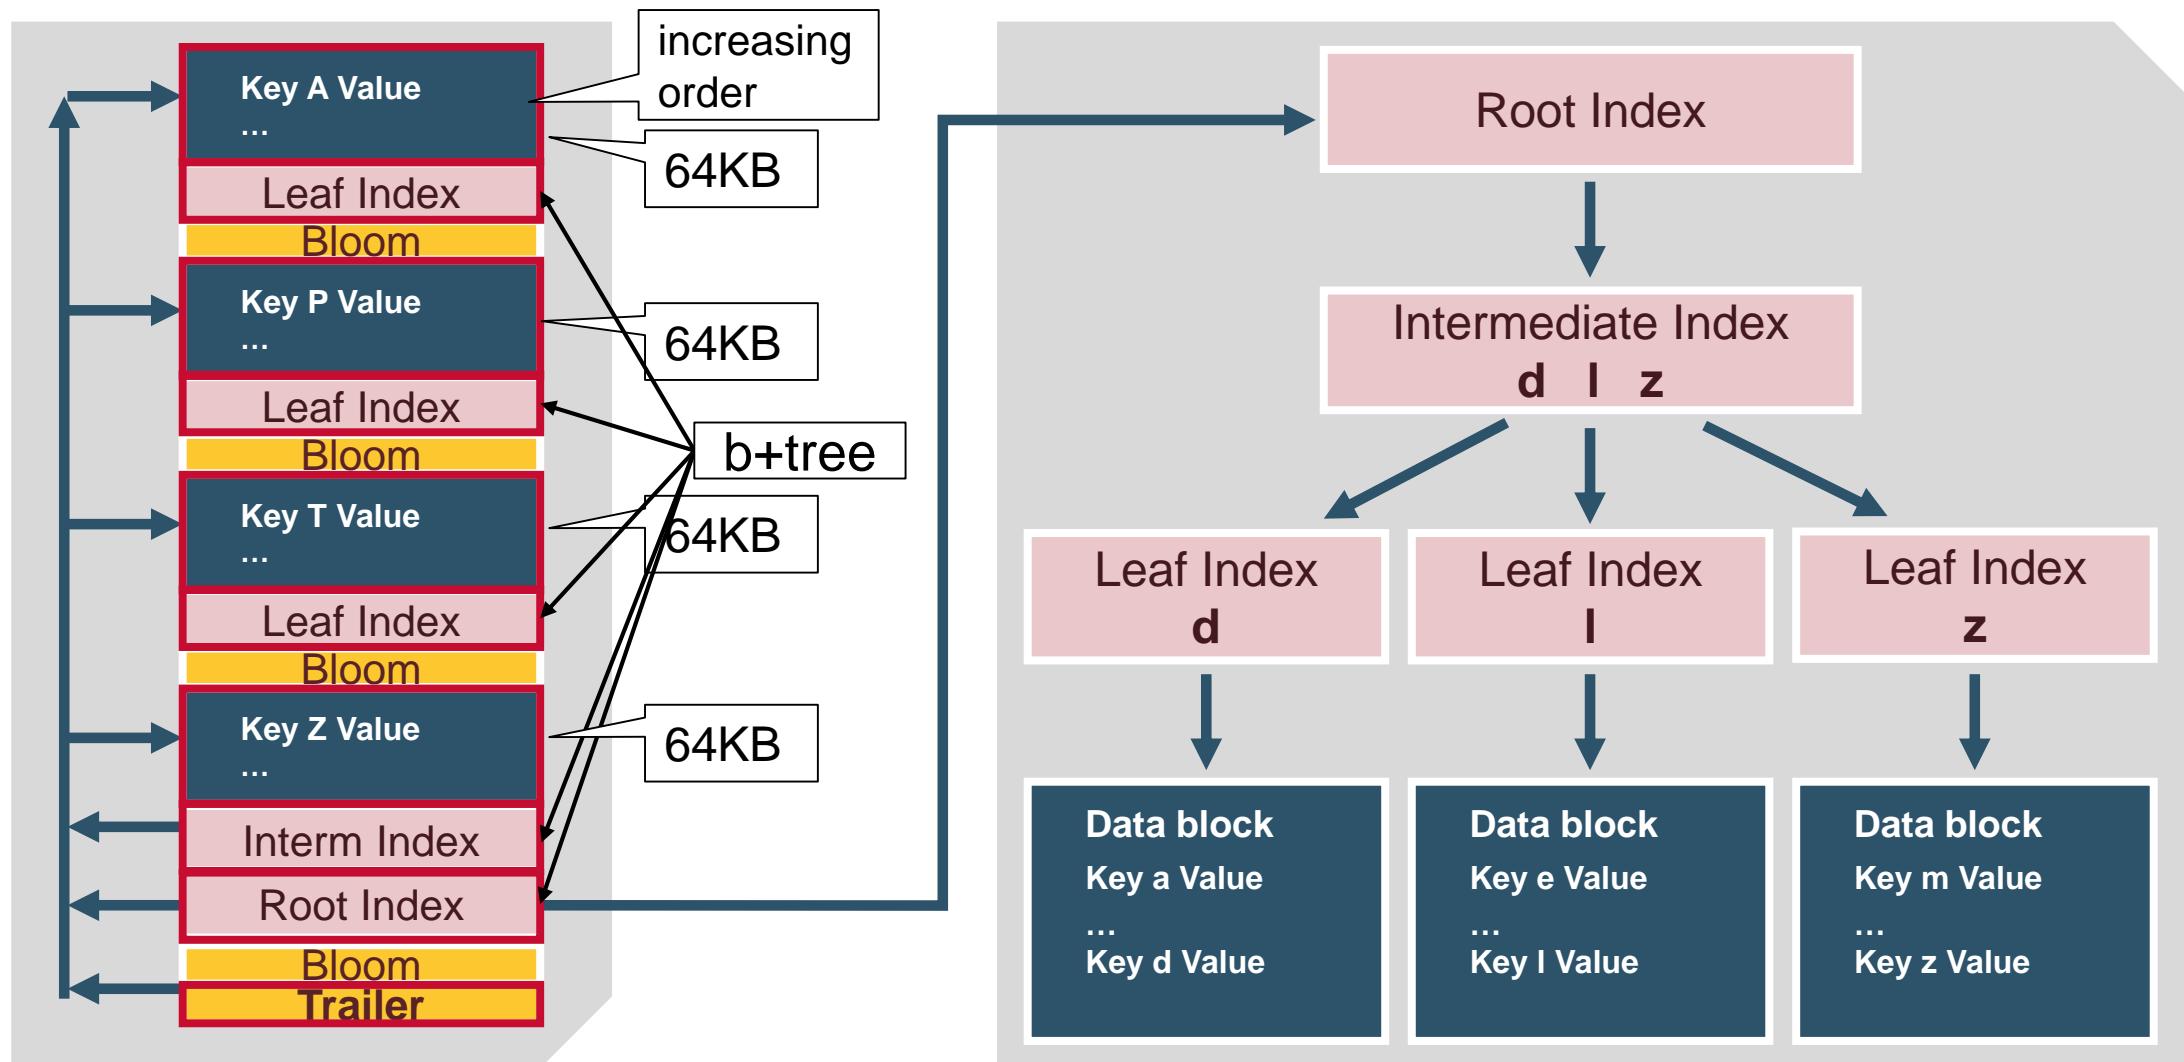
HBase HFile



Key		Value	
Key	CF1:Col	version	value
ra	cf1:ca	v1	1
rb	cf1:cb	v2	4
rb	cf1:cb	v1	3
rc	cf1:ca	v1	5
Key		Value	
Key	CF2:Col	version	value
ra	cf2:ca	v1	2
rc	cf2:ca	v2	7
rc	cf2:ca	v1	6
rc	cf2:cd	v1	8

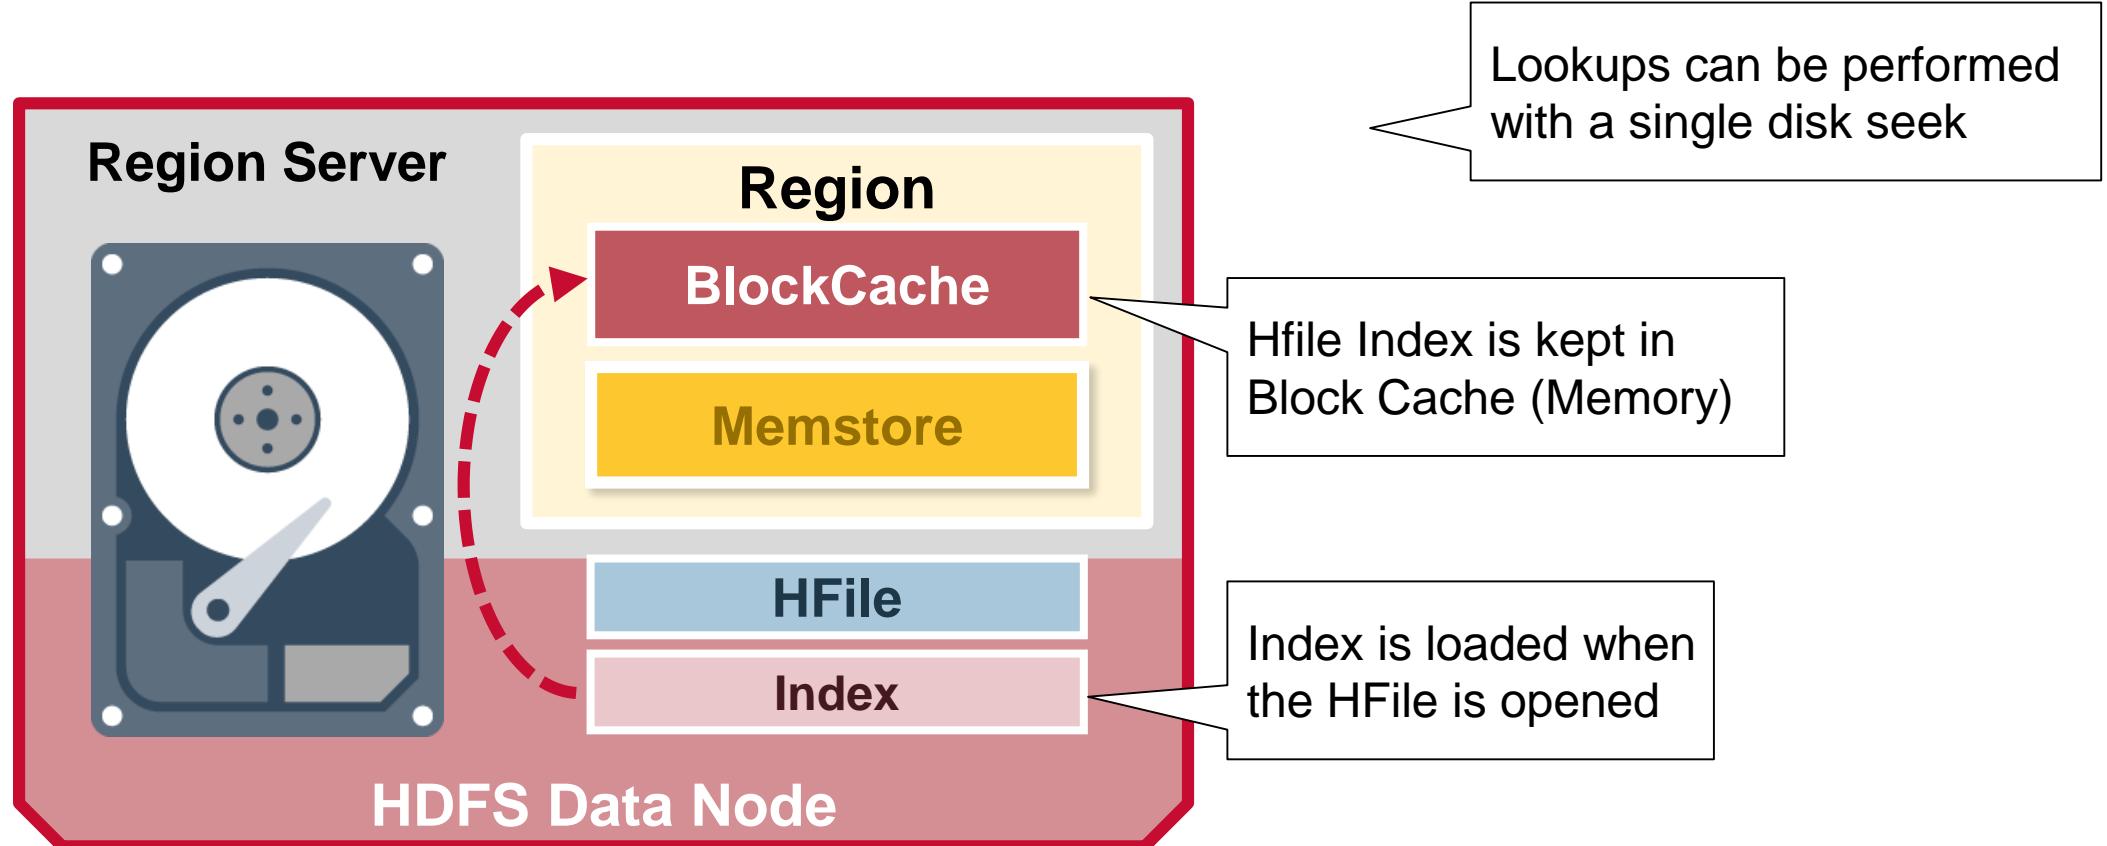


HBase HFile Structure





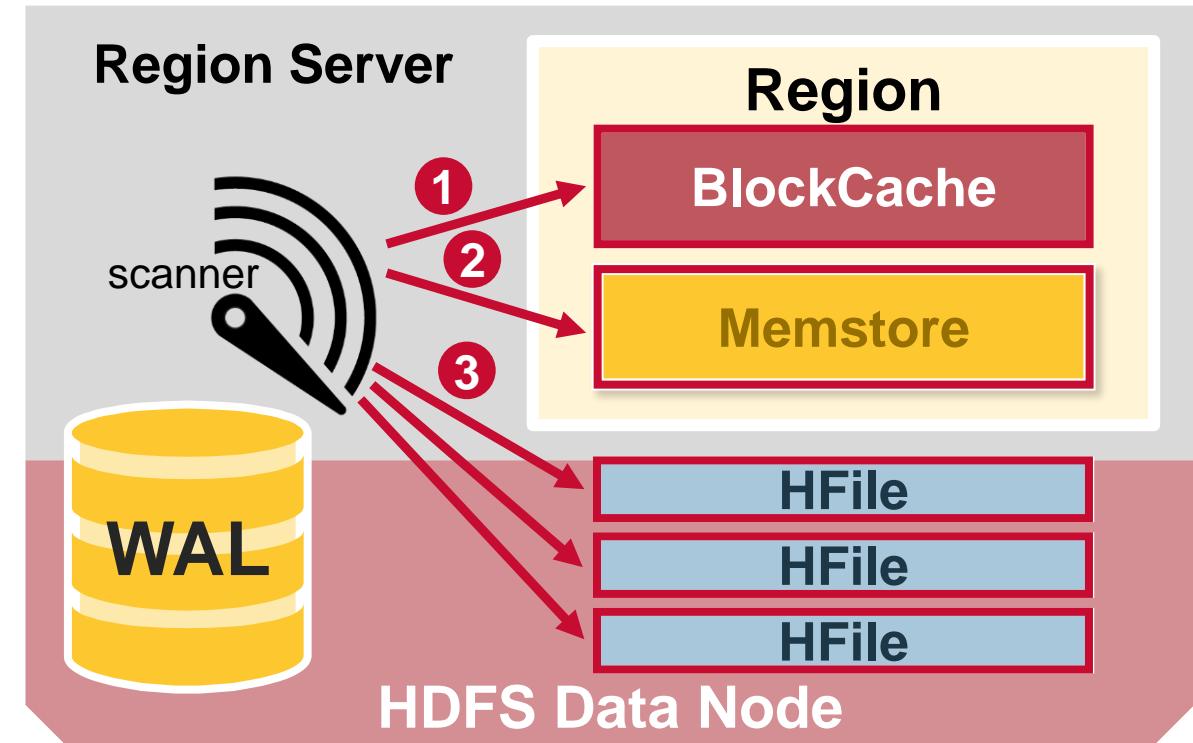
HFile Index





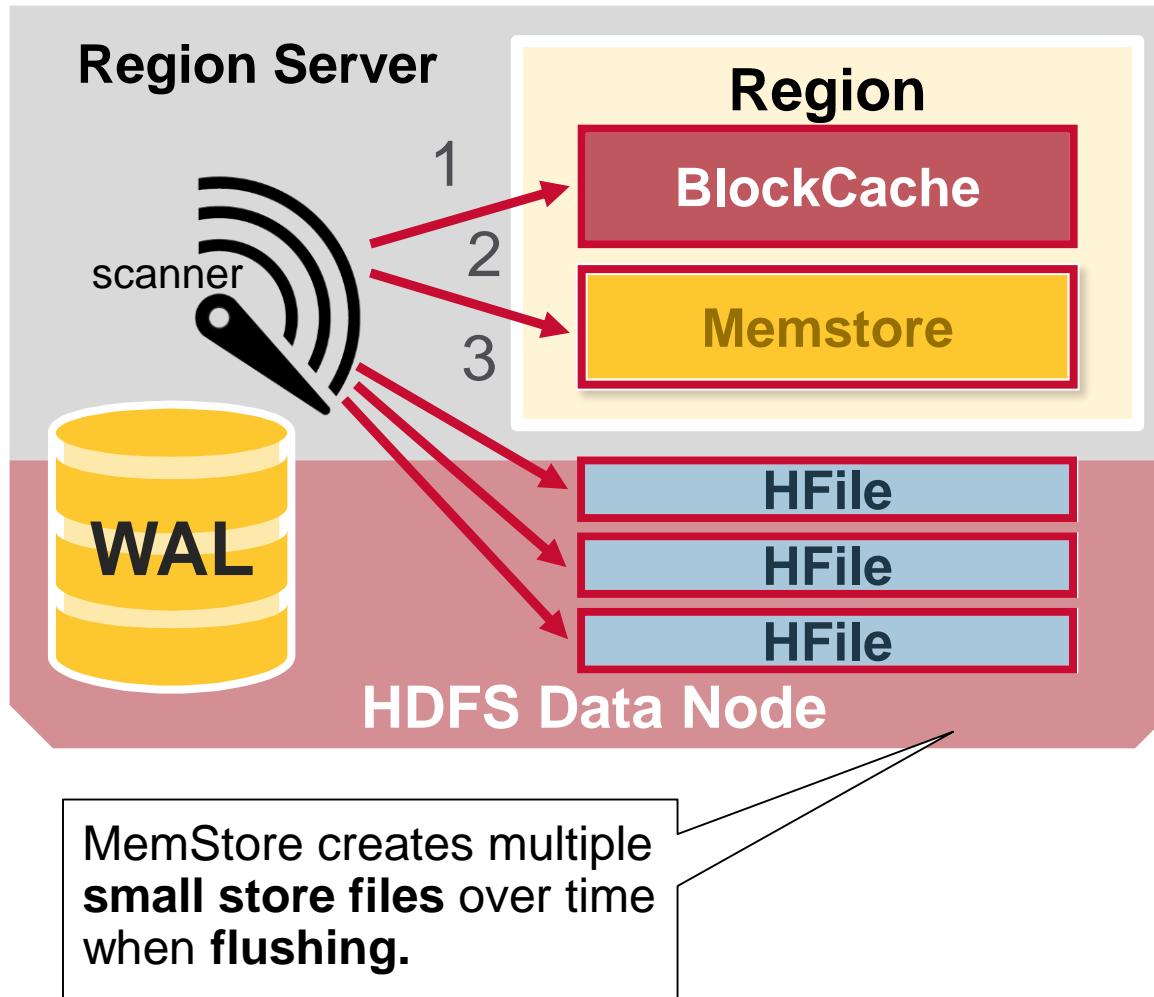
HBase Read Merge

- 1 First the scanner looks for the Row KeyValues in the Block cache
- 2 Next the scanner looks in the Memstore
- 3 If all row cells not in memstore or blockCache, look in HFiles





HBase Read Merge

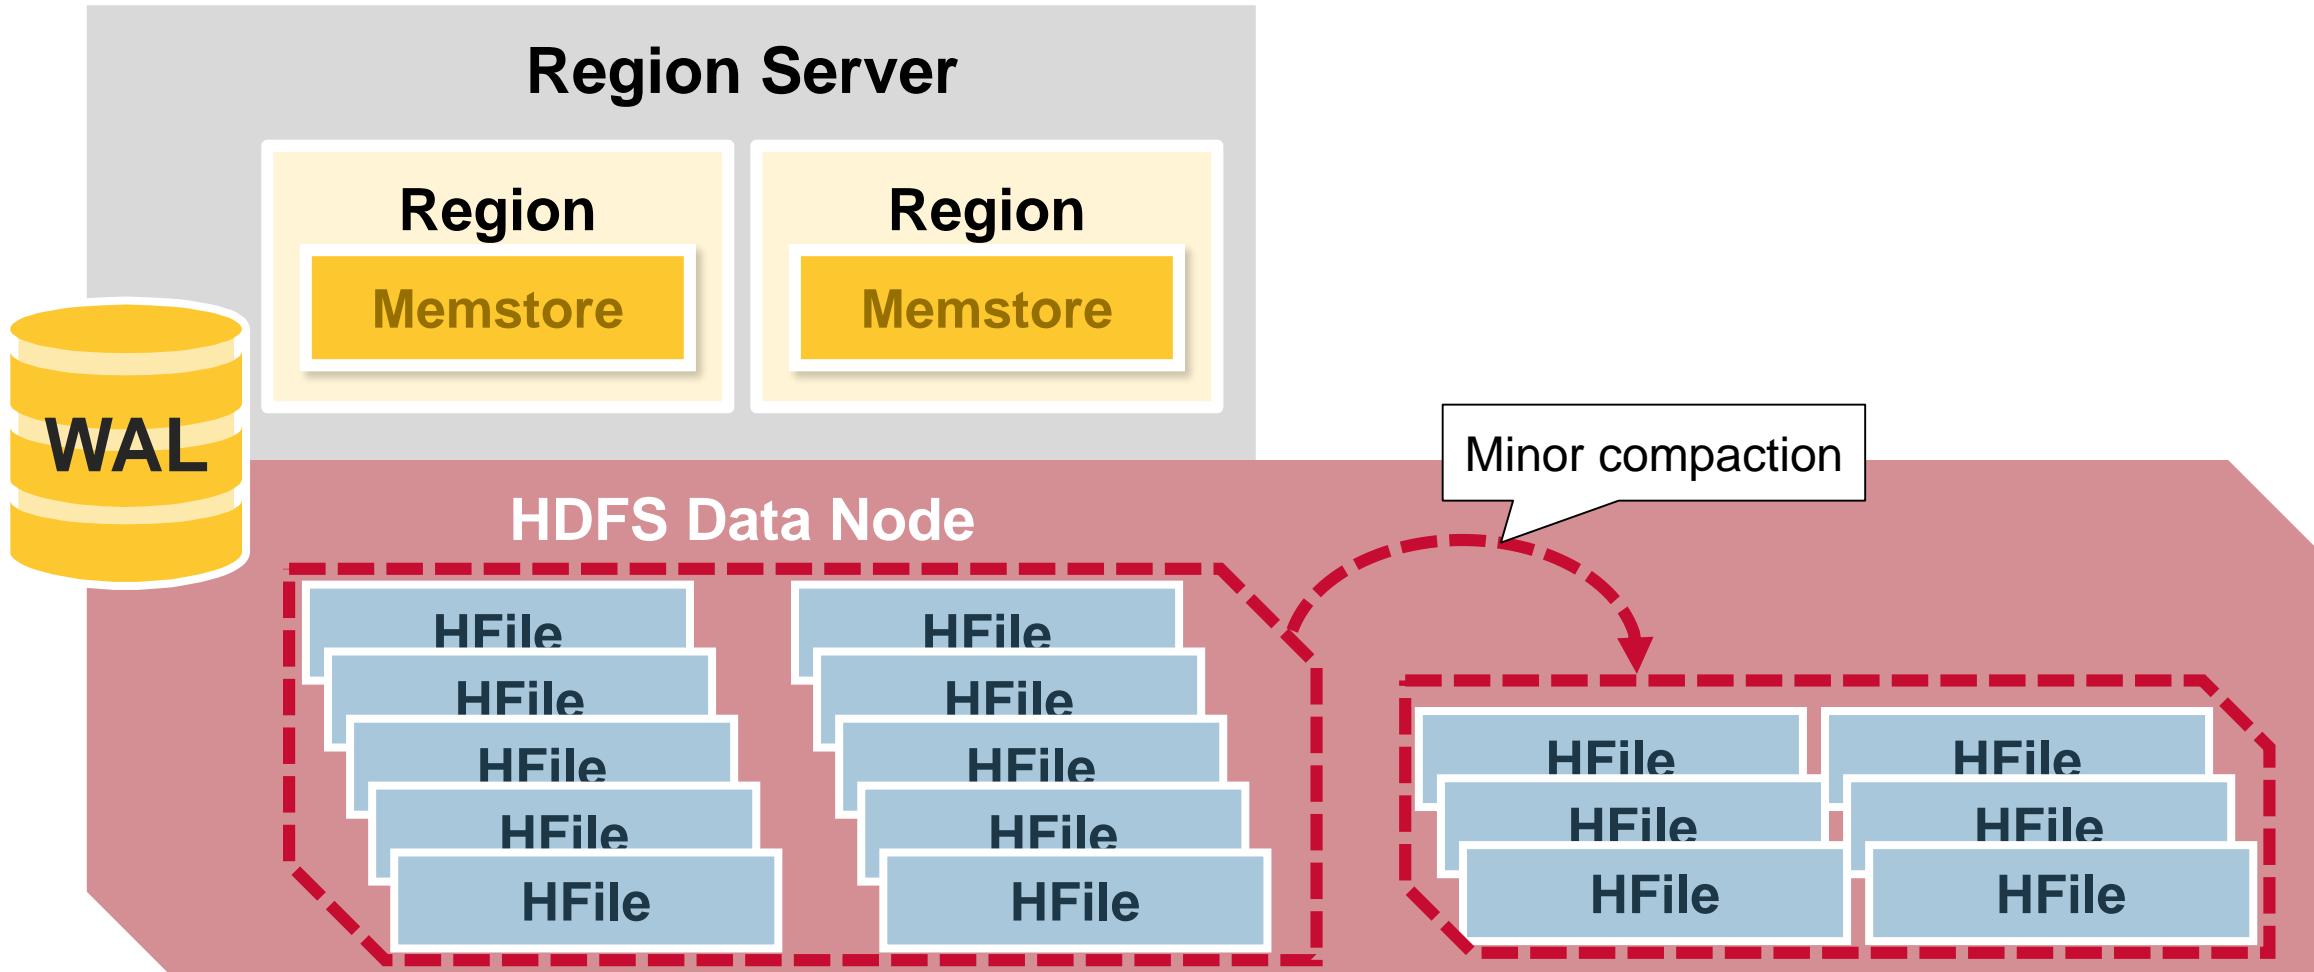


Read Amplification

- multiple files have to be examined

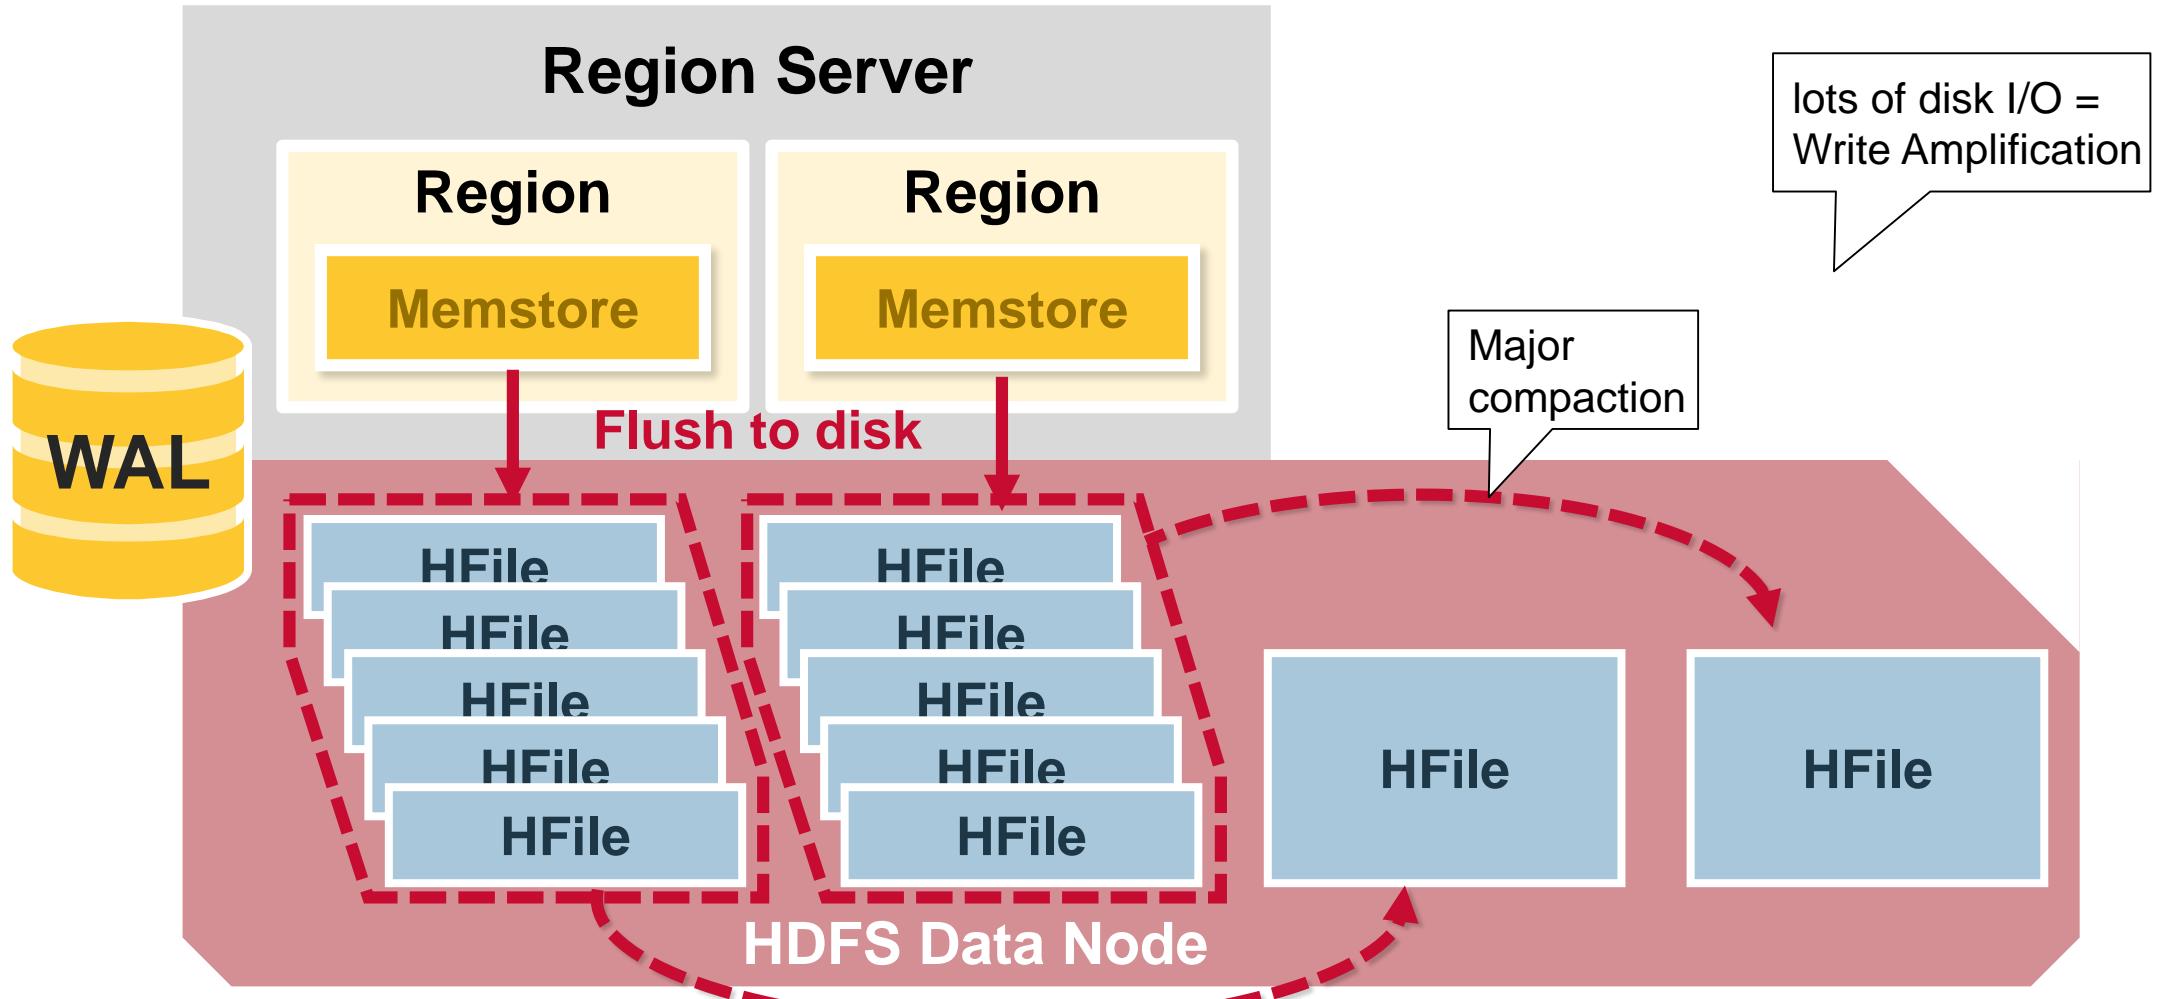


HBase Minor Compaction

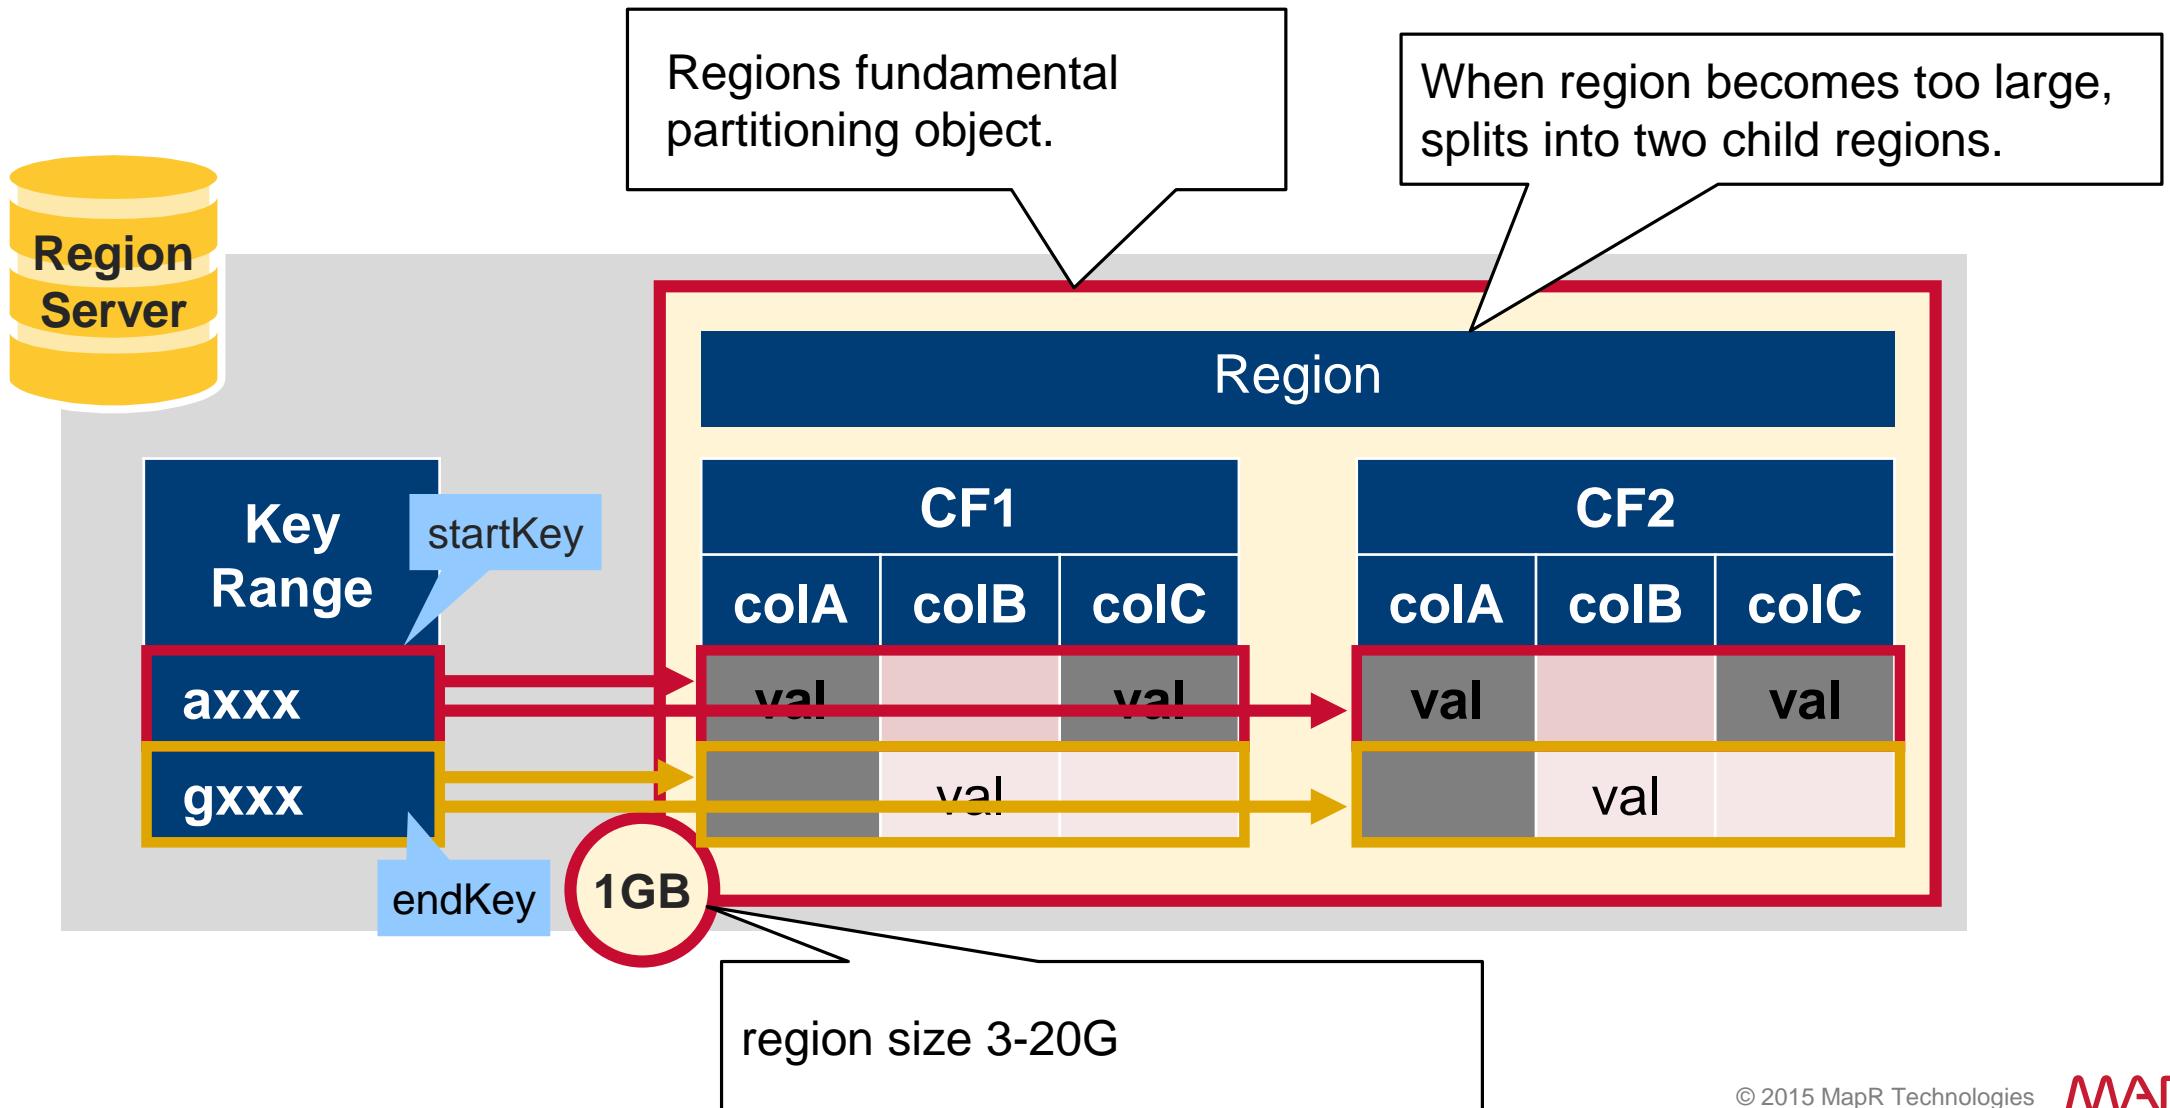




HBase Major Compaction

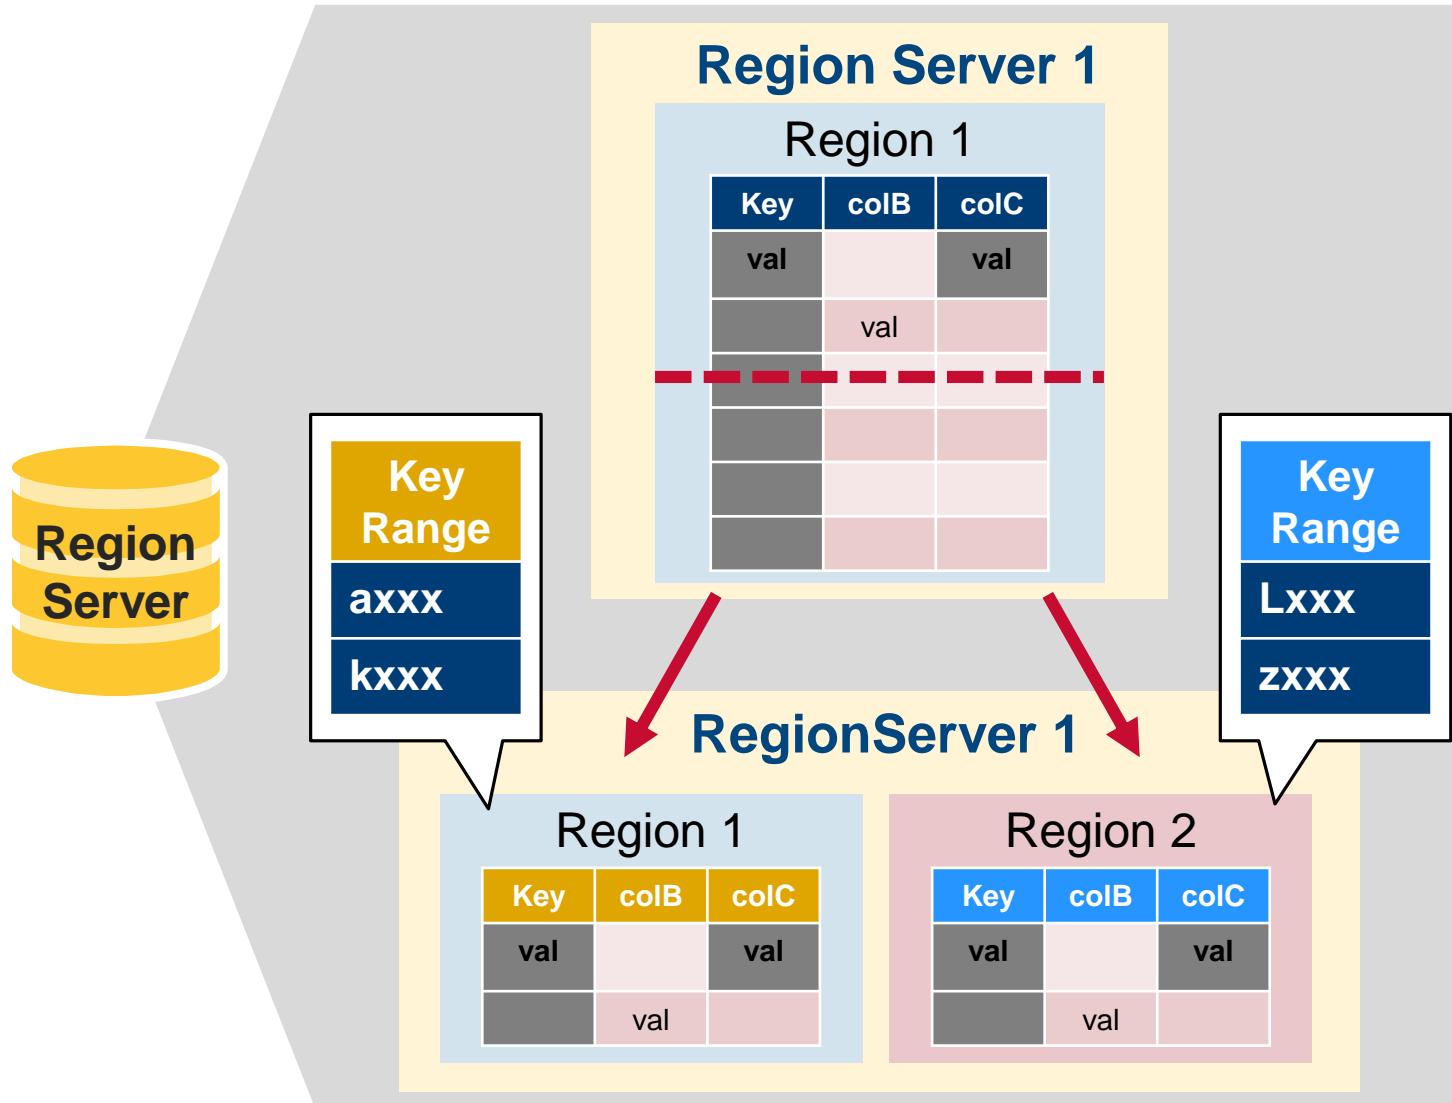


Region = contiguous keys





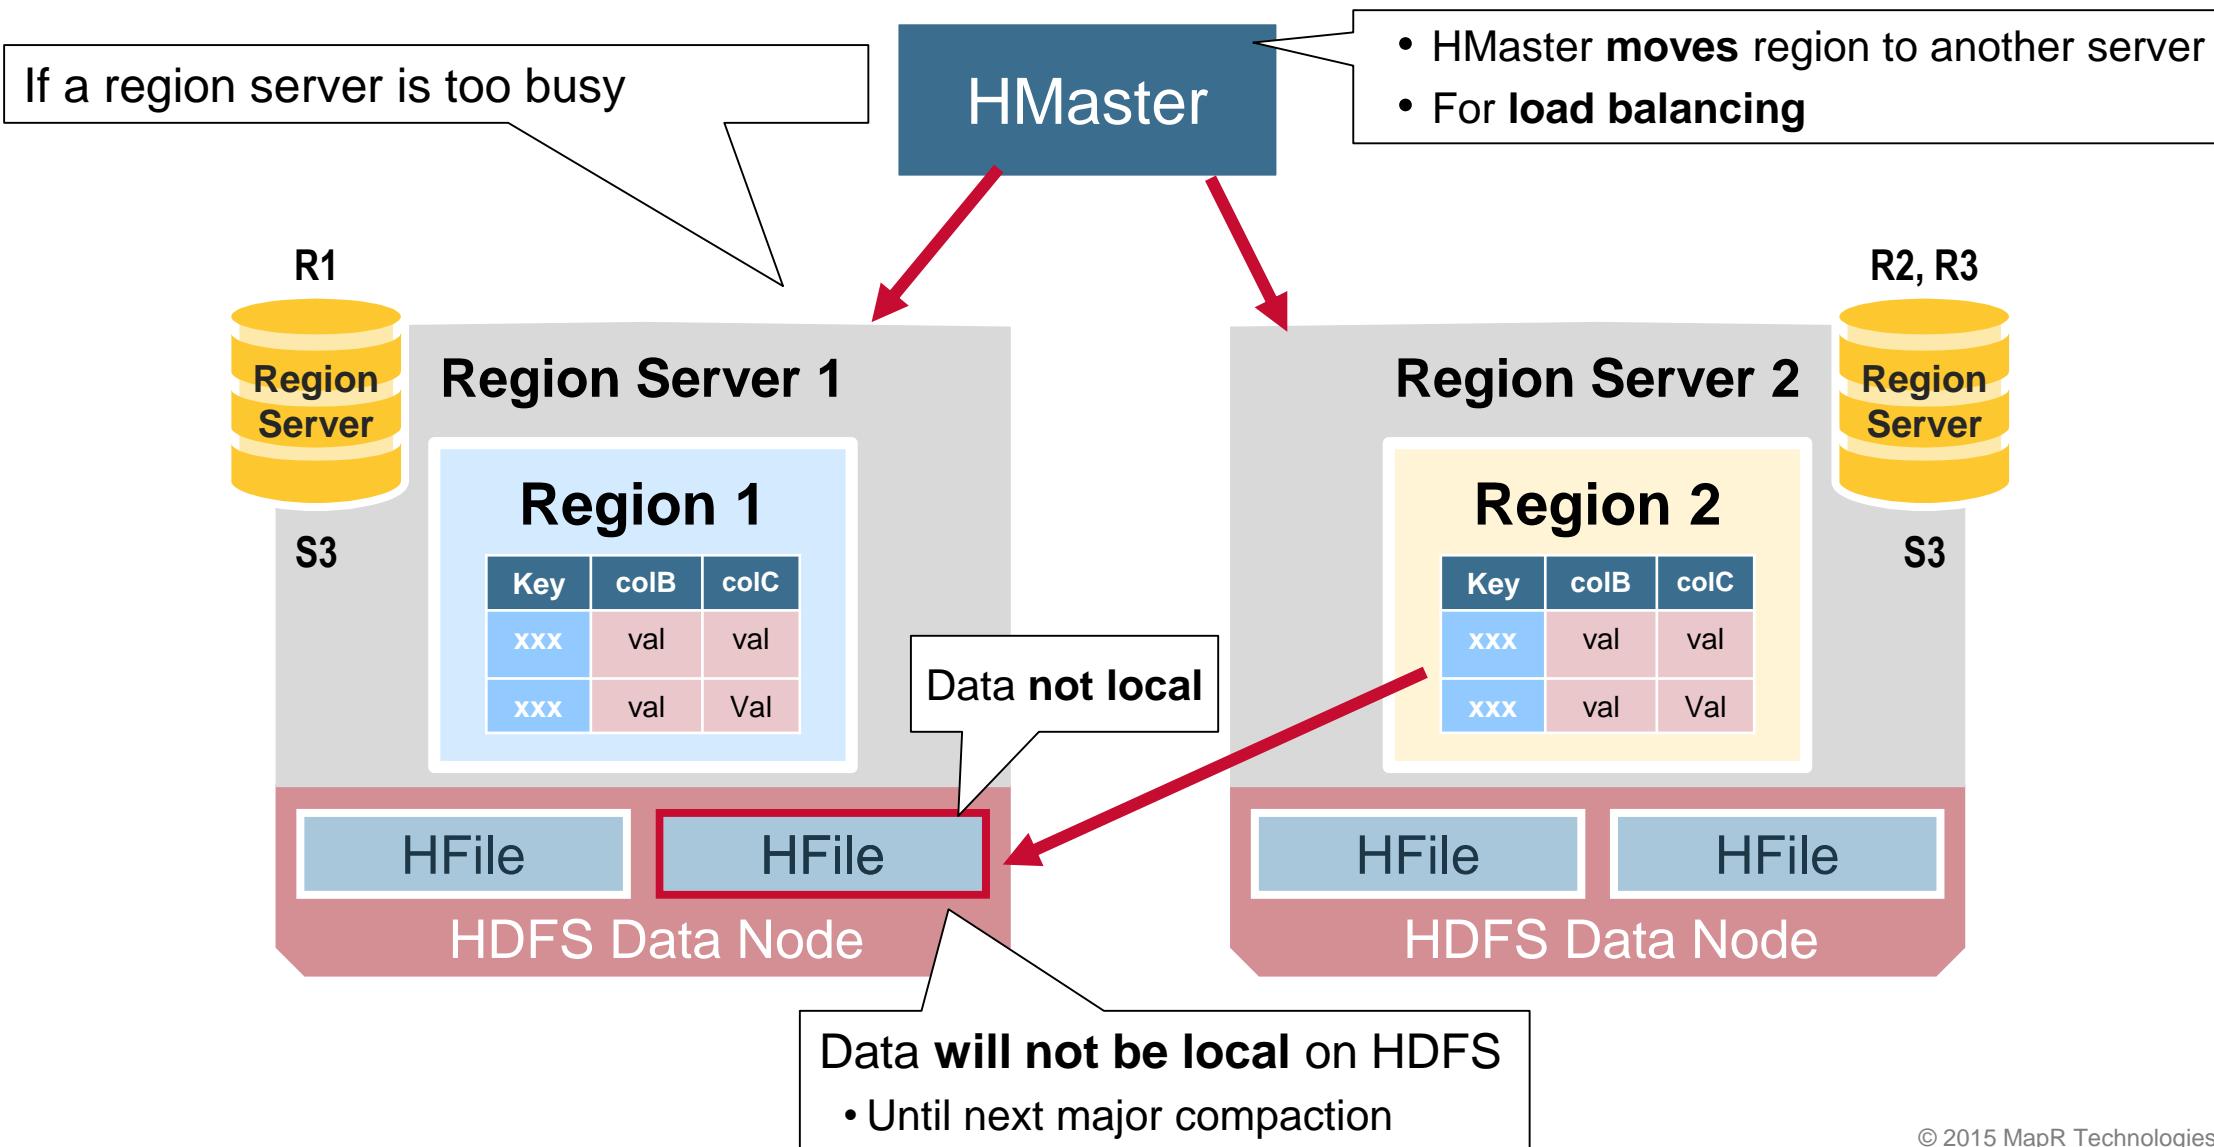
Region Split



when region size >
hbase.hregion.max.
filesize → split

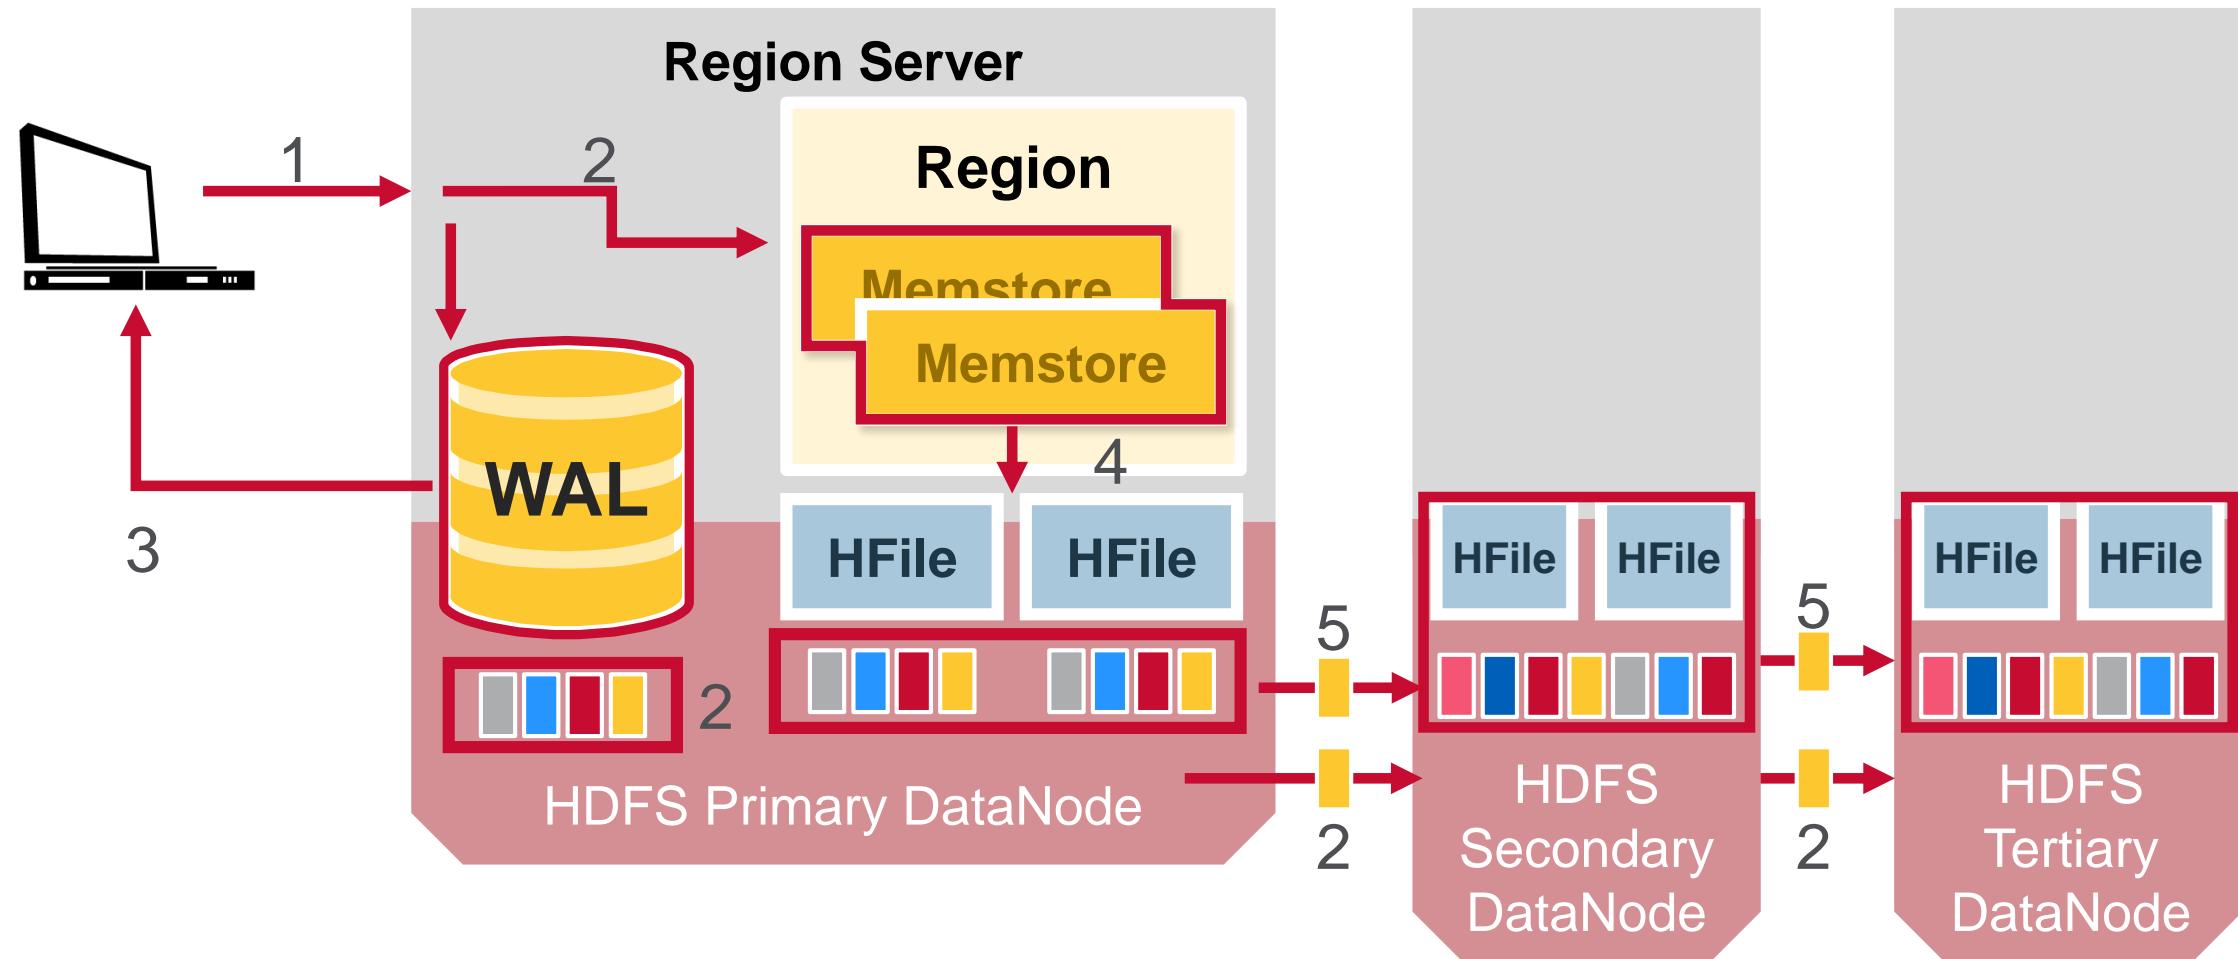


Region Load Balancing



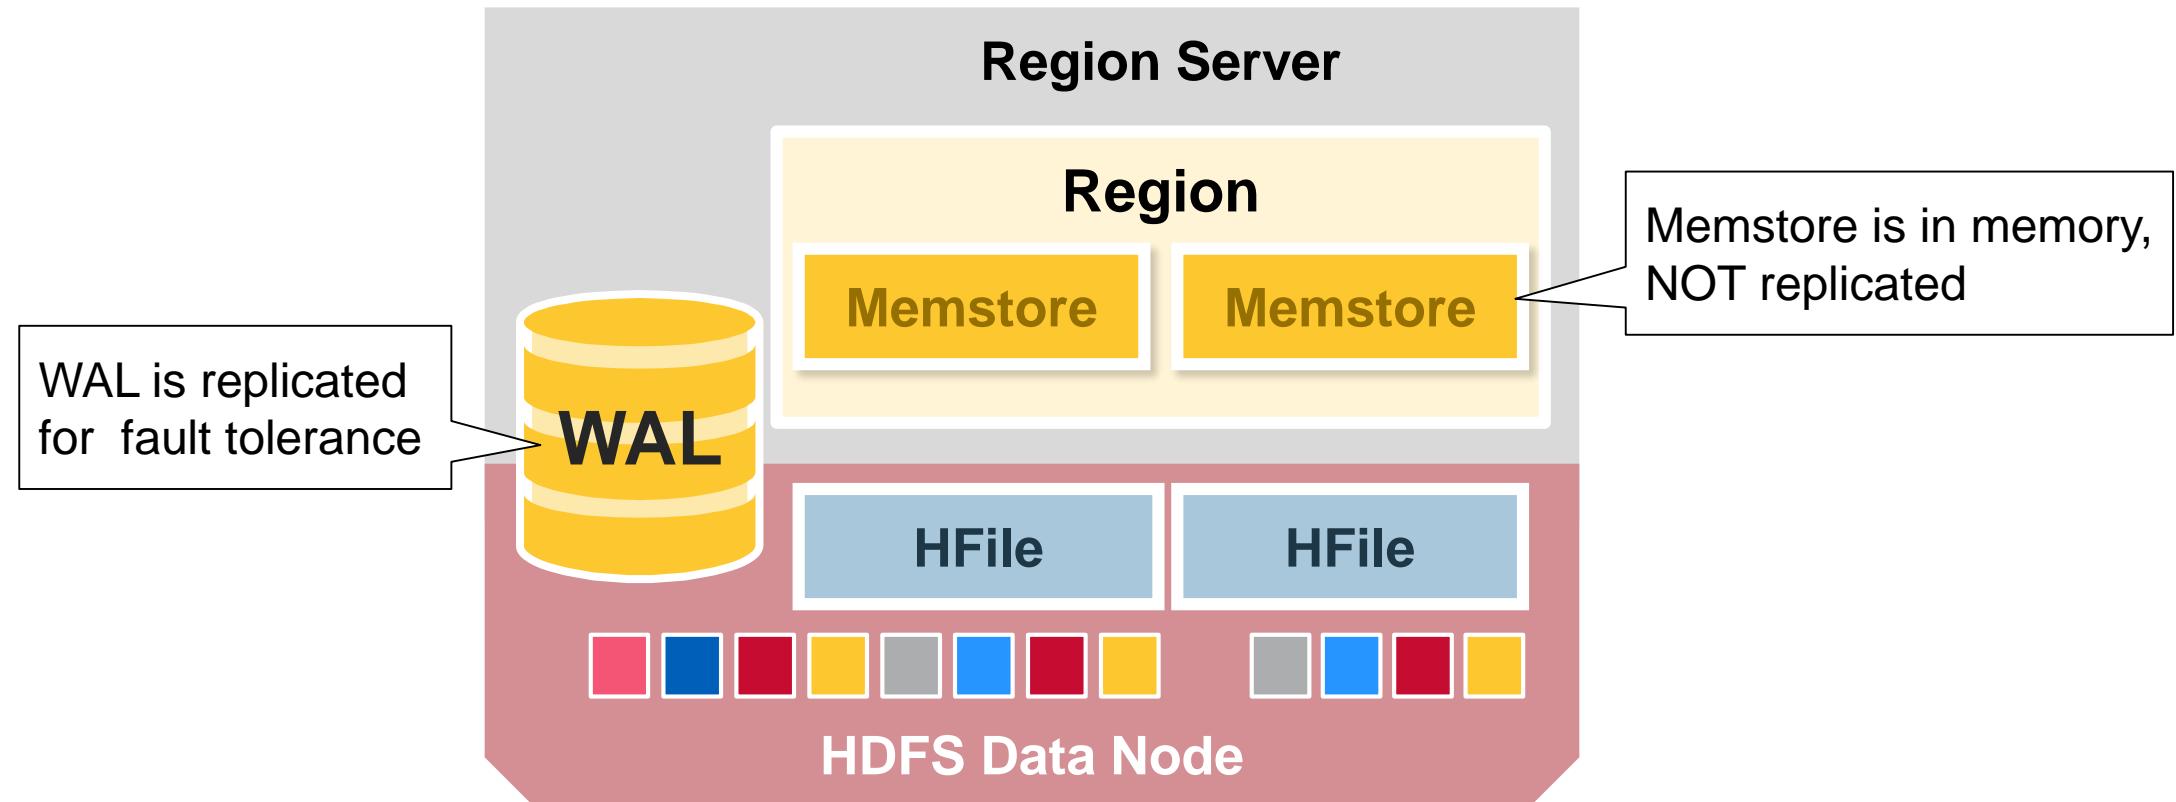


HDFS Data Replication





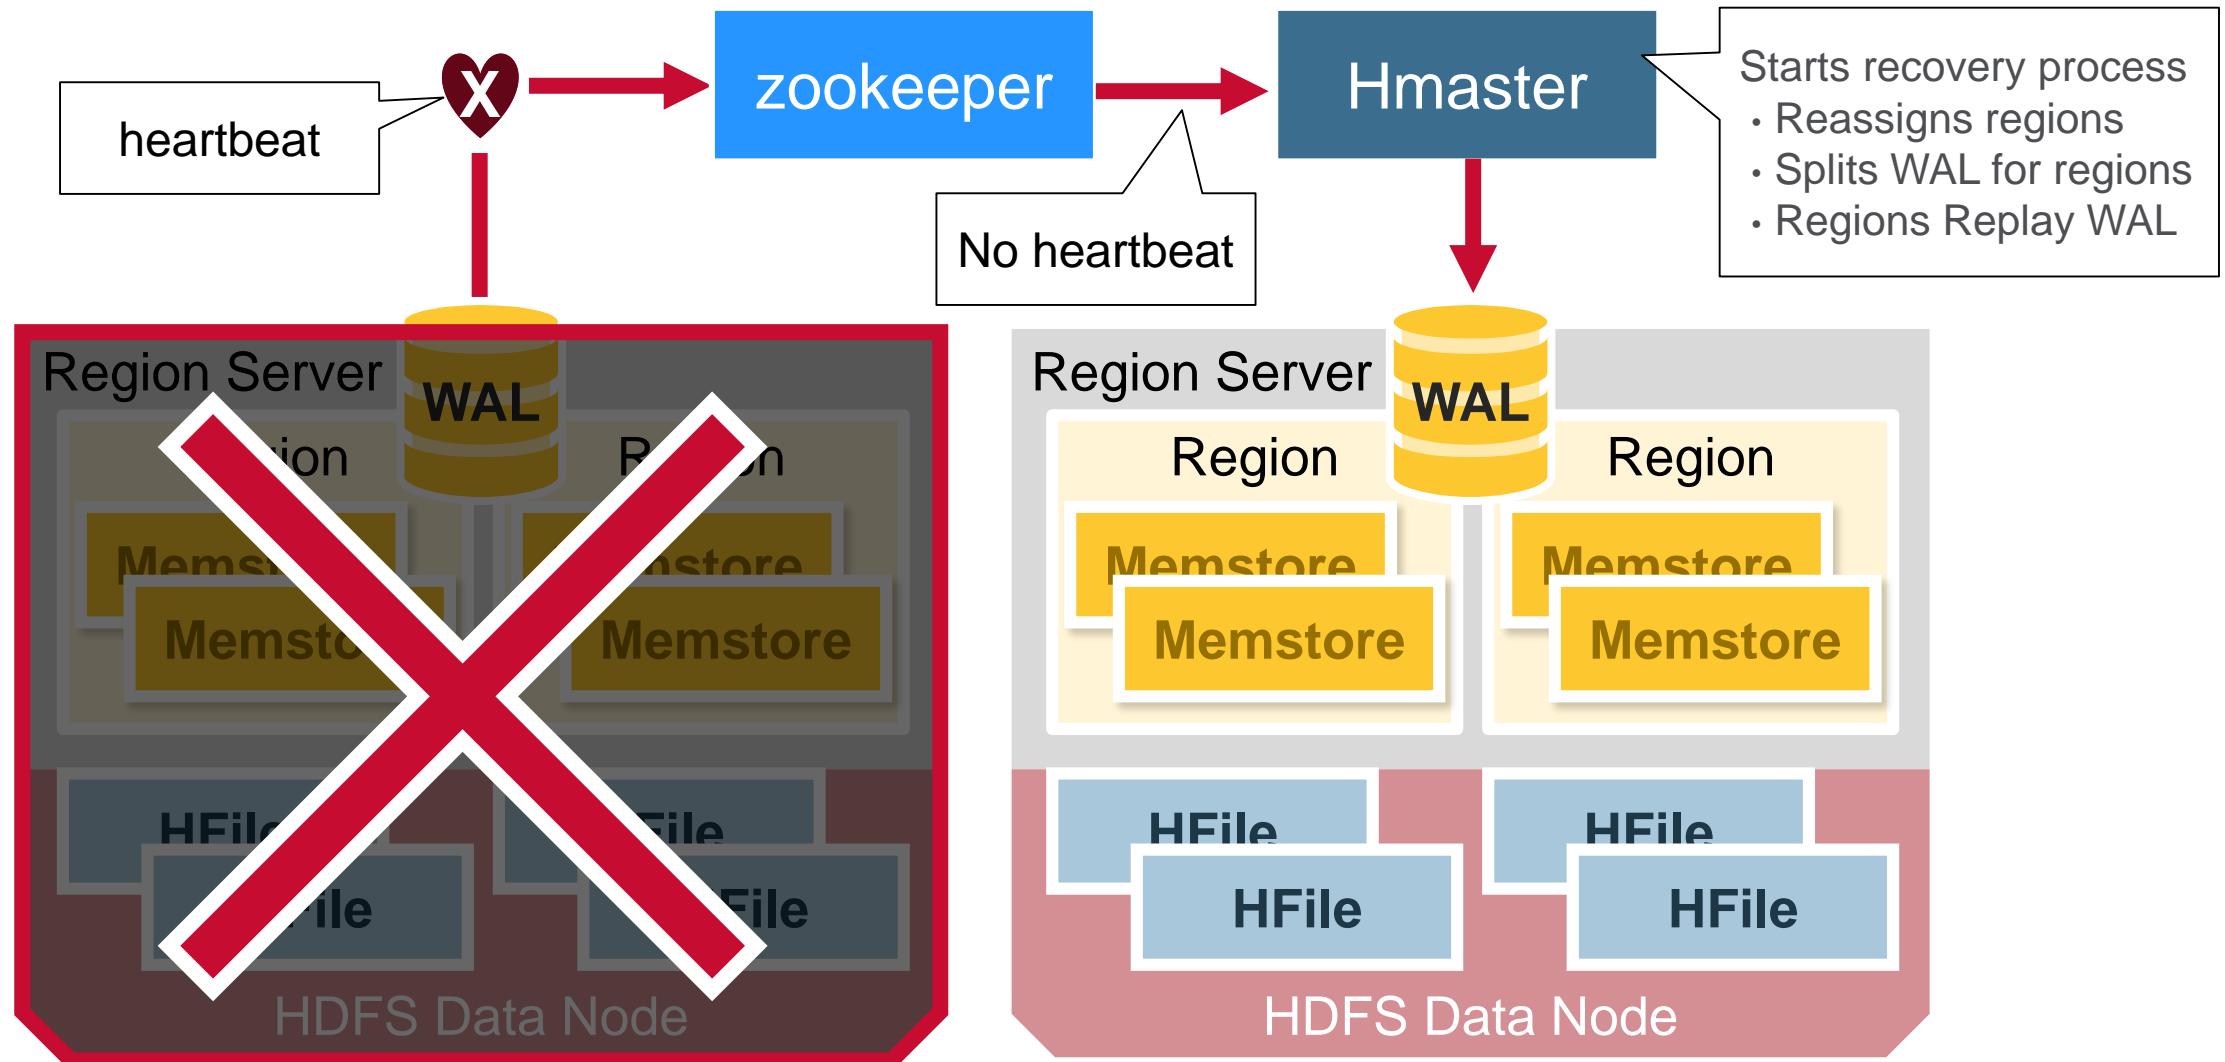
HDFS Data Replication – (2)



How does HBase recover updates not persisted to HFiles?

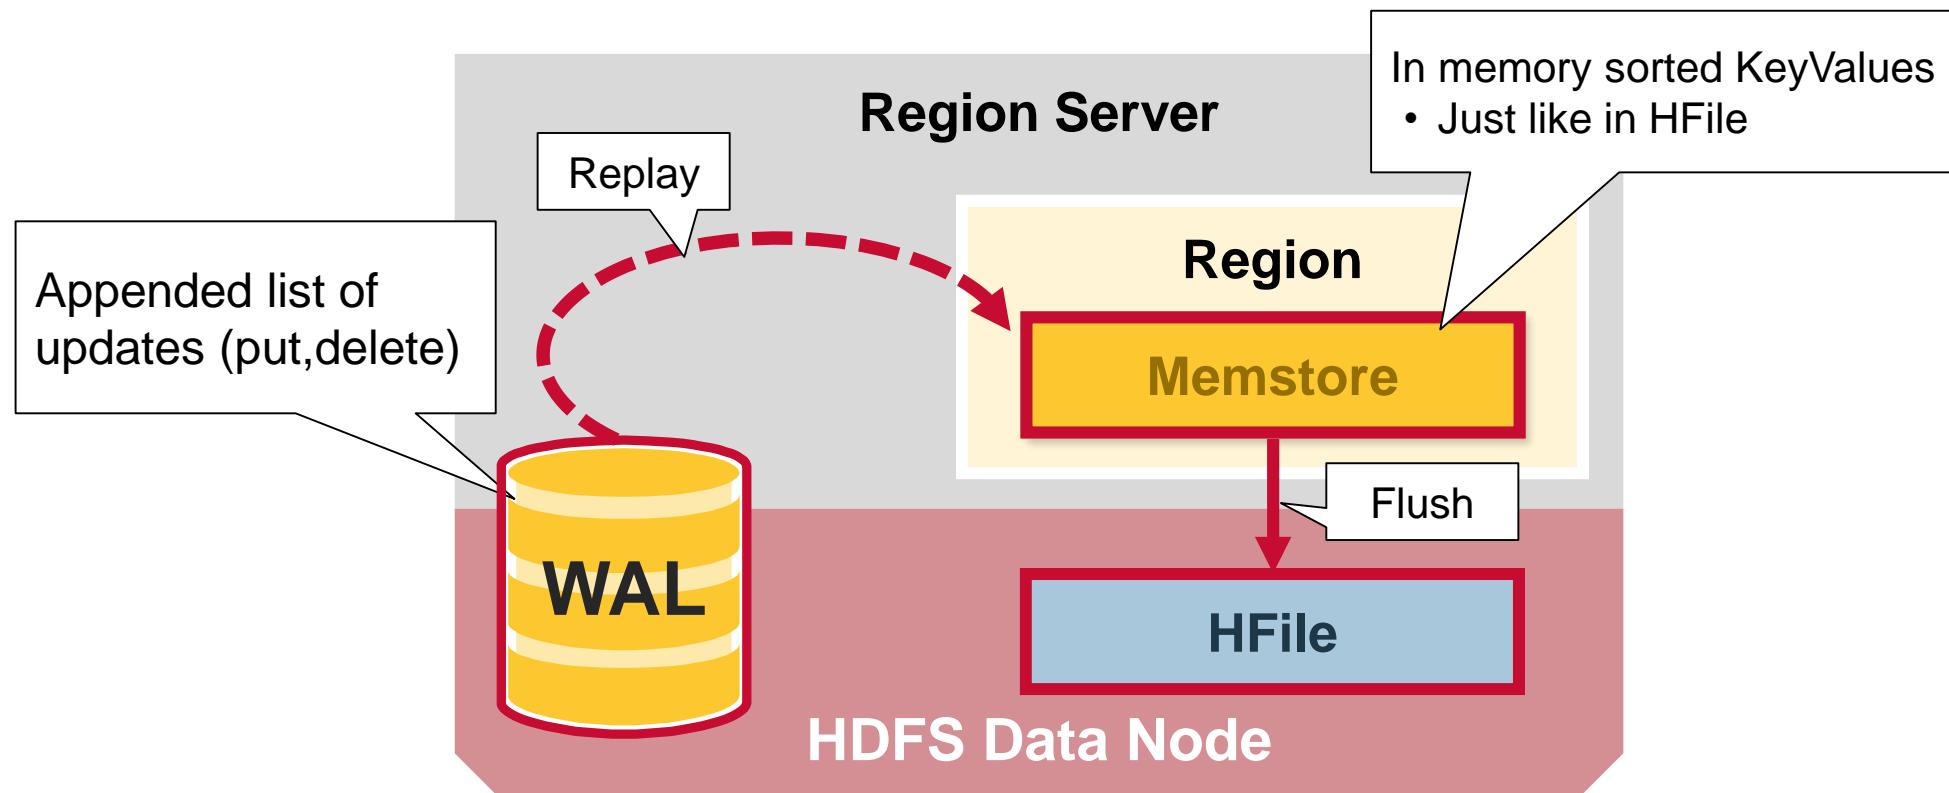


HBase Crash Recovery





Data Recovery





Better than many NoSQL data store solutions, hence its popularity

- **Strong consistency model**
 - When a write returns, all readers will see same value
- **Scales automatically**
 - Regions split when data grows too large
 - Uses HDFS to spread and replicate data
- **Built-in recovery**
 - Using Write Ahead Log (similar to journaling on file system)
- **Integrated with Hadoop**
 - MapReduce on HBase is straightforward

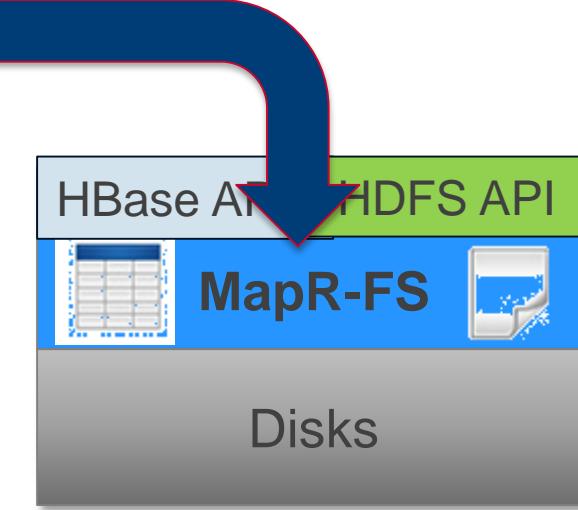
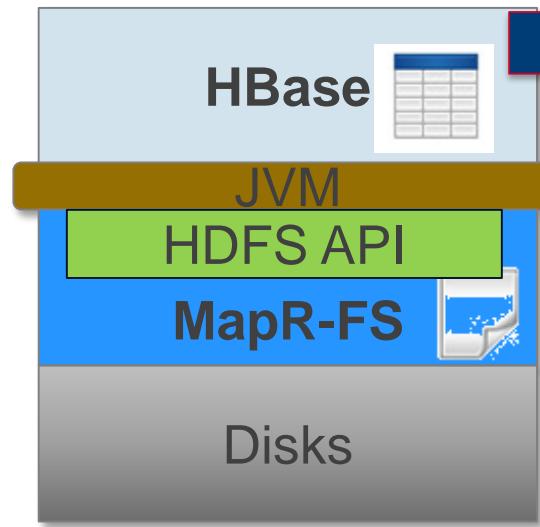
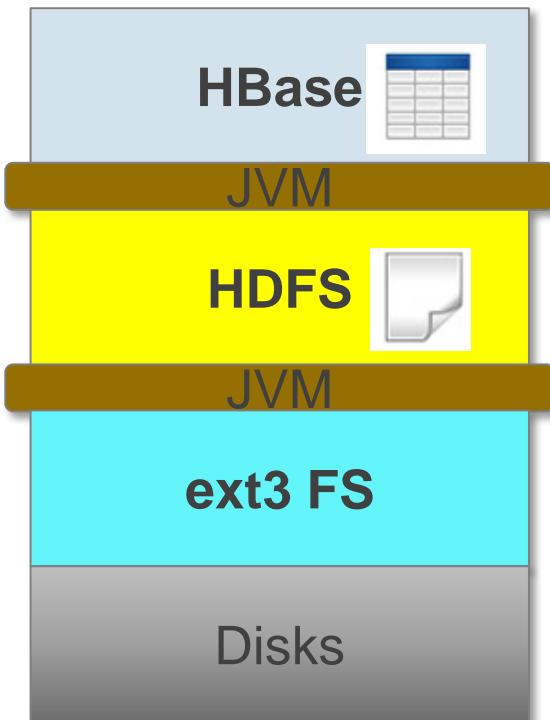
MapR-DB





MapR-DB and Files in a Unified Storage Layer

- MapR-DB Tables and MapR-FS Files in a unified read/write filesystem.



Apache
HBase on Hadoop

Apache HBase on
MapR Filesystem

Mapr-DB Integrated
into Filesystem



Tables Integrated into MapR read/write File System:

- **MapR-DB tables use the HBase data model and API**
- **Key differences between MapR tables and Apache HBase**
 - **Tables** part of the MapR **Read/Write** File system
 - **Guaranteed data locality**
 - **Smarter** load balancing
 - Uses container **Replicas**
 - **Smarter** fail over
 - Uses container **replicas**
 - Multiple small WALs
 - **Faster** recovery
 - **No compaction !**



HBase Use Cases



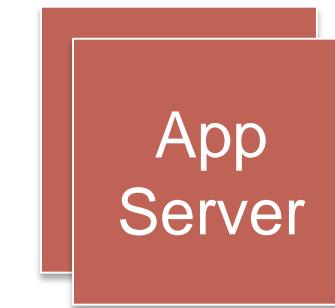
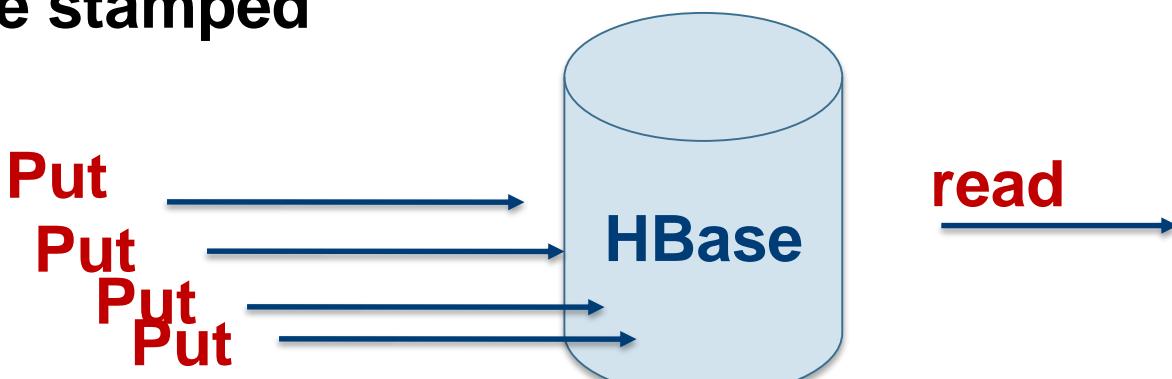
3 Main Use Case Categories

- Time Series Data, Stuff with a Time Stamp
 - Sensor, System Metrics, Events, log files
 - Stock Ticker, User Activity
 - Hi Volume, Velocity Writes

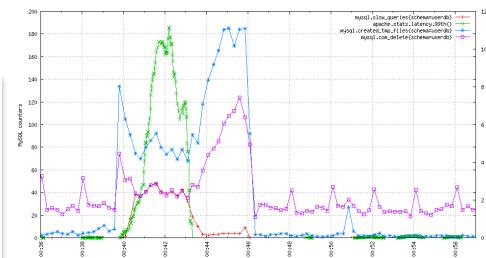
OpenTSDB



Event time stamped
data

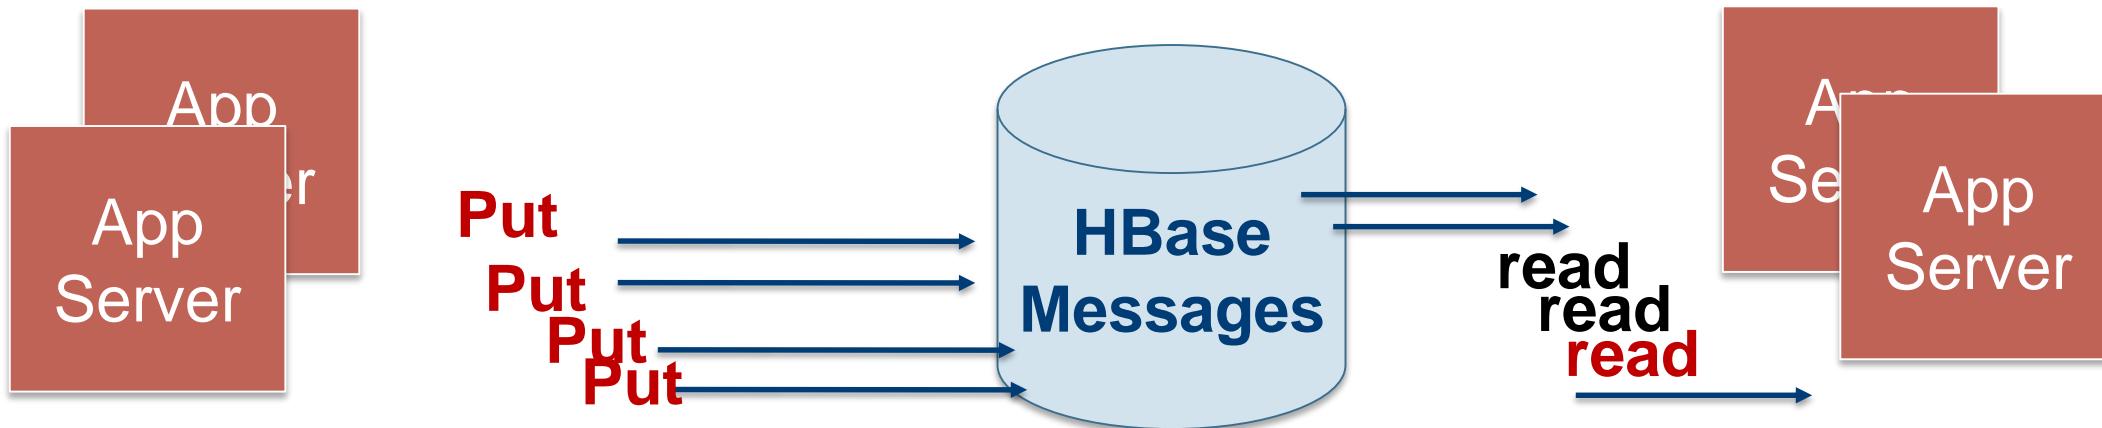


Data for real-time monitoring.



3 Main Use Case Categories

- Information Exchange
 - email, Chat, Inbox: **Facebook**
 - Hi Volume, Velocity Write/Read

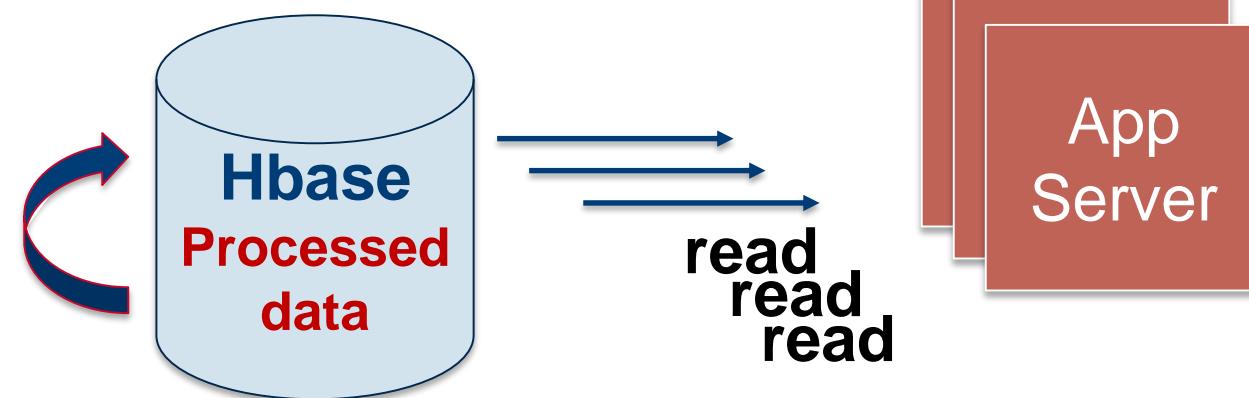


3 Main Use Case Categories

- Content Serving, Web Application Backend
 - Online Catalog: Gap, World Library Catalog.
 - Search Index: ebay
 - Online Pre-Computed View: Groupon, Pinterest
 - Hi Volume, Velocity Reads



Bulk Import
Pre-Computed
Materialized View



Agenda

- Why do we need NoSQL / HBase?
- Overview of HBase & HBase data model
- HBase Architecture and data flow
- **Demo/Lab using HBase Shell**
 - Create tables and CRUD operations using MapR Sandbox
- Design considerations when migrating from RDBMS to HBase
- HBase Java API to perform CRUD operations
 - Demo / Lab using Eclipse, HBase Java API & MapR Sandbox
- How to work around transactions

Schema Design Guidelines

- HBase tables ≠ Relational tables!
 - HBase Design for Access Patterns



Use Case Example



Record Stock Trade Information in a Table

Trade data

Trade

- Timestamp
- Stock symbol
- Price per share
- Volume of trade

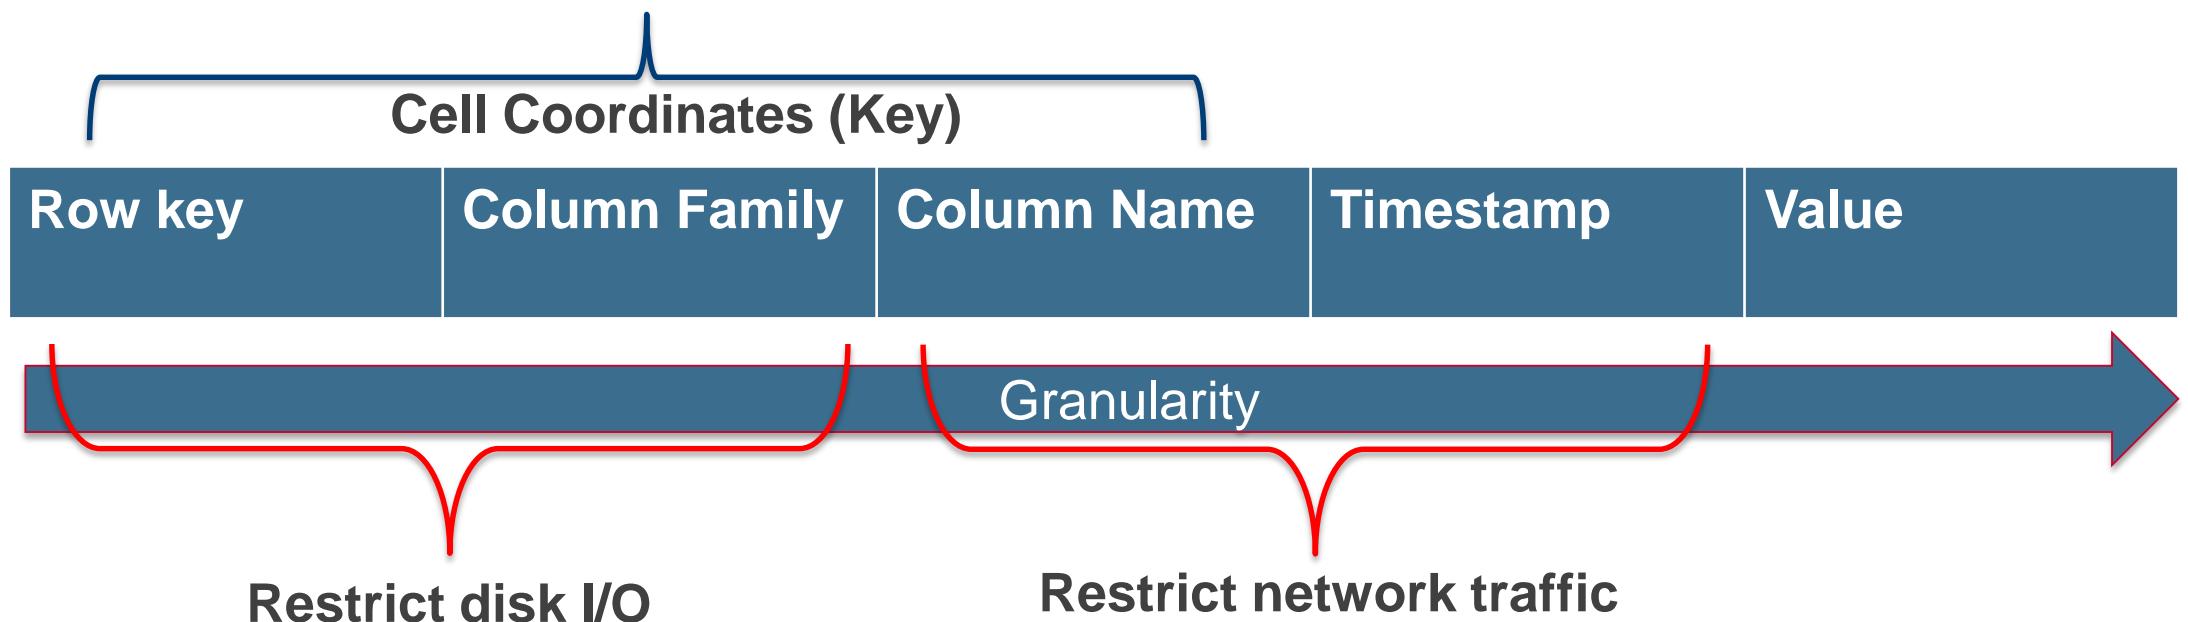
Example

- 1381396363000
(epoch timestamp with millisecond granularity)
- AMZN
- \$304.66
- 1333 shares



Intelligent keys

- Only the row keys are indexed
- Compose the key with attributes used for searching
 - Composite key : 2 or more identifying attributes
 - Like multi-column index design in RDB



 Composite Keys

- Use composite row key:
 - Include multiple elements in the row key
 - Use a separator or fixed length
- Example row key format:

SYMBOL	+	DATE (YYYYMMDD)
---------------	----------	------------------------

 - Ex: GOOG_20131012
- Scans can use partial keys
 - Ex: "GOOG" or "GOOG_2014"

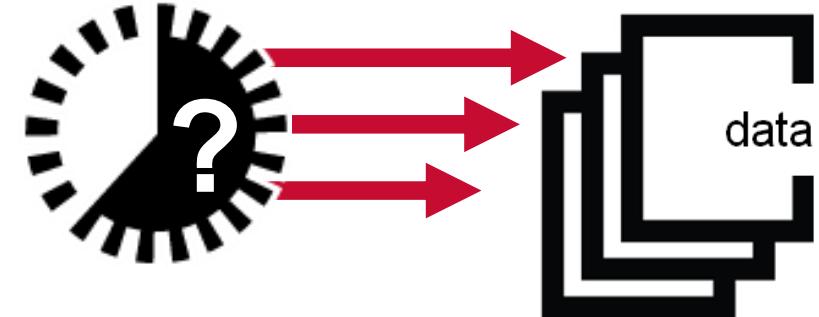




Consider Access Patterns for Application

How will data be retrieved?

- By date? By hour? By companyId?
 - *Rowkey design*

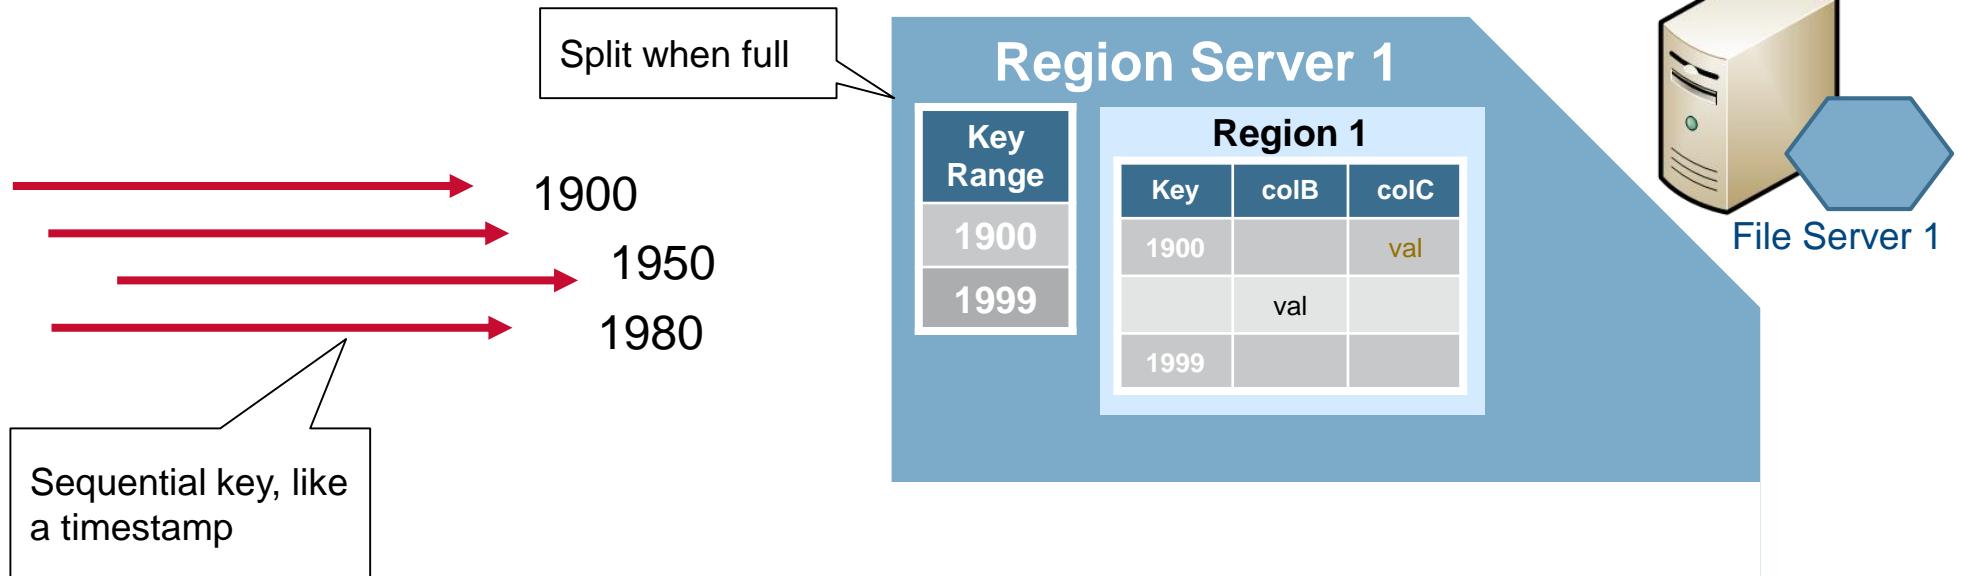


What if the Date/Timestamp is leftmost ?

Key				
1391813876369_AMZN				
1391813876370_AMZN				
1391813876371_GOOG				

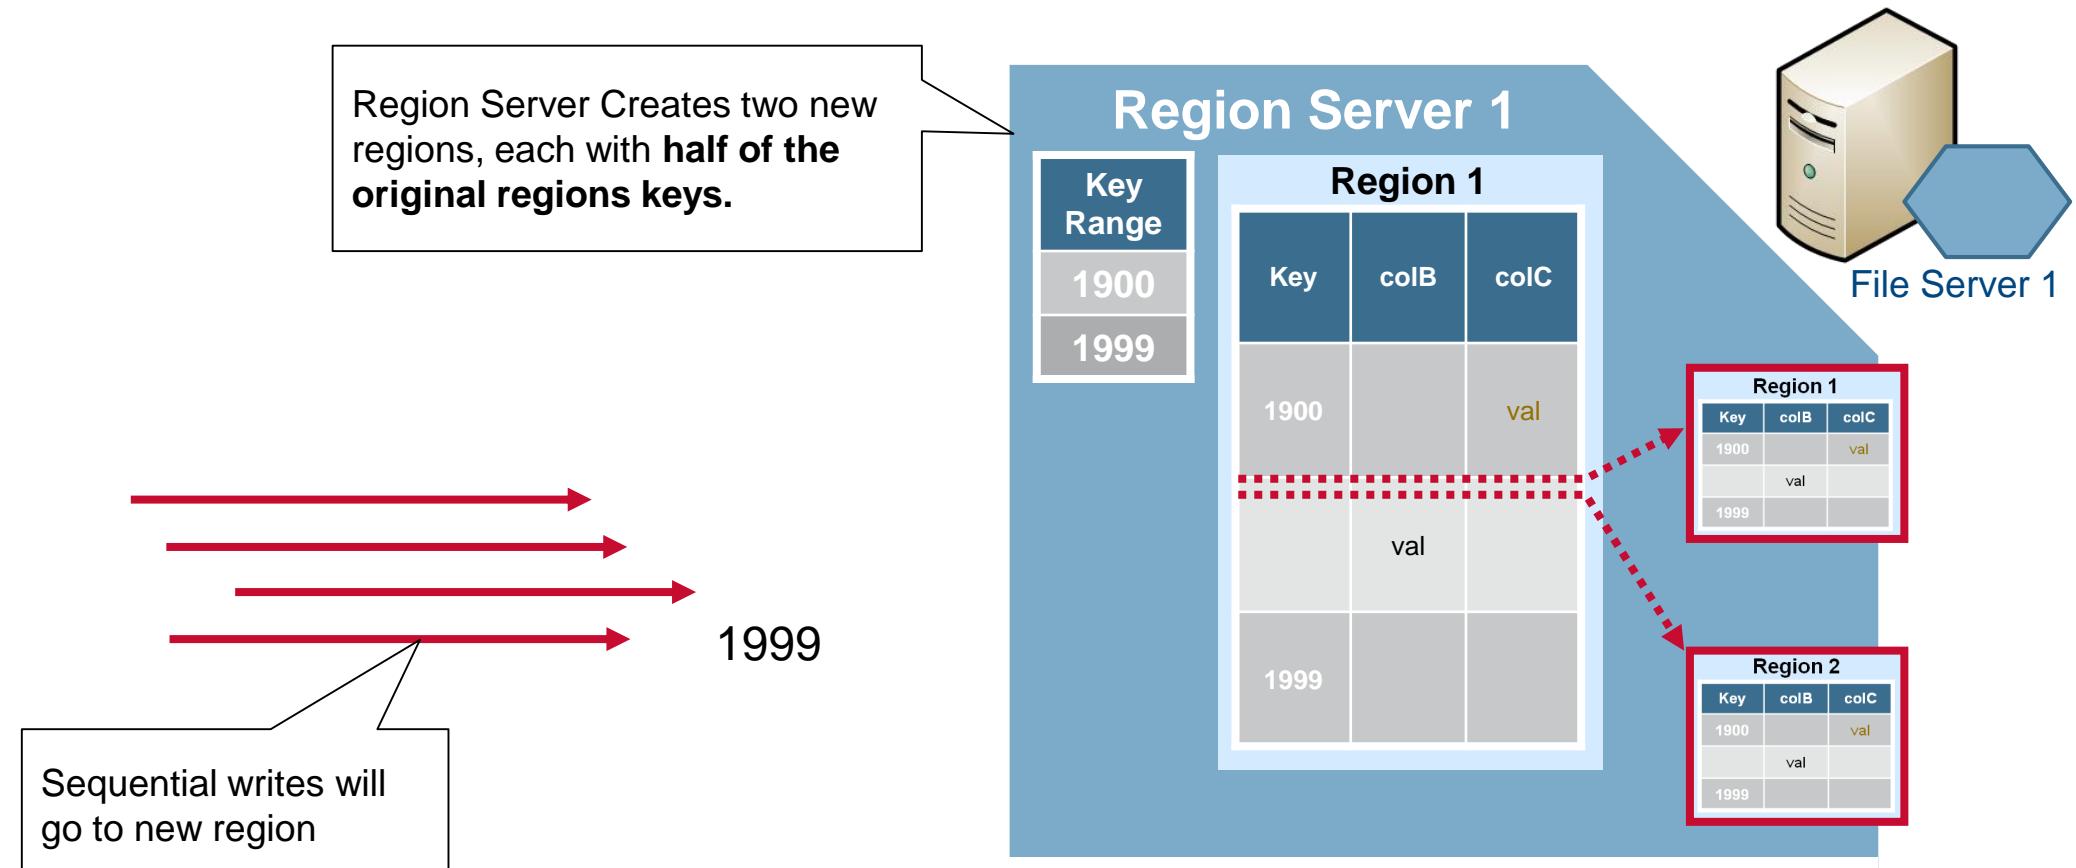


Hot-Spotting & Region Splits



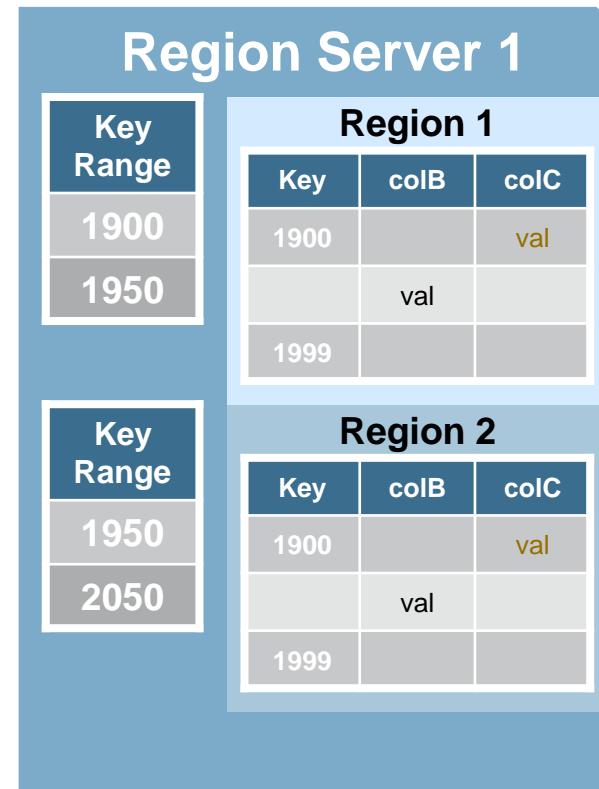
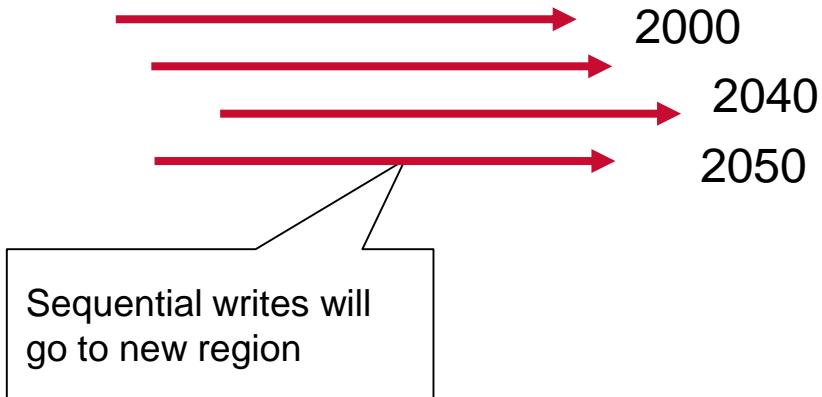


Hot-Spotting & Region Splits





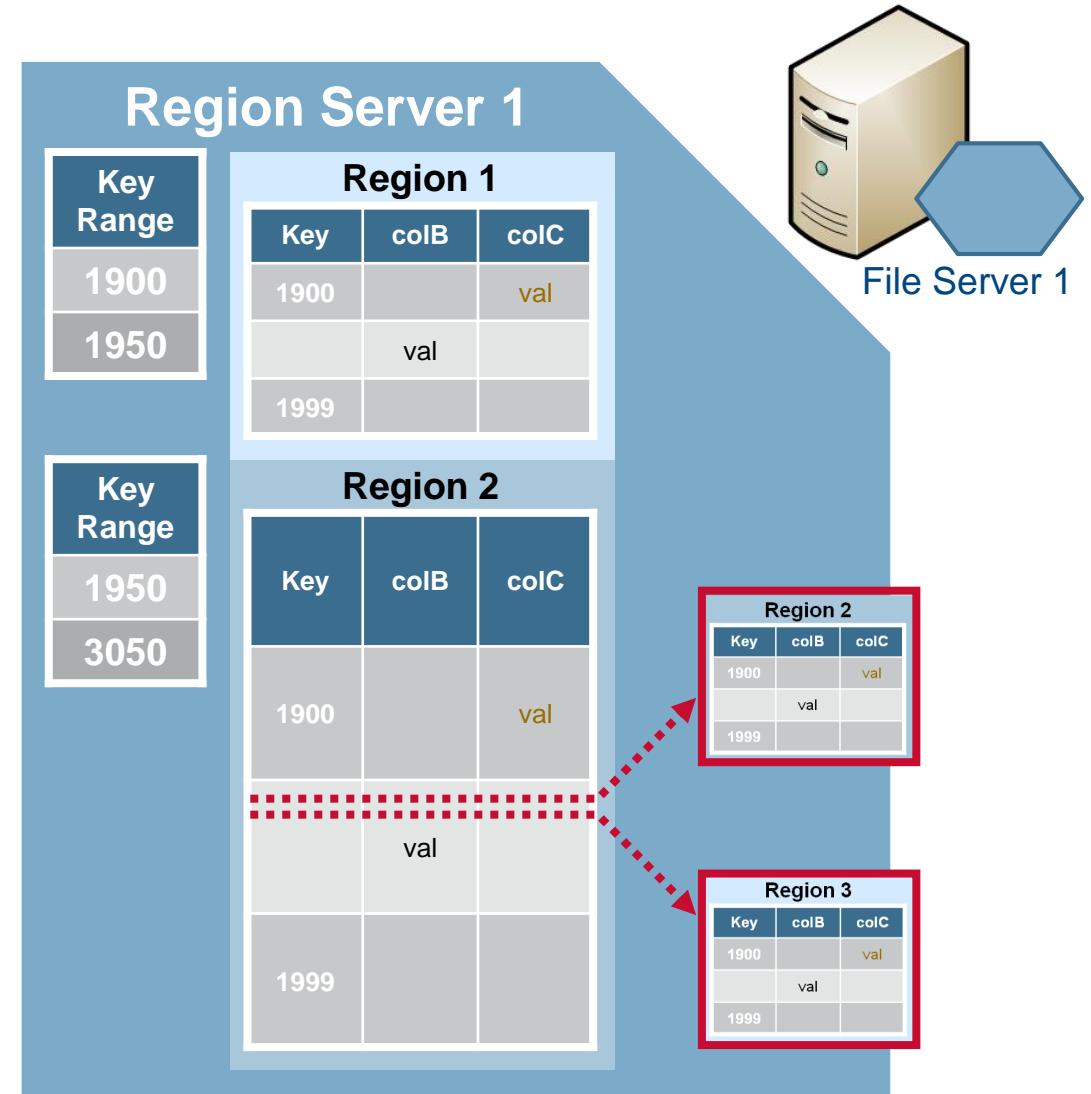
Hot-Spotting & Region Splits





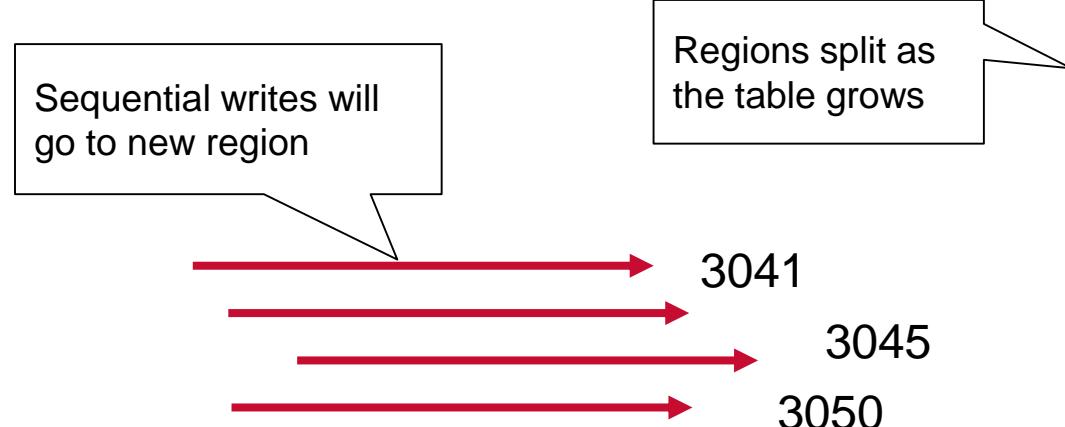
Hot-Spotting & Region Splits

Sequential writes will go to new region





Hot-Spotting & Region Splits



Region Server 1		
Region 1		
Key Range	colB	colC
1900		val
1950	val	
1999		

Region 2		
Key Range	colB	colC
1950		val
2050	val	
1999		

Region 3		
Key Range	colB	colC
2051		val
3050	val	
1999		





Hot-Spotting Summary

- Caused by **row keys** that are *written* in **sequential** order
 - Ex: row keys written in order, 0000, 0001, 0002, 0003...
- All writes go to only one server at a time
 - **Bottlenecks** write performance
- Results in **Inefficient splitting**
 - Regions fill to half the maximum, but never more.

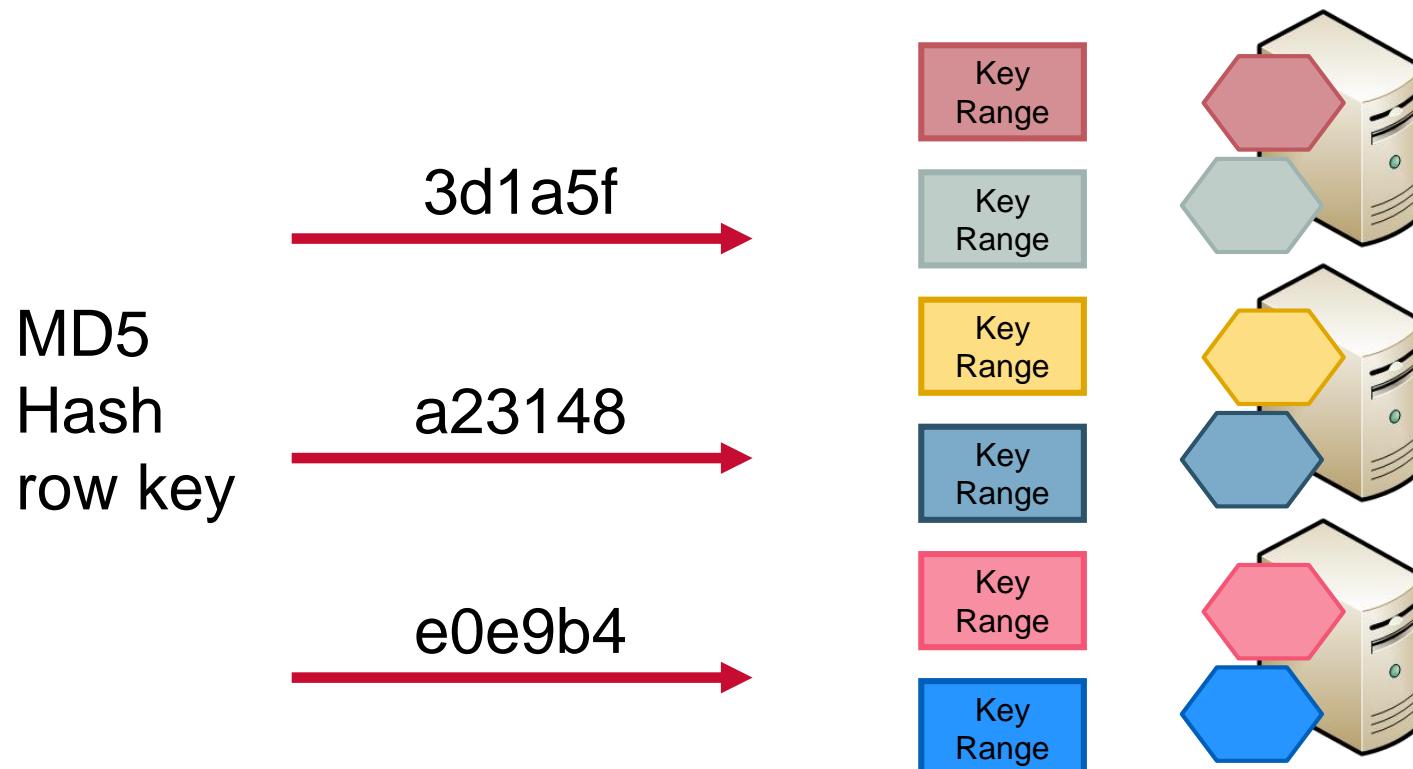


Random Keys

Random writes will go to different regions

If table was pre-split or **big enough** to have split

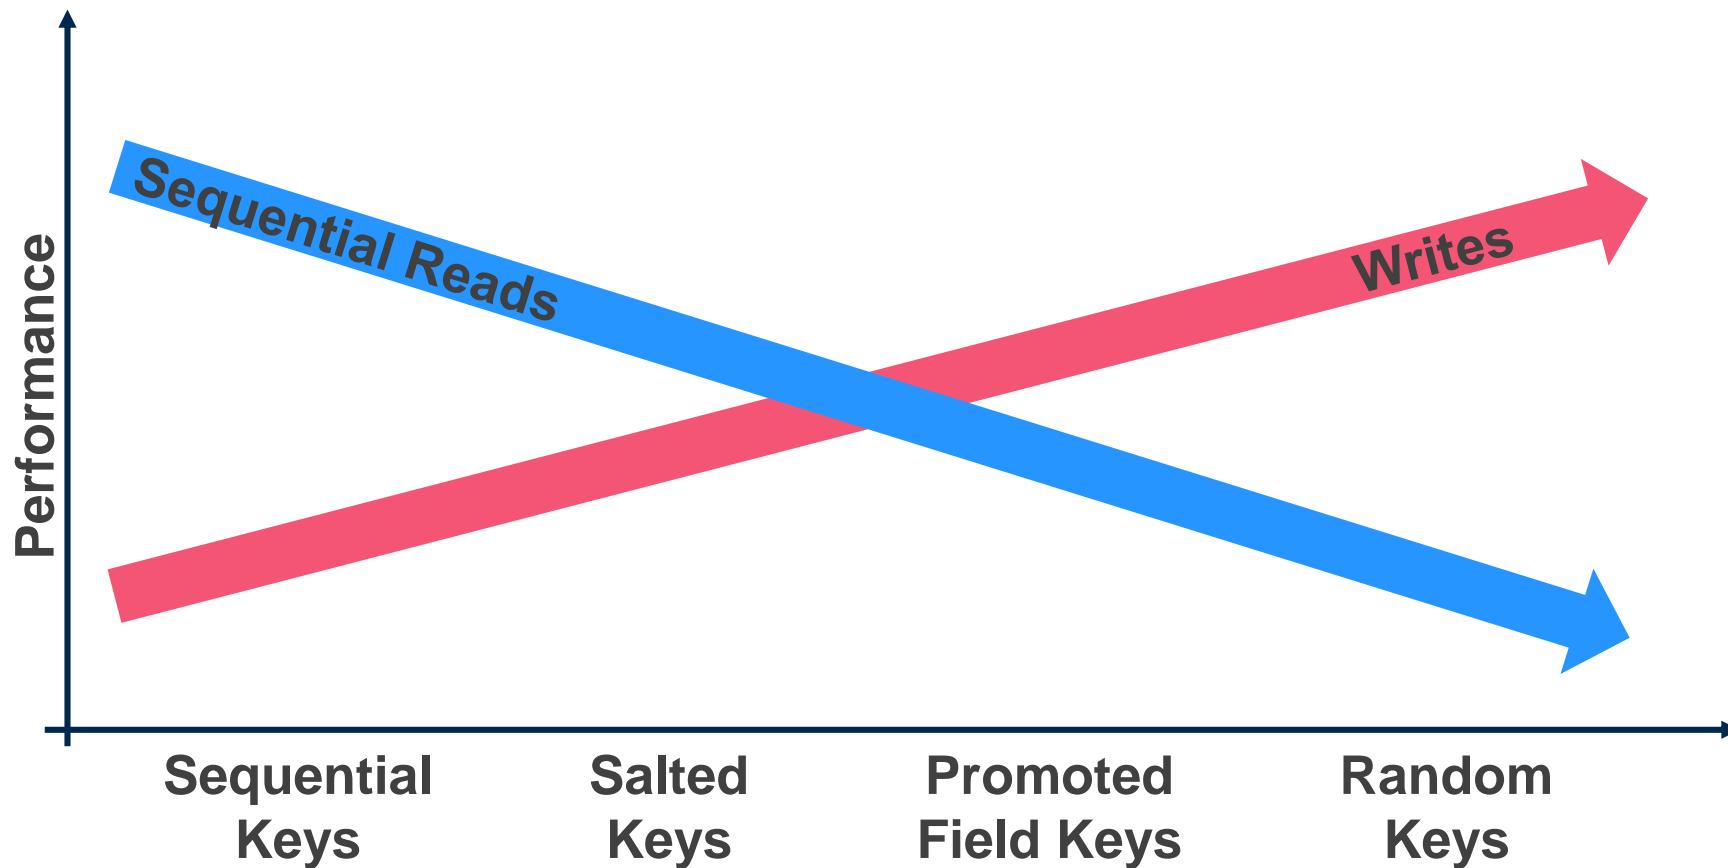
```
d = MessageDigest.getInstance ("MD5");  
byte[] prefix = d.digest(Bytes.toBytes(s));
```





Sequential vs. Random keys

Random is better for **writing** , but **sequential** is better for **scanning** row keys

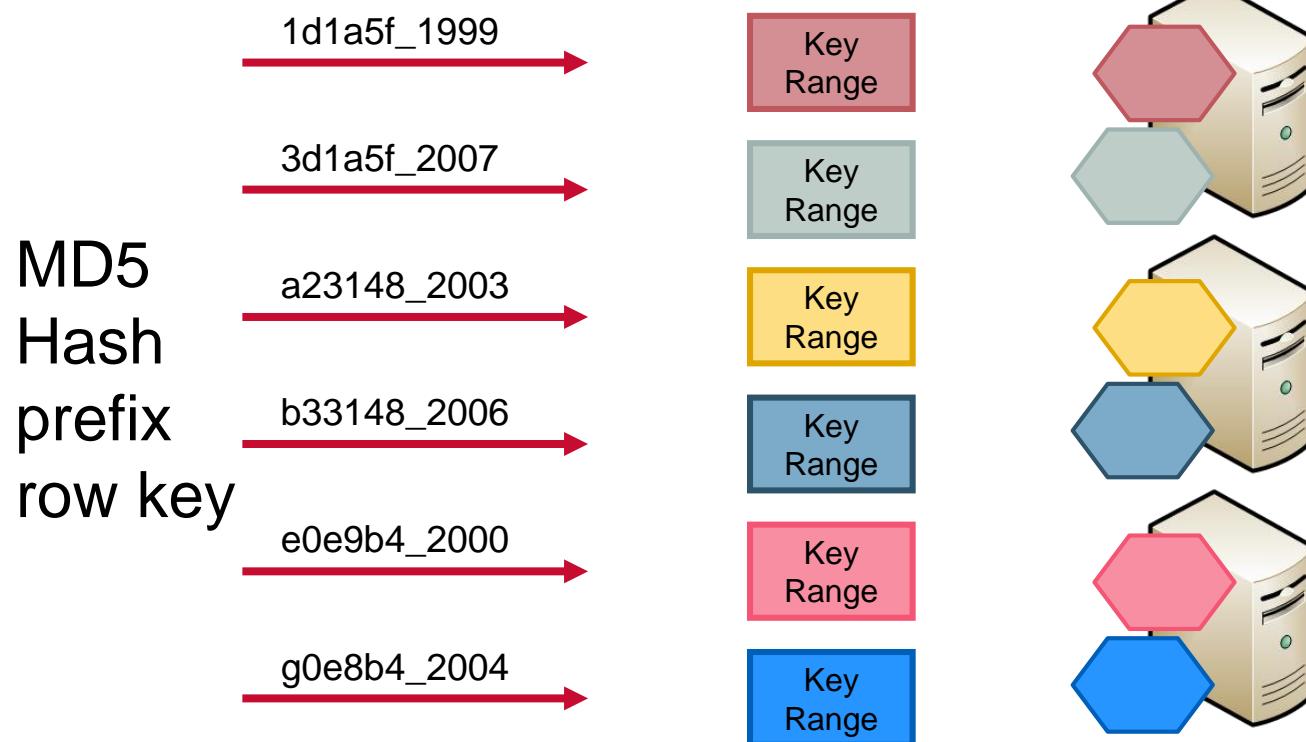




Prefix with a Hashed Field Key

- Prefix the row key with a (shortened) hash:

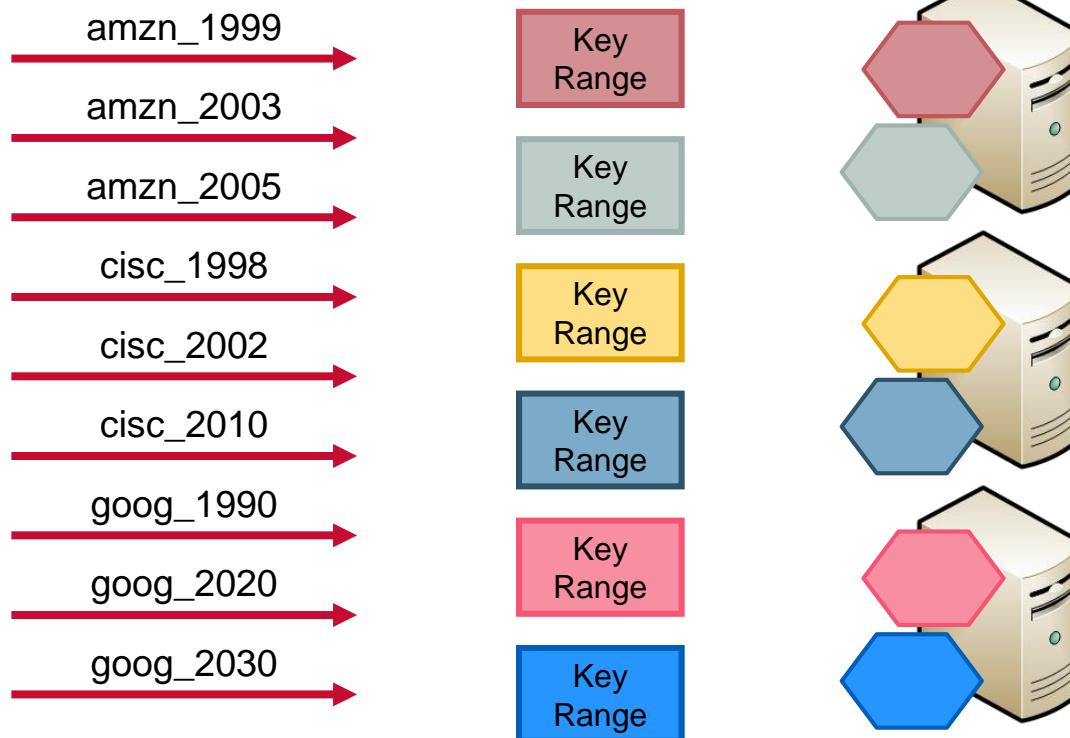
```
byte[] hash = d.digest(Bytes.toBytes(fieldkey));  
Bytes.putBytes(rowkey, 0, hash, 0, length);
```





Prefix, Promote a Field Key

- Prefix or promote identifying/searchable value to front of key

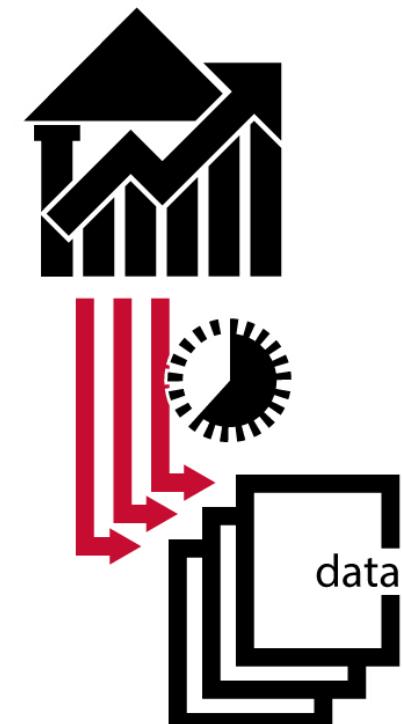




Consider Access Patterns for Application

- Which trade data needs fastest access (or most frequent)?
 - **Row key ordering**
- What if you want to retrieve the stocks by symbol & date?
 - Scan by: STOCKSYMBOL_TIMESTAMP

SYMBOL	+	timestamp			
Row Key					
AMZN_1391813876369					
AMZN_1391813876370					
GOOG_1391813876371					



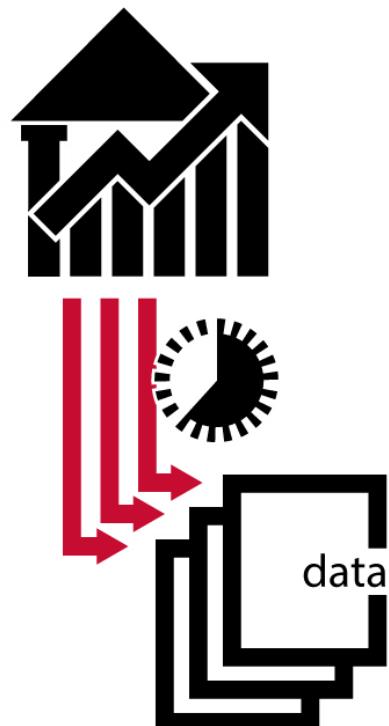
- What if you usually want to retrieve the most recent?



Last In First Out Access: Use Reverse-Timestamp

- Row keys are sorted in increasing order
- For fast access to **most-recent** writes:
 - Design composite row key with **reverse-timestamp** that **decreases** over time.
 - Scan by row key prefix **Decreasing**: **[MAXTIME-TIMESTAMP]**
 - Ex: `Long.MAX_VALUE-date.getTime()`

SYMBOL	+	Reverse timestamp			
Key					
AMZN_98618600666					
AMZN_98618600777					
GOOG_98618608888					



*`Long.MAX_VALUE = 263 - 1`



Consider Access Patterns for Application

How will data be retrieved?

- What are the needs for atomicity of transactions?
 - ***Column design***
 - **More Values in a single row**
 - Works well to get or **update multiple values**



Tall or Flat Tables

Tall Narrow

Term	Percentage
GMOs	100%
Organic	100%
Natural	100%
Artificial	100%
Organic	~95%
Natural	~95%
Artificial	~95%
Organic	~85%
Natural	~85%
Artificial	~85%

Flat Wide

A 10x10 grid of squares. The first column contains four dark blue squares, while the other nine columns contain light blue squares. A thick black border surrounds the grid.



Consider Access Patterns for Application



- Are Price & Volume data typically accessed together, or are they unrelated?
 - *Column family structure*
- Column Families
 - Group data that will be read & stored together



Tall Table for Stock Trades

Row key format:

SYMBOL

+

Reverse timestamp

Ex: AMZN_98618600888

Group data that will be
read & stored together



Row key	CF: CF1	
	CF1:price	CF1:vol
AMZN_98618600666	12.34	2000
AMZN_98618600777	12.41	50
AMZN_98618600888	12.37	10000
CSCO_98618600777	23.01	1000



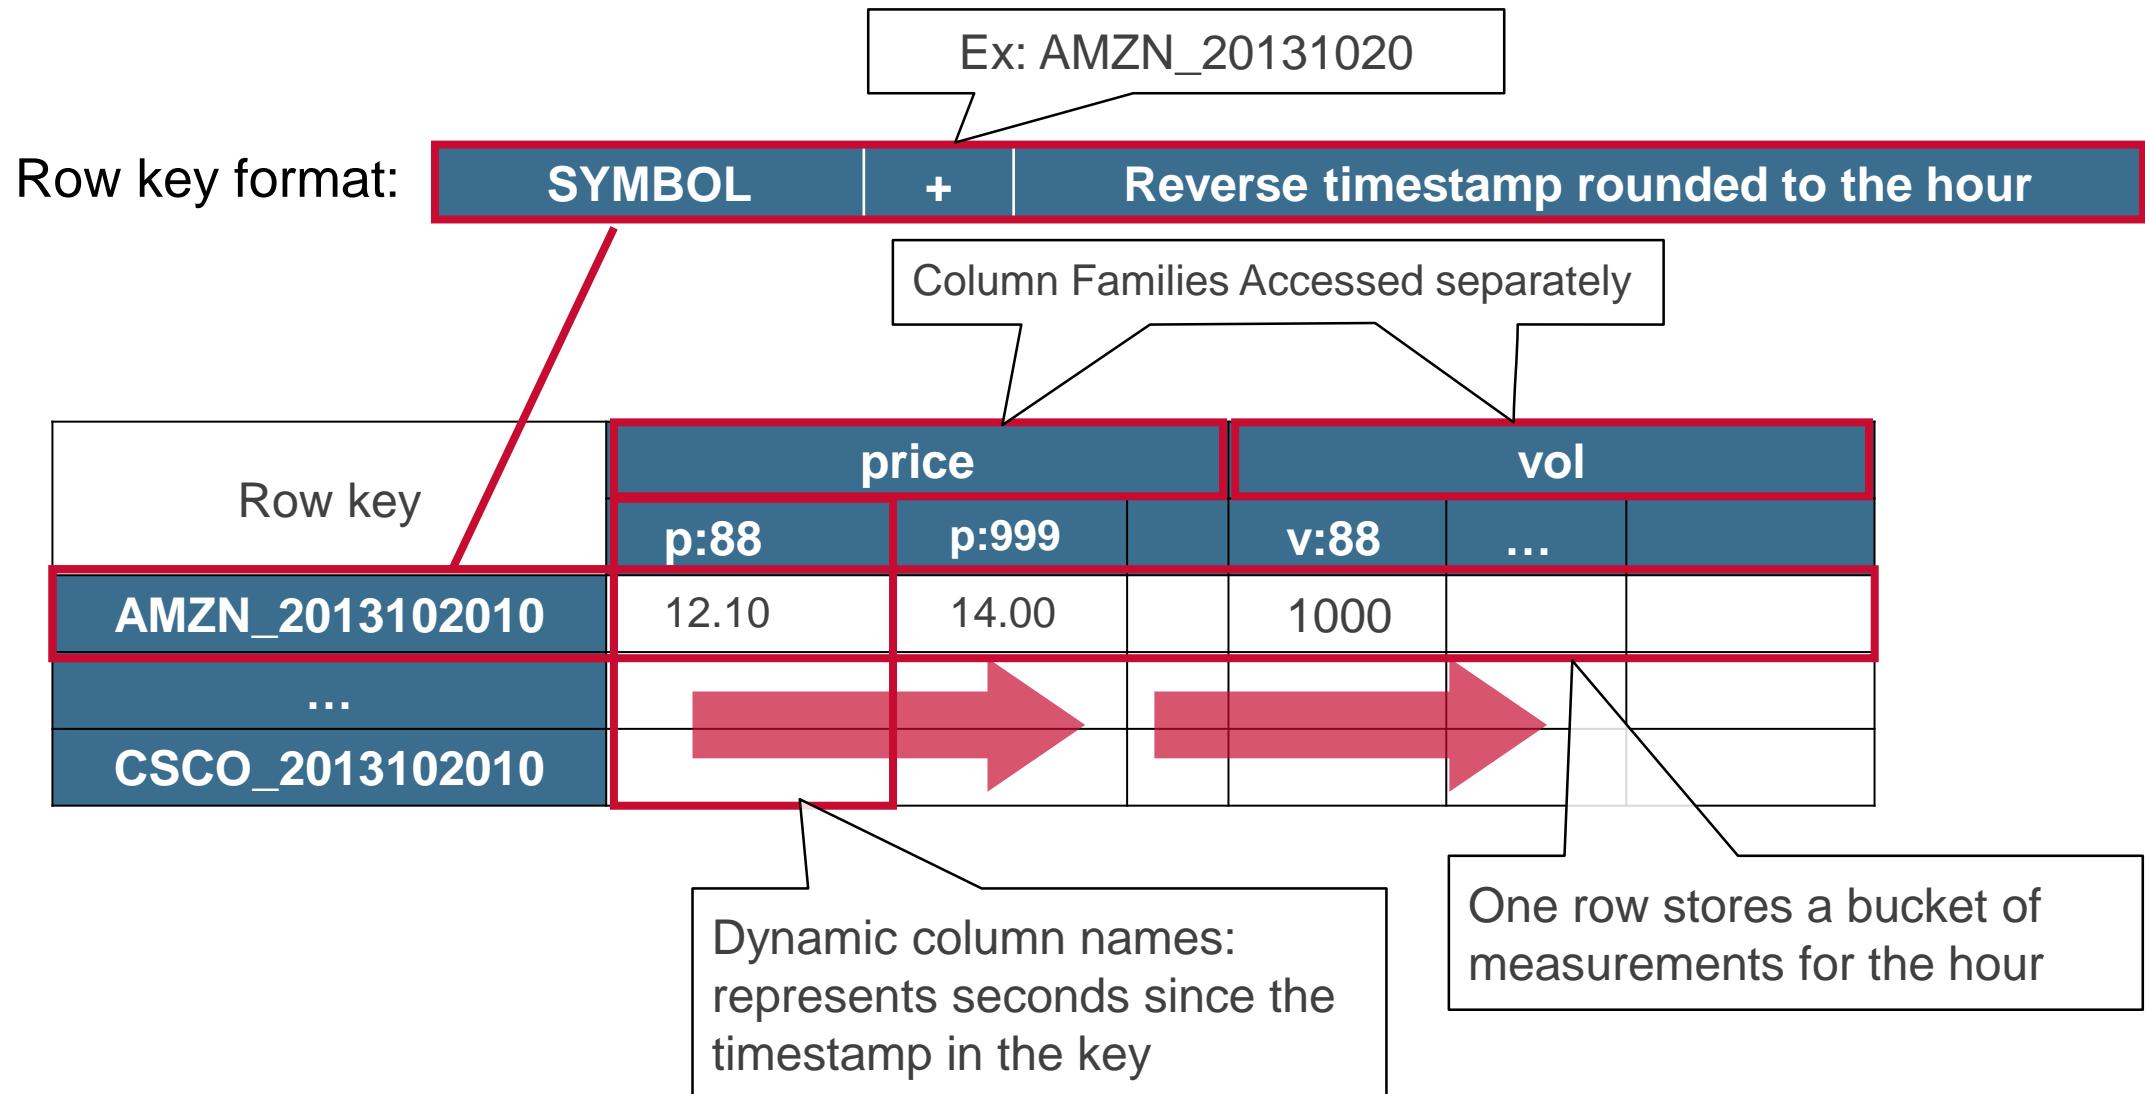
Consider Access Patterns for Application



- Are Price & Volume data typically accessed separately ?
 - *Column family structure*
- Column Families
 - Separate data that will be **not** be read together
- Columns
 - Column names are **dynamic**



Medium Wide Table for Stock Trades





Consider Access Patterns for Application



- Column Families
 - How many Versions?
 - *Max Versions*



Wide Table for Stock Trades

Row key format: **SYMBOL | + | date YYYYMMDD**

Ex: **AMZN_20131020**

Date in the row key	Hour in the column name	Separate price & volume data into column families	Set Column Family to store Max Versions, timestamp in the version
Row key		CF price	CF vol
AMZN_20131020	price:00	...	vol:00
...	12.37	12.34	10000
CSCO_20130817	23.01		1000



Flat-Wide Vs. Tall-Narrow Tables

- Tall-Narrow provides better **query granularity**
 - Finer grained Row Key
 - Works well with scan
- Flat-Wide supports built-in row **atomicity**
 - More Values in a single row
 - Works well to update multiple values (row atomicity)
 - Works well to get multiple associated values

Column names can be dynamic, every row does not need to have same columns

Lesson: Schemas can be very flexible and can even change on the fly

Lesson: Have to know the queries to design in performance



Lab – Import airline data into HBase

```
$ tail ontime.csv
```

```
2014,1,1,31,5,2014-01-31,WN,N7704B,228,TUS,LAS,1946,46.00,1958,43.00,0.00,,  
75.00,60.00,365.00,11.00,0.00,0.00,0.00,32.00,
```

```
$ wc -l ontime.csv
```

```
471949 lines of flight information
```

Import command:

```
-Dimporttsv.columns=$CF:year,$CF:qtr,$CF:month,$CF:dom,$CF:dow,HBASE_ROW_KEY,  
$CF:carrier,$CF:tailnum,$CF:flightnumber,$CF:origin,$CF:dest,$CF:deptime,$CF:depdelay,$CF  
:arrtime,$CF:arrdelay,$CF:cncl,$CF:cnclcode,$CF:elaptime,$CF:airtime,$CF:distance,$CF:carri  
erdelay,$CF:weatherdelay,$CF:nasdelay,$CF:securitydelay,$CF:aircraftdelay\
```



Lab – Import airline data into HBase

2014,1,1,31,5,**2014-01-31**,WN,N7704B,228,TUS,LAS,1946,46.00,1958,43.00,0.00,,75.00,60.00,365.00,11.00,0.00,0.00,0.00,32.00,

HBASE_ROW_KEY

Table: **airline**

Row-key	cf1											
	Date	year	Qtr	month	dom	dow	carrier	tailnum	Flight number	origin	dest	...
2014-01-31	2014	1	1		31	5	WN	N7704B	228	TUS	LAX	...



Why does the table only have 31 rows ??

\$ hbase shell

hbase(main):003:0> count '/tables/airline'
31 row(s) in 0.0560 seconds

**hbase(main):003:0> scan '/tables/airline',
{LIMIT => 1}**

ROW	COLUMN+CELL
2014-01-01	column=cf1:aircraftdelay, timestamp=1424960635661, value=0.00
2014-01-01	column=cf1:airtime, timestamp=1424960635661, value=201.00
2014-01-01	column=cf1:arrdelay, timestamp=1424960635661, value=26.00
2014-01-01	column=cf1:arrtime, timestamp=1424960635661, value=1311
2014-01-01	column=cf1:carrier, timestamp=1424960635661, value=WN
2014-01-01	column=cf1:carrierdelay, timestamp=1424960635661, value=0.00
2014-01-01	column=cf1:cndl, timestamp=1424960635661, value=0.00
2014-01-01	column=cf1:cndlcode, timestamp=1424960635661, value=
2014-01-01	column=cf1:depdelay, timestamp=1424960635661, value=1.00
2014-01-01	column=cf1:deptime, timestamp=1424960635661, value=0931
2014-01-01	column=cf1:dest, timestamp=1424960635661, value=MCO
2014-01-01	column=cf1:distance, timestamp=1424960635661, value=1142.00
2014-01-01	column=cf1:dom, timestamp=1424960635661, value=1
2014-01-01	column=cf1:dow, timestamp=1424960635661, value=3
2014-01-01	column=cf1:dummy, timestamp=1424960635661, value=
2014-01-01	column=cf1:elaptime, timestamp=1424960635661, value=195.00
2014-01-01	column=cf1:flightnumber, timestamp=1424960635661, value=1147
2014-01-01	column=cf1:month, timestamp=1424960635661, value=1
2014-01-01	column=cf1:nasdelay, timestamp=1424960635661, value=26.00
2014-01-01	column=cf1:origin, timestamp=1424960635661, value=MHT
2014-01-01	column=cf1:qtr, timestamp=1424960635661, value=1
2014-01-01	column=cf1:securitydelay, timestamp=1424960635661, value=0.00
2014-01-01	column=cf1:tailnum, timestamp=1424960635661, value=N264LV
2014-01-01	column=cf1:weatherdelay, timestamp=1424960635661, value=0.00
2014-01-01	column=cf1:year, timestamp=1424960635661, value=2014

1 row



Lab – Import airline data into HBase

mapping to Row Key and Columns:

```
timing:year=2014,  
timing:qtr=1,  
timing:month=1,  
timing:dom=31,  
timing:dow=5,  
HBASE_ROW_KEY_3=2014-01-31, // date  
HBASE_ROW_KEY_1=WN, // carrier  
info:tailnum=N7704B,  
HBASE_ROW_KEY_2=228, // flight number  
HBASE_ROW_KEY_4=TUS, // orig  
HBASE_ROW_KEY_5=LAS, // dest  
timing:deptime=1946,  
delay:depdelay=46.00,
```

```
timing:arrtime=1958,  
delay:arrdelay=43.00,  
info:cncl=0.00,  
info:cnclcode="",  
stats:elaptime=75.00,  
stats:airtime=60.00,  
stats:distance=365.00,  
delay:carrierdelay=11.00,  
delay:weatherdelay=0.00,  
delay:nasdelay=0.00,  
delay:securitydelay=0.00,  
delay:aircraftdelay=32.00,
```



Scans are limited

\$ hbase shell

```
scan '/tables/airline'
```

```
scan '/tables/airline', { STARTROW => '2014'}
```

```
scan '/tables/airline', { STARTROW => '2014-01-20', STOPROW => '2014-01-21'}
```

1 row



Lab – Import airline data into HBase

Import mapping to Row Key and Columns:

Row-key	delay			info			stats		timing	
Carrier-Flightnumber-Date-Origin-destination	Air Craft delay	Arr delay	Carrier delay	cncl	Cncl code	tailnum	distance	elaptime	arrtime	Dep time
AA-1-2014-01-01-JFK-LAX		13		0		N7704	2475	385.00	359	...



Do you see the advantages to this row key, CFs ?

\$ hbase shell

hbase(main):003:0> count '/tables/airline'

471,949 row(s)

```
hbase(main):003:0>scan '/tables/airline',
{STARTROW => 'WN-228-2014-01-31-TUS-LAS',
LIMIT => 1}
```

ROW	COLUMN+CELL
WN-228-2014-01-31-TUS-LAS	column=delay:aircraftdelay, timestamp=1425775096289, value=32.00
WN-228-2014-01-31-TUS-LAS	column=delay:arrdelay, timestamp=1425775096289, value=43.00
WN-228-2014-01-31-TUS-LAS	column=delay:carrierdelay, timestamp=1425775096289, value=11.00
WN-228-2014-01-31-TUS-LAS	column=delay:depdelay, timestamp=1425775096289, value=46.00
WN-228-2014-01-31-TUS-LAS	column=delay:nasdelay, timestamp=1425775096289, value=0.00
WN-228-2014-01-31-TUS-LAS	column=delay:securitydelay, timestamp=1425775096289, value=0.00
WN-228-2014-01-31-TUS-LAS	column=delay:weatherdelay, timestamp=1425775096289, value=0.00
WN-228-2014-01-31-TUS-LAS	column=info:cncl, timestamp=1425775096289, value=0.00
WN-228-2014-01-31-TUS-LAS	column=info:cnclcode, timestamp=1425775096289, value=
WN-228-2014-01-31-TUS-LAS	column=info:dummy, timestamp=1425775096289, value=
WN-228-2014-01-31-TUS-LAS	column=info:tailnum, timestamp=1425775096289, value=N7704B
WN-228-2014-01-31-TUS-LAS	column=stats:airtime, timestamp=1425775096289, value=60.00
WN-228-2014-01-31-TUS-LAS	column=stats:distance, timestamp=1425775096289, value=365.00
WN-228-2014-01-31-TUS-LAS	column=stats:elaptime, timestamp=1425775096289, value=75.00
WN-228-2014-01-31-TUS-LAS	column=timing:arptime, timestamp=1425775096289, value=1958
WN-228-2014-01-31-TUS-LAS	column=timing:deptime, timestamp=1425775096289, value=1946
WN-228-2014-01-31-TUS-LAS	column=timing:dom, timestamp=1425775096289, value=31
WN-228-2014-01-31-TUS-LAS	column=timing:dow, timestamp=1425775096289, value=5
WN-228-2014-01-31-TUS-LAS	column=timing:month, timestamp=1425775096289, value=1
WN-228-2014-01-31-TUS-LAS	column=timing:qtr, timestamp=1425775096289, value=1
WN-228-2014-01-31-TUS-LAS	column=timing:year, timestamp=1425775096289, value=2014

1 row



Comparing Relational vs HBase Schema design And more Complex Schema design





\$ hbase shell

```
scan '/tables/airline', { STARTROW => 'AA-1-2014-01-01-JFK-LAX',  
STOPROW => 'AA-10'}
```

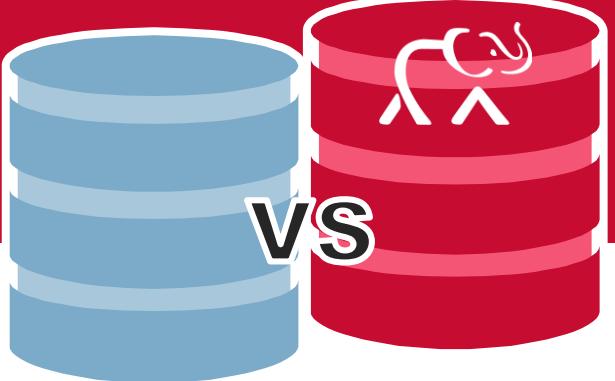
```
scan '/tables/airline', {COLUMNS=>['delay'], STARTROW => 'DL'}
```

```
scan '/tables/airline', FILTER=>"ValueFilter(=,'binary:239.00')"
```

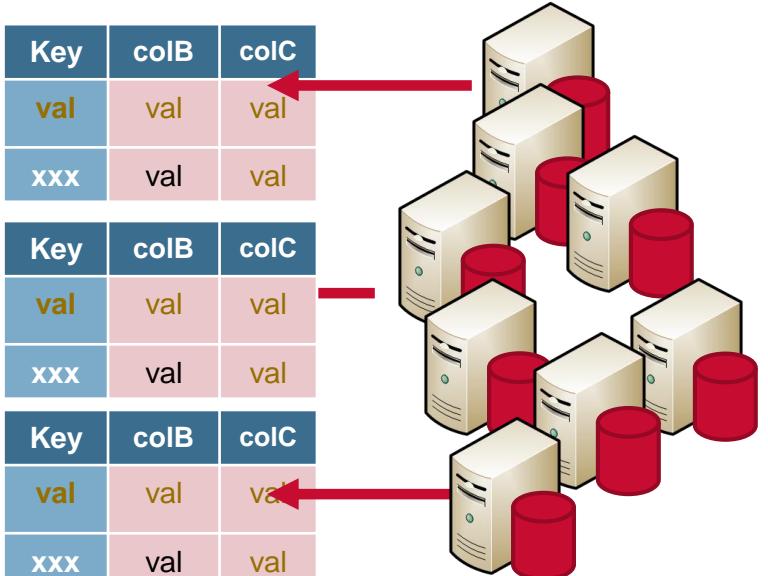
1 row



Relational vs. HBase Schemas



- Relational design
 - Data centric, focus on entities & relations
 - Query joins
 - New views of data from different tables easily created
 - Does not scale across cluster
- HBase is designed for clustering:
 - Distributed data is stored and accessed together
 - Query centric, focus on how the data is read
 - Design for the questions



Data that is accessed together is stored together



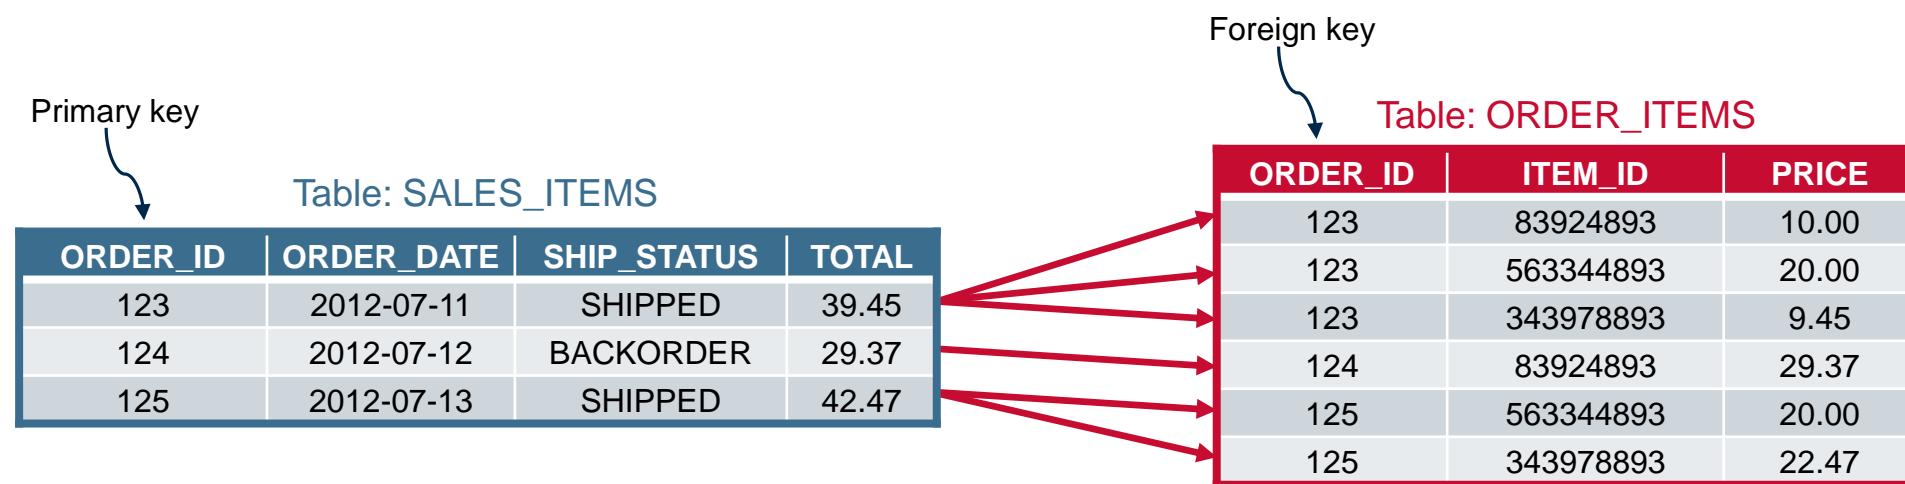
Normalization

Goal of Normalization

- Eliminate redundant data
- Put repeating information in its own table

Normalized database

- Causes joins
 - Data has to be retrieved from more tables
 - Queries can take more time





Denormalization

OrderId	Data:date	item:id1	item:id2	item:id3
123	20131010	\$10	\$20	\$9.45

Order & Order_Items in **same table**

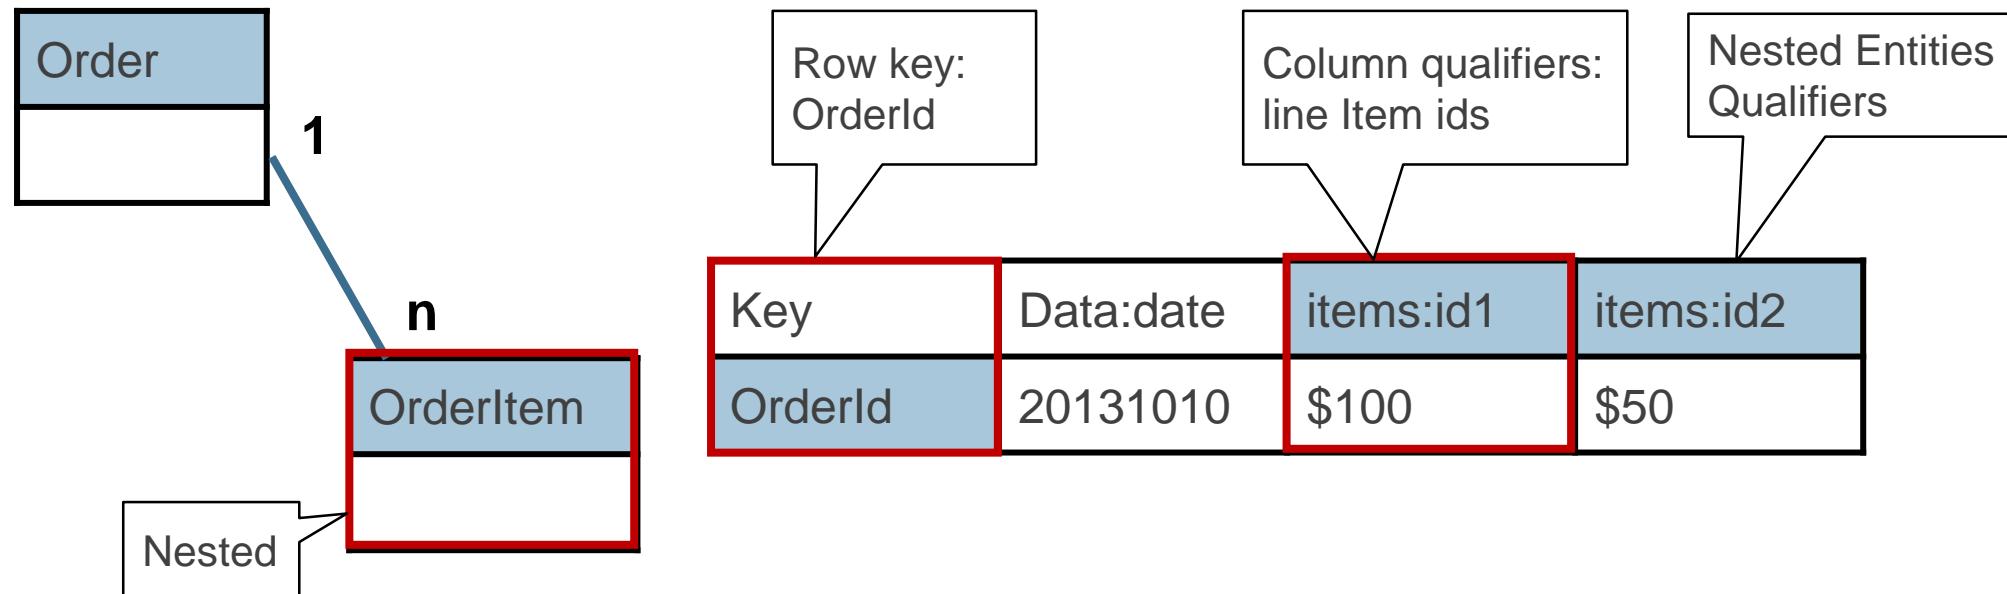
Duplicating data in more than one table

Replacement for JOINs

The table illustrates denormalization by showing an Order row and three corresponding Order_Items rows. The Order row contains fields: OrderId (123), Data:date (20131010). The Order_Items row contains fields: item:id1 (\$10), item:id2 (\$20), and item:id3 (\$9.45). A large watermark 'denormalization' is overlaid across the table cells.



Parent-Child Relationship –Nested Entity



A one-to-many relationship can be modeled as a **single** row
Embedded, Nested Entity

Reads are faster across a cluster
retrieve data about entity and related entities in **one read**

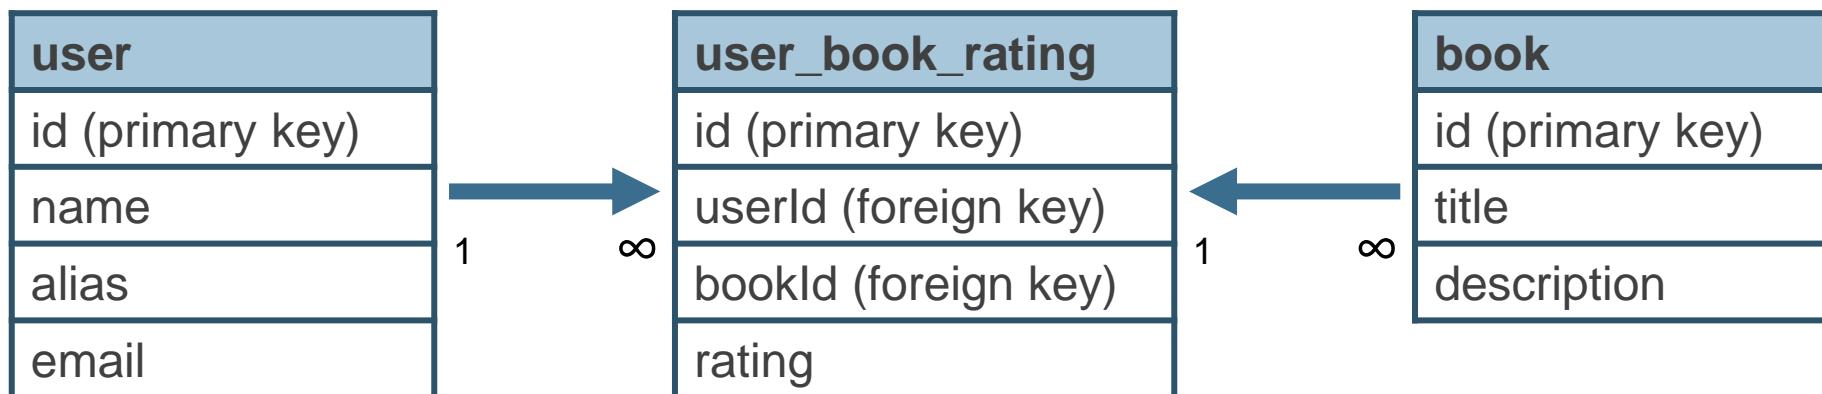


Many to Many Relationship - RDBMS

Queries

- Get name for user x
- Get title for book x
- Get books and corresponding ratings for userId x
- Get all userids and corresponding ratings for book y

Online book store





Many to Many Relationship - HBase

Queries

- Get books and corresponding ratings for userId x
- Get all userids and corresponding ratings for book y

User table Column family for book ratings by userid for bookids

Key	data:fname	...	rating:bookid1	rating:bookid2
userid1			5	4

Book table Column family for ratings for bookid by userid

Key	data:title	...	rating:userid1	rating:userid2
bookid1			5	4



Generic Data, Event Data, Entity-Attribute-Value

- **Generic table:** Entity, Attributes, Values
 - Event Id, Event attributes, Values
 - Object-property-value , name-value pairs , **schema-less**

patientXYZ-ts1, Temperature , "102"

patientXYZ-ts1, Coughing, "True"

patientXYZ-ts1, Heart Rate, "98"

- **Advantage of HBase**
 - Define columns on the fly, put attribute name in column qualifier
 - Group data by column families

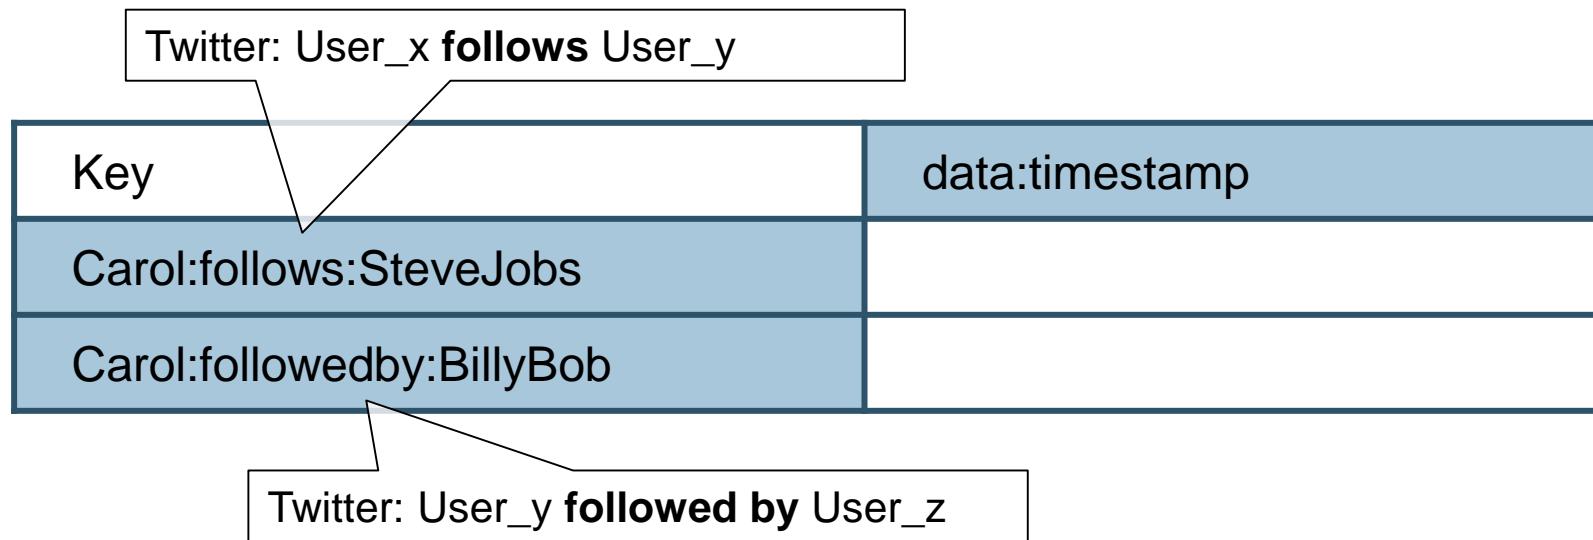
Event id=row key	Event type name=qualifier	Event measurement=value
Key	event:heartrate	event:coughing
Patientxyz-ts1	98	true
	event:temperature	102



Self Join Relationship - HBase

Queries

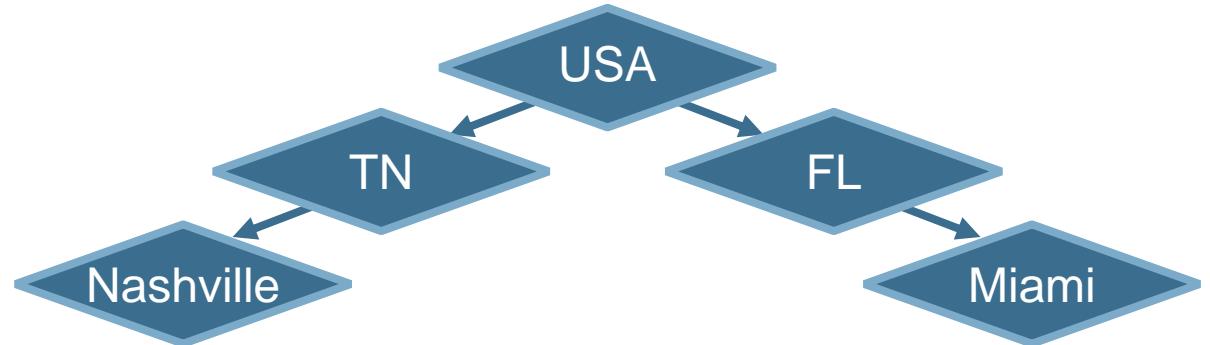
- Get all users who Carol follows
- Get all users following Carol





Tree, Graph Data

- Row=node
- Row key=node id
- Parents, children in columns
 - Col Qualifier=edge id



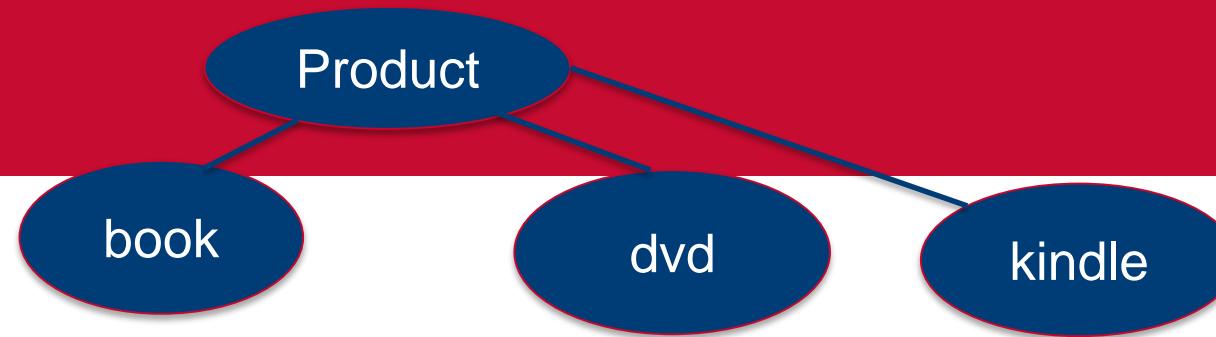
parent

children

Key	P:USA	P:TN	P:FL	C:TN	C:FL	C:Nashvl	C:Miami
USA				state	state		
TN	country					city	
FL	country						city
Nashville		state					
Miami			state				



Inheritance Mapping

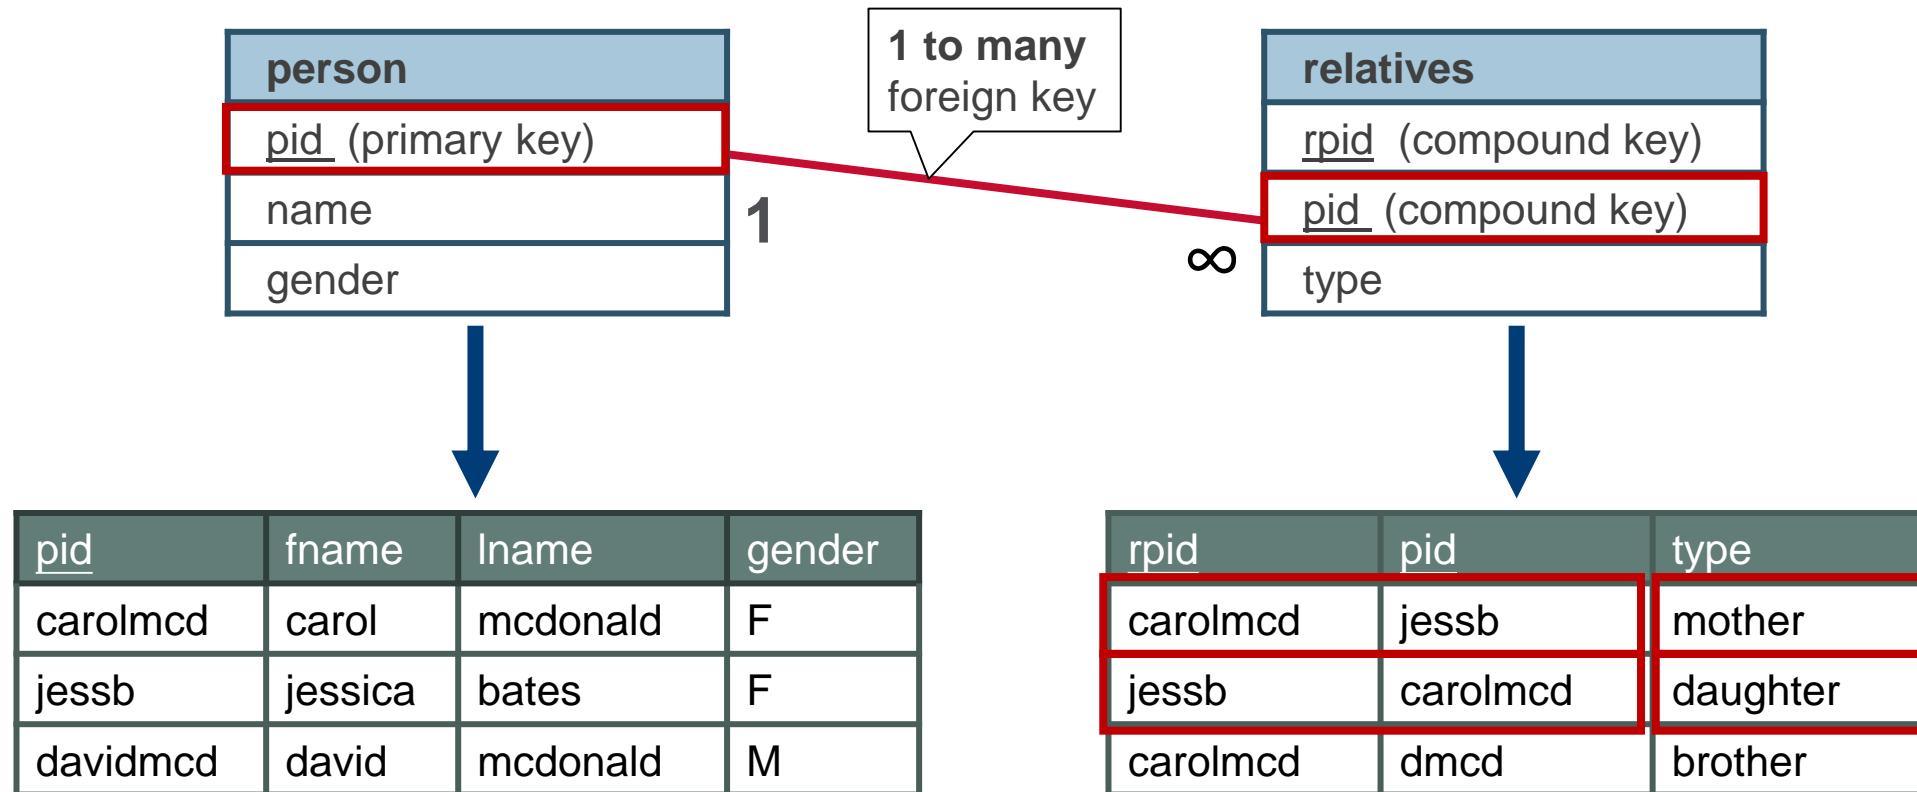


- Online Store Example Product table
 - Put **sub class type abbreviation** in **key prefix** for searching
 - Columns do not all have to be the same for different types

Key	price	title	details	model
Bok+id1	10	HBase		
Dvd+Id2	15	stones		
Kin+Id3	100			fire



Use Case 1 – Person's Relatives



get all the relatives for a userid ?



HBase Modeling Concepts

- 1 Identify Entities
- 2 Identify Queries
- 3 Identifying Attributes
- 4 Non-identifying attributes
- 5 Relationships
- 6 Secondary index (foreign keys)



1

Identify Entities

2

Identify Queries

- What information accessed together in one get
- What information needs scanning

3

Identifying Attributes

4

Non-identifying attributes

5

Relationships

6

Secondary index (foreign keys)

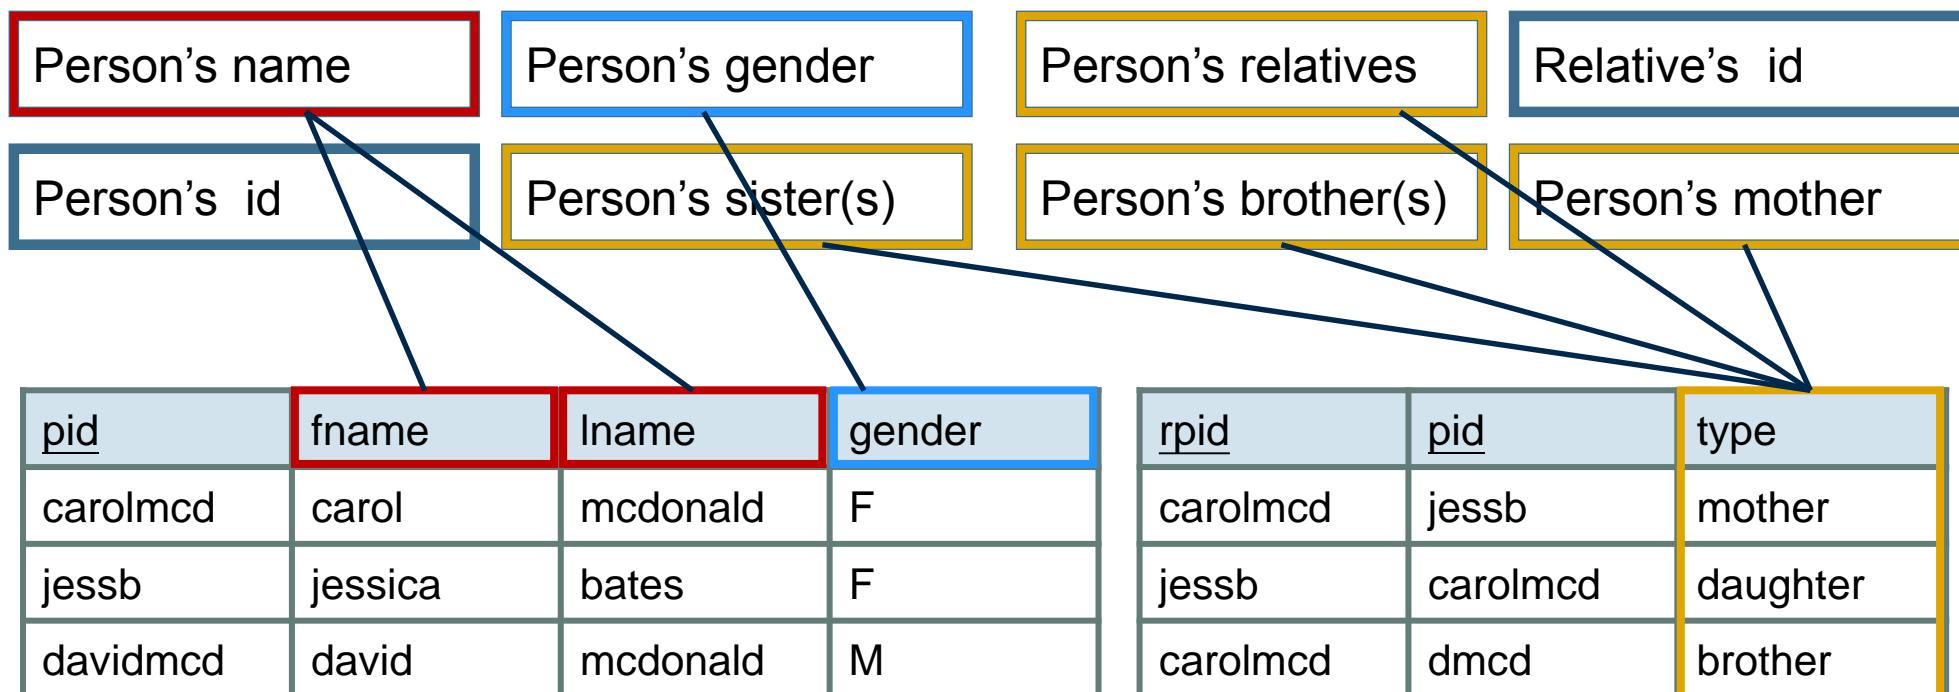


Use Case 1 - Queries

2

Identify queries

get all the relatives for a userid ?





HBase Modeling Concepts

- 1 Identify Entities
- 2 Identify Queries
- 3 Identifying Attributes
- 4 Non-identifying attributes
- 5 Relationships
- 6 Secondary index (foreign keys)



Use Case 1 – Identifying & Non-Identifying Attributes

3

Identifying attributes used in queries

Identify which attributes can be used to identify unique instances of the entity (entity = Person):

A. pid

B. pid + lname + fname

The diagram illustrates a database schema for a 'person' entity. At the top is a conceptual table labeled 'person' with four columns: 'pid' (primary key), 'name', and 'gender'. Below it is a detailed table with the same columns: 'pid', 'fname', 'lname', and 'gender'. A dashed vertical line connects the 'pid' column of the conceptual table to the 'pid' column of the detailed table, indicating that 'pid' is the primary key.

person			
pid (primary key)			
name			
gender			
carolmcd	carol	mcdonald	F
jessb	jessica	bates	F
davidmcd	david	mcdonald	M



Identifying & Non Identifying Attributes – Option A

3

Identifying attributes used in queries

Identify which attributes can be used to identify unique instances of the entity (entity = Person):

A. pid

Entity: Person

Identifying attribute: pid

4

Non-identifying attributes: fname, lname, gender

Person table

rowkey pid	Info: fname	Info: lname	Info: gender
caroljmcd	carol	mcdonald	F
jessicab	jessica	bates	F
...			

Row key = pid Column Family = Info Columns = fname, lname, gender get, scan by pid



Identifying & Non Identifying Attributes – Option B

3

Identifying attributes used in queries

Identify which attributes can be used to identify unique instances of the entity (entity = Person):

B. pid + lname + fname

Entity: Person

Identifying attribute: **pid, fname, lname**

4

Non-identifying attributes: **gender**

Person table			Depends on if you want to be able to scan by first name and last name Or just pid
Row key = pid_fname_lname	Column Family = Info	Columns = gender	
rowkey pid_fname_lname	Info: gender		
caroljmcd_carol_mcdonald	F		
jessicab_jessica_bates	F		
...			

The diagram illustrates a Person table in a column-oriented database. The table has three columns: Row key, Column Family, and Columns. The Row key column contains the value "rowkey pid_fname_lname". The Column Family column contains the value "Info: gender". The Columns column contains the value "F". A red box highlights the Row key and Column Family columns. A yellow box highlights the "Info: gender" cell. Arrows point from the text labels above the table to their corresponding columns. A red box also surrounds the text "Depends on if you want to be able to scan by first name and last name Or just pid".



HBase Modeling Concepts

- 1 Identify Entities
- 2 Identify Queries
- 3 Identifying Attributes
- 4 Non-identifying attributes
- 5 Relationships
 - One to many -> nested or embedded entities in column family
- 6 Secondary index (foreign keys)



Use Case 1 - Relationships

- 5 One to many → nested or embedded entities in column family

How to model?

Parent identifying attribute(s) = row key

Child identifying attributes = column qualifiers



Use Case 1 – Relationships (2)

5

What are the child identifying attributes?

(Click the column heading(s) from the table on the right.)

Based on which attribute you use as column qualifier, you have different options.

relatives
<u>rpid</u> (compound key)
<u>pid</u> (compound key)
type



rpid	pid	type
carolmcd	jessb	mother
jessb	carolmcd	daughter
carolmcd	dmcd	brother



Use Case 1 – Relationships (3)

5

Option A

Child identifying attribute = type

Option B

Child identifying attribute = rpid



Relationships #A

5

Option A

Child identifying attribute = type

Person table

rowkey	info: fname	info: lname	info: gender	relation: brother1	relation: daughter1	relation: daughter2	relation: mother	relation: sister1	...
pid									
caroljmcd	carol	mcdonald	F	davidmcd	jessicab	sarahb			
jessicab	jessica	bates	F				caroljmcd	sarahb	
sarahb							caroljmcd	jessicab	
davidmcd	david							carol	

Info column family

relation column family

Requires keeping a counter
for number of sisters,
brothers, uncles...



Relationships #2

5

Option B

Child identifying attribute = rpid

Person table

rowkey	info: fname	info: lname	info: gender	...	relation: jessicab	relation: sarahb	relation: caroljmcd	relation: davidmcd	...
pid									
caroljmcd	carol	mcdonald	F		daughter	daughter		brother	
jessicab	jessica	bates	F			sister	mother	uncle	
sarahb	sarah		F		sister		mother	uncle	
davidmcd	david		M		uncle	uncle	brother		

Info column family

relation column family

Dynamic column name



HBase Modeling Concepts

- 1 Identify Entities
- 2 Identify Queries
- 3 Identifying Attributes
- 4 Non-identifying attributes
- 5 Relationships
- 6 Secondary index (foreign keys)
 - Put identifying data in lookup table



Use Case 1 – Secondary index

6

Secondary index (foreign keys)

- Put identifying data in lookup table

Person table

rowkey	info: fname	Info: lname	Info: gender
caroljmcd	carol	mcdonald	F
jessicab	jessica	bates	F
...			

Relatives lookup table

rowkey	Info:ts
caroljmcd_mother_jessicab	125666
caroljmcd_mother_sarahb	125675
caroljmcd_brother_davidmcd	
jessicab_daughter_caroljmcd	
jessicab_sister_sarahb	

Requires querying 2 tables
Good solution if relationships (relatives) need frequent updating



Summary – Use Case 1

1 Entity = Person

2 Queries – get person information
Get all relatives for person
Get all sisters for person, etc.

3 Identifying attributes → row key
pid, lname, fname

4 Non-identifying attributes: gender

5 Relationships:

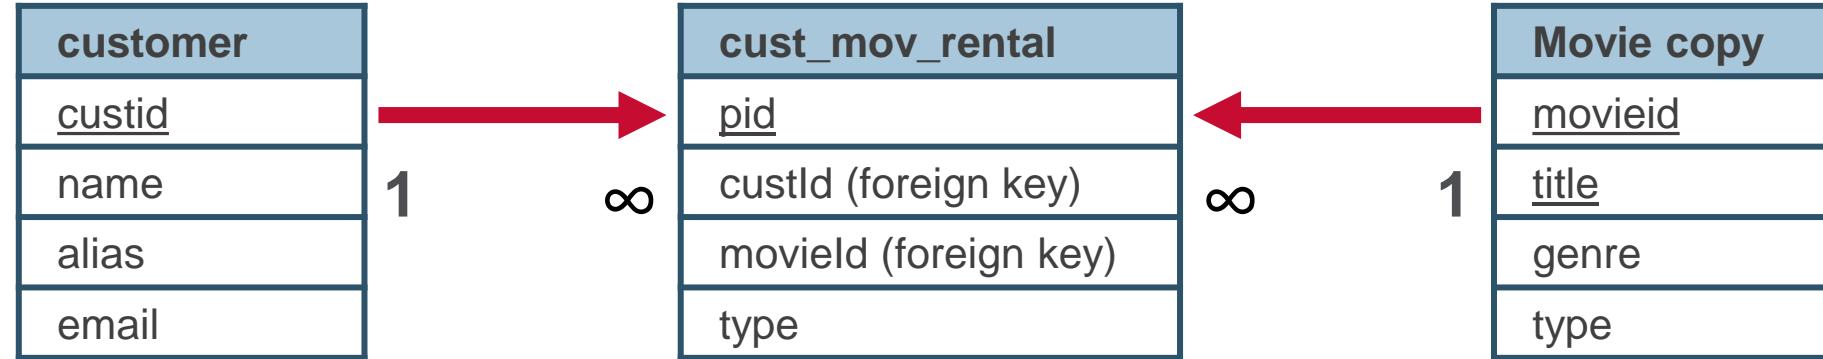
1. Use type → need counter
2. Use rpid → dynamic

6 Secondary index

3. Use foreign key → lookup table



Use Case 2 - Customers Movies Rental



<u>pid</u>	Name	email
carolj	carol	cm@y
jessb	jessica	
dmcd	david	

<u>pid</u>	rid	type
carolj	TS1	DVD
jessb	BC1	online

<u>movieid</u>	<u>title</u>	genre
TS1	Toy story	drama
BC1	Big Chill	drama



Identifying, Non Identifying Attributes #1

Customer table			
rowkey	info:	Info:	...
custid	name	gender	
Movie table			
rowkey	info:	Info:	...
movieid	title	description	

Info column family

Info column family



- Identifying Attributes used in queries-> row key
- Non identifying attributes -> columns
- **Put identifying data for relationships:**
 - in a **column family column qualifier**
 - Or put it in a lookup table
- **Relationships:**
 - **Many to many -> 2 tables:**
 - ?



Relationships #1

Customer table		Info column family	rental column family		
rowkey	info:	Info: name	... rental: movieid1	rental: movieid2	...
custid1			dvd		
custid2				online	

Movie table		Info column family	rental column family		
rowkey	info:	Info: description	... rental: custid1	rental: custid2	...
movieid1			dvd		
movieid2				online	



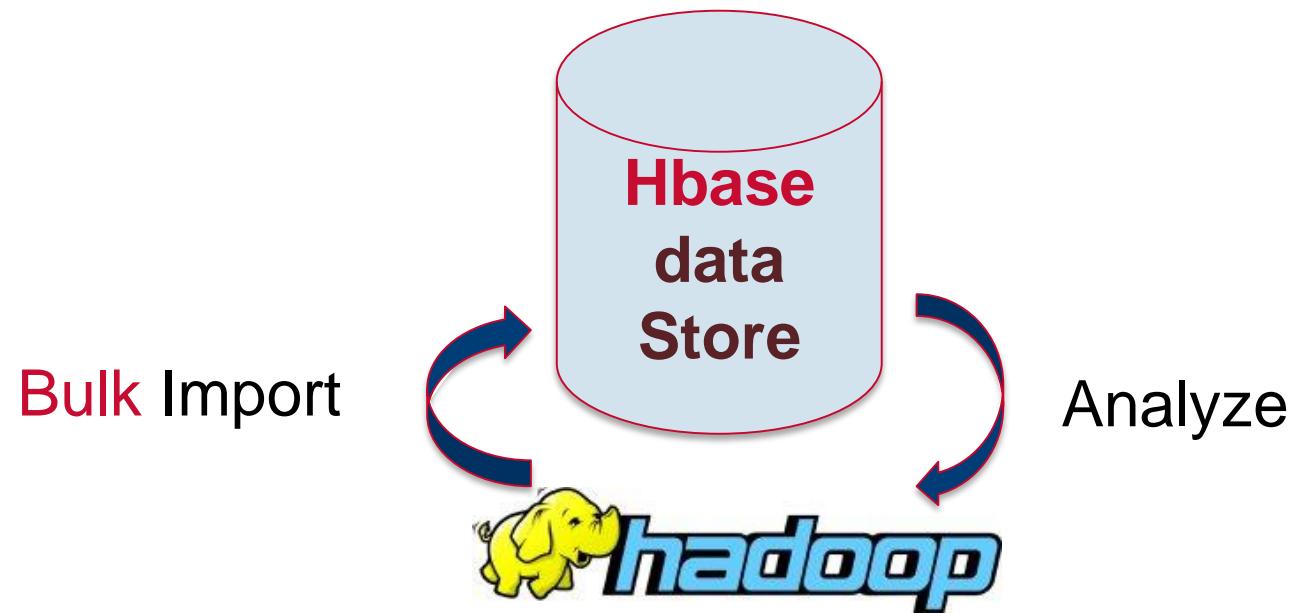
Learning Goals

- ▶ Using MapReduce or Hive with HBase



Use Cases:

- Large scale **offline** ETL analytics, Generating derived data
 - Bulk Import

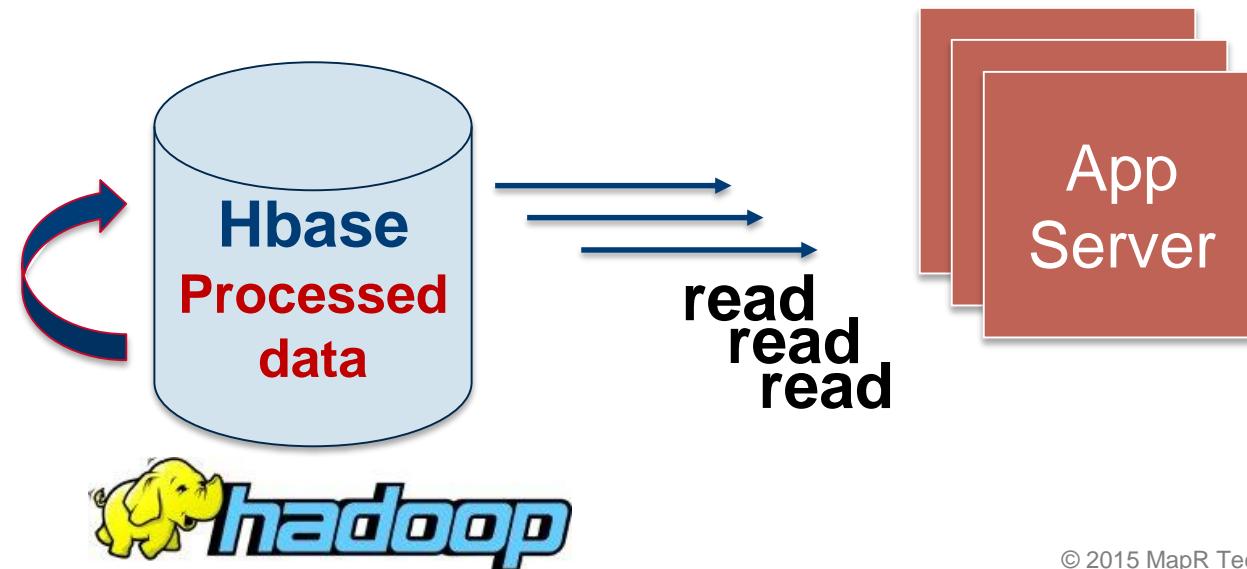




Use Cases:

- **Materialized View, Pre-Calculated Summaries**
 - Map Reduce to **update schema offline**
 - **Online Catalog, Online Dashboard**

Bulk Import

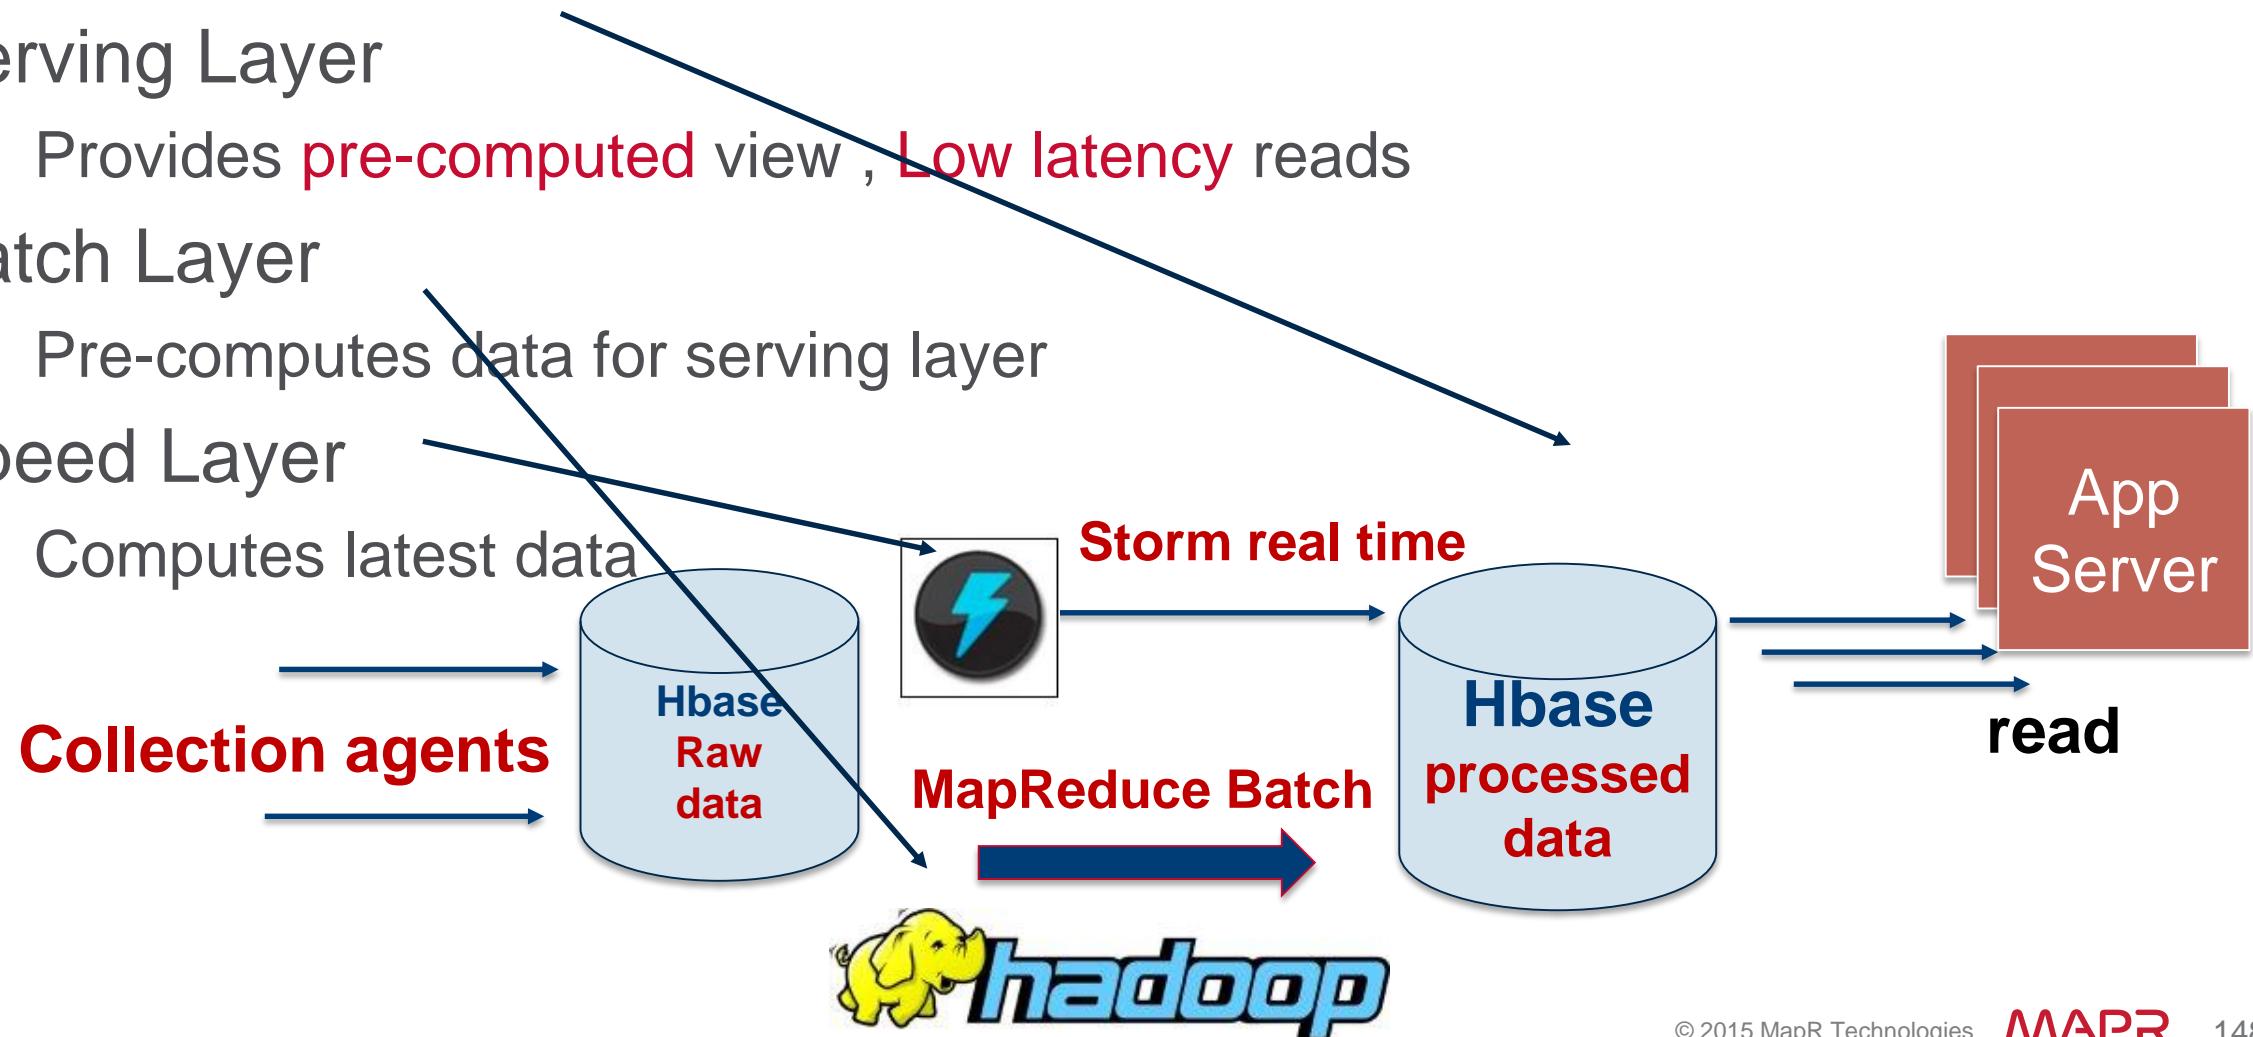




Data access patterns

Lambda architecture

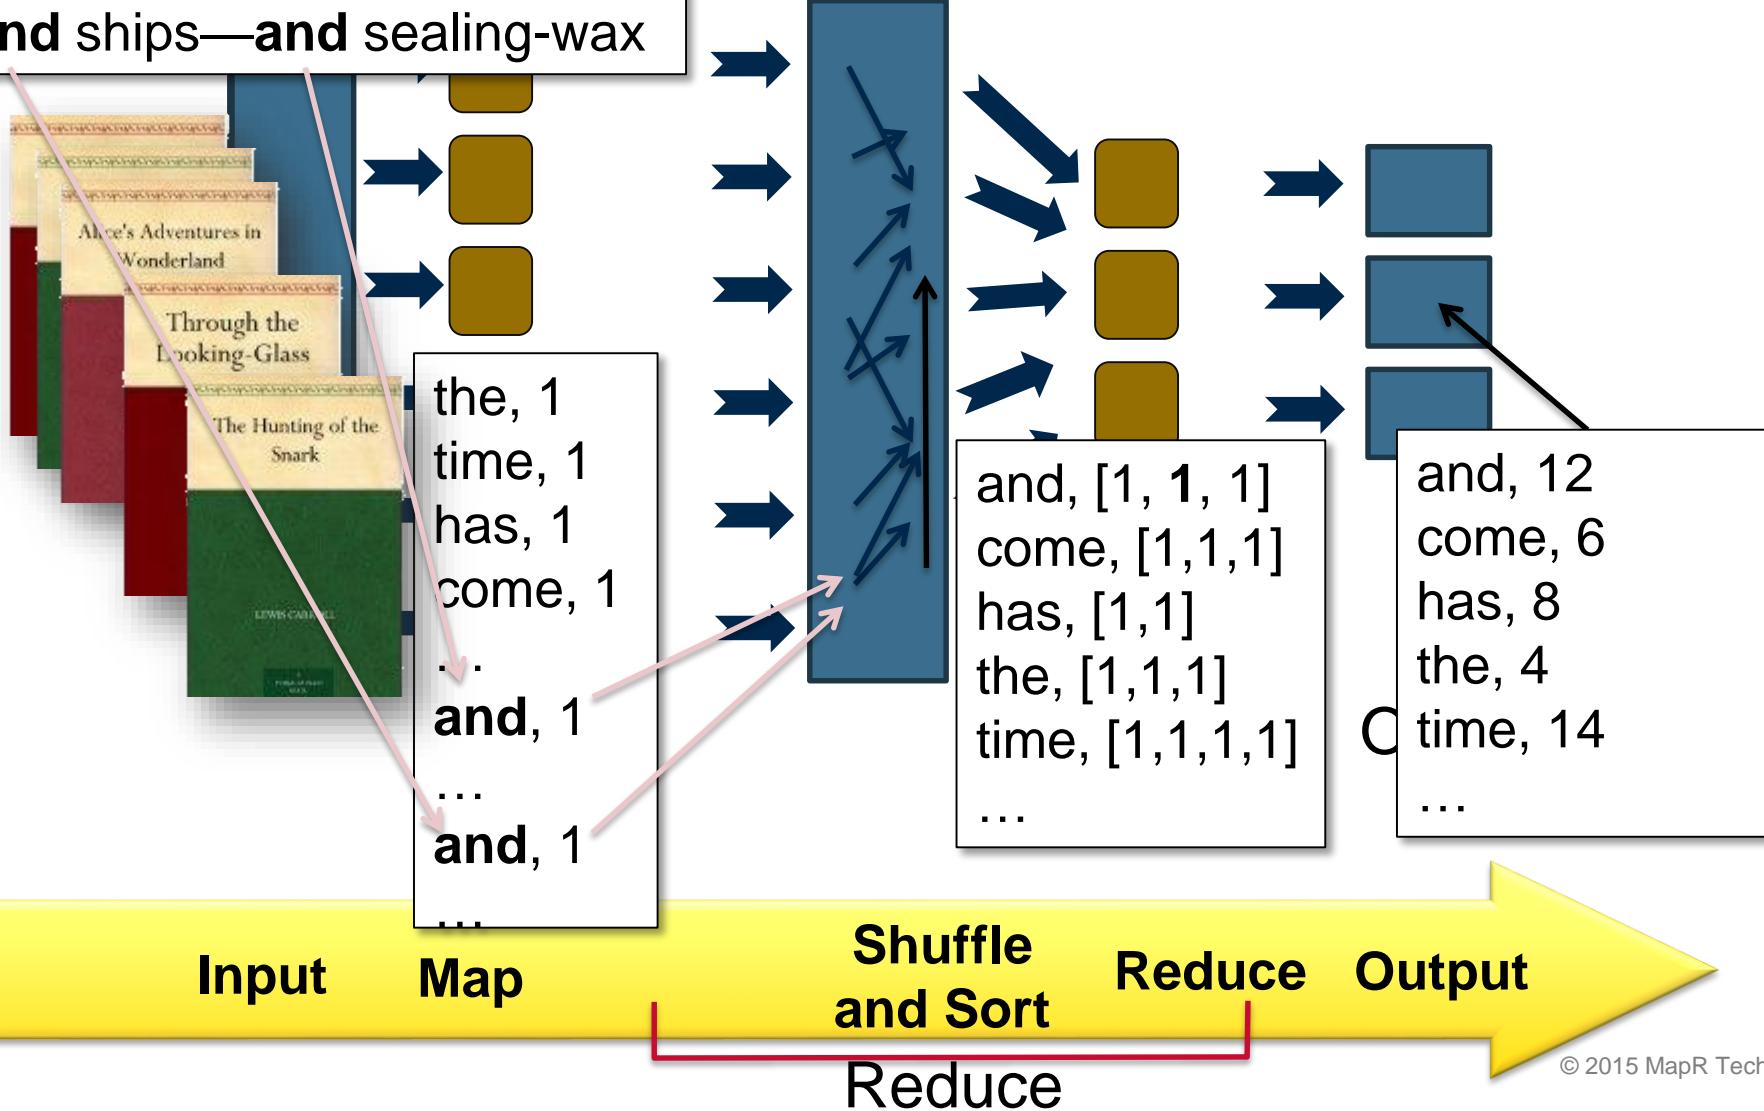
- Serving Layer
 - Provides pre-computed view , **Low latency** reads
- Batch Layer
 - Pre-computes data for serving layer
- Speed Layer
 - Computes latest data





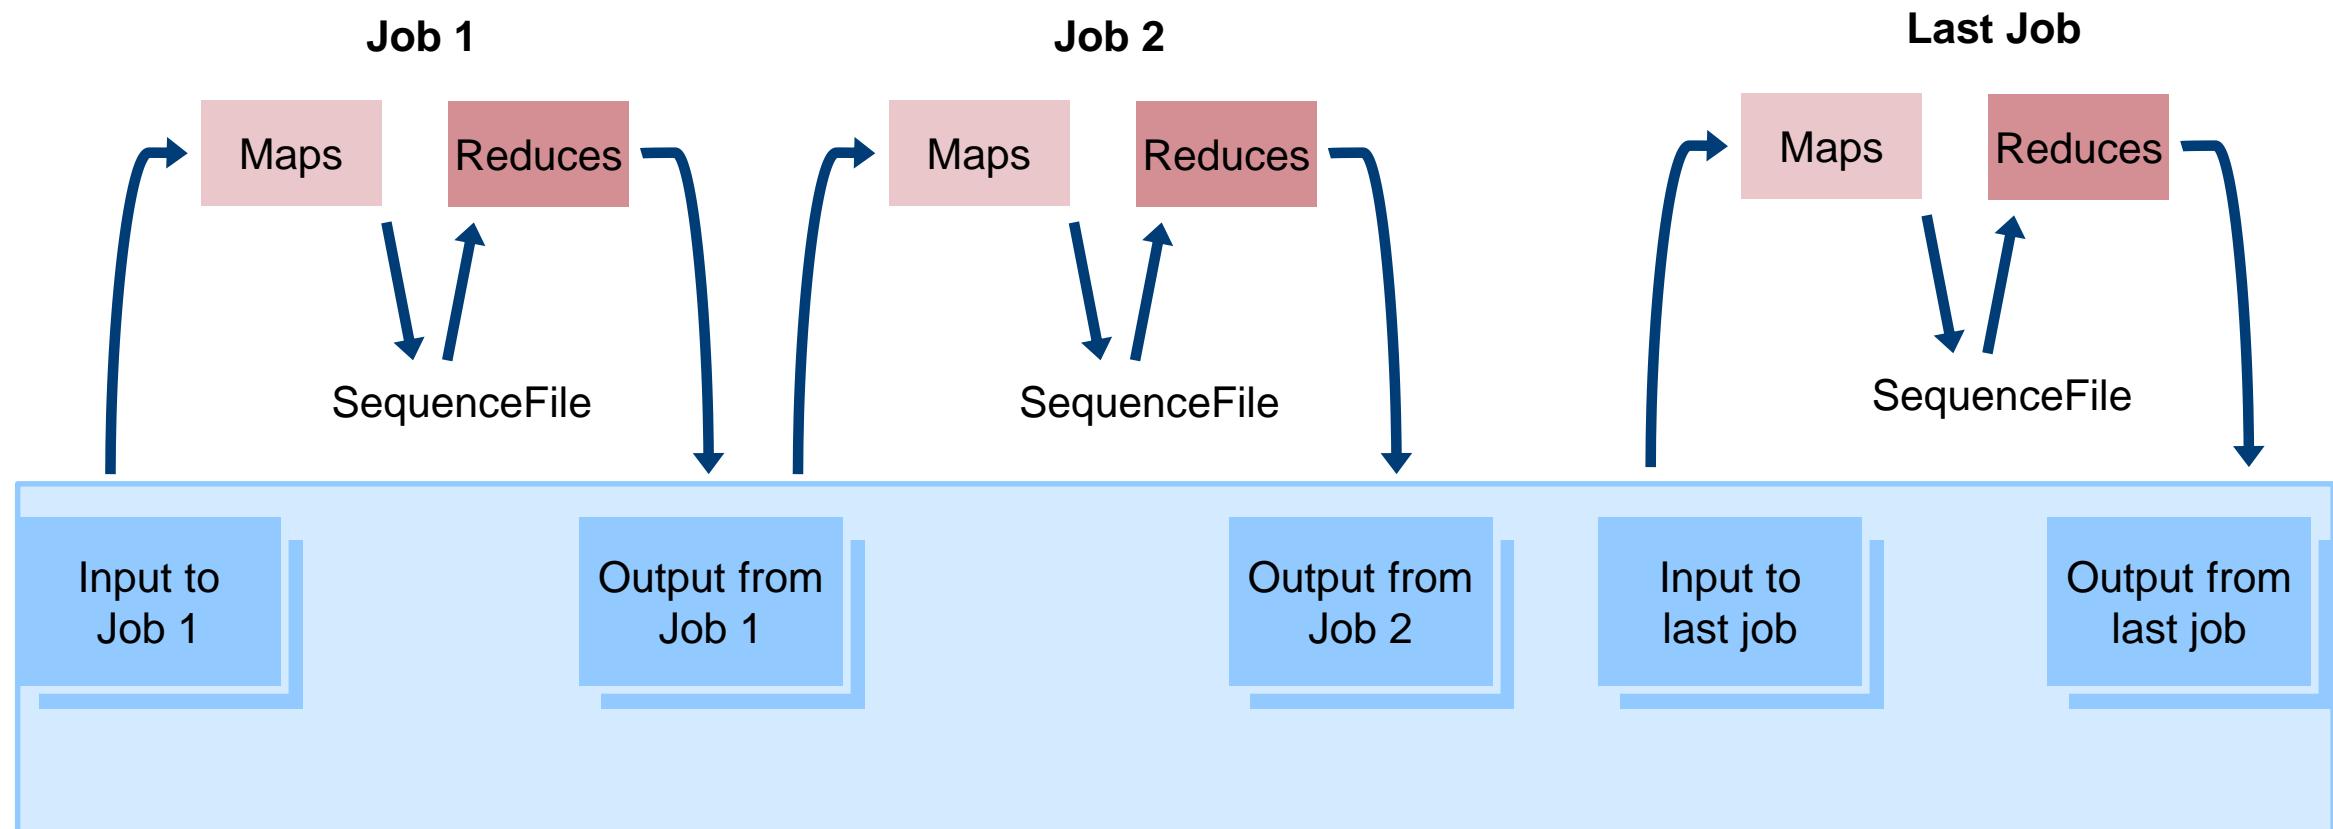
Example: Word Count

"The time has come," the Walrus said,
"To talk of many things:
Of shoes—and ships—and sealing-wax





Typical MapReduce Workflows





MapReduce Design Patterns

- Summarization
 - Inverted index, counting
- Filtering
 - Top ten, distinct
- Aggregation
- Data Organization
 - partitioning
- Join
 - Join data sets
- Metapattern
 - Job chaining





MapReduce Can be Complex

```
public void reduce(Text key, Iterable<Text> values, Context context) throws
IOException, InterruptedException {
for (Text value: values) {
    compositeString = value.toString();
    compositeStringArray = compositeString.split("_");
    tempYear = new Text(compositeStringArray[0]);
    tempValue = new Long(compositeStringArray[1]).longValue();
    if(tempValue < min) {
        min=tempValue;
        minYear=tempYear;
    }
}
Text keyText = new Text("min" + "(" + minYear.toString() + ") : ");
context.write(keyText, new FloatWritable(min));
}
```



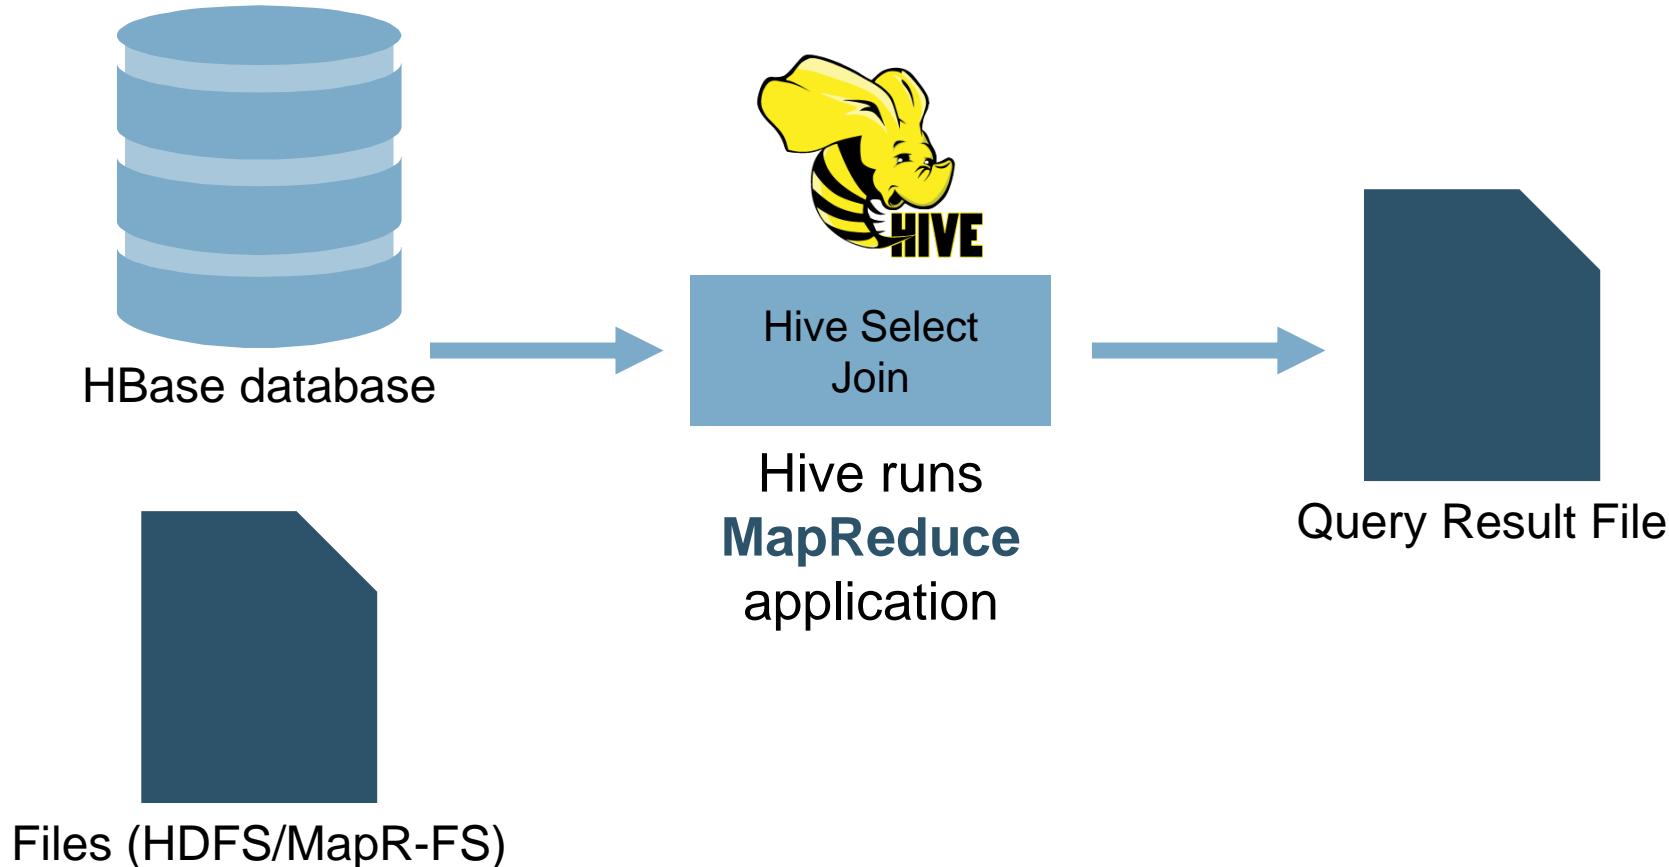
What is Hive?

- Data Warehouse on top of Hadoop
 - Gives ability to query without programming
 - Used for analytical querying of data
- SQL like execution for Hadoop
- SQL evaluates to MapReduce code
 - Submits jobs to your cluster





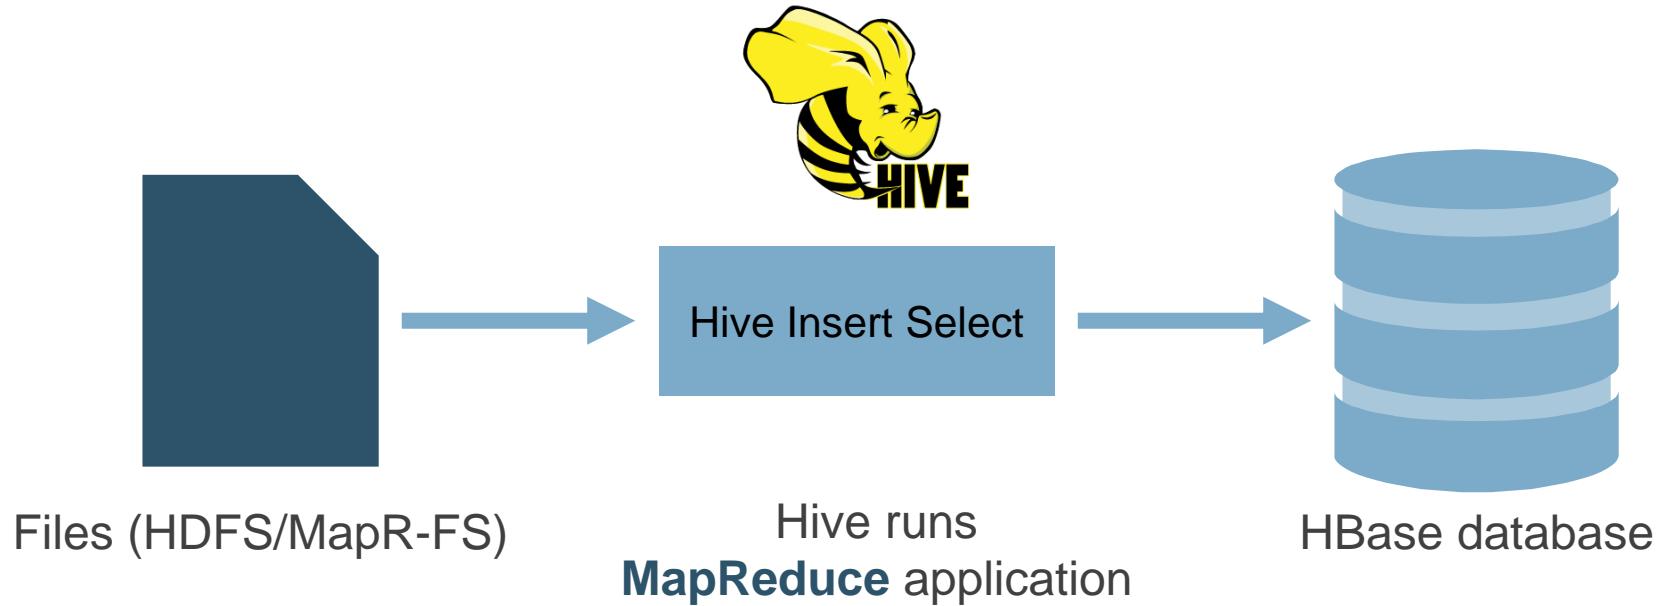
Using HBase as a MapReduce/Hive Source



EXAMPLE: Data Warehouse for Analytical Processing queries



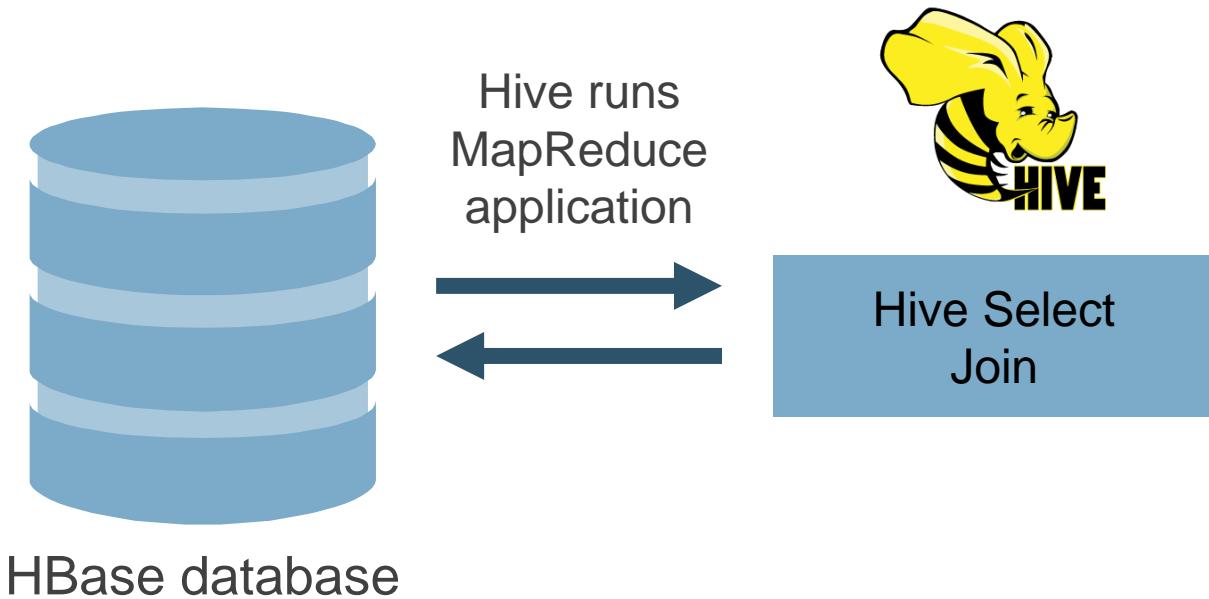
Using HBase as a MapReduce or Hive Sink



EXAMPLE: bulk load data into a table



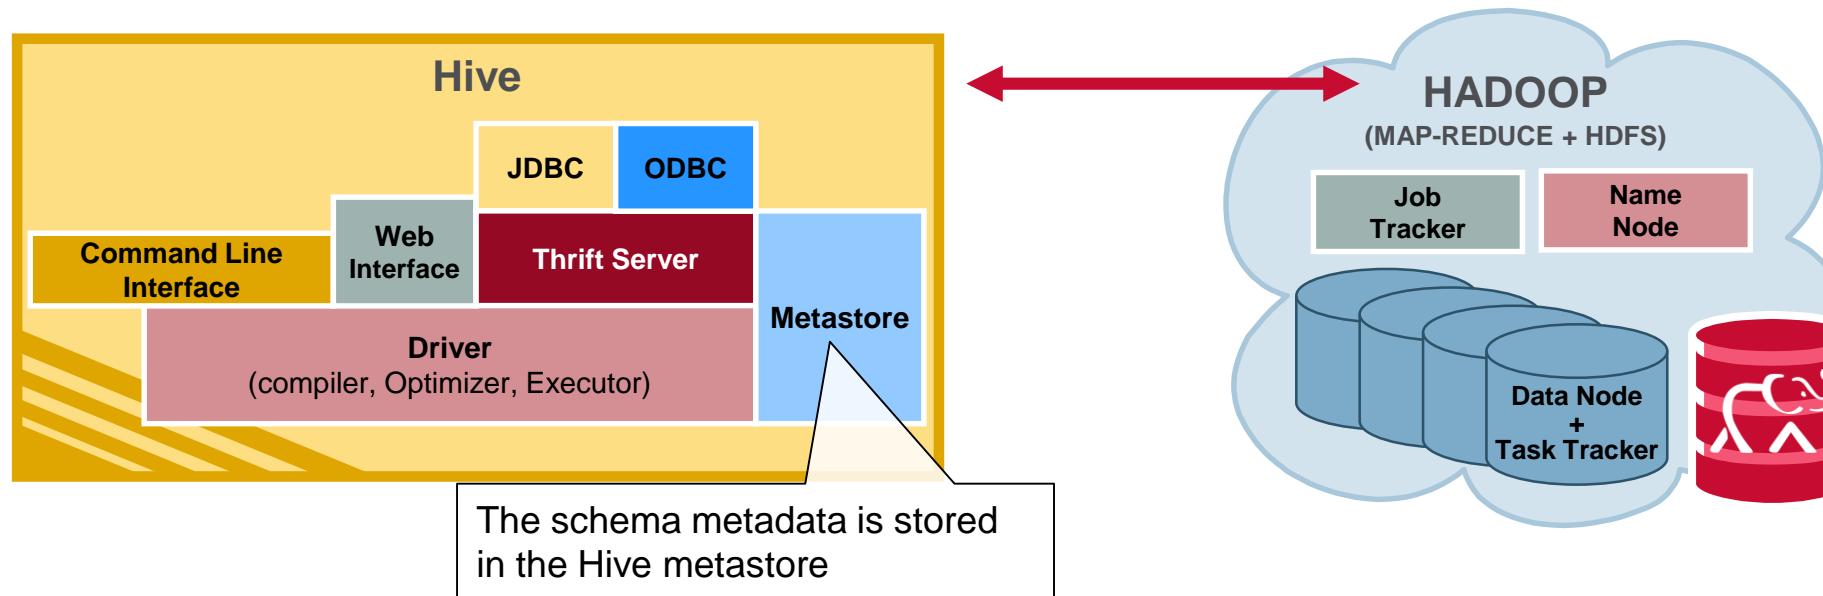
Using HBase as a Source & Sink



*EXAMPLE: calculate and store summaries,
Pre-Computed, Materialized View*



Hive Metastore



Hive Table definition

key string	price bigint	vol bigint

HBase trades_tall Table

key	cf1:price	cf1:vol
AMZN_986186008	12.34	1000
AMZN_986186007	12.00	50



Hive HBase





Hive HBase – External Table

```
CREATE EXTERNAL TABLE trades(key string, price bigint, vol bigint)
STORED BY 'org.apache.hadoop.hive.hbase.HBaseStorageHandler'
WITH SERDEPROPERTIES ("hbase.columns.mapping"= ":key,cf1:price#b,cf1:vol#b")
TBLPROPERTIES ("hbase.table.name" = "/usr/user1/trades_tall");
```

trades		
key string	price bigint	vol bigint

Hive Table definition

Points to
External

key	cf1:price	cf1:vol
AMZN_986186008	12.34	1000
AMZN_986186007	12.00	50

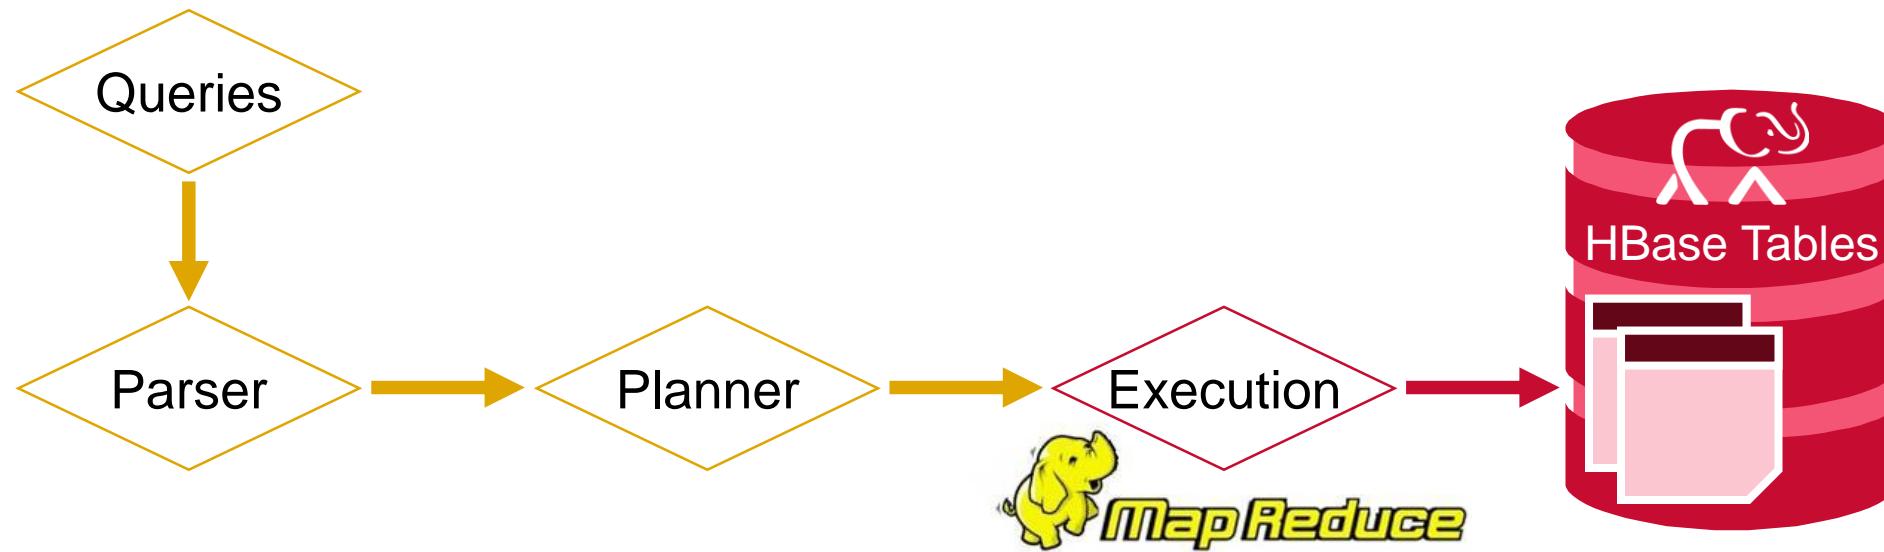
HBaseTable



Hive HBase – Hive Query

SQL evaluates to MapReduce code

```
SELECT AVG(price) FROM trades WHERE key LIKE "GOOG" ;
```

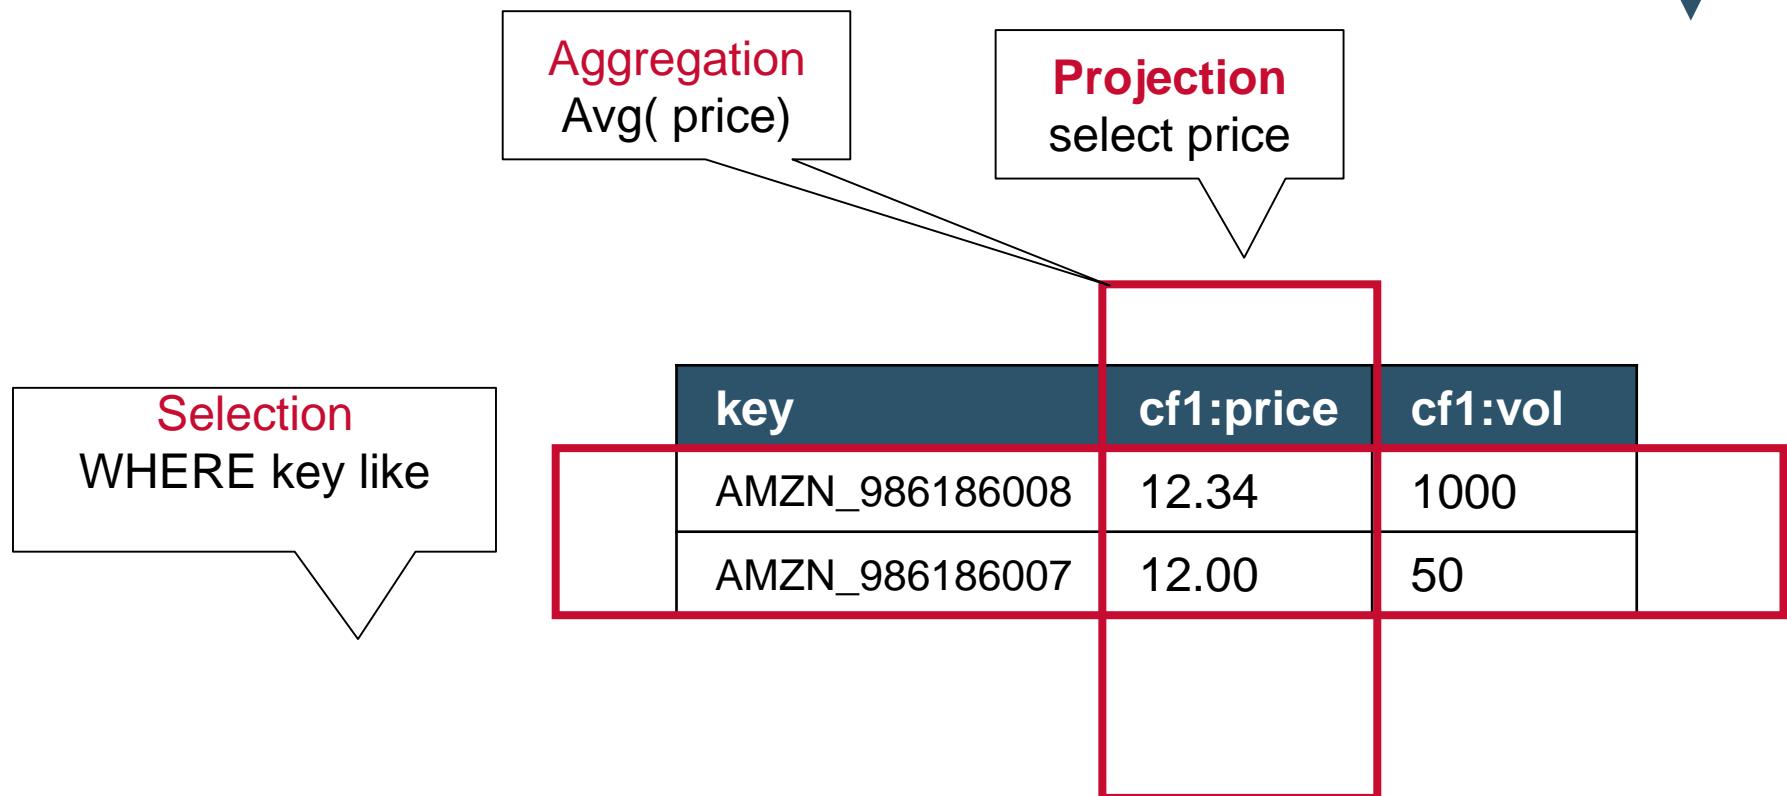




Hive HBase – External Table

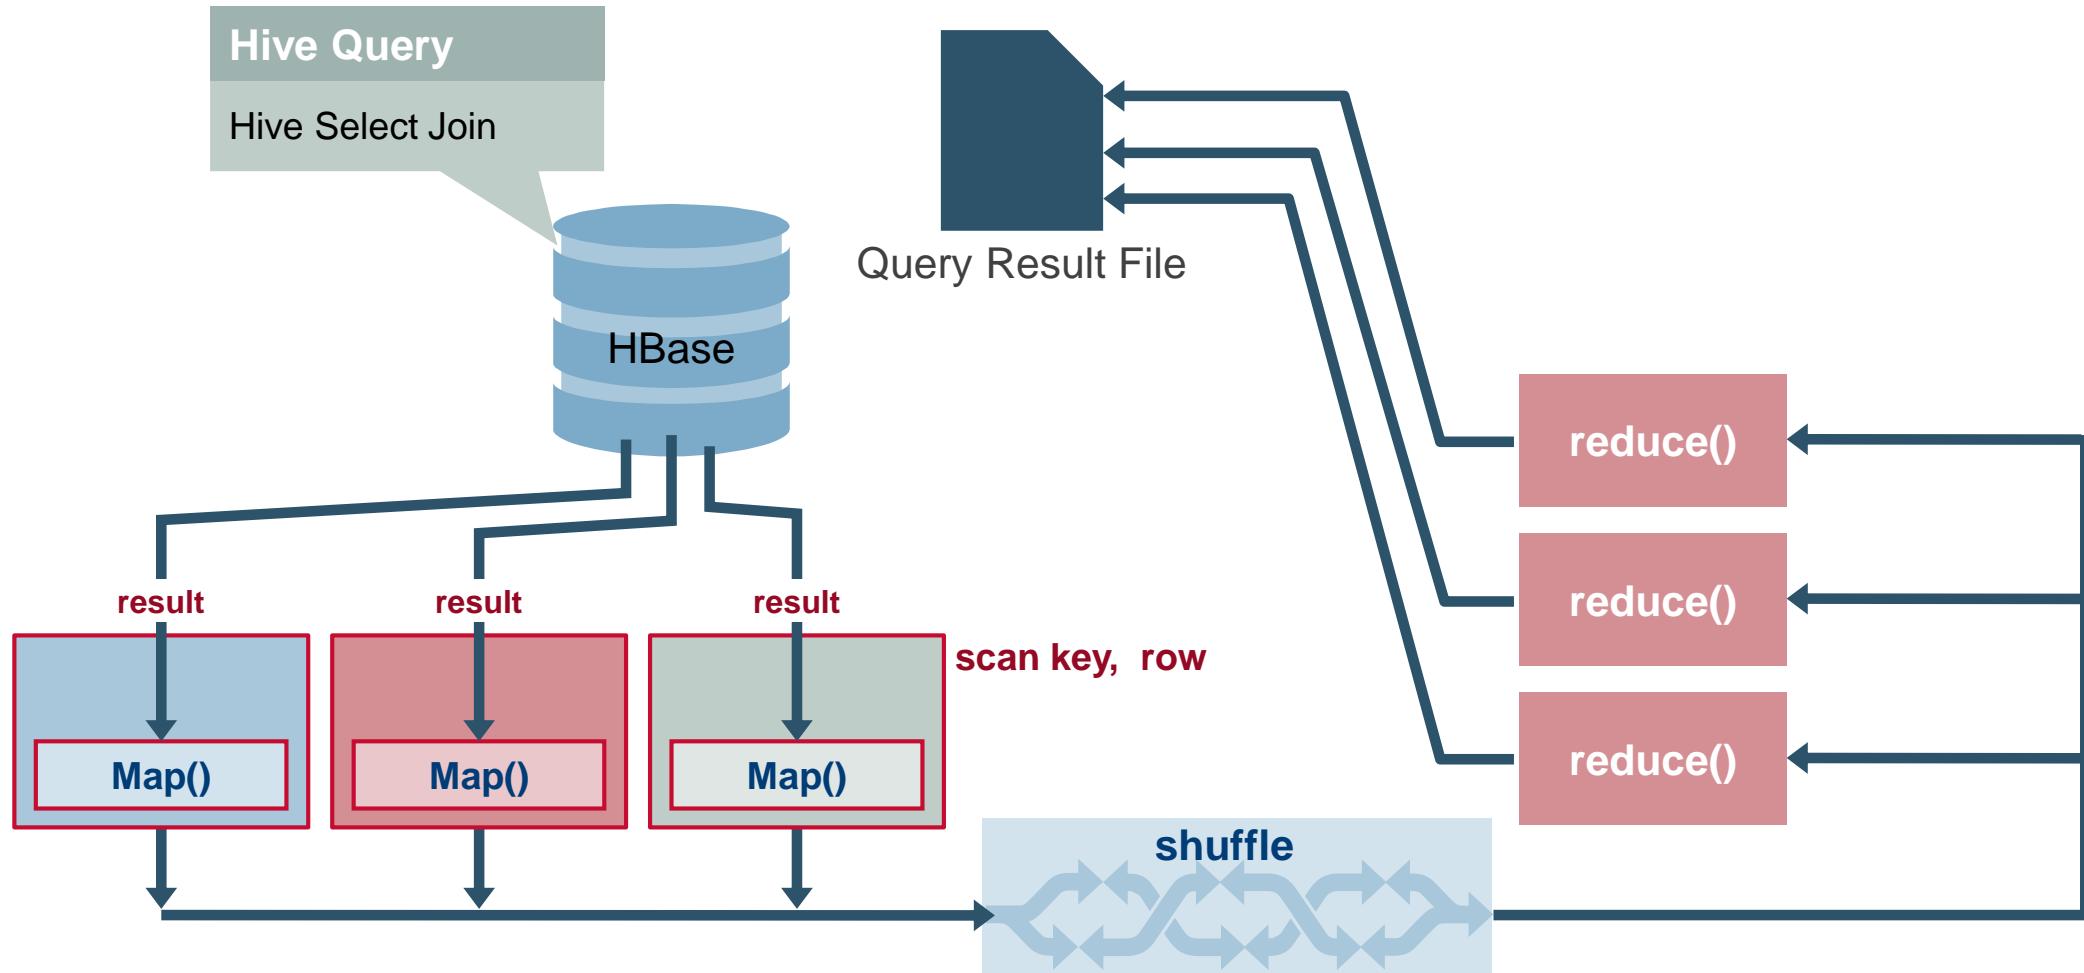
SQL evaluates to MapReduce code

```
SELECT AVG(price) FROM trades WHERE key LIKE "AMZN" ;
```





Hive Map Reduce





Hive Query Plan

```
EXPLAIN SELECT AVG(price) FROM trades WHERE key LIKE "GOOG%";
```

STAGE PLANS:

Stage: Stage-1

Map Reduce

Map Operator Tree:

TableScan

Filter Operator

 predicate: (key like 'GOOG%') (type: boolean)

Select Operator

Group By Operator

Reduce Operator Tree:

Group By Operator

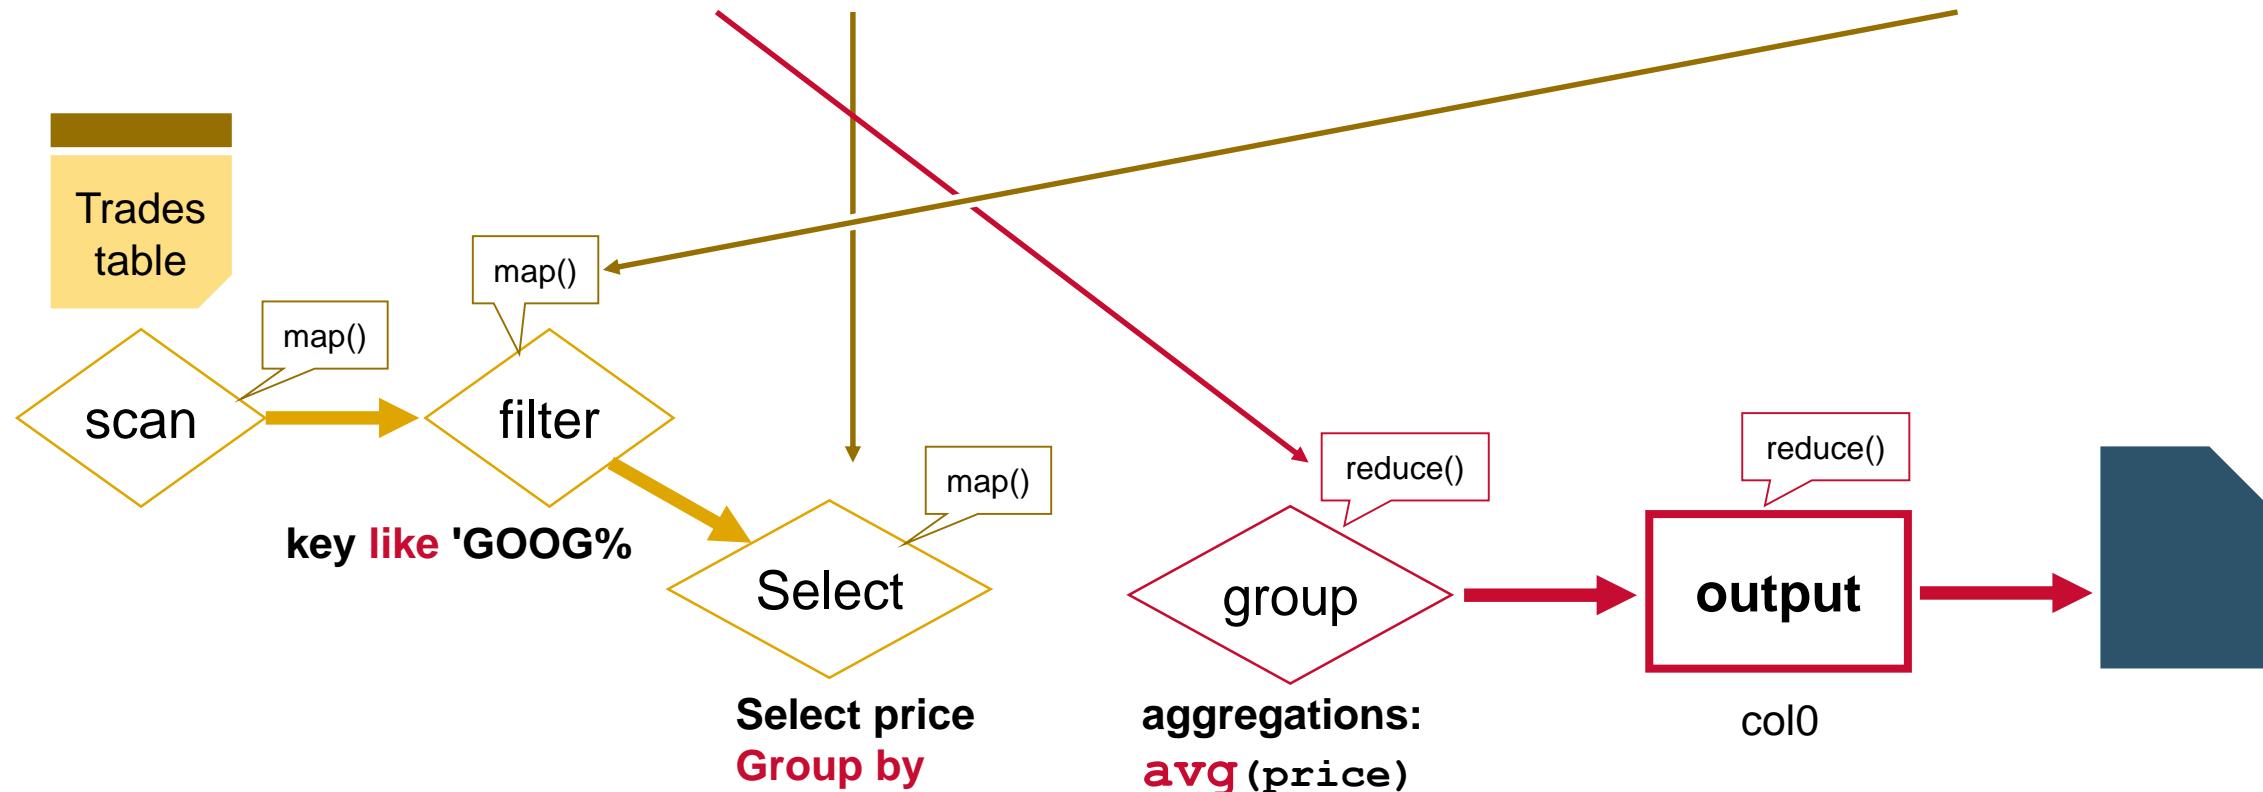
Select Operator

File Output Operator



Hive Query Plan – (2)

```
hive> SELECT AVG(price) FROM trades WHERE key LIKE "GOOG%";
```





Some Hive Design Patterns

- Summarization
 - **Select min(delay), max(delay), count(*) from flights group by carrier;**
- Filtering
 - `SELECT * FROM trades WHERE key LIKE "GOOG%" ;`
 - `SELECT price FROM trades DESC LIMIT 10 ;`
- Join

```
SELECT tableA.field1, tableB.field2 FROM tableA  
JOIN tableB  
ON tableA.field1 = tableB.field2;
```



Lab – Query HBase airline data with Hive

Import mapping to Row Key and Columns:

Row-key	delay			info			stats		timing	
Carrier-Flightnumber-Date-Origin-destination	Air Craft delay	Arr delay	Carrier delay	cncl	Cncl code	tailnum	distance	elaptime	arrtime	Dep time
AA-1-2014-01-01-JFK-LAX		13		0		N7704	2475	385.00	359	...



Count number of cancellations by reason (code)

\$ **hive**

```
hive> explain select count(*) as
cancellations, cnclcode from flighttable
where cncl=1 group by cnclcode order by
cancellations asc limit 100;
```

OK

STAGE DEPENDENCIES:

Stage-1 is a root stage

Stage-2 depends on stages: Stage-1

Stage-0 is a root stage

STAGE PLANS:

Stage: Stage-1

Map Reduce

Map Operator Tree:

TableScan

Filter Operator

Select Operator

Group By Operator

aggregations: count()

Reduce Output Operator

Reduce Operator Tree:

Group By Operator

aggregations: count(VALUE._col0)

Select Operator

File Output Operator

Stage: Stage-2

Map Reduce

Map Operator Tree:

TableScan

Reduce Output Operator

Reduce Operator Tree:

Extract

Statistics: Num rows: 0 Data size: 0 Basic stats: NONE Column stats: NONE

Limit

File Output Operator

Stage: Stage-0

Fetch Operator

limit: 100



2 MapReduce jobs

\$ **hive**

```
hive> select count(*) as cancellations, cnclcode from flighttable where  
cncl=1 group by cnclcode order by cancellations asc limit 100;
```

Total jobs = 2

MapReduce Jobs Launched:

Job 0: Map: 1 Reduce: 1 Cumulative CPU: 13.3 sec MAPRFS Read: 0

MAPRFS Write: 0 SUCCESS

Job 1: Map: 1 Reduce: 1 Cumulative CPU: 1.52 sec MAPRFS Read: 0

MAPRFS Write: 0 SUCCESS

Total MapReduce CPU Time Spent: 14 seconds 820 msec

OK

4598 C

7146 A



Find the longest airline delays

\$ **hive**

```
hive> select arrdelay,key from flighttable where arrdelay > 1000 order by arrdelay desc limit 10;
```

MapReduce Jobs Launched:

Map: 1 Reduce: 1

OK

1530.0 AA-385-2014-01-18-BNA-DFW

1504.0 AA-1202-2014-01-15-ONT-DFW

1473.0 AA-1265-2014-01-05-CMH-LAX

1448.0 AA-1243-2014-01-21-IAD-DFW

1390.0 AA-1198-2014-01-11-PSP-DFW

1335.0 AA-1680-2014-01-21-SLC-DFW

1296.0 AA-1277-2014-01-21-BWI-DFW

1294.0 MQ-2894-2014-01-02-CVG-DFW

1201.0 MQ-3756-2014-01-01-CLT-MIA

1184.0 DL-2478-2014-01-10-BOS-ATL

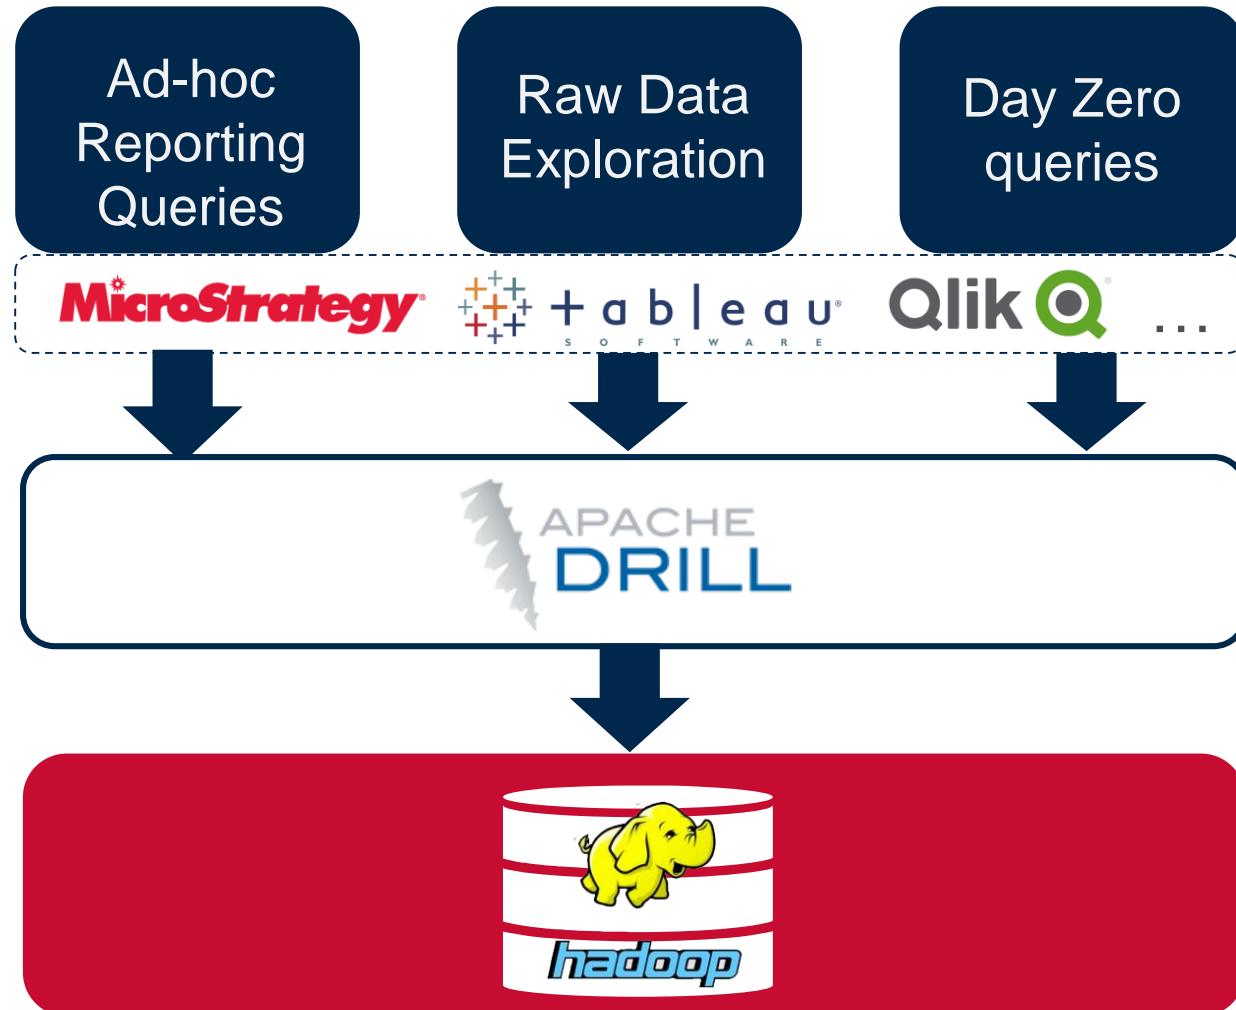


- Hive



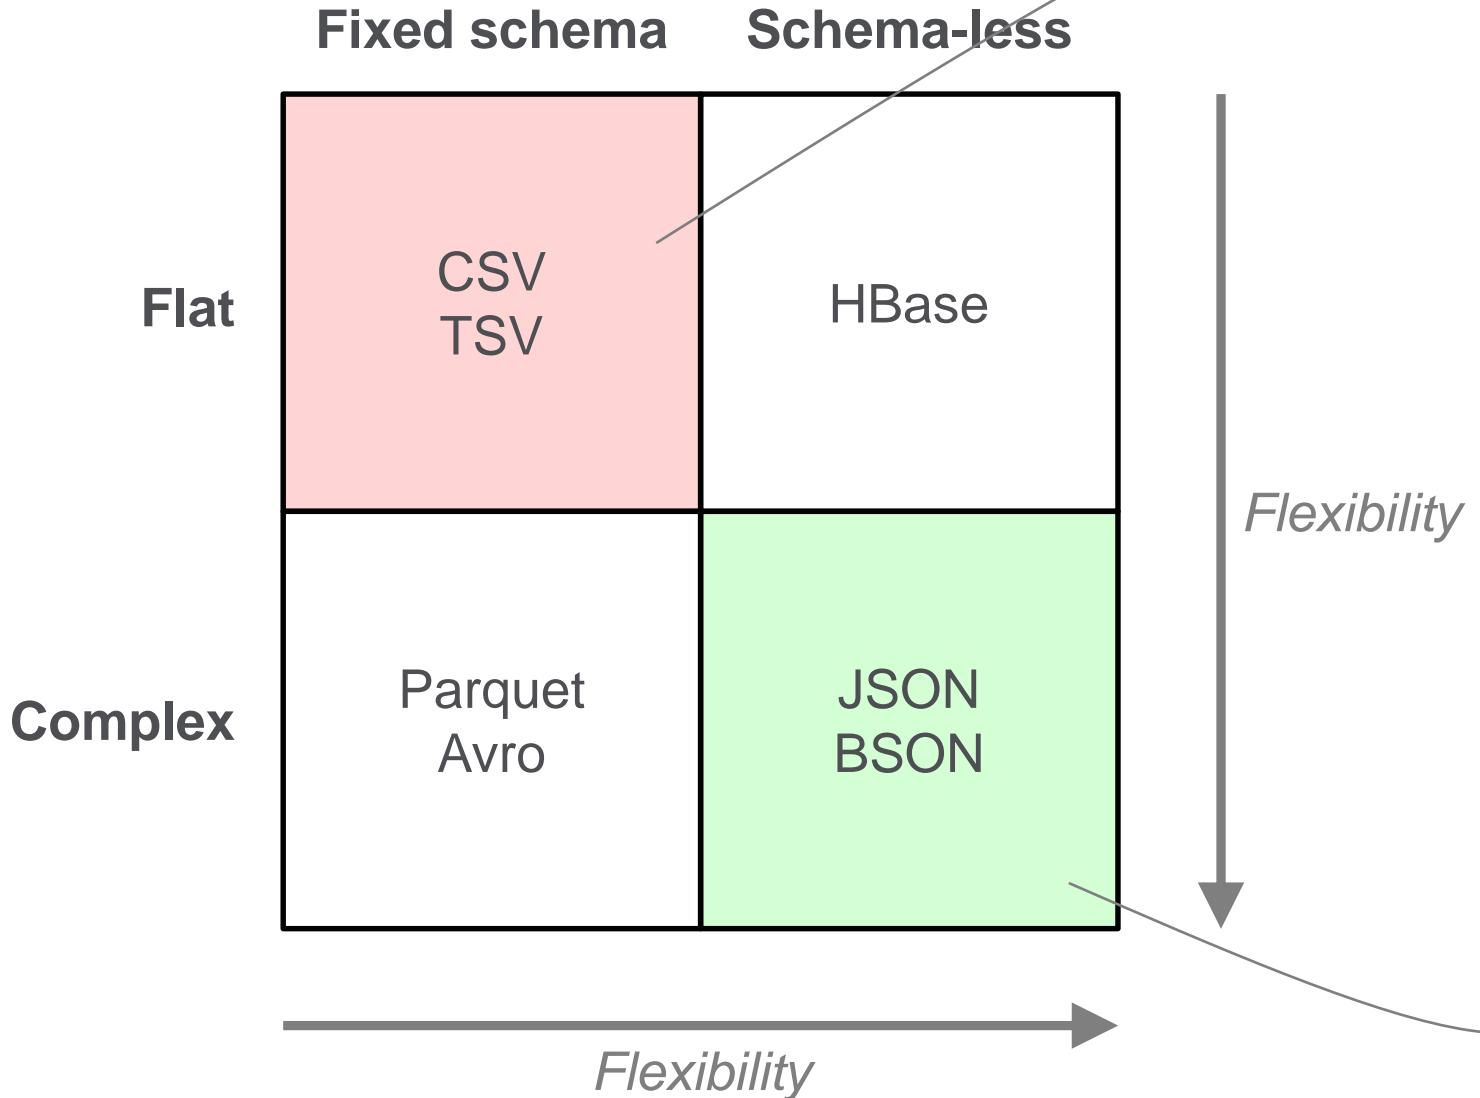
Self-Service Data Exploration

Direct access to Hadoop data from familiar BI / Analytics tools- ANSI SQL compatible





Drill's Data Model is Flexible



RDBMS/SQL-on-Hadoop table

Name	Gender	Age
Michael	M	6
Jennifer	F	3

Apache Drill table

{ name: { first: Michael, last: Smith }, hobbies: [ski, soccer], district: Los Altos } { name: { first: Jennifer, last: Gates }, hobbies: [sing], preschool: CCLC }
--

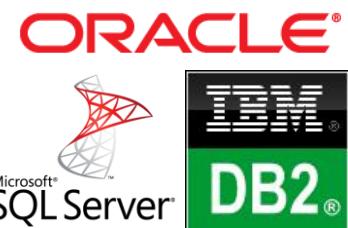


Drill Supports *Schema Discovery On-The-Fly*

Schema Declared In Advance

- Fixed schema
- Leverage schema in centralized repository (Hive Metastore)

SCHEMA ON
WRITE



SCHEMA
BEFORE READ



Schema Discovered On-The-Fly

- Fixed schema, evolving schema or schema-less
- Leverage schema in centralized repository or self-describing data

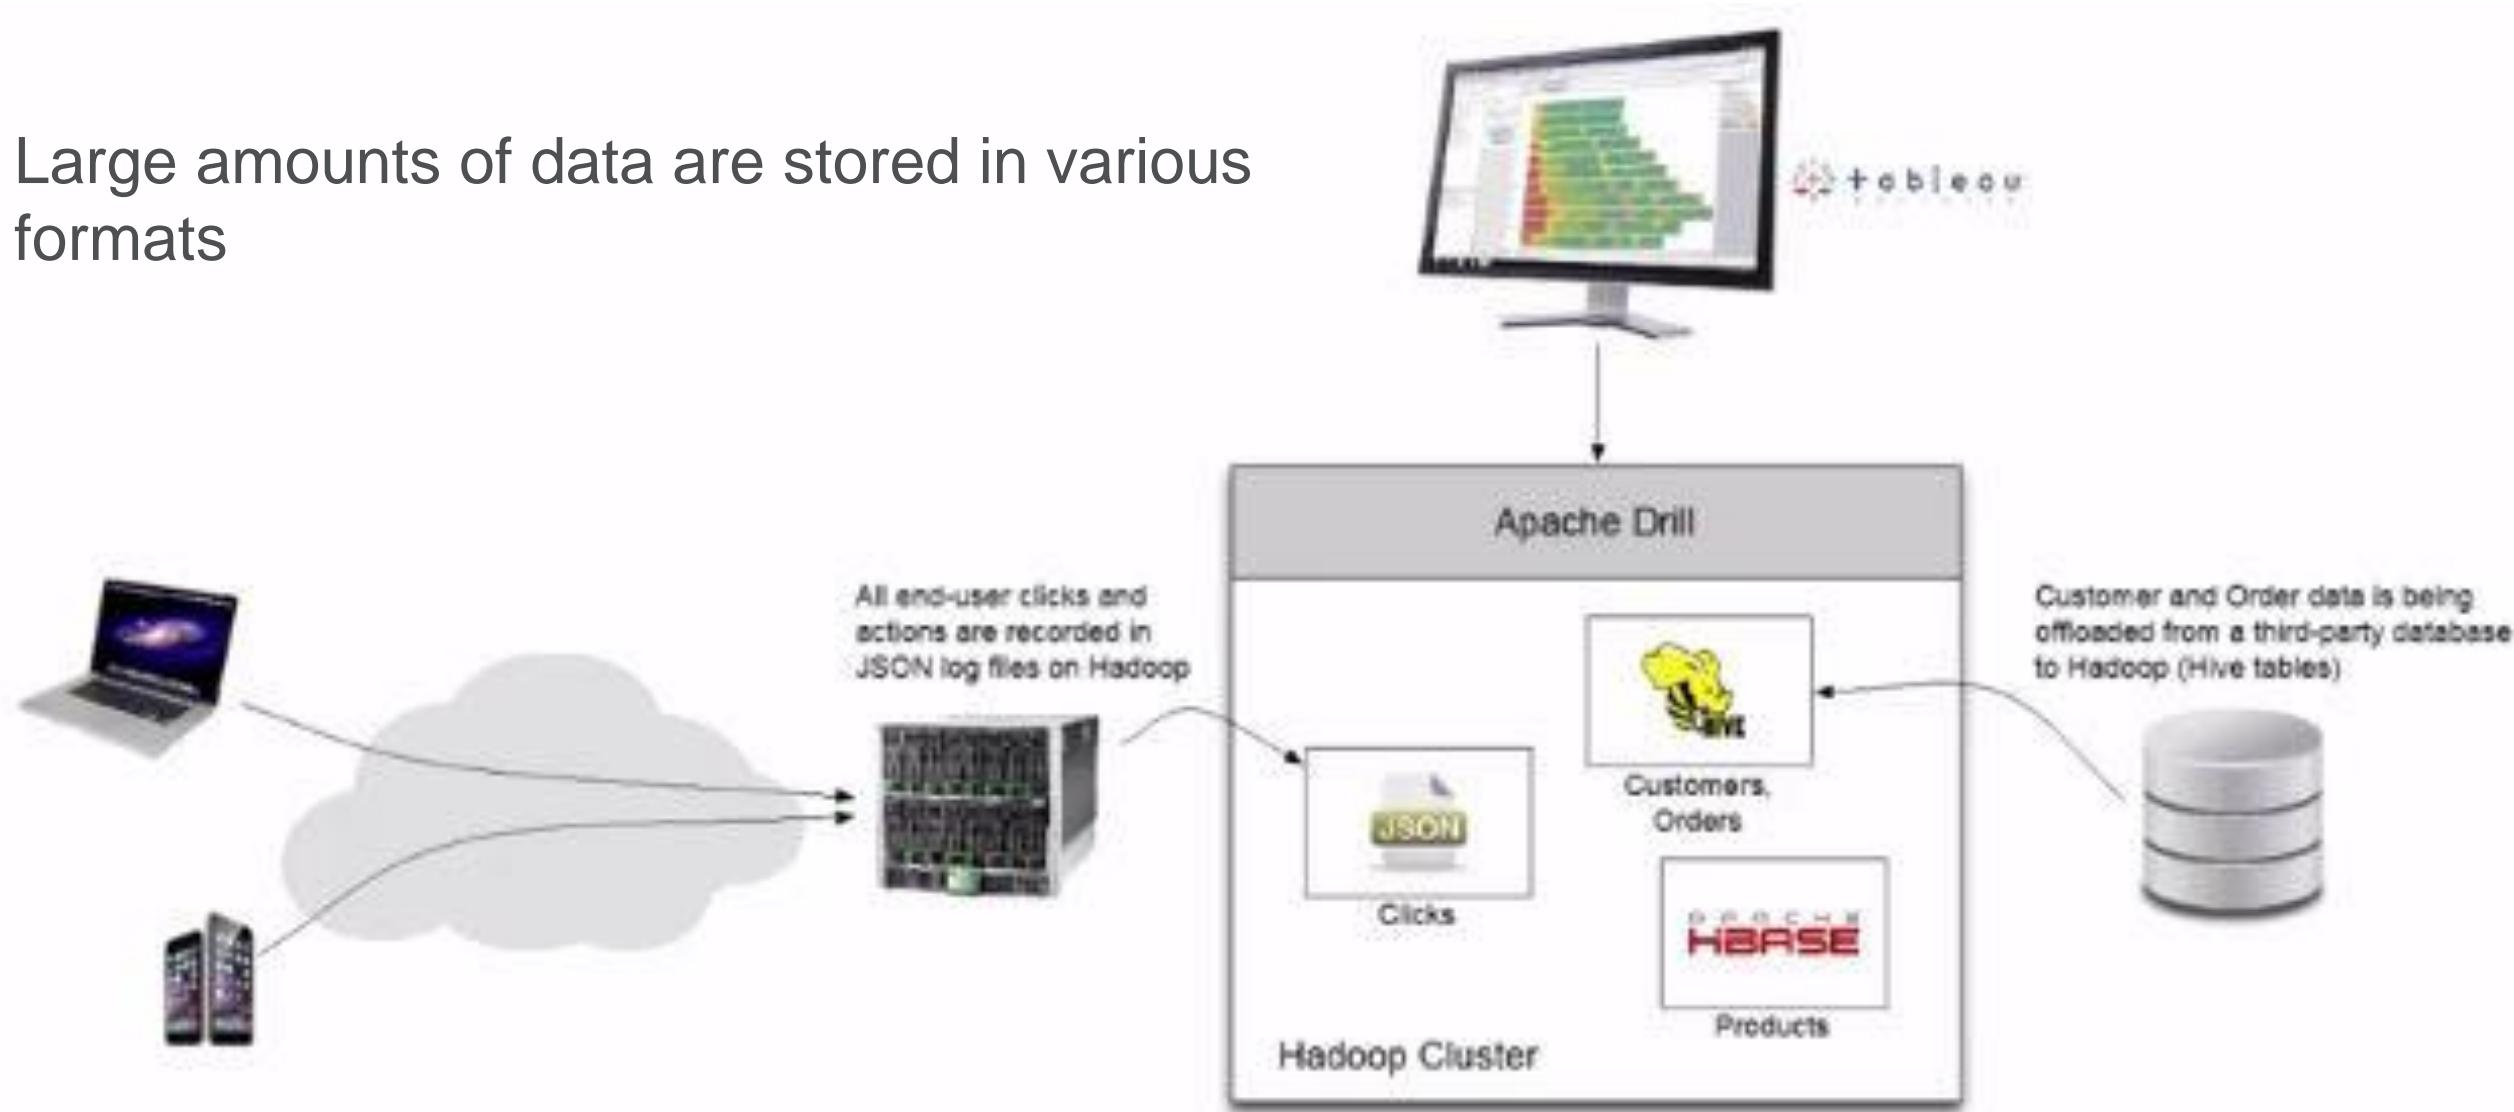
SCHEMA ON THE
FLY





Use Case: Online Retail Business

Large amounts of data are stored in various formats



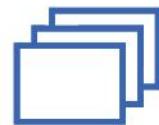


Use Case: Online Retail Business

Data Analyst



Analyst explores various data sources using Apache Drill



Web, mobile log Files of User clicks on online retail web Site, Mobile app

Orders

Customer Profile
Products



- Orders
 - Hive



Order schema

order_id	month	cust_id	state	Prod_id	Order_total
67212	June	10001	ca	909	13



HBase Data

- Products Data
 - HBase
- Customer Data
 - Hbase

rowkey	CF: details		CF: pricing
	category	name	price
10	laptop	sony	1000

rowkey	CF: address		CF: loyalty		CF:personal	
	state	agg_rev	membership	age	gender	
1	"va"	197.00	"Silver"	"15-20"	"FEMALE"	

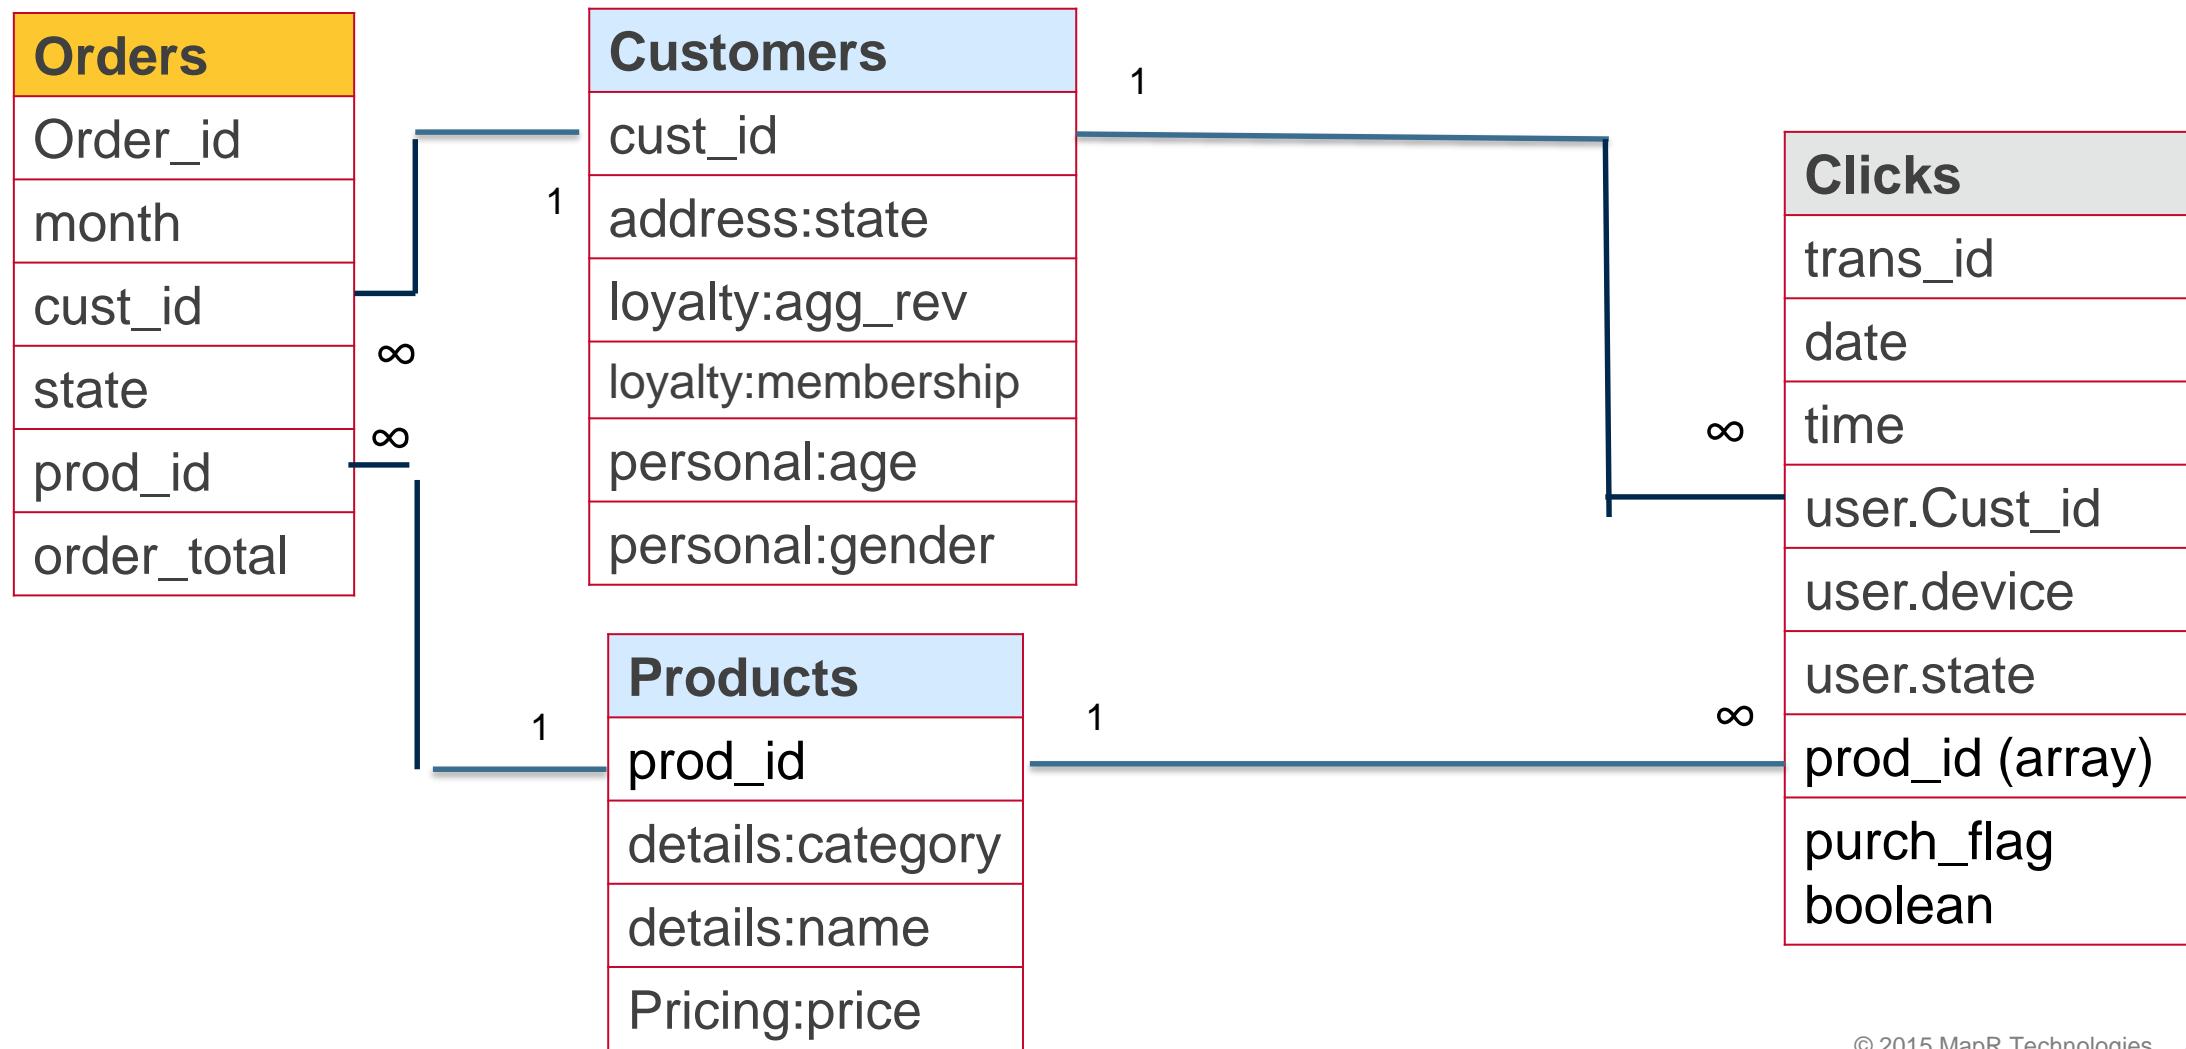


- User Web clicks logs
 - JSON files

trans_id	date	user_info.cust_id	user_info.device	user_info.state	trans_info.prod_id	trans_info.purch_flag
67212	June	10001	“iOS5”	“ca”	[174,2]	false



High Level “Relations” between Retail data





Use Case

- ▶ Explore the Use Case Data

Hive

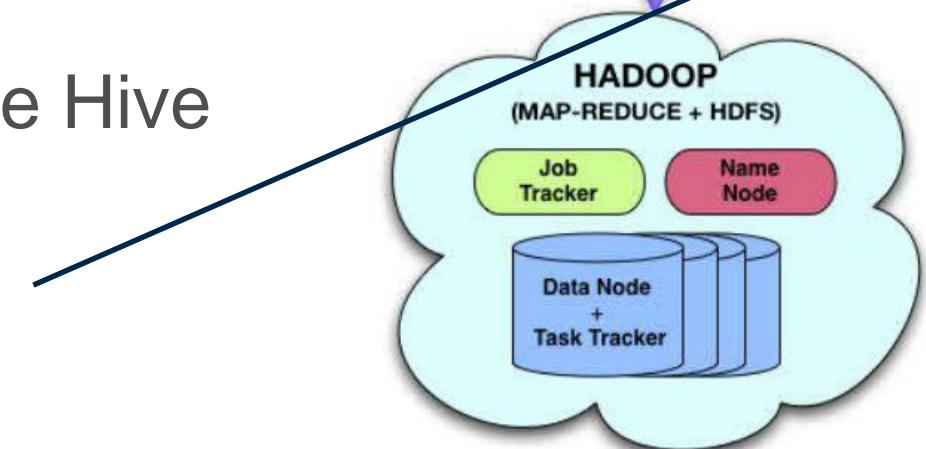
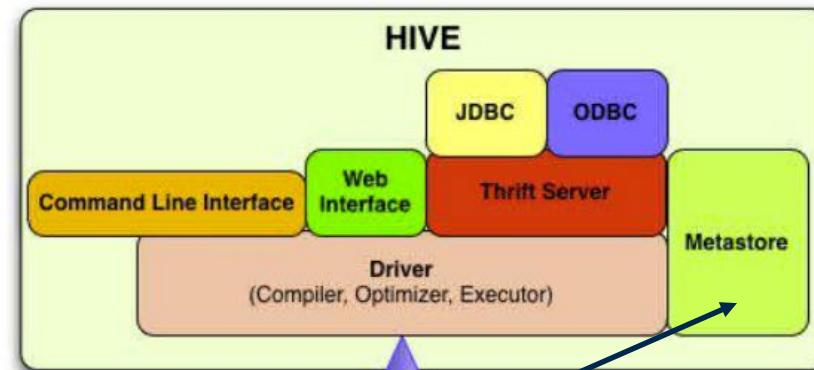
HBase

Files



Drill can read Hive Metastore

- With Hive you have to define a schema for files or Hbase tables
- The schema metadata is stored in the Hive metastore
- Drill can read the Hive metastore



Order schema

order_id	month	cust_id	state	Prod_id	Order_total

Metadata for file

/data/orders/month2.log.csv

15552	February	10005	CA	909	10



Exploring the Hive Order Tables

Drill Explorer

New Snippet Ctrl+N

Browse SQL

Schemas:

- + cp.default
- + dfs.clicks
- + dfs.data
- + dfs.default
- + dfs.logs
 - logs
 - 2012
 - 2013
 - 2014
- + dfs.root
- + dfs.tmp
- + dfs.views
- + hive.default
 - orders
- + maprdb
 - customers
 - address
 - loyalty
 - personal
 - embeddedclicks
 - products

Metadata: `hive.default`orders``

	COLUMN_NAME	DATA_TYPE	IS_NULLABLE
1	order_id	BIGINT	YES
2	month	VARCHAR	YES
3	cust_id	BIGINT	YES
4	state	VARCHAR	YES
5	prod_id	BIGINT	YES

Data Preview:

	order_id	month	cust_id	state	prod_id	order_total
1	67212	June	10001	ca	909	13
2	70302	June	10004	ga	420	11
3	69090	June	10011	fl	44	76
4	68834	June	10012	ar	0	81
5	71220	June	10018	az	411	24
6	61287	June	1001	nj	104	134
7	68553	June	10021	ca	117	67
8	68109	June	10022	tx	337	10
9	68526	June	10025	mi	11	63
10	69362	June	10028	tx	430	65
11	68624	June	10030	fl	808	51

Refresh

Connect...

Sample

100

Rows

Close



HBase is “schema less”

- Drill allows direct queries on HBase Tables, without mapping to a schema (like with Hive or Impala)

RowKey	CF1			CF2				...
	colA	colB	colC	colA	colB	colC	colD	
aXXX	val			val	val		val	
...								



Exploring the HBase Tables

Browse SQL

Schemas:

- + cp.default
- + dfs.clicks
- + dfs.data
- + dfs.default
- + dfs.logs
- + dfs.root
- + dfs.tmp
- + dfs.views
- + hive.default
- maprdb
 - customers
 - address
 - loyalty
 - personal
 - embeddedclicks
 - products
 - details
 - pricing

Metadata: `maprdb`.`products`

	Column_Name	Data_Type	Value
▶ 1	category	Varchar(40)	40
2	name	Varchar(128)	128

Data Preview:

	Row_Key	category	name
▶ 1	0	laptop	"Sony notebook"
2	1	Envelopes	#10-4 1/8 x 9 1/2 Pre
3	10	Storage & Organization	24 Capacity Maxi Da
4	100	Labels	Avery 498
5	101	Labels	Avery 49
6	102	Labels	Avery 501

 Query the Clicks File

```
> use dfs.clicks;  
> SELECT * FROM `clicks/clicks.json` limit 2;
```

Browse SQL

View Definition SQL:

```
select * from `clicks/clicks.json` limit 2;
```

Total Number of Records: 2

	trans_id	date	time	user_info	trans_info
▶ 1	31920	2014-04-26	12:17:12	{ "cust_id": 22526, "device": "IOS5", "state": "il"}	{ "prod_id": [174, 2], "purch_flag": "false" }
2	31026	2014-04-20	13:50:29	{ "cust_id": 16368, "device": "AOS4.2", "state": "nc"}	{ "prod_id": [], "purch_flag": "false" }



Explore the logs directory

Drill Explorer

Browse SQL

Schemas:

- + cp.default
- + dfs.clicks
- + dfs.data
- + dfs.default
- dfs.logs
 - logs
 - 2012
 - 1
 - {...} log.json
 - 10
 - 11
 - 12
 - 2
 - 3
 - 4
 - 5
 - 6
 - 7
 - 8
 - 9
 - 2013
 - 2014
 - dfs.root

Metadata:

	COLUMN_NAME	DATA_TYPE	IS_NULLABLE
▶ 1	Metadata does not exist		

Data Preview:

	trans_id	date	time	cust_id	device	state	camp_id	keyword
▶ 1	100	01/30/2012	17:14:14	4	AOS4.2	ms	16	i
2	120	01/29/2012	06:52:05	0	IOS5	va	4	crime
3	138	01/30/2012	10:03:10	0	IOS5	ky	5	laughing
4	154	01/07/2012	20:24:28	2	AOS4.2	ky	10	that
5	160	01/25/2012	11:43:53	9107	AOS4.2	oh	9	me
6	169	01/17/2012	12:14:27	1119	IOS5	ny	17	the
7	180	01/13/2012	17:57:47	5145	IOS5	nc	7	you
8	182	01/09/2012	17:05:30	1	AOS4.4	ny	13	sir
9	184	01/04/2012	14:59:08	7927	AOS4.2	ar	3	one
10	186	01/24/2012	14:14:48	24254	IOS5	ca	10	make



Query the logs directory

use dfs.logs;

SELECT *

FROM logs limit 100;

Browse SQL

View Definition SQL:

```
SELECT * FROM dfs.logs.`logs` limit 400
```

Total Number of Records: 400

	dir0	dir1	trans_id	date	time	cust_id	device	state	camp_id	keywords
219	2014	8	36017	08/02/2014	08:44:33	6	AOS4.2	ak	13	people
220	2014	8	36044	08/02/2014	18:48:36	29	IOS5	ms	7	oh
221	2014	8	36061	08/04/2014	18:53:31	1	IOS5	la	2	it's
222	2013	2	12115	02/23/2013	19:48:24	3	IOS5	az	5	who's
223	2013	2	12127	02/26/2013	19:42:03	11459	IOS5	wa	10	for
224	2013	2	12138	02/09/2013	05:49:01	1	IOS6	ca	7	minutes
225	2013	2	12139	02/23/2013	06:58:20	1	AOS4.4	ms	7	i



Data Sources and Storage Plugins Configuration





Data Source is in the Query

```
SELECT *  
FROM dfs.clicks.`clicks/clicks.json`
```

where you want it from

<storage plugin>.<workspace>.<tablename>

↑
OPTIONAL

A *storage plugin name*

- DFS (default)
- HBase
- Hive

A *workspace*

- Sub-directory
- Hive database

A *table*

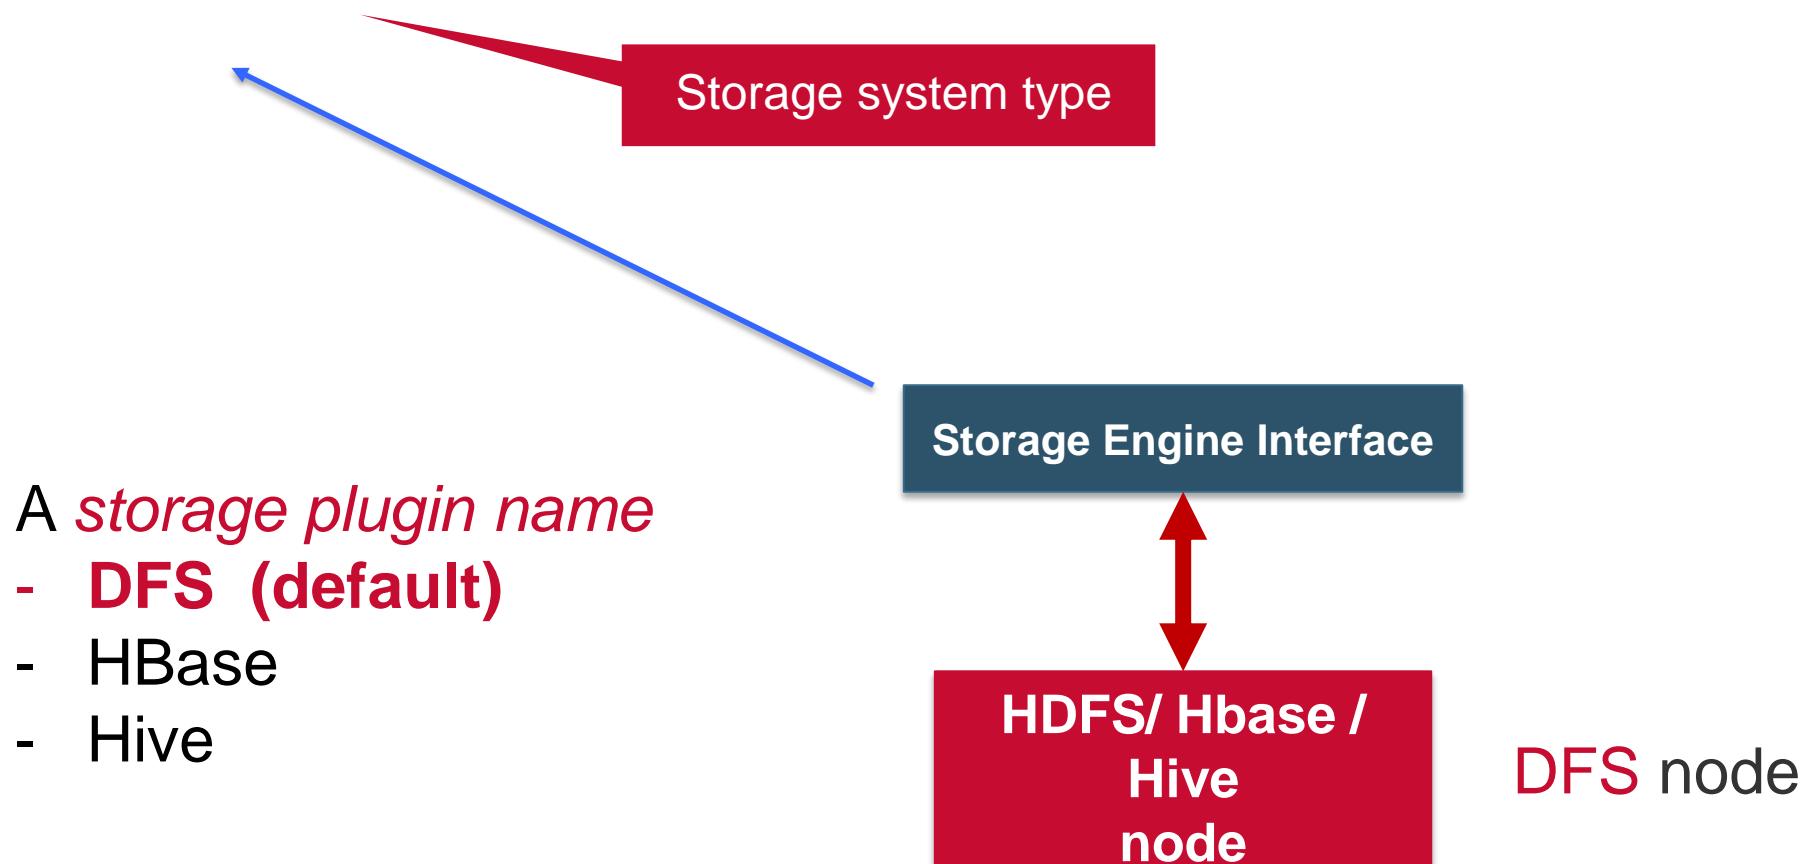
- pathname
- HBase table
- Hive table



Data Source is in the Query

```
SELECT *  
FROM dfs.clicks.`clicks/clicks.json`
```

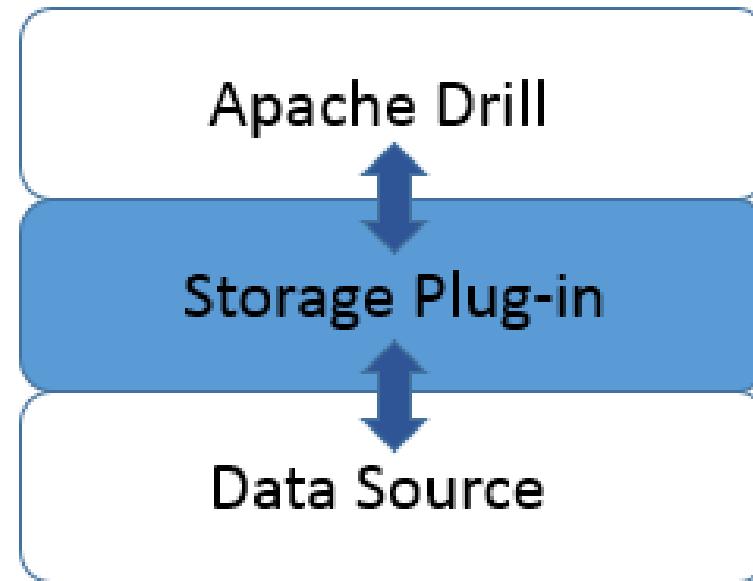
where you want it from





Storage plugins

- Storage plugins provide:
 - Metadata from data source
- Storage plugins perform scanner and writer functions

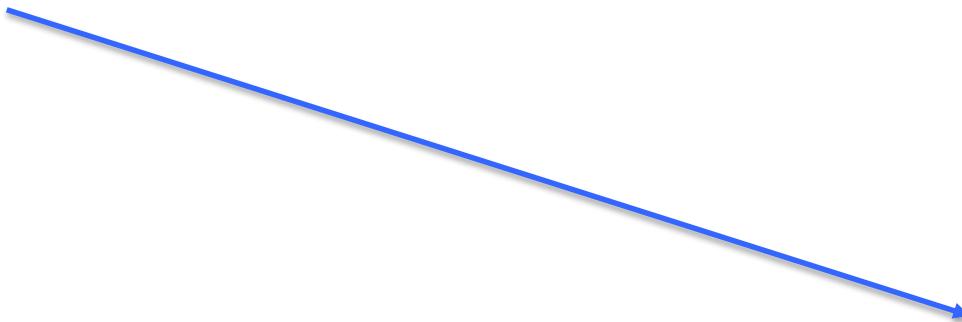




Specifying the data source

Storage Plugin Registration

FROM **dfs.clicks.`clicks/clicks.json`**



The screenshot shows the Apache Drill interface with the URL 192.168.110.133:8047/storage. The 'Apache Drill' tab is selected. The 'Enabled Storage Plugins' section lists four plugins: cp, dfs, hive, and maprdb. Each plugin has 'Update' and 'Disable' buttons next to it.

Storage Plugin	Action	Action
cp	Update	Disable
dfs	Update	Disable
hive	Update	Disable
maprdb	Update	Disable

<https://cwiki.apache.org/confluence/display/DRILL/Storage+Plugin+Registration>



Define Workspaces in the DFS Storage Plugin

The screenshot shows the Apache Drill storage configuration interface. At the top, there are buttons for 'Update' and 'Disable' for the 'cp' plugin. Below this, for the 'dfs' plugin, the 'Update' button is circled in red. The main area displays a list of disabled storage plugins: hbase, hive, and mongo, each with 'Update' and 'Enable' buttons.

Apache Drill

cp Update Disable

dfs **Update** Disable

Disabled Storage Plugins

hbase	Update	Enable
hive	Update	Enable
mongo	Update	Enable

The screenshot shows the Apache Drill storage configuration interface for the 'dfs' plugin. The title is 'Configuration'. The JSON configuration for workspaces is displayed:

```
"workspaces": {  
    "root": {  
        "location": "/",  
        "writable": false,  
        "defaultInputFormat": null  
    },  
    "tmp": {  
        "location": "/tmp",  
        "writable": true,  
        "defaultInputFormat": null  
    },  
    "demo": {  
        "location": "/Users/tshiran/Development/demo/data",  
        "writable": true,  
        "defaultInputFormat": null  
    }  
}
```

At the bottom, there are buttons for 'Back', 'Update', 'Disable', and 'Delete', with 'Delete' being highlighted in red. A large red oval encloses the 'demo' workspace configuration.

Apache Drill

Configuration

```
"workspaces": {  
    "root": {  
        "location": "/",  
        "writable": false,  
        "defaultInputFormat": null  
    },  
    "tmp": {  
        "location": "/tmp",  
        "writable": true,  
        "defaultInputFormat": null  
    },  
    "demo": {  
        "location": "/Users/tshiran/Development/demo/data",  
        "writable": true,  
        "defaultInputFormat": null  
    }  
}
```

Back Update Disable **Delete**



Query Basics Data Source Workspace

```
SELECT *\nFROM dfs.clicks.`clicks/clicks.json`
```

OPTIONAL

A workspace
- Sub-directory

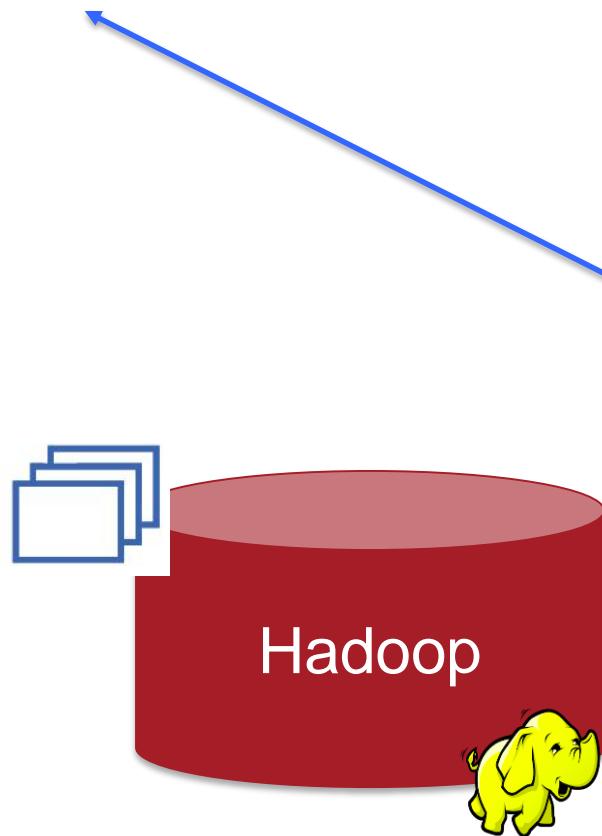
workspace defines a schema to query

```
{\n    "type" : "file",\n    "enabled" : true,\n    "connection" : "maprfs:///",\n    "workspaces" : {\n        "clicks" : {\n            "location" : "/mapr/data",\n            "writable" : false,\n            "storageformat" : json\n        }\n    }\n}
```



Data Source table

```
SELECT *\nFROM dfs.clicks.`clicks/clicks.json`
```



where you want it from

- A *table*
- **pathname**
- HBase table
- Hive table



Setting the Workspace

- ▶ You can set a workspace:

```
drill:> use dfs.clicks;
```

```
drill:> SELECT * FROM `clicks/clicks.json`
```

```
{  
  "type" : "file",  
  "enabled" : true,  
  "connection" : "maprfs:/// ",  
  "workspaces" : {  
    "clicks" : {  
      "location" : "/mapr/data",  
      "writable" : false,  
      "storageformat" : json  
    }  
  }  
}
```

Browse SQL

View Definition SQL:

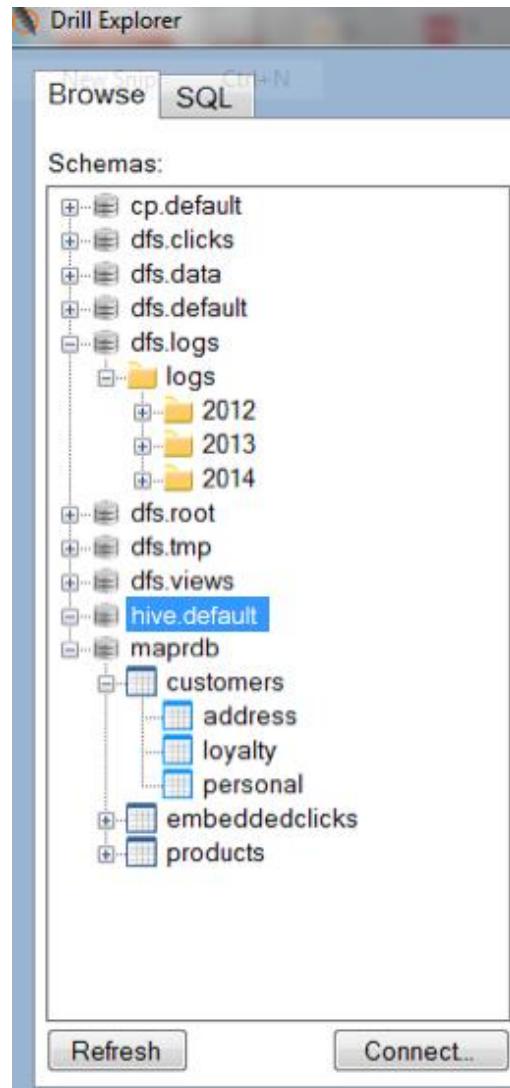
```
select * from `clicks/clicks.json` limit 2;
```

Total Number of Records: 2

	trans_id	date	time	user_info	trans_info
▶ 1	31920	2014-04-26	12:17:12	{ "cust_id": 22526, "device": "IOS5", "state": "il"}	{ "prod_id": [174, 2], "purch_flag": "false" }
2	31026	2014-04-20	13:50:29	{ "cust_id": 16368, "device": "AOS4.2", "state": "nc"}	{ "prod_id": [], "purch_flag": "false" }



Define Data Interaction & Schema Detection

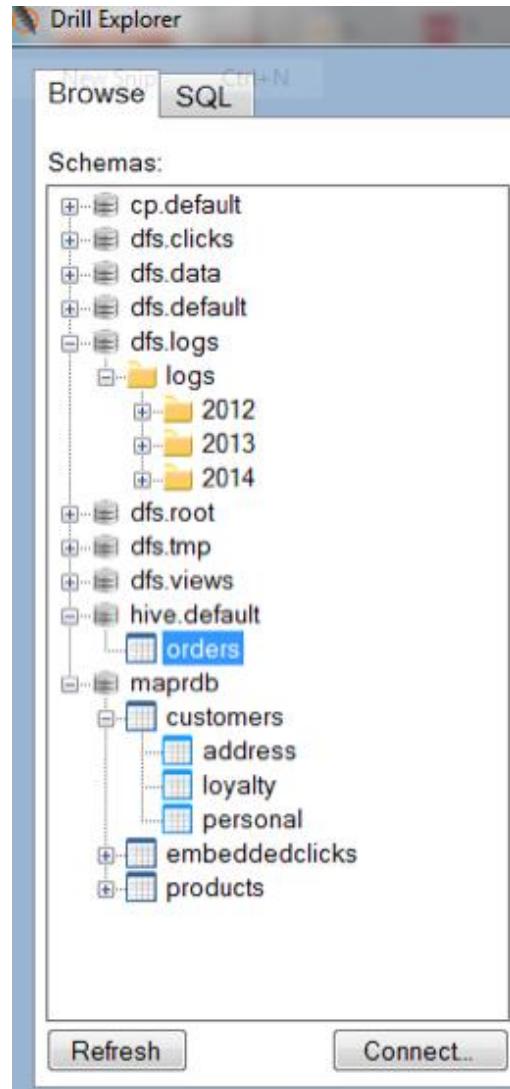


Storage system type

```
SELECT `month`, SUM(order_total)  
as sales FROM hive.orders GROUP  
BY `month` ORDER BY sales desc;
```



Define Data Interaction & Schema Detection



Storage system type

```
SELECT `month`, SUM(order_total)  
as sales FROM hive.orders GROUP  
BY `month` ORDER BY sales desc;
```

Table
Name

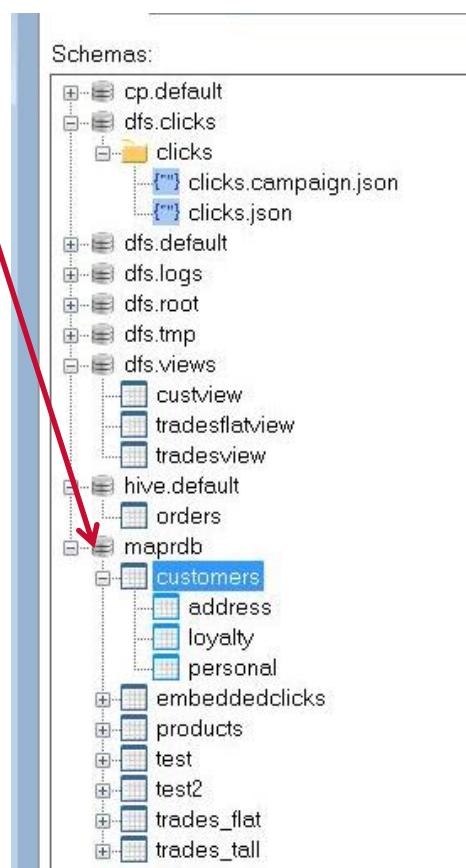


Complex Query Examples

```
SELECT cast(cust.personal.name as varchar(40)) as cust_name,  
clk.prod_id, clk.`date`  
FROM maprdb.customers as cust
```

Storage system type

Table Name



Metadata: `maprdb`.`customers`

	COLUMN_NAME	DATA_TYPE
1	row_key	ANY
2	address	(VARCHAR(1), ANY) MAP
3	loyalty	(VARCHAR(1), ANY) MAP
4	personal	(VARCHAR(1), ANY) MAP

Data Preview:

	row_key	address	loyalty	personal
1	0x3130303031	{ "state" : "InZhlg=="} { "agg_rev" : "MTk3", "membership" : "InNpbHZlcil="}	{ "agg_rev" : "MjE1LTlwlglg==", "membe	{ "age" : "ijE1LTlwlglg==", "in
2	0x3130303035	{ "state" : "ImIulg=="}	{ "agg_rev" : "MjMw", "membership" : "InNpbHZlcil="}	{ "age" : "ijI2LTM1lg==", "in
3	0x3130303036	{ "state" : "ImNhlg=="} { "agg_rev" : "MjUw", "membership" : "InNpbHZlcil="}	{ "agg_rev" : "MjUw", "membership" : "InNpbHZlcil="}	{ "age" : "ijI2LTM1lg==", "in
4	0x3130303037	{ "state" : "Im11lg=="} { "agg_rev" : "MjYz", "membership" : "InNpbHZlcil="}	{ "agg_rev" : "MjYz", "membership" : "InNpbHZlcil="}	{ "age" : "ijUxLTEwMCI=", "in
5	0x3130303130	{ "state" : "Im1ulg=="} { "agg_rev" : "MjAy", "membership" : "InNpbHZlcil="}	{ "agg_rev" : "MjAy", "membership" : "InNpbHZlcil="}	{ "age" : "ijUxLTEwMCI=", "in
6	0x3130303131	{ "state" : "Imhplg=="} { "agg_rev" : "MTY5", "membership" : "InNpbHZlcil="}	{ "agg_rev" : "MTY5", "membership" : "InNpbHZlcil="}	{ "age" : "ijUxLTEwMCI=", "in
7	0x3130303132	{ "state" : "Im9olg=="} { "agg_rev" : "MTU1", "membership" : "InNpbHZlcil="}	{ "agg_rev" : "MTU1", "membership" : "InNpbHZlcil="}	{ "age" : "ijIxLTl1lg==", "in
8	0x3130303134	{ "state" : "Im5qlg=="} { "agg_rev" : "MTk5", "membership" : "InNpbHZlcil="}	{ "agg_rev" : "MTk5", "membership" : "InNpbHZlcil="}	{ "age" : "ijE1LTlwlglg==", "in
9	0x3130303138	{ "state" : "Im55lg=="} { "agg_rev" : "MjQ4", "membership" : "InNpbHZlcil="}	{ "agg_rev" : "MjQ4", "membership" : "InNpbHZlcil="}	{ "age" : "ijUxLTEwMCI=", "in
10	0x31303032	{ "state" : "ImN0lg=="} { "agg_rev" : "MTA5MQ==", "membership" : "ImdvbGQi"}	{ "agg_rev" : "MTA5MQ==", "membership" : "ImdvbGQi"}	{ "age" : "ijIxLTl1lg==", "in

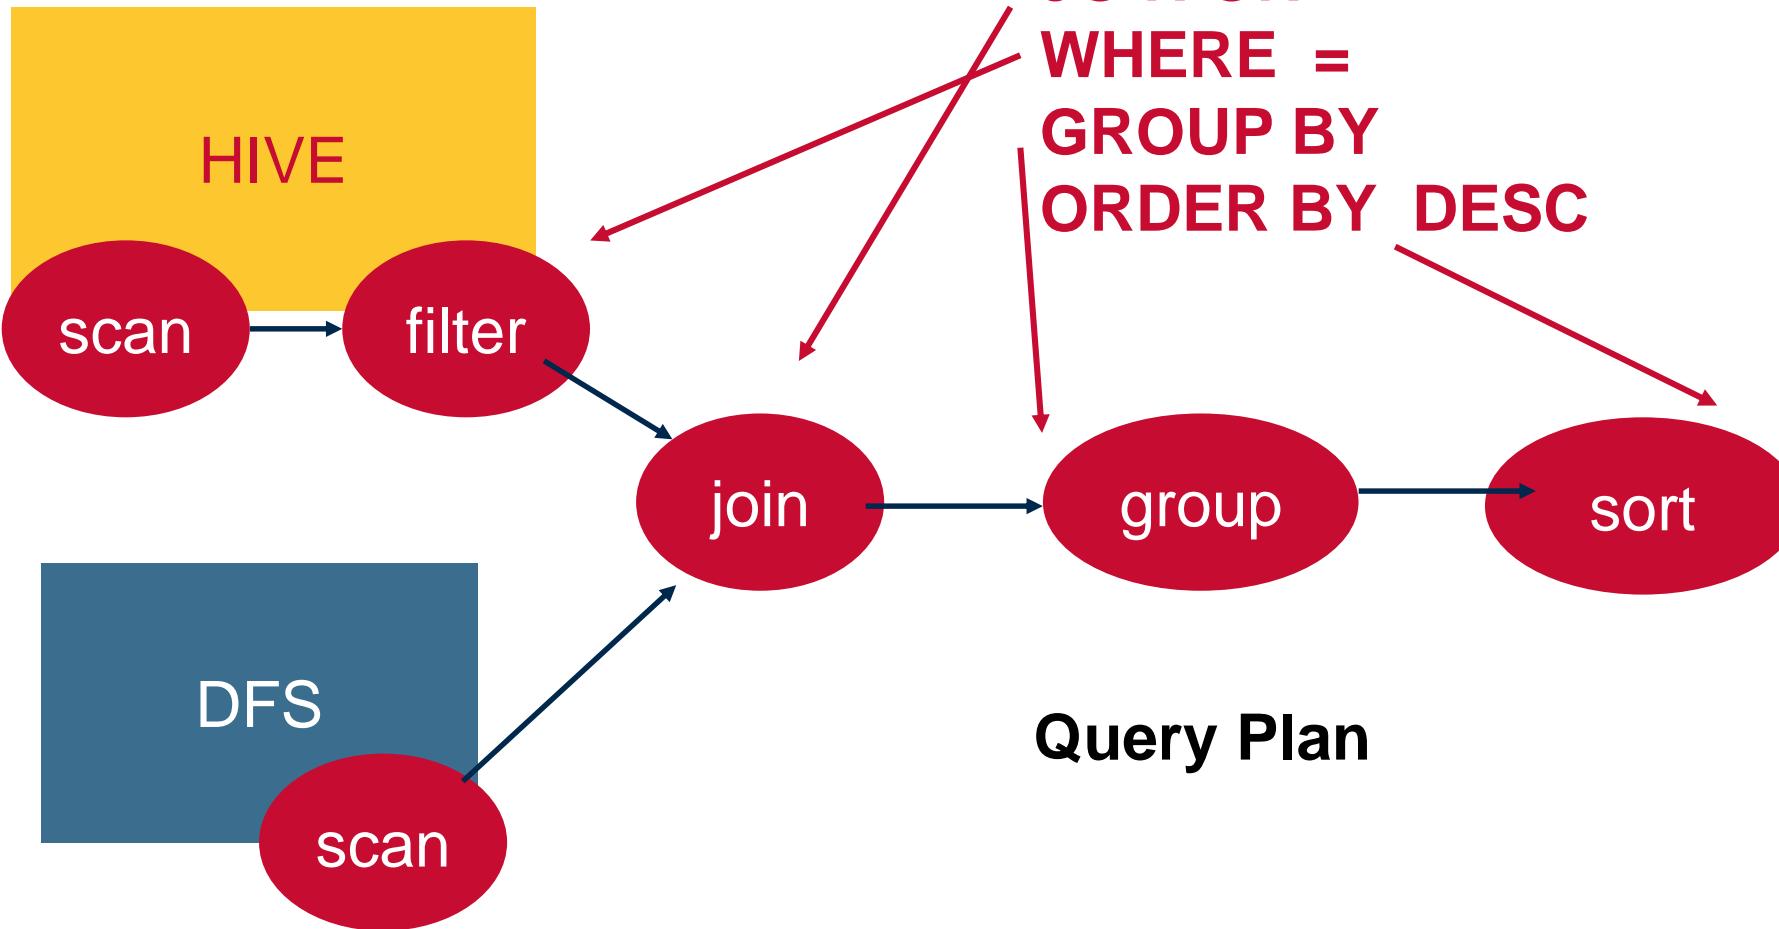


Drill Architecture and flow





SQL is the WHAT, now we will look at the HOW





Understanding the Explain plan

Business Scenario

- Understanding the Explain plan for this query
 - Which device do customers who made orders use ?

```
SELECT t.user_info.device device, t.user_info.cust_id cust_id, o.order_id
FROM dfs.clicks.`clicks/clicks.json` t, hive.orders o
WHERE t.user_info.cust_id=o.cust_id ;
```

JOIN ON

id	date	cust_id	device	state	prod_id	purch_flag
67212	June	10001	“iOS5”	“ca”	[174,2]	false

Web log clicks

order_id	month	cust_id	state	Prod_id	total
67212	June	10001	ca	909	13
70302	May	10004	ga	420	11

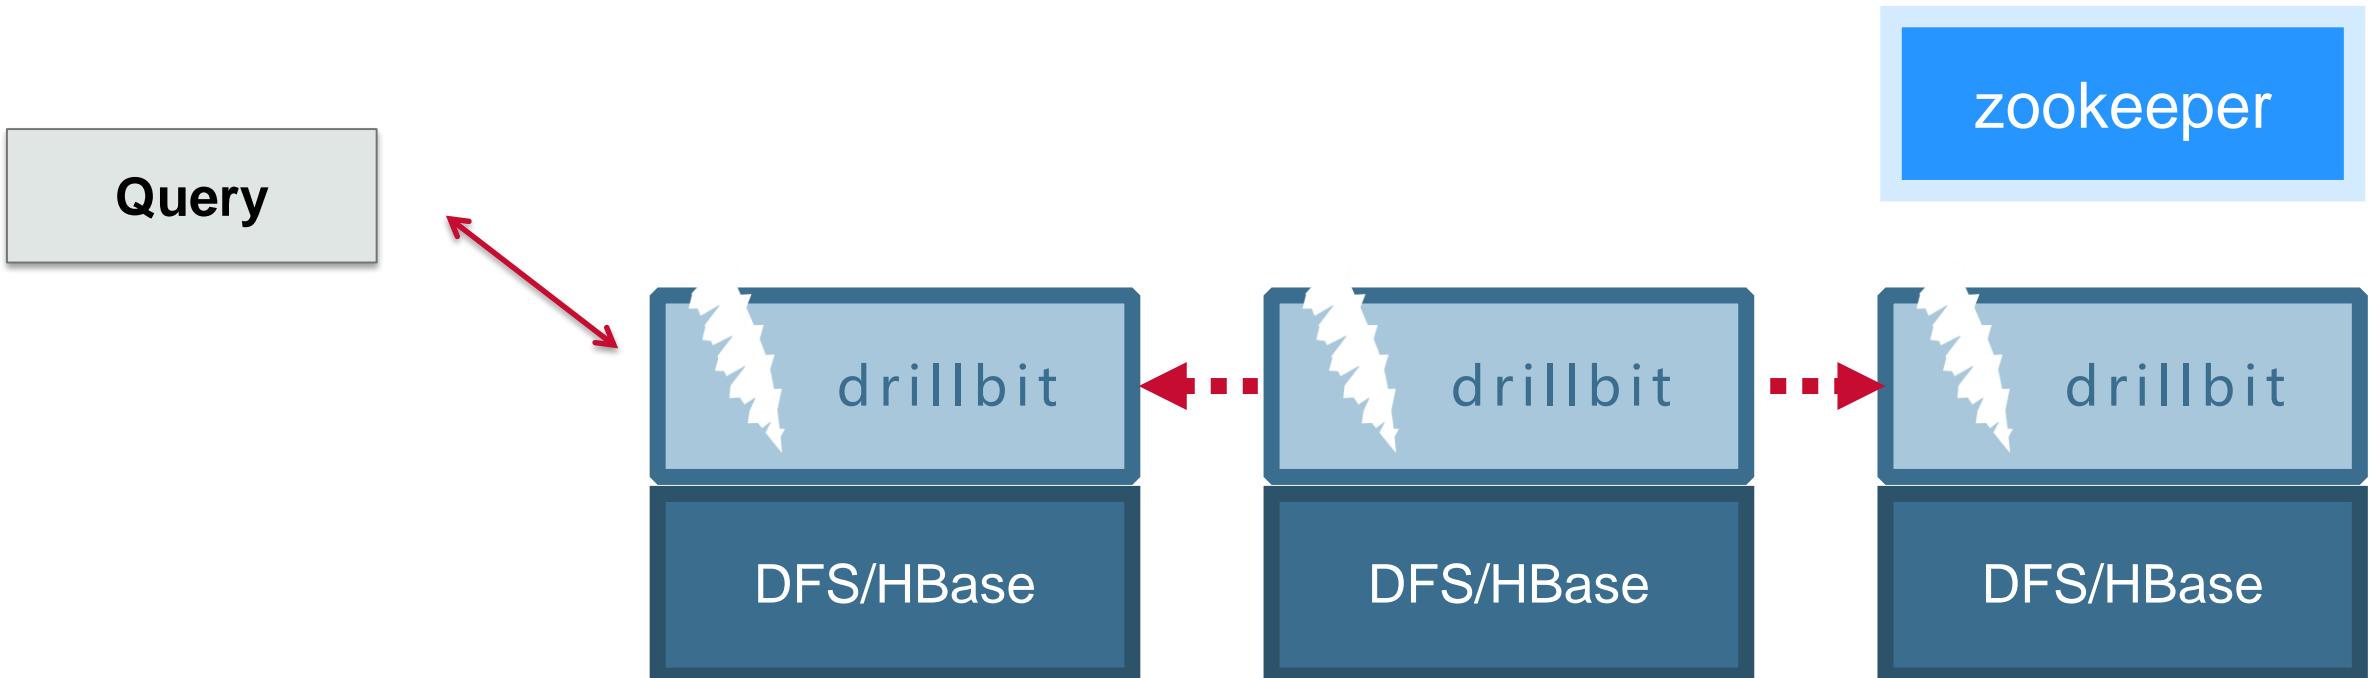
HIVE orders



How Does It Work?

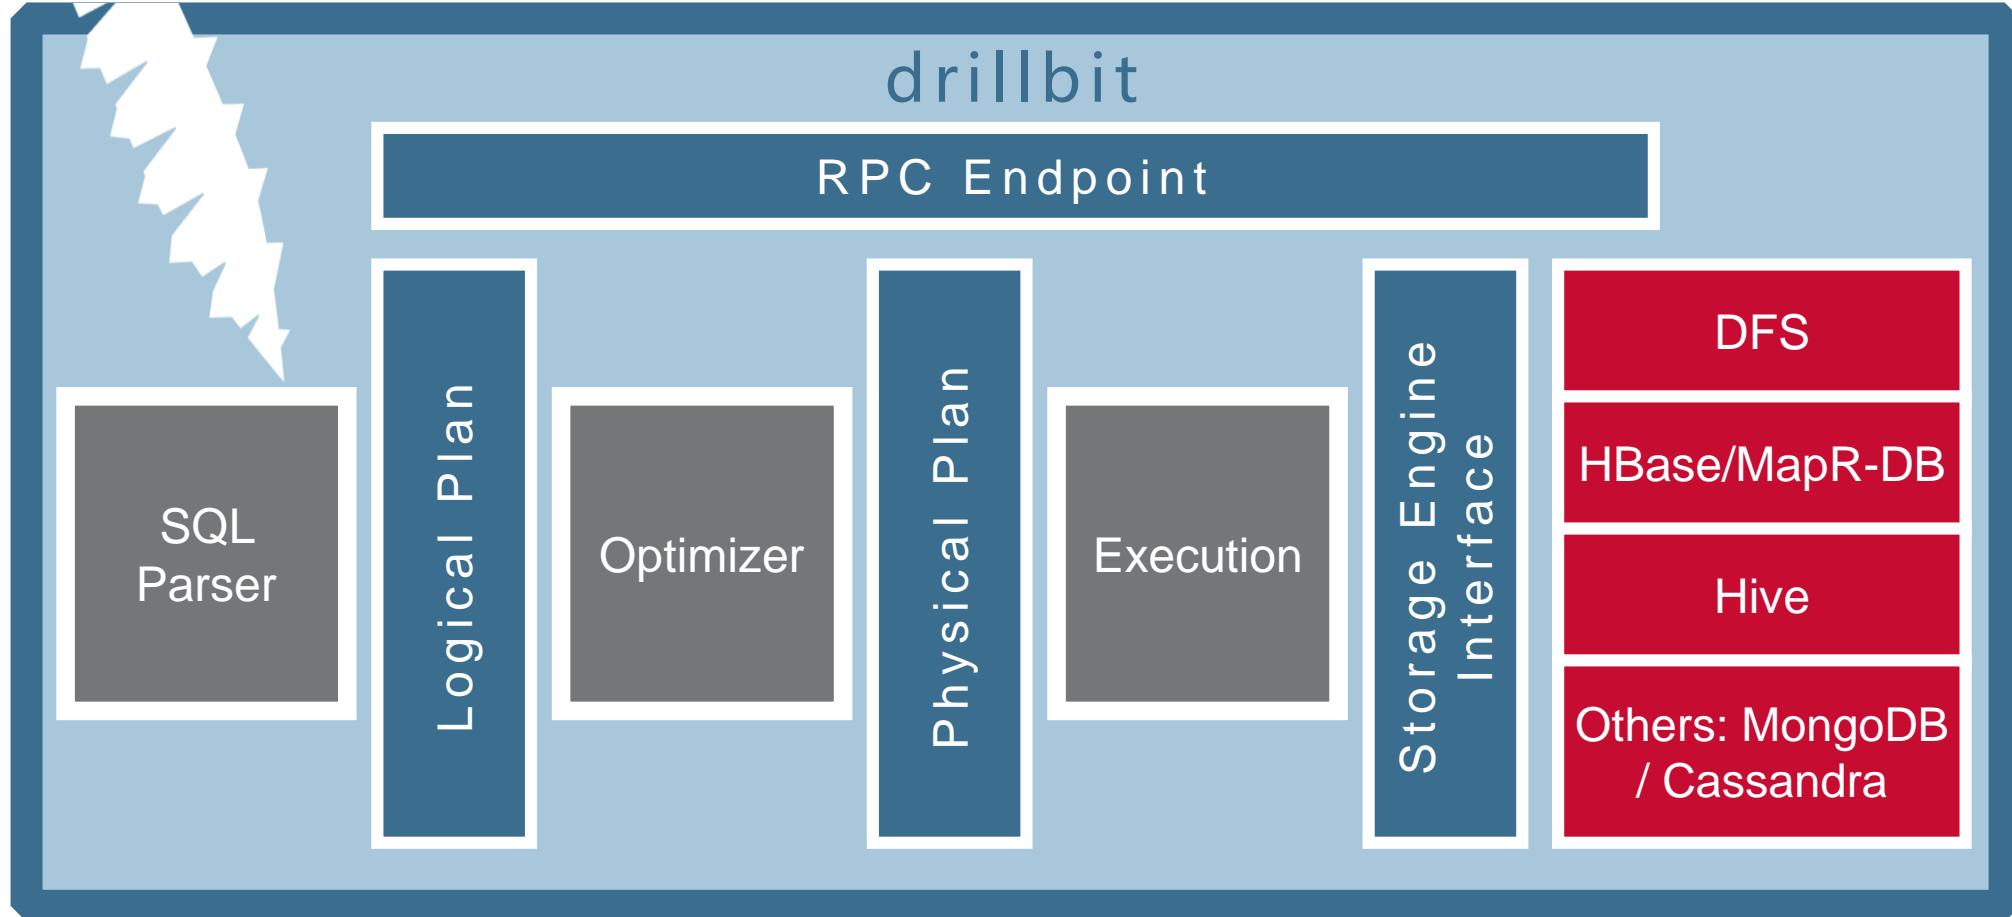
```
SELECT device, cust_id, order_id  
FROM clicks.json t, hive.orders o  
WHERE t.cust_id=o.cust_id
```

- “Drillbits” run on each node
 - designed to **maximize data locality**
- Coordination, query planning, optimization, scheduling, and execution are **distributed**





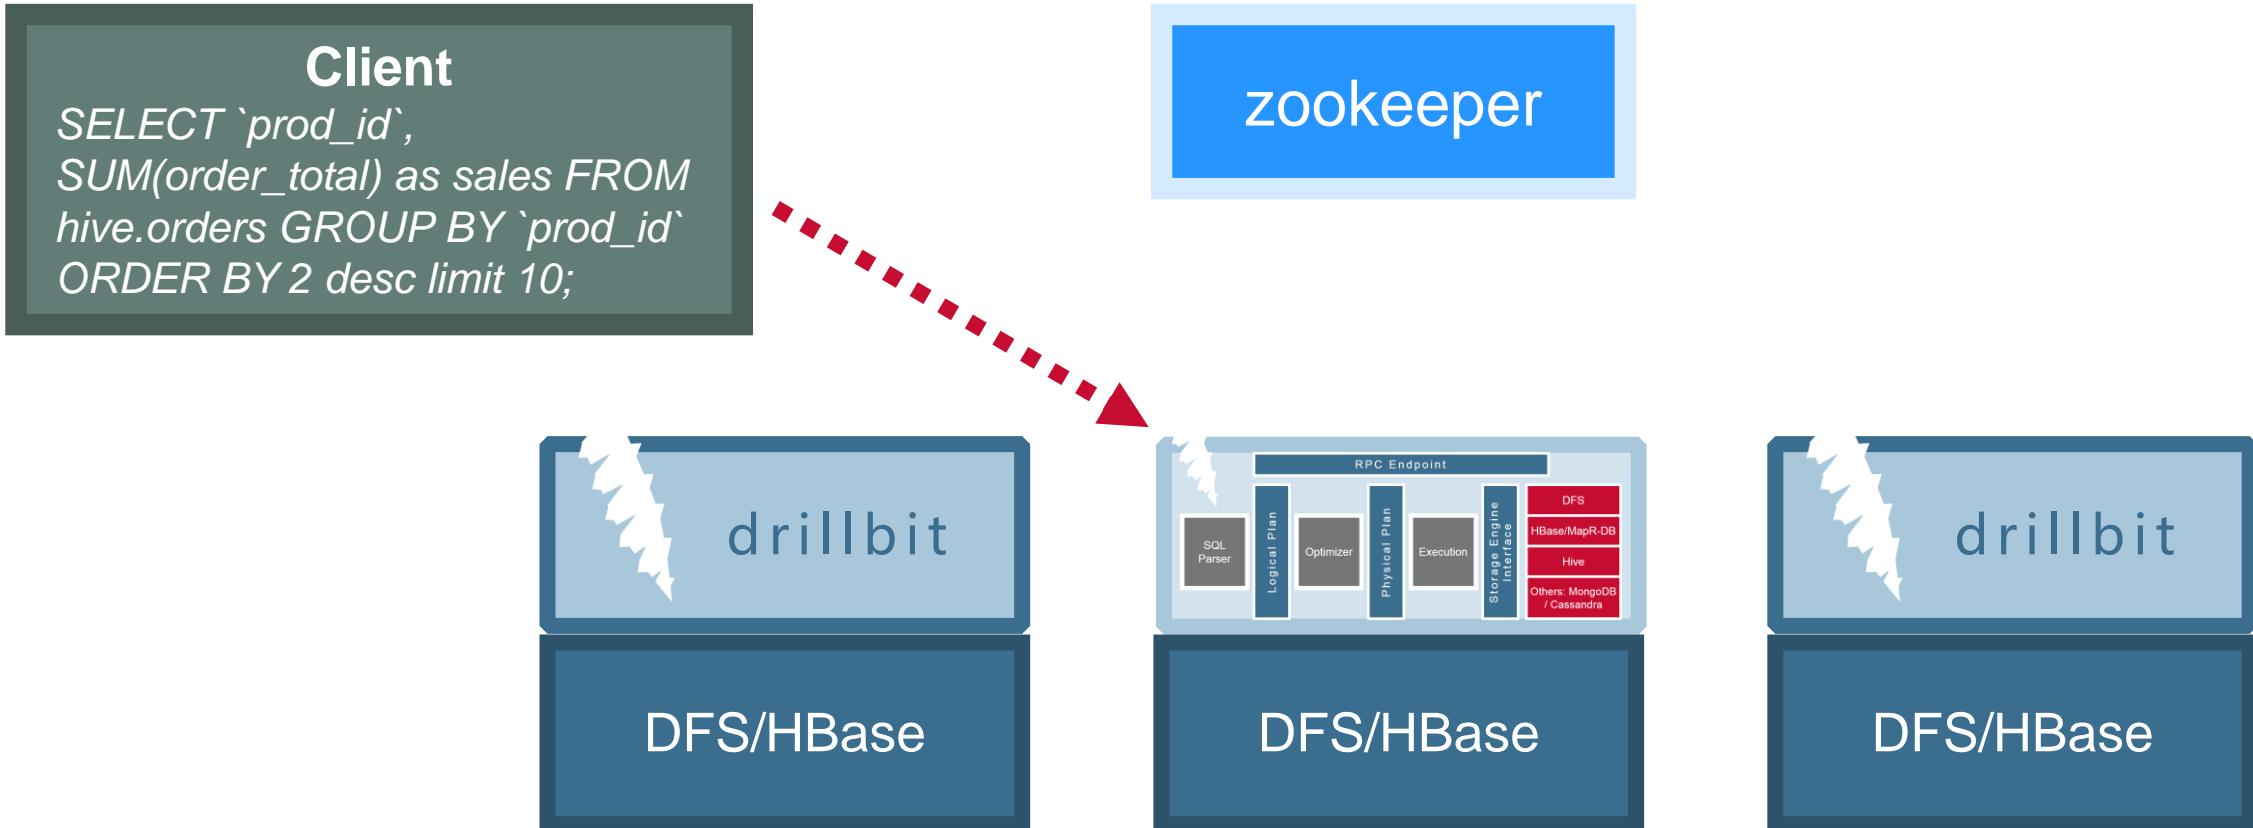
Core Modules within a Drillbit





Query Execution

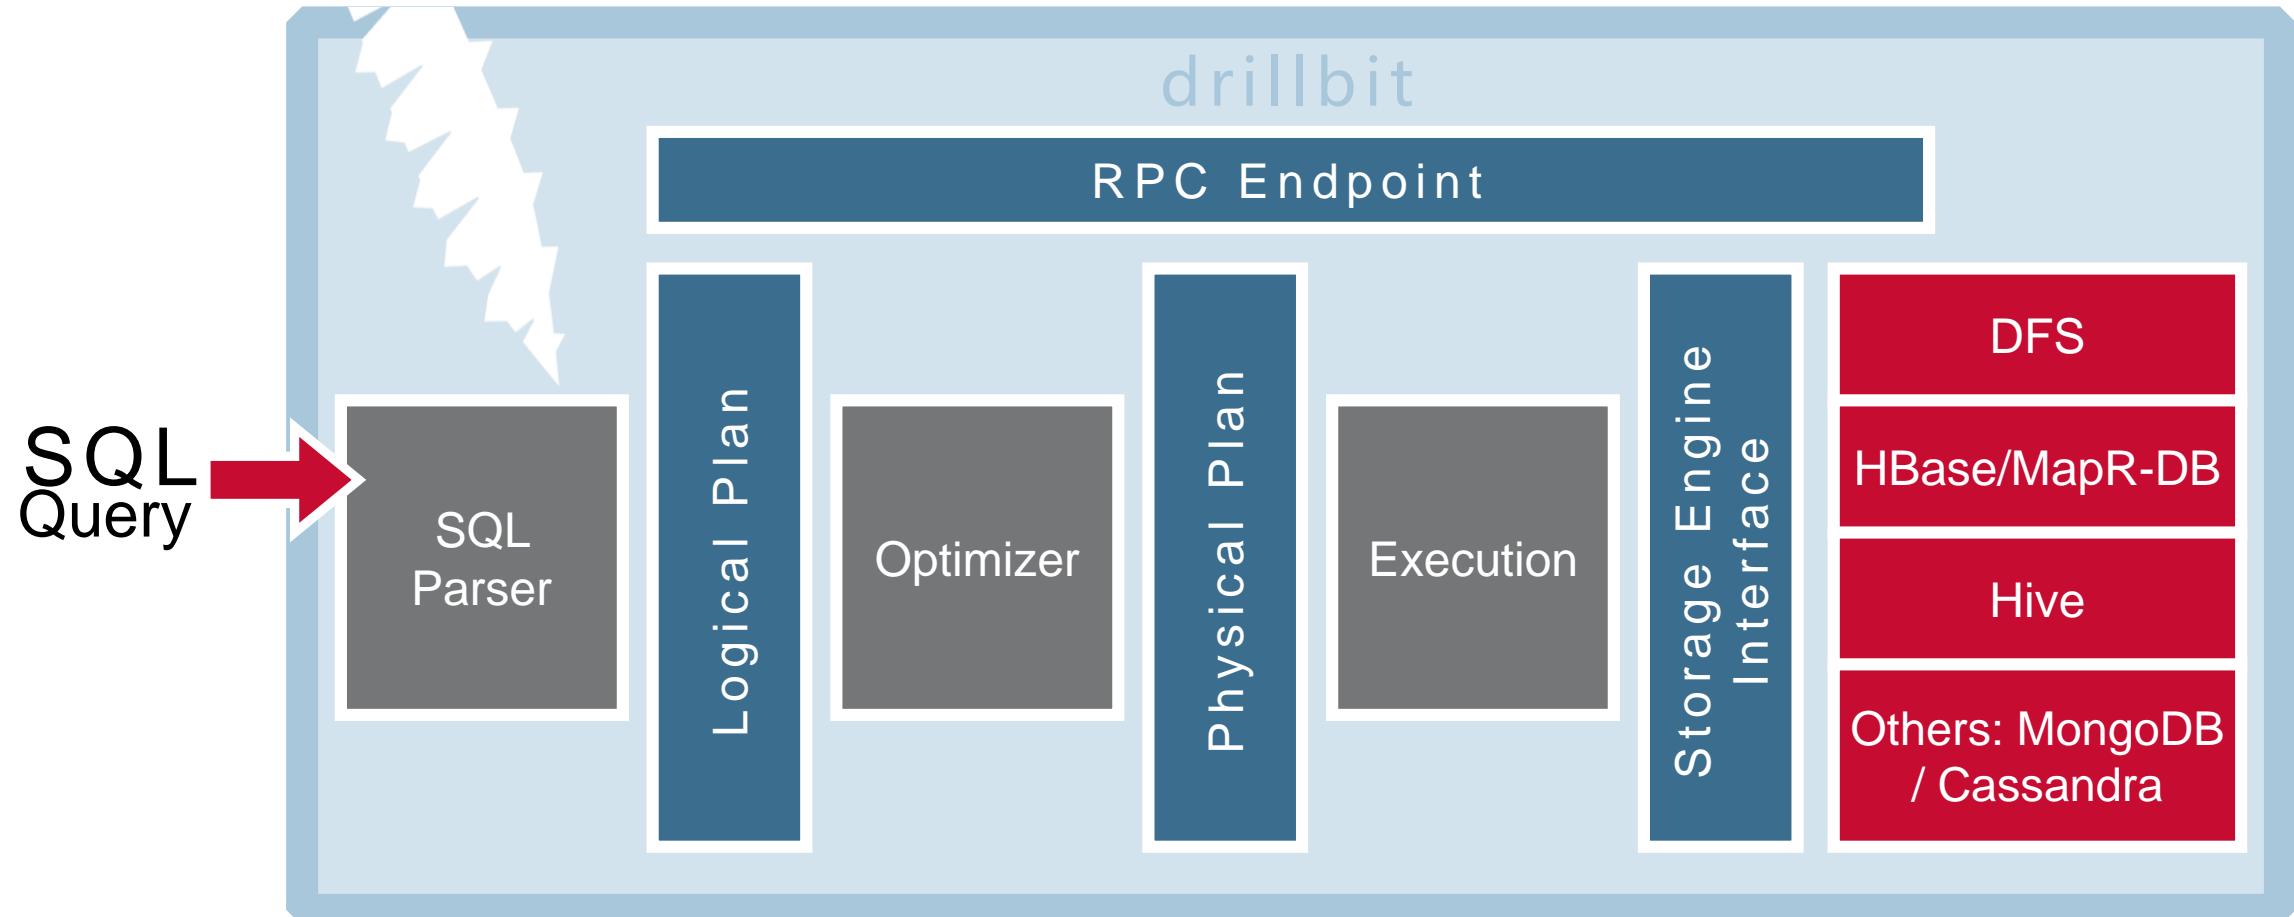
1. Query comes to any Drillbit (JDBC, ODBC, CLI)





Query Execution Plan

2. Drillbit generates execution plan based on query optimization & locality

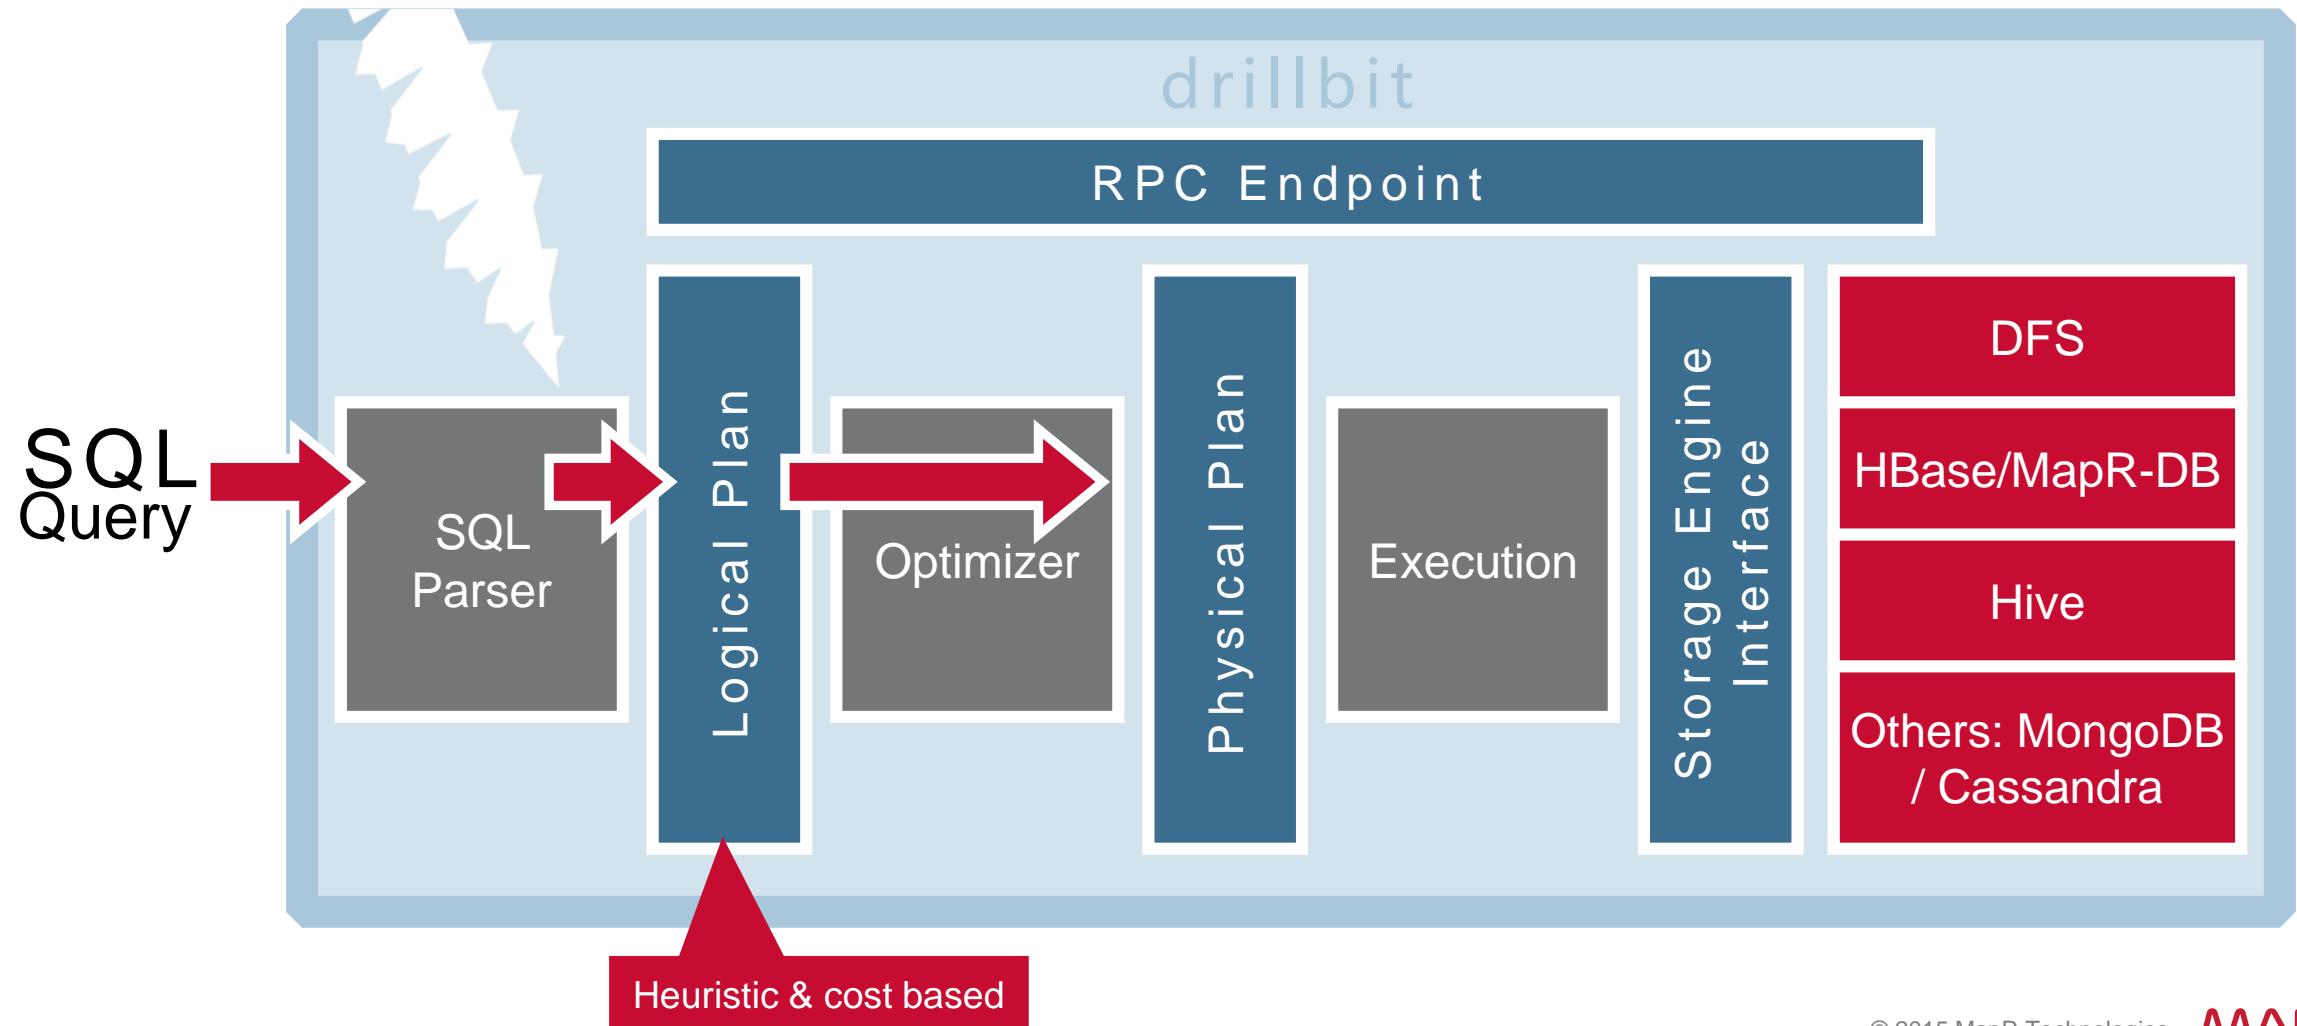


SQL What we want to do (analyst friendly)



Query Execution Plan

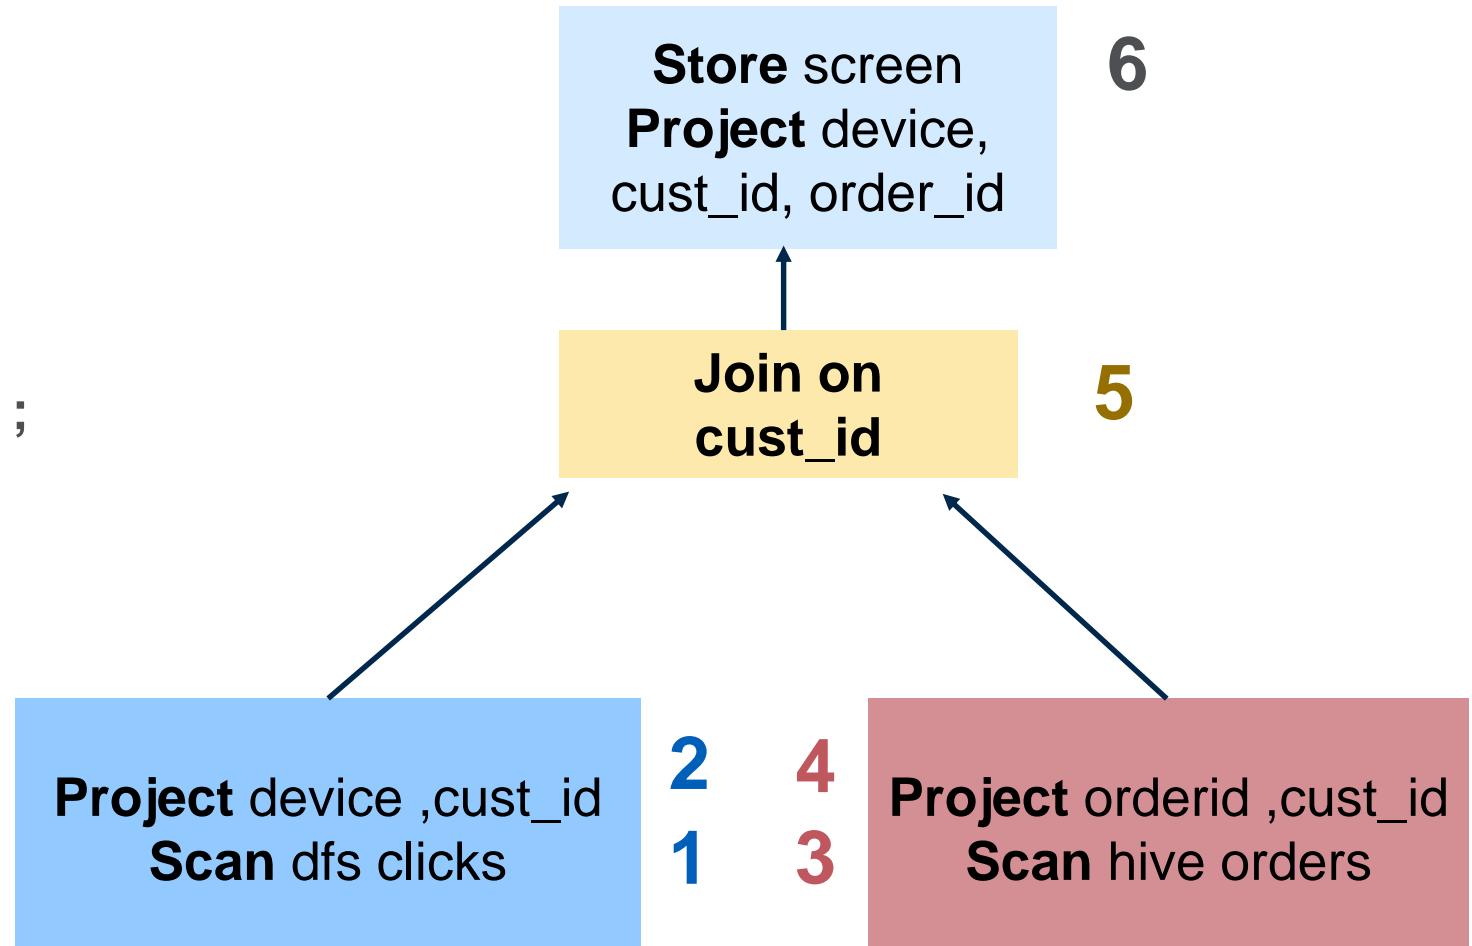
Logical Plan: What we want to do (language agnostic, computer friendly)





Logical Plan

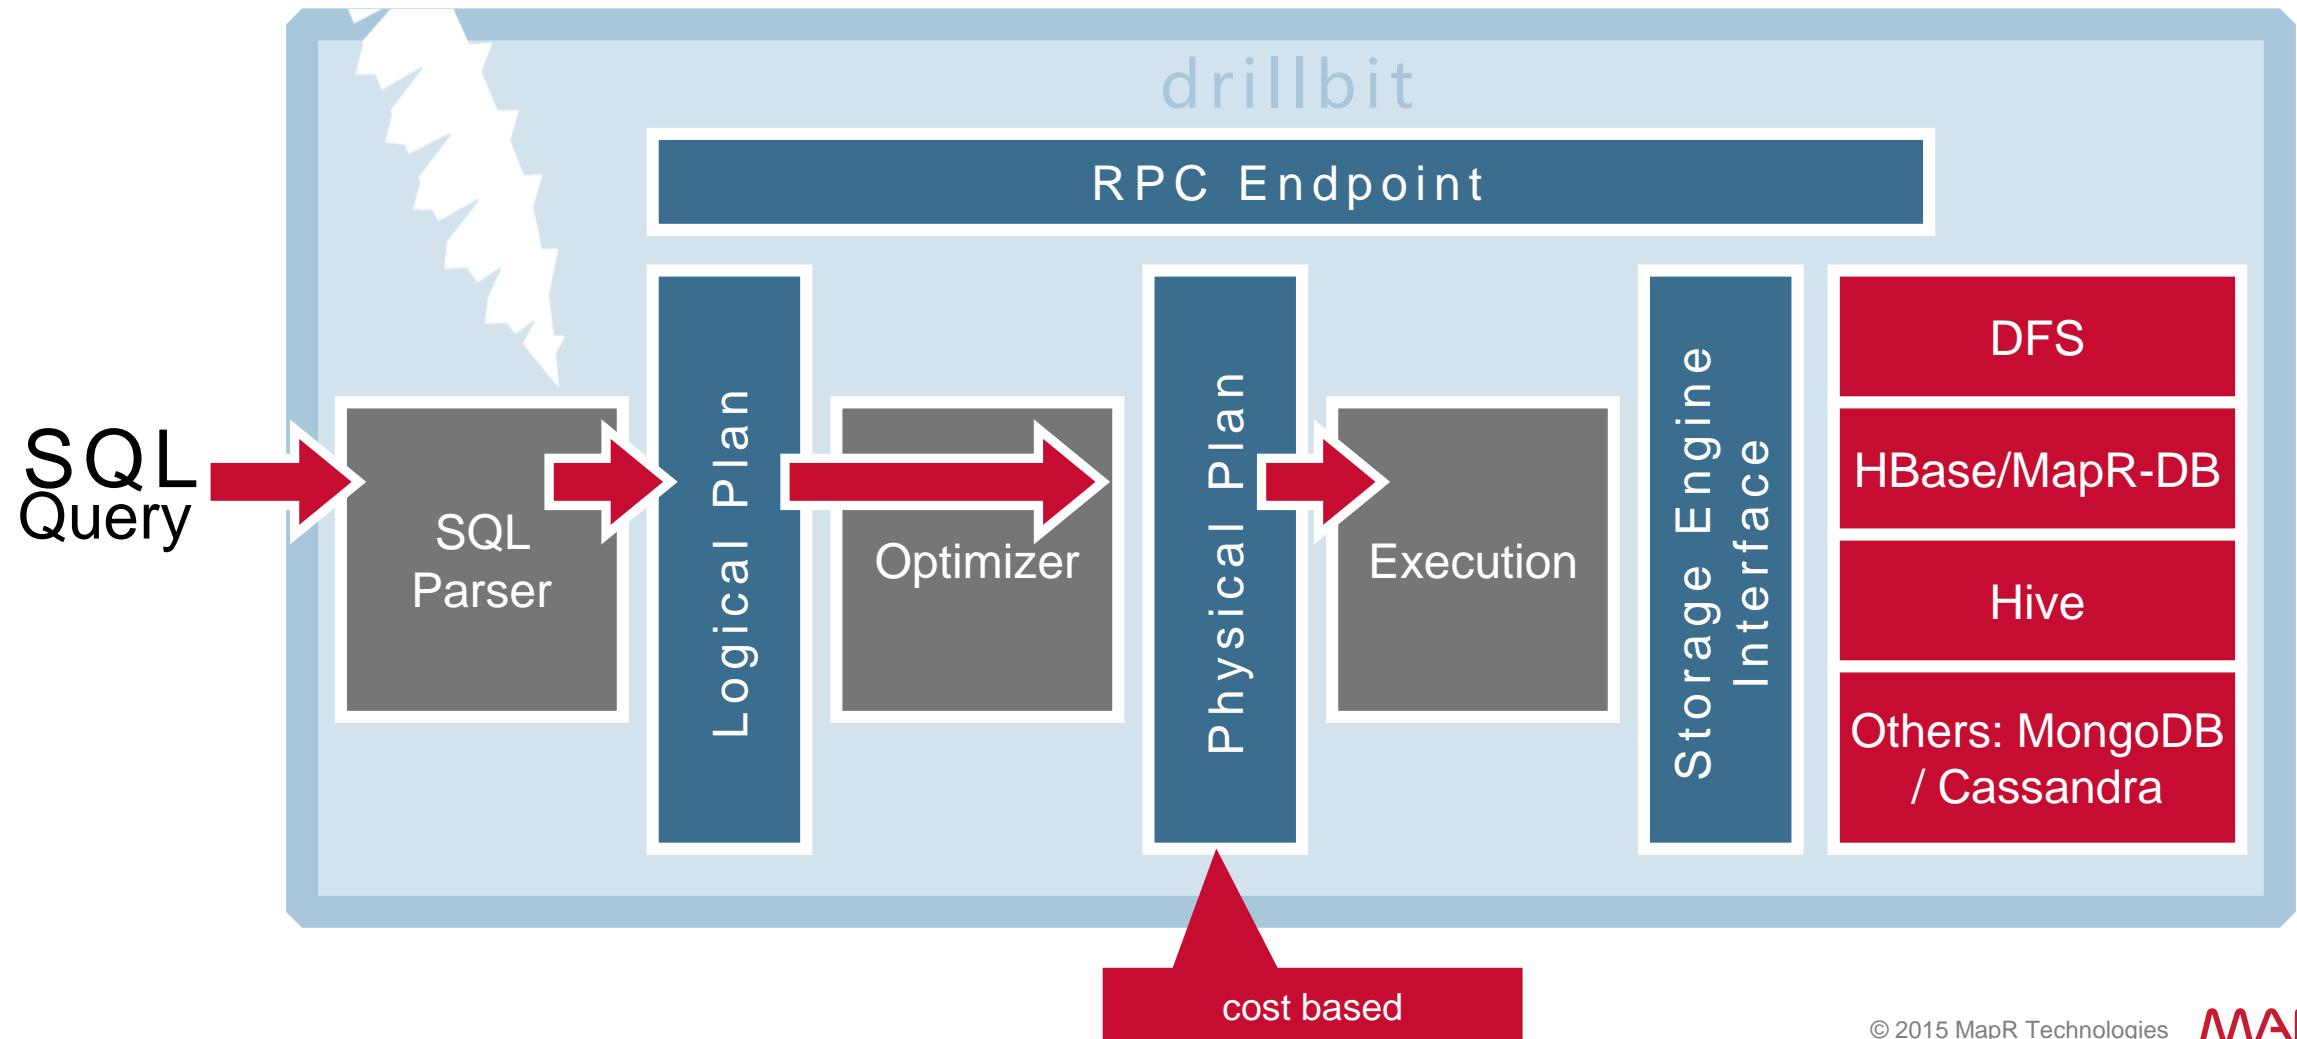
```
SELECT t.user_info.device  
device, t.user_info.cust_id  
cust_id, o.order_id  
FROM  
dfs.clicks.`clicks/clicks.json`  
t, hive.orders o  
WHERE  
t.user_info.cust_id=o.cust_id ;
```





Query Execution Plan

Physical Plan: How we want to do it (the best way we can tell)





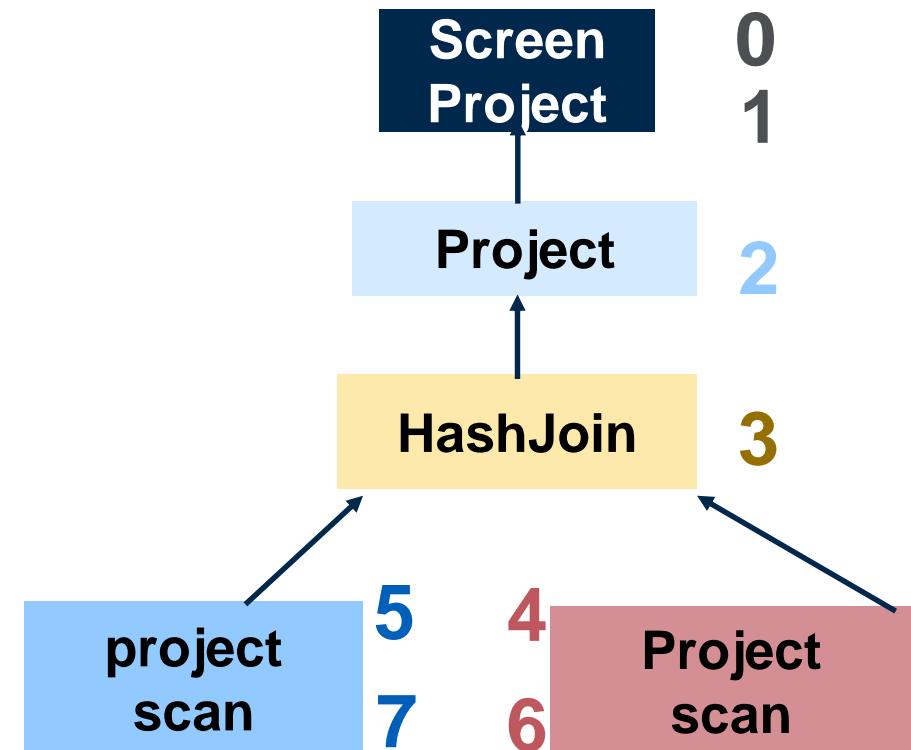
Drill Generates a Physical Plan, based on available operations

- A Physical Plan is composed of major Fragments
- Each operator has a particular **OperatorId**.
- You can capture a JSON representation of this plan, modify manually and submit

```
> !set maxwidth 10000
> explain plan for select * from customer limit 5;
...
00-00  Screen
00-01  Project
00-02  Project
00-03  HashJoin
00-04  Project
00-05  Scan
00-06  Project
00-07  Scan
```

Major Fragment Id

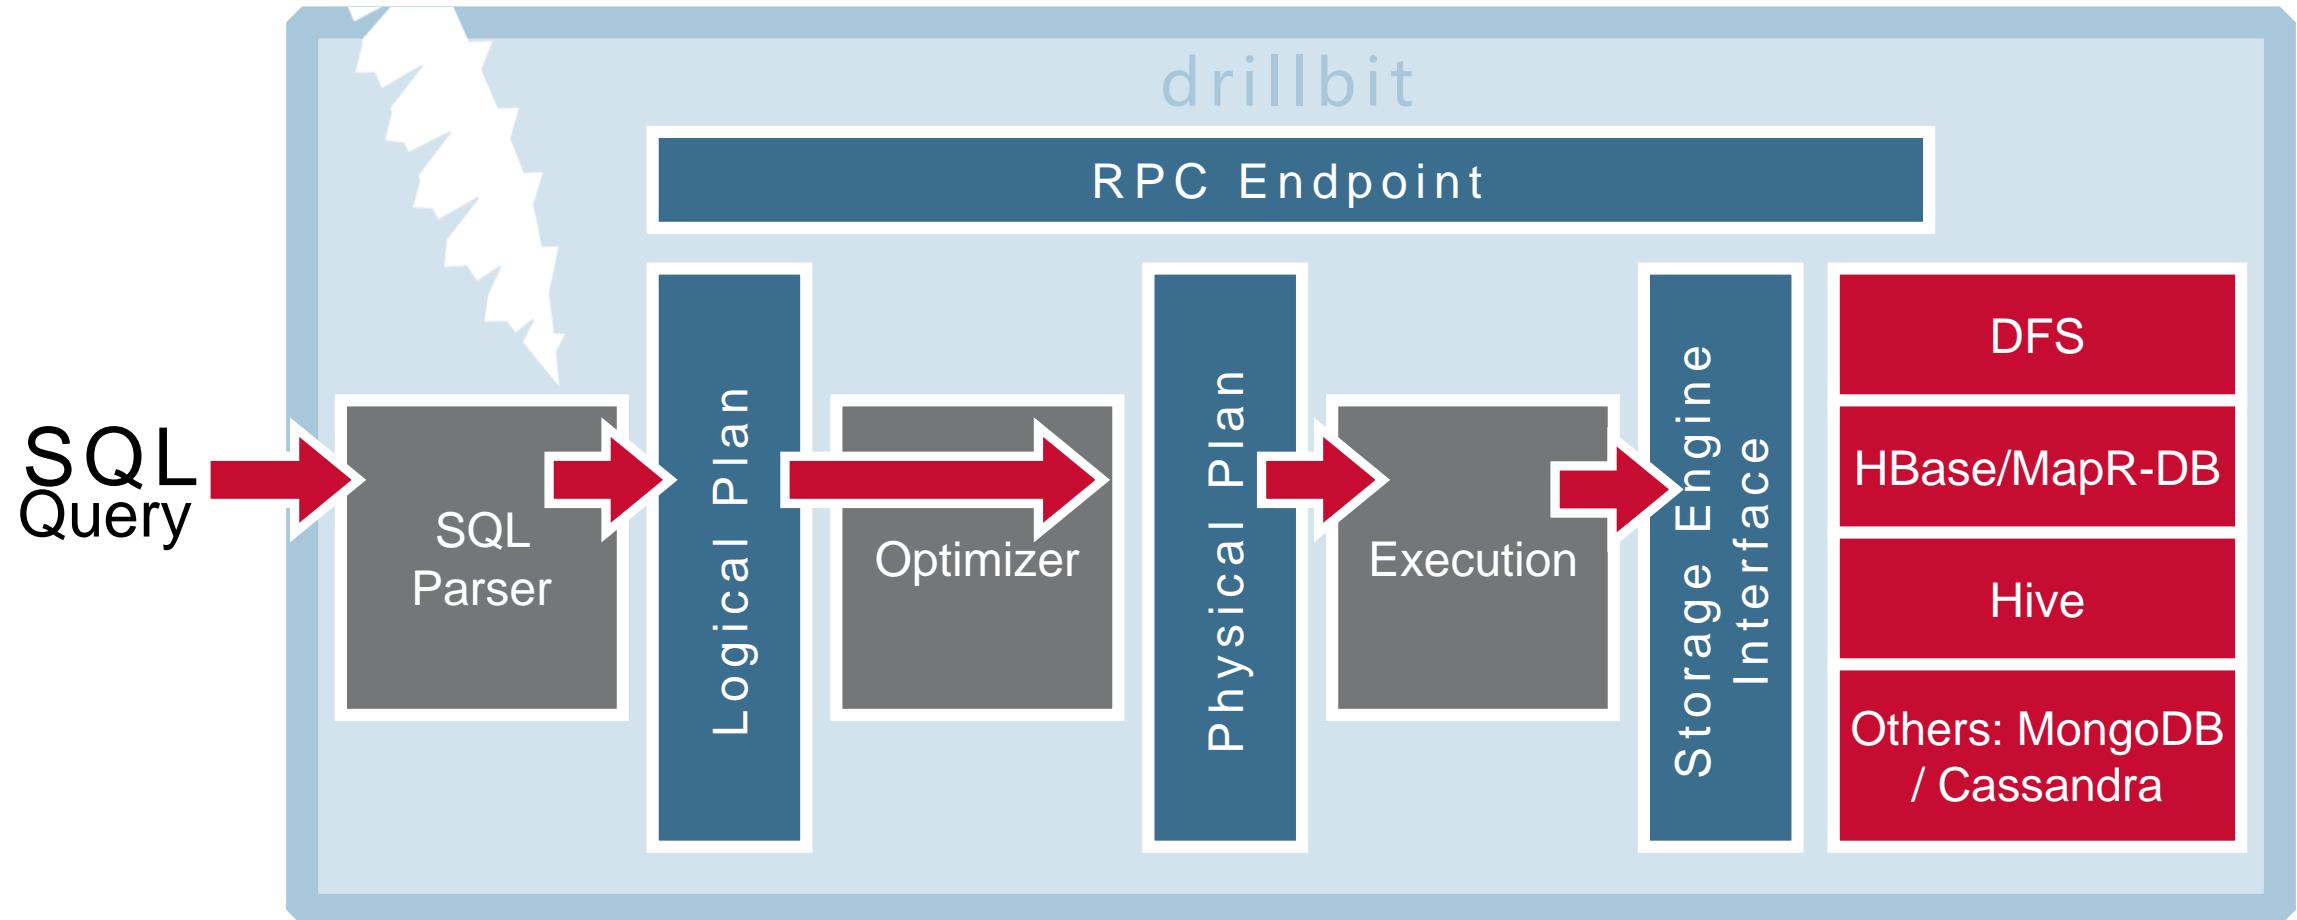
Operator Id





Query Execution Plan

Execution Plan (fragments): **Where** we want to do it



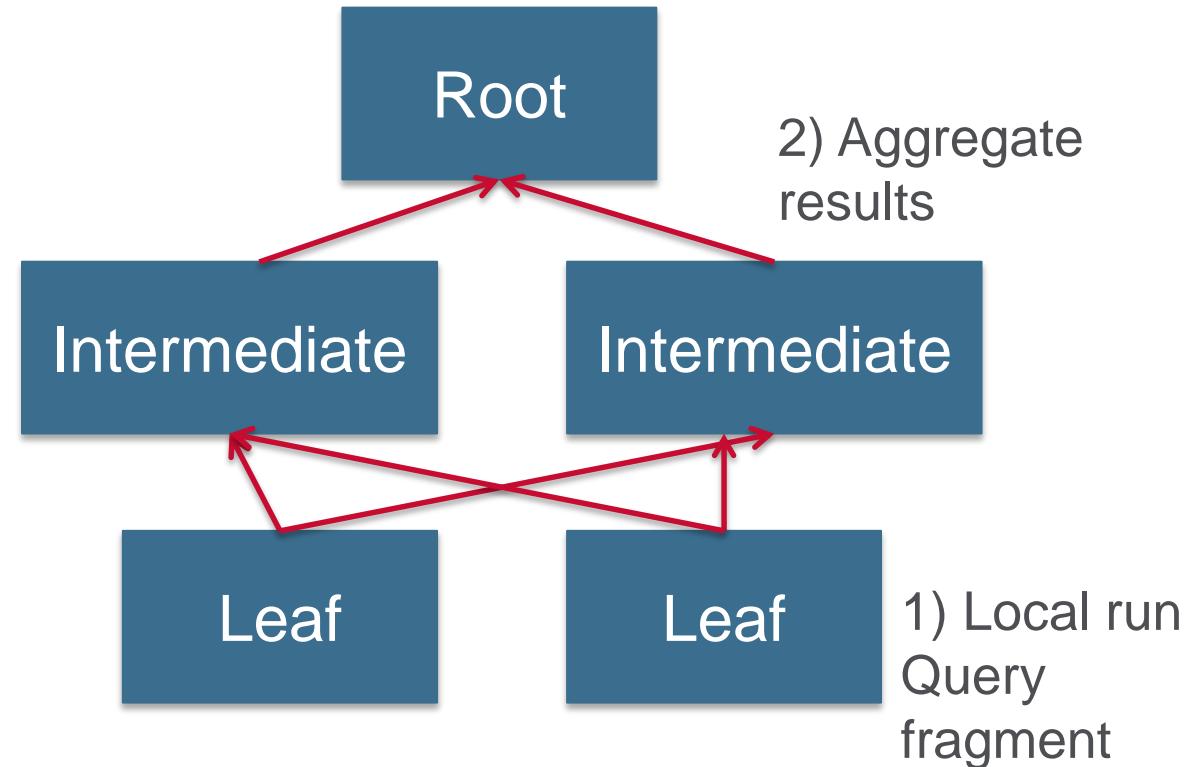
Plan fragments sent to drill bits



Query Execution Plan tree

Each execution plan has:

- One **root** fragment
- **1. Leaf fragments**
 - 1st to run
 - Leaf nodes **run query** fragment
- Intermediate fragments



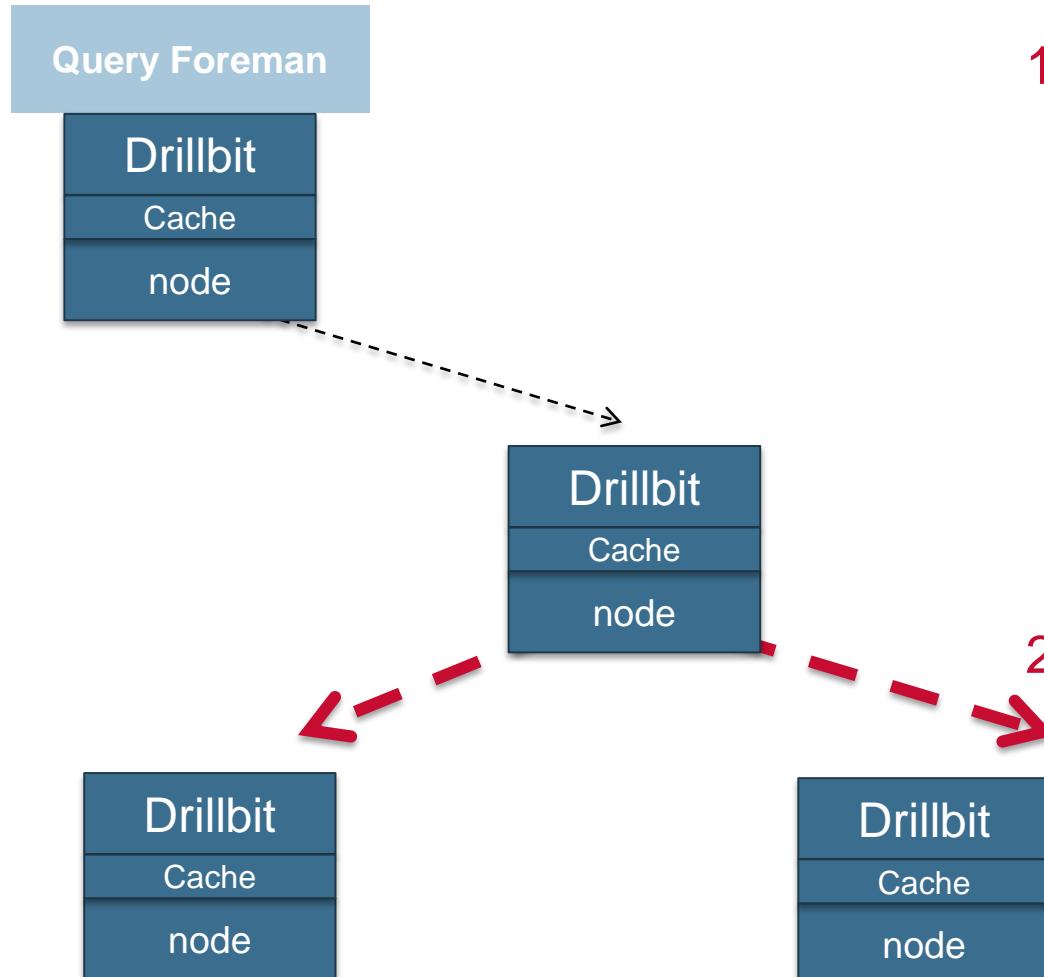


Foreman manages execution

SQL
Query

```
SELECT device, cust_id, order_id  
FROM clicks.json t, hive.orders o  
WHERE t.cust_id=o.cust_id ;
```

cust_id	device
10001	“iOS5”
10004	AOS4.2



1. Originating Drillbit manages query execution

2. leafs run query.

order_id	cust_id
67212	10001
70302	10004



Leaf fragments pass back results

Leaf nodes run query fragment on data source and send results up the tree

```
SELECT device, cust_id, order_id  
FROM clicks.json t, hive.orders o  
WHERE t.cust_id=o.cust_id ;
```

Project device ,cust_id
Scan dfs clicks

cust_id	device
10001	“iOS5”
10004	AOS4.2



**run
Query
fragment
On data
source**



Project orderid ,cust_id
Scan hive orders

order_id	cust_id
67212	10001
70302	10004

 Columnar Storage

- Drill **optimizes for** Columnar Storage
 - Parquet <http://parquet.incubator.apache.org/>
- Only retrieve columns participating in query

Logical table representation

a	b	c
a1	b1	c1
a2	b2	c2
a3	b3	c3
a4	b4	c4
a5	b5	c5

Row layout

a1	b1	c1	a2	b2	c2	a3	b3	c3	a4	b4	c4	a5	b5	c5
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

Column layout

a1	a2	a3	a4	a5	b1	b2	b3	b4	b5	c1	c2	c3	c4	c5
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

 Columnar Execution

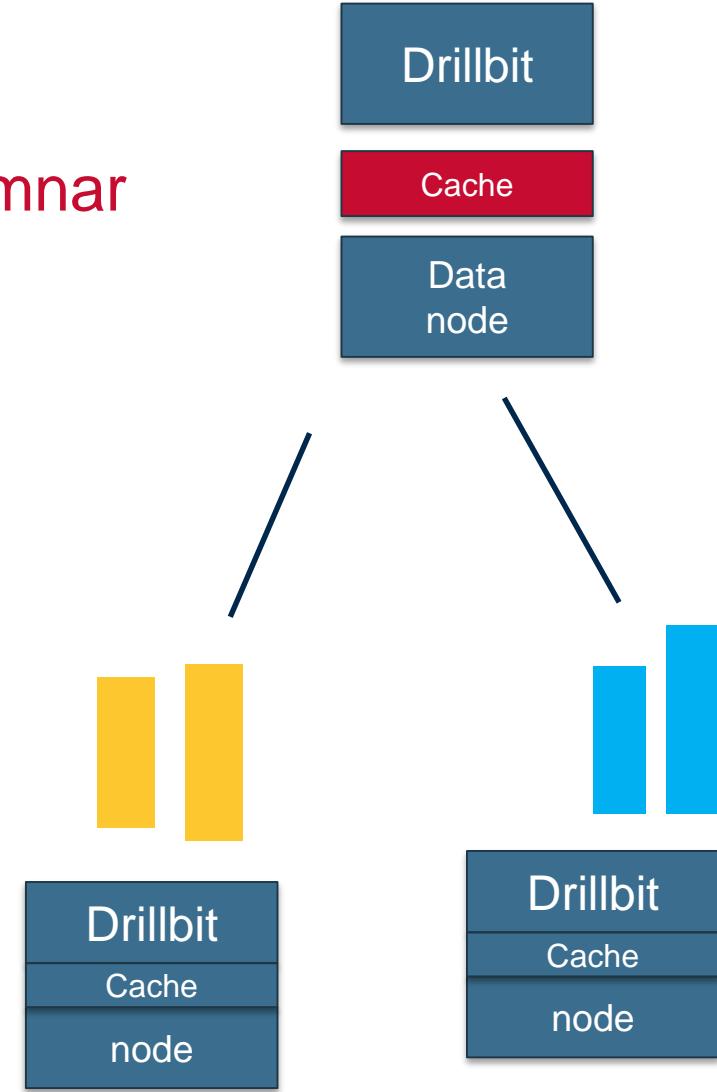
- Drill optimizes for Columnar **Execution**
 - **in-memory** data model that is hierarchical and **columnar**

order id	Cust_id
101	abc
102	def
104	ghi
105	jkl
108	mno
112	pqr

Data node

devic	Cust_id
101	abc
102	def
104	ghi
105	jkl
108	mno
112	pqr

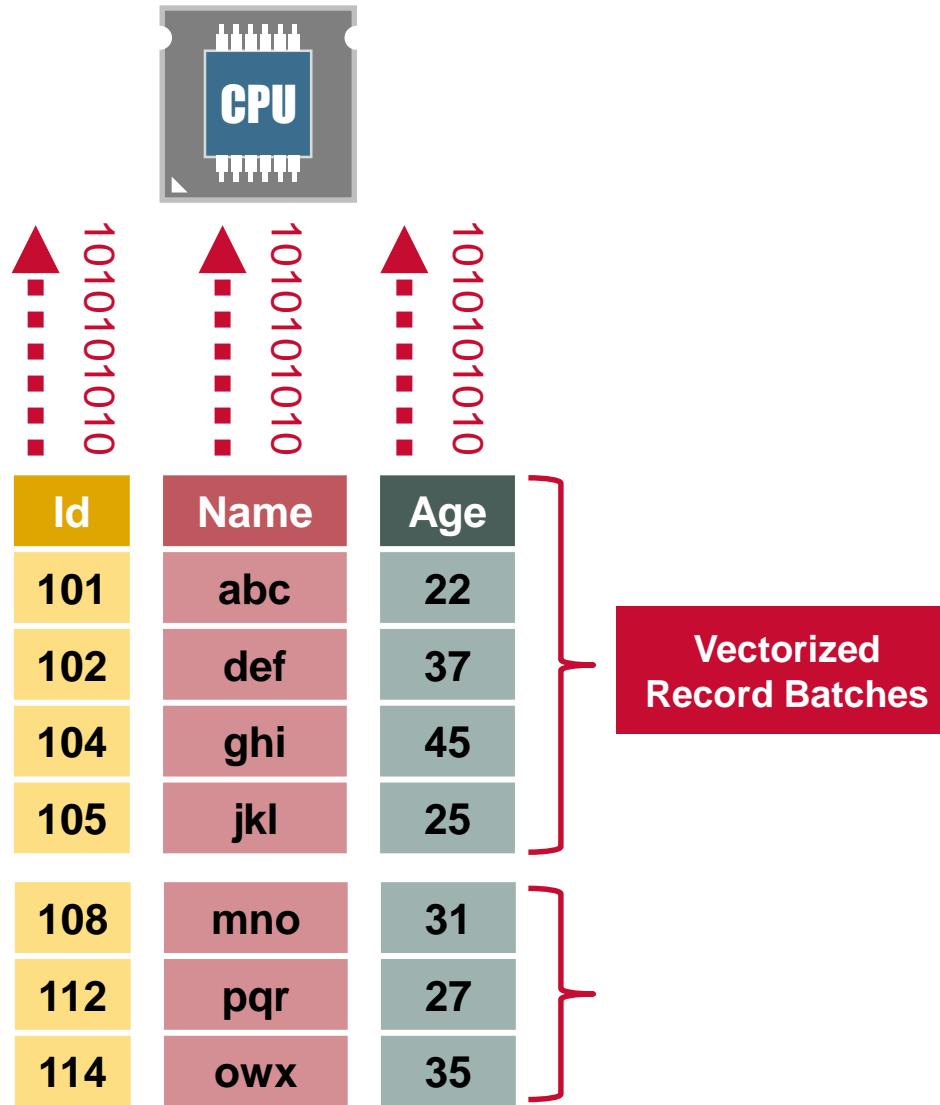
Data node





Vectorization

sets of columns values



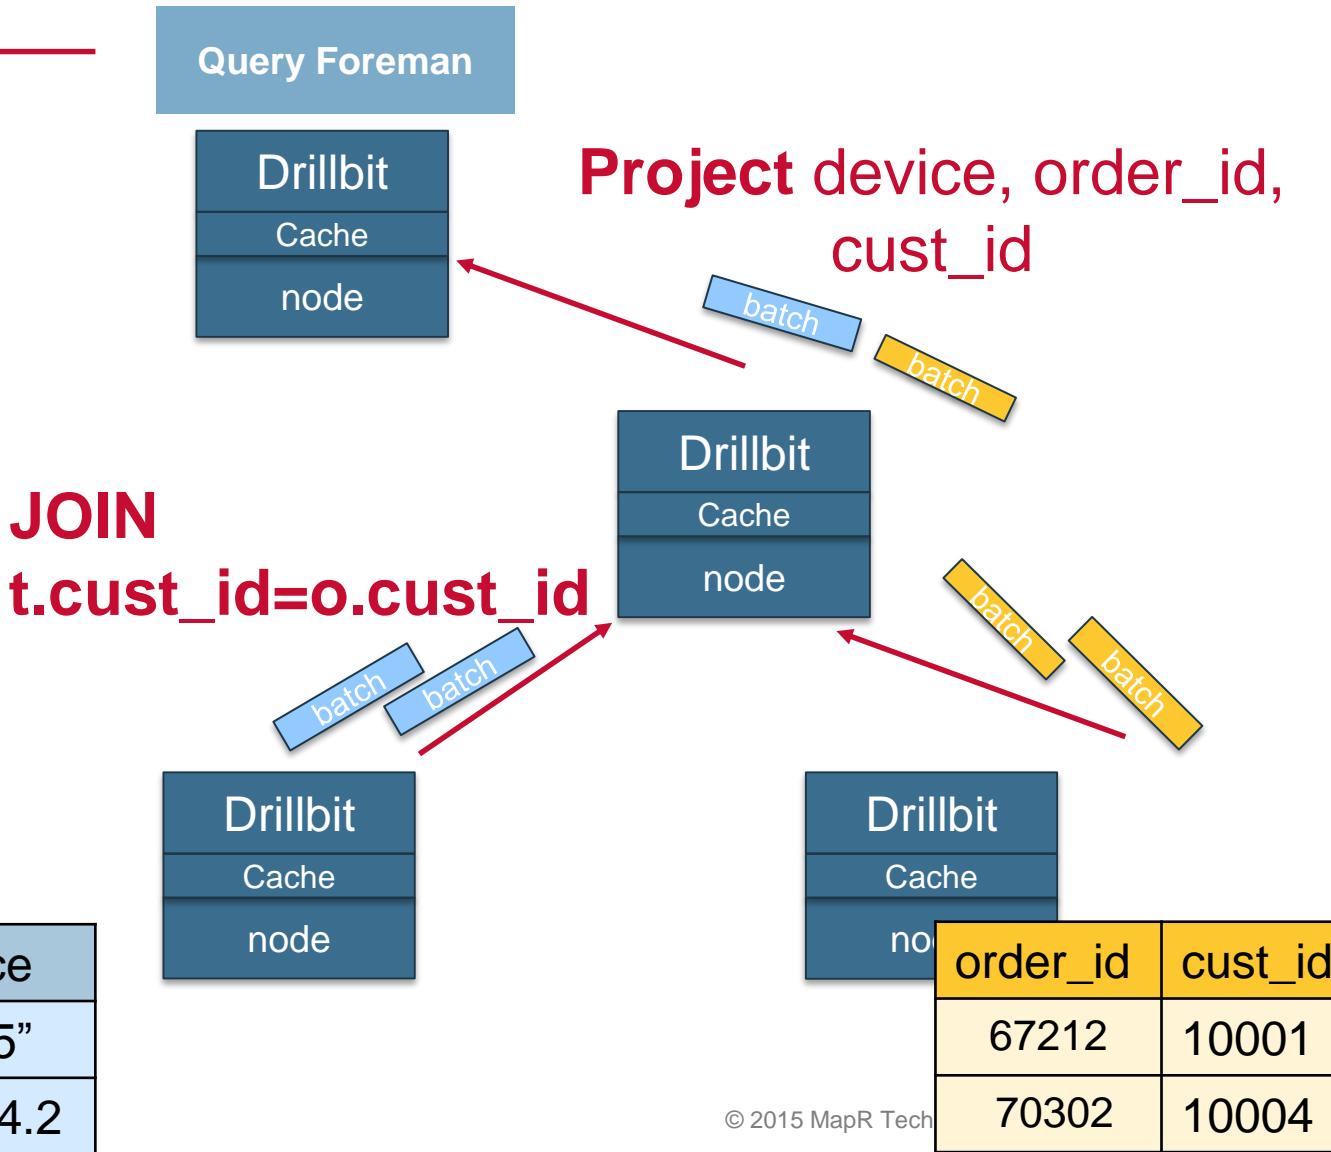


Data is returned, filtered

- Intermediate fragments **start** when they **receive** data
- **aggregates results** from leaf nodes and
 - send results up the tree

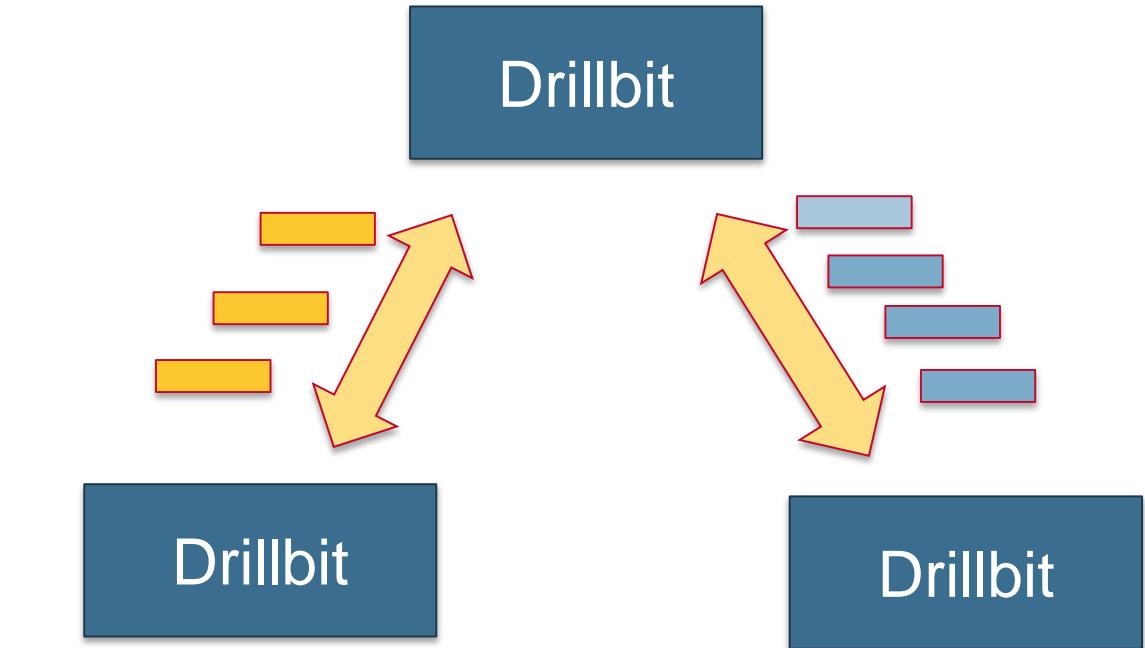
```
SELECT device, cust_id, order_id  
FROM clicks.json t, hive.orders o  
WHERE t.cust_id=o.cust_id ;
```

cust_id	device
10001	“iOS5”
10004	AOS4.2



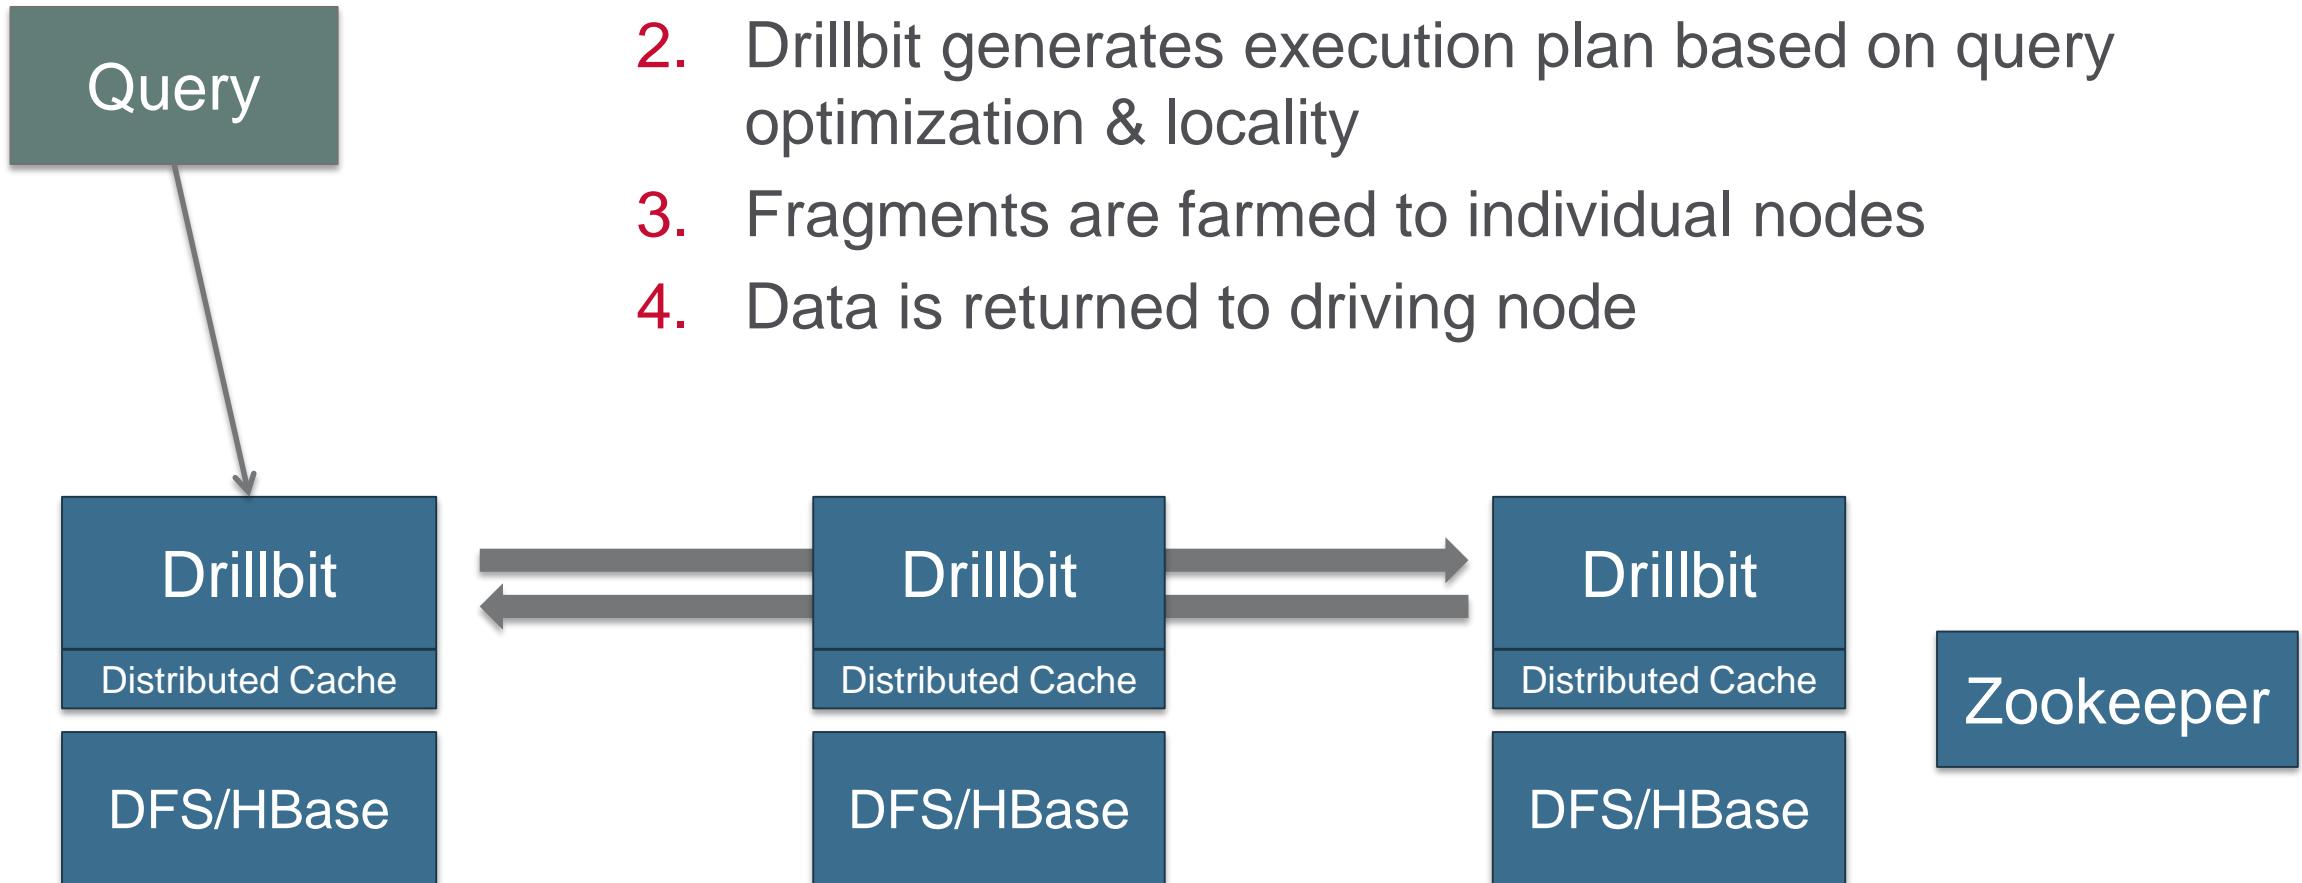
 Pipelining

- Record batches are pipelined between nodes
 - Operators works on a batch
- query execution happens in-memory
 - (as much as possible)





Drill query execution at a glance





Drill Query Profiles

- For each query, Drill generates a unique **query id**
- Profiles are saved under this identifier

The screenshot shows the Apache Drill web interface at `127.0.0.1:8047/profiles`. The top navigation bar includes links for Apache Drill, Query, Profiles, Storage, Metrics, and Threads. Below the navigation, a light blue box displays the message "No running queries.". Under the heading "Completed Queries", there is a table with two columns: "Time" and "Query". A single entry is listed: "04/10/2015 16:56:46" followed by a query string. A red oval has been drawn around the query string.

Time	Query
04/10/2015 16:56:46	<code>SELECT t.user_info.device device, t.user_info.cust_id cust_id, o.order_id FROM dfs.clicks.'clicks/clicks.json' t, hive.orders o WHERE t.user_info.</code>



WebUI: Query and Physical Plan

Query and Planning

Query Physical Plan Visualized Plan Edit Query

```
SELECT t.user_info.device device, t.user_info.cust_id cust_id, o.order_id FROM
```

Query Physical Plan Visualized Plan Edit Query

```
00-00      Screen: rowcount = 5097.0, cumulative cost = {35834.7 rows, 112730.0 memory}, id = 7735
00-01      Project(device=[$0], cust_id=[$1], order_id=[$2]): rowcount = 5097.0, cumulative cost = {35834.7 rows, 112730.0 memory}, id = 7735
00-02      Project(device=[ITEM($0, 'device')], cust_id=[ITEM($0, 'cust_id')], order_id=[ITEM($0, 'order_id'))]: rowcount = 5097.0, cumulative cost = {35834.7 rows, 112730.0 memory}, id = 7734
00-03      HashJoin(condition=[=(#1, #3)], joinType=[inner]): rowcount = 3280.0, cumulative cost = {35834.7 rows, 112730.0 memory}, id = 7733
00-04      Project(user_info=[#0], $f2=[CAST(ITEM($0, 'cust_id')):BIGINT]): rowcount = 3280.0, cumulative cost = {35834.7 rows, 112730.0 memory}, id = 7730
00-05      Scan(groupscan=[EasyGroupScan [selectionRoot=/mapr/demo/`orders`/`customer`/`order`]]): rowcount = 5097.0, cumulative cost = {5097.0 rows, 5097.0 memory}, id = 7729
00-06      Scan(groupscan=[HiveScan [table=Table(tableName:orders, partitionKeys:[])]]): rowcount = 5097.0, cumulative cost = {5097.0 rows, 5097.0 memory}, id = 7728
```

Major Fragment Id

Operator Id



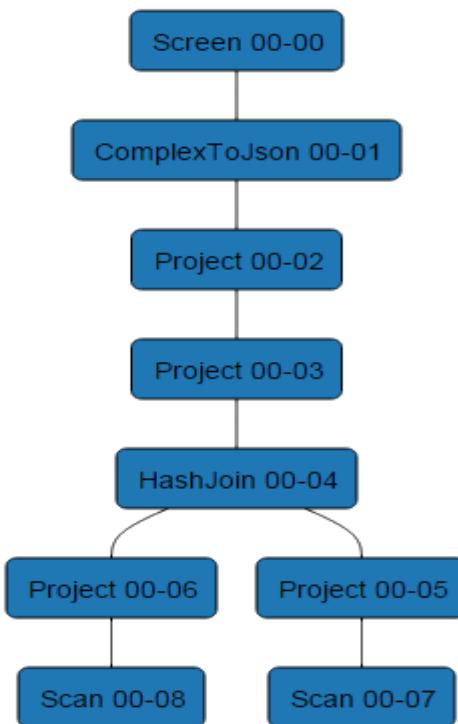
WebUI: Visualized Plan



Query and Planning

Query Physical Plan Visualized Plan

[Edit Query](#)



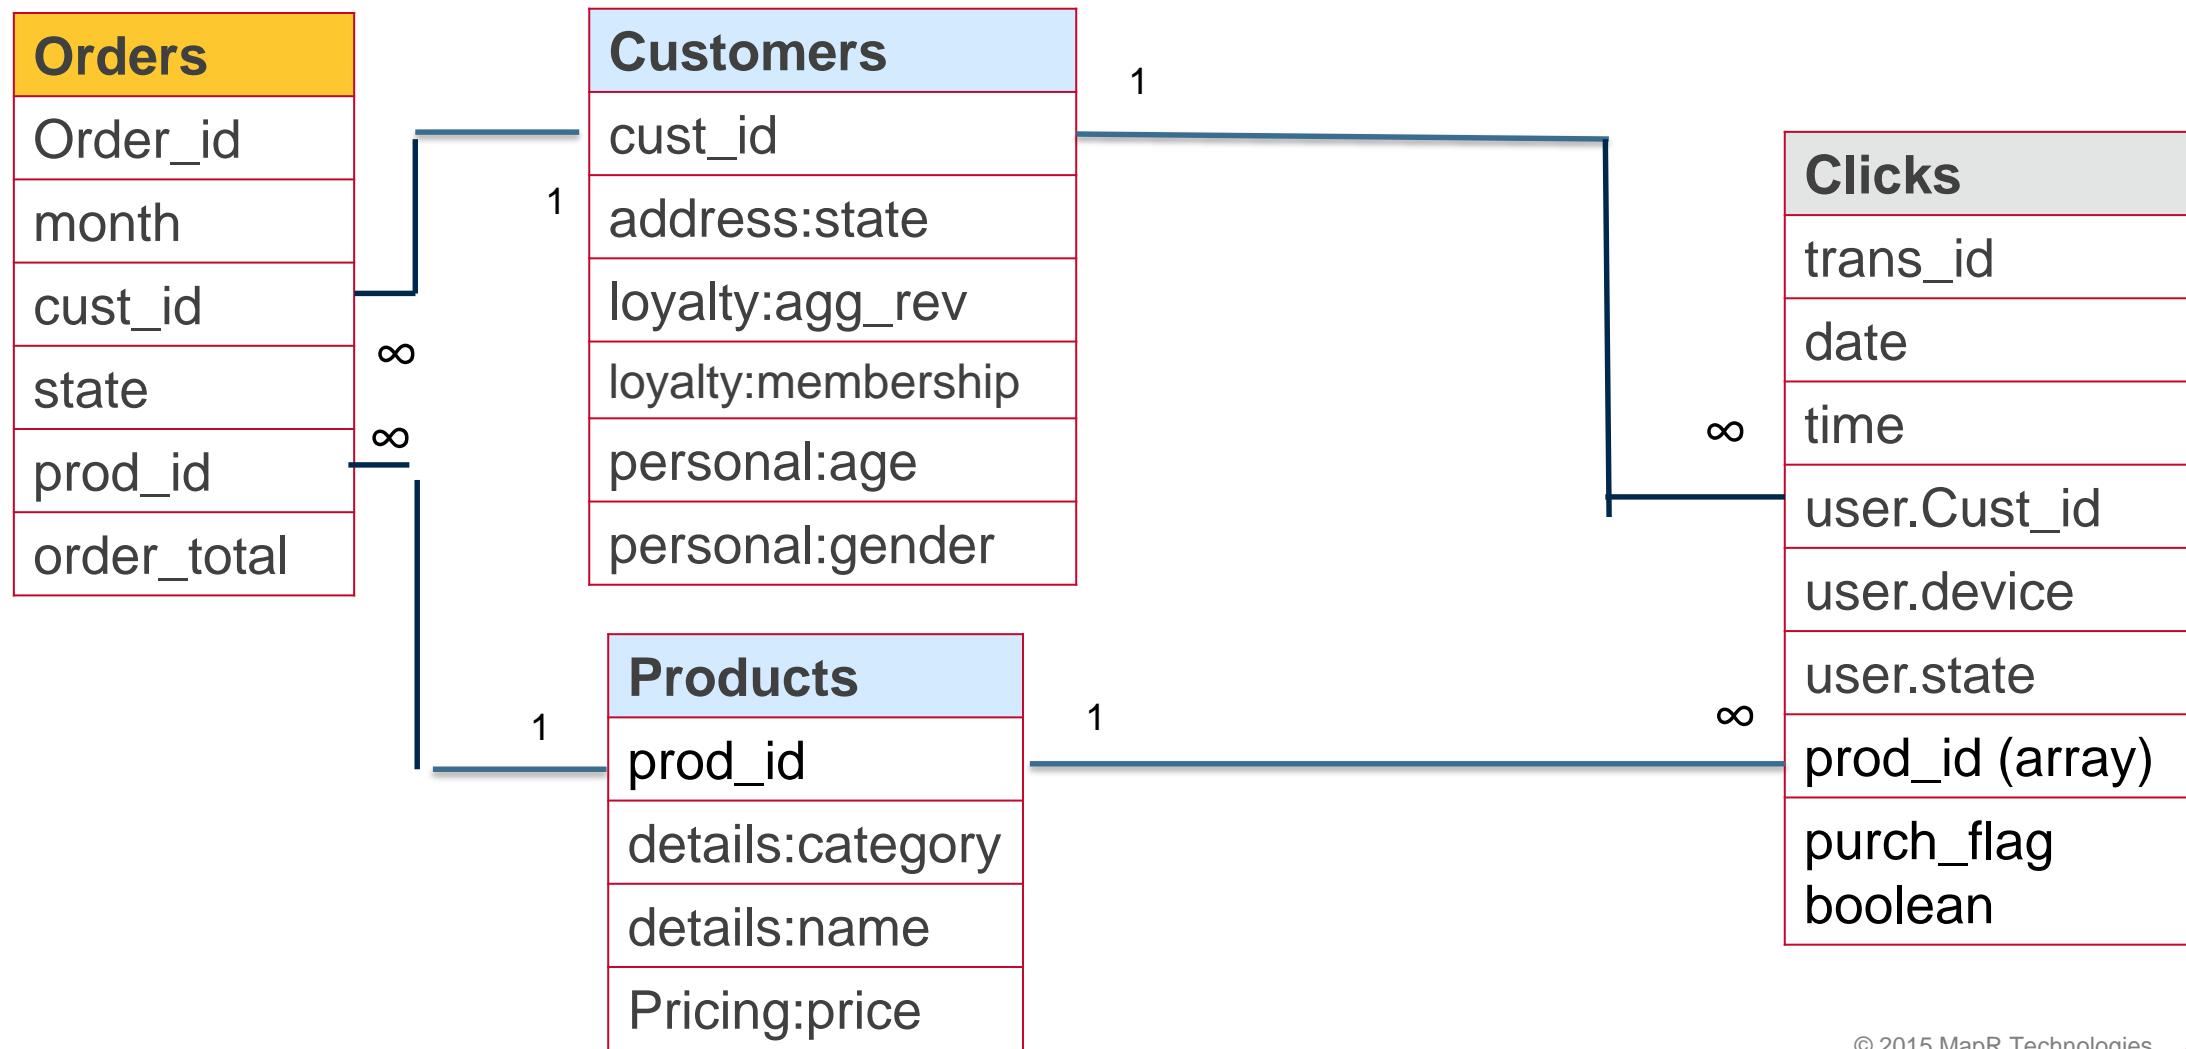


Combining and Aggregating Data from Files, Hive, and HBase using Drill and ANSI SQL

- ▶ GROUP BY
- ▶ UNION
- ▶ JOIN
- ▶ SUBQUERY
- ▶ CAST
- ▶ CREATING A VIEW
- ▶ CREATING A TABLE



High Level “Relations” between Retail data





SQL Basic Query Syntax

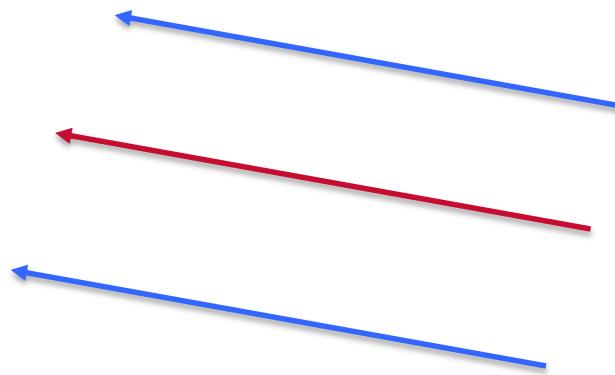
```
SELECT select_list
[ FROM table_source ]
[ WHERE search_condition ]
[ GROUP BY group_by_expression ]
[ HAVING search_condition ]
[ ORDER BY order_expression [ ASC | DESC ] ]
```

ANSI SQL with no modifications



Query Basics Select

```
SELECT cust_id  
FROM hive.`orders`  
WHERE state='ca'
```



Selects columns

where you want it from

Filters records to
return



Aggregating Data from Files, Hive, and HBase using Drill and ANSI SQL

- ▶ Aggregate Functions, GROUP BY

- UNION

- JOIN

- SUBQUERY

- CAST

- CREATING A VIEW

- CREATING A TABLE



Summarizing Data with Aggregate Functions, GROUP BY

Business Scenario

- What are the Sales total by month for California ?

SELECT GROUP BY WHERE STATE='CA' SUM

order_id	month	cust_id	state	Prod_id	Order_total
67212	June	10001	ca	909	13
70302	May	10004	ga	420	11



Query Basics Group by

```
SELECT `month`, SUM(order_total) AS sales  
FROM hive.`orders`  
WHERE state='ca'  
GROUP BY `month`  
ORDER BY 2 desc
```

Query from orders for **sales totals** by **month**

ALIAS

Total Number of Records:: 10

	month	sales
1	May	119586
2	June	116322
3	April	101363
4	March	99540
5	July	90285
6	October	80090
7	August	71255
8	February	63527
9	September	45552
10	January	38799



Query Basics: Aggregate Functions

```
avg(expression)
count(*)
count([DISTINCT] expression)
max(expression)
min(expression)
sum(expression)
```

Standard SQL clauses work in the same way in Drill queries as in relational database queries



Use Drill with ANSI SQL

- GROUP BY
- UNION
- ▶ Combining Data from Files, Hive, and HBase using Drill and ANSI SQL
- ▶ JOIN
- SUBQUERY
- CAST
- CREATING A VIEW
- CREATING A TABLE



Combining Data from multiple tables with JOIN

Business Scenario

- Which device do customers who made orders use ?

JOIN ON



order_id	cust_id
67212	10001
70302	10004

HIVE orders



user_info.cust_id	user_info.device
10001	“iOS5”
10004	AOS4.2

Web log clicks



Query Basics JOIN

```
SELECT t.user_info.device device, t.user_info.cust_id cust_id, o.order_id  
FROM dfs.clicks.`clicks/clicks.json` t, hive.`orders` o  
WHERE t.user_info.cust_id=o.cust_id
```

Query to get the device,
and order id from the click
log and hive orders, join
on customer id

Drill Explorer

Browse SQL

View Definition SQL:

```
SELECT t.user_info.device device
FROM dfs.clicks.`clicks/clicks.json` t
WHERE t.user_info.cust_id=o.cust_id
```

Total Number of Records:: 50

	device	cust_id	order_id
1	IOS5	22526	104718
2	IOS5	22526	21418
3	IOS5	22526	87760
4	IOS5	10269	80070



Cast function with HBase

Business Scenario

- Select an HBase column in a query

rowkey	CF: address		CF: loyalty		CF:personal		
	state	agg_rev	membership	age	gender	name	
1	"va"	197.00	"silver"	"15-20"	"FEMALE"	Brittany Park	



Selecting a Column with HBase

```
>SELECT t.personal.name FROM maprdb.customers t  
LIMIT 1;
```

```
+-----+  
| EXPR$0 |  
+-----+  
| [B@6780874d |  
+-----+
```

table
ALIAS

rowkey	CF: address		CF: loyalty		CF:personal		
	state	agg_rev	membership	age	gender	name	
1	"va"	197.00	"silver"	"15-20"	"FEMALE"	Brittany Park	



CONVERT_FROM function with HBase

Business Scenario

- Convert HBase bytes to datatypes in a query

```
+-----+  
| EXPR$0 |  
+-----+  
| [B@6780874d |  
+-----+
```

Data Preview:

	Row_Key	age	gender	name
▶ 1	10001	"15-20"	"FEMALE"	"Corrine Mecham"
2	10005	"26-35"	"MALE"	"Brittany Park"

rowkey	CF: address	CF: loyalty		CF:personal		
	state	agg_rev	membership	age	gender	name
1	"va"	197.00	"silver"	"15-20"	"FEMALE"	Brittany Park



Using convert_from with HBase

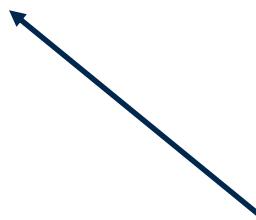
```
> SELECT convert_from(t.personal.name, 'UTF8') as pName  
> FROM customers t LIMIT 2;
```

pName
"Corrine Mecham"
"Brittany Park"

2 rows selected

Table
ALIAS

column
ALIAS





convert_from with HBase

```
use maprdb;
select convert_from(row_key , 'UTF8') as cust_id,
convert_from(t.personal.name , 'UTF8') as name,
convert_from(t.personal.gender , 'UTF8') as gender,
convert_from(t.personal.age , 'UTF8') as age,
convert_from(t.address.state , 'UTF8') as state,
convert_from(t.loyalty.agg_rev, 'UTF8') as agg_rev,
convert_from(t.loyalty.membership, 'UTF8') as membership
from customers t limit 5;
```

cust_id	name	gender	age	state	agg_rev
10001	"Corrine Mecham"	"FEMALE"	"15-20"	"va"	197
10005	"Brittany Park"	"MALE"	"26-35"	"in"	230
10006	"Rose Lokey"	"MALE"	"26-35"	"ca"	250
10007	"James Fowler"	"FEMALE"	"51-100"	"me"	263
10010	"Guillermo Koehler"	"OTHER"	"51-100"	"mn"	20



Combining and Aggregating Data from Files, Hive, and HBase using Drill and ANSI SQL

- GROUP BY
- UNION
- JOIN
- SUBQUERY
- CONVERT_FROM
- ▶ CREATING A VIEW
- CREATING A TABLE



Simplify future Queries with a View

Business Scenario

- Make the **Converted HBase data easily available in a view**

```
+-----+  
| EXPR$0 |  
+-----+  
| [B@6780874d |  
+-----+
```

View Definition SQL:

```
select convert_from(row_key , 'UTF8') as cust_id,  
convert_from(t.personal.name , 'UTF8') as name,  
convert_from(t.personal.gender , 'UTF8') as gender,  
convert_from(t.personal.age , 'UTF8') as age,
```

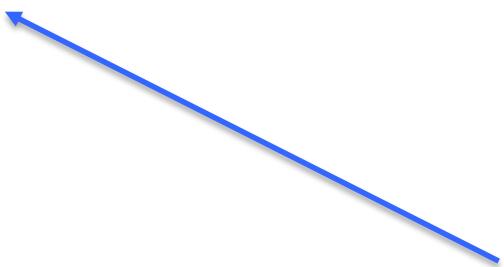
Total Number of Records::5

	cust_id	name	gender	age	state	agg_rev	membership
▶ 1	10001	"Corrine Mecham"	"FEMALE"	"15-20"	"va"	197	"silver"
2	10005	"Brittany Park"	"MALE"	"26-35"	"in"	230	"silver"



Mutable Workspace for Views and/or Tables

use **dfs.mydata**



Mutable workspace

```
"type": "file",
"enabled": true,
"connection": "maprfs:///",
"workspaces": {
  "mydata": {
    "location": "/mapr/data/views",
    "writable": true,
    ...
  },
}
```



Create a View

```
use dfs.mydata;
```

```
create or replace view custview as
select convert_from(row_key , 'UTF8') as cust_id,
convert_from(t.personal.name , 'UTF8') as name,
convert_from(t.personal.gender , 'UTF8') as gender,
convert_from(t.personal.age , 'UTF8') as age,
convert_from(t.address.state , 'UTF8') as state,
convert_from(t.loyalty.agg_rev, 'UTF8') as agg_rev,
convert_from(t.loyalty.membership, 'UTF8') as membership
from maprdb.customers t;
```



View created

Browse SQL

Schemas:

- cp.default
- dfs.clicks
- dfs.csv
- dfs.data
- dfs.default
- dfs.logs
- dfs.mydata
 - custview
 - cust_orders
 - student.parquet
 - voter.parquet
- dfs.root
- dfs.tmp
- hive.default
- maprdb

Metadata: dfs.mydata.custview

	COLUMN_NAME	DATA_TYPE	IS_NULLABLE
1	cust_id	ANY	NO
2	name	ANY	NO
3	gender	ANY	NO
4	age	ANY	NO
5	state	ANY	NO

Data Preview:

	cust_id	name	gender	age	state	agg_rev	members
1	10001	"Corrine Mecham"	"FEMALE"	"15-20"	"va"	197	"silver"
2	10005	"Brittany Park"	"MALE"	"26-35"	"in"	230	"silver"
3	10006	"Rose Lokey"	"MALE"	"26-35"	"ca"	250	"silver"
4	10007	"James Fowler"	"FFMALE"	"51-100"	"me"	263	"silver"



Drill View Creates metadata in the filesystem

A view is simply a **special file with a specific extension (.drill)**:

```
$ cat /mapr/demo.mapr.com/data/views/custview.view.drill
```

metadata which was created for this view:

```
{
  "name" : "custview",
  "sql" : "SELECT CAST(`row_key` AS INTEGER) AS `cust_id`, CAST(`t`.`personal`['name'] AS VARCHAR(20)) AS `name`,
CAST(`t`.`personal`['gender'] AS VARCHAR(10)) AS `gender`, CAST(`t`.`personal`['age'] AS VARCHAR(10)) AS `age`,
CAST(`t`.`address`['state'] AS VARCHAR(4)) AS `state`, CAST(`t`.`loyalty`['agg_rev'] AS DECIMAL(7, 2)) AS `agg_rev`,
CAST(`t`.`loyalty`['membership'] AS VARCHAR(20)) AS `membership`\nFROM `maprdb`.`customers` AS `t`",
  "fields" : [ {
    "name" : "cust_id",
    "type" : "INTEGER"
  }, {
    "name" : "name",
    "type" : "VARCHAR",
    "precision" : 20
  }, {
    "name" : "gender",
  ...
}
```



What is a Drill View?

- A Drill view is a JSON file
- Stored metadata for a query
- no actual data
 - Accesses latest data
- View Can be used :
 - to simplify complex queries
 - Like a table
 - with BI tools like Tableau



How to Use a View ?

```
use dfs.mydata;
select * from custview limit 1;
+-----+-----+-----+-----+-----+
| cust_id | name          | gender       | age        | state       | agg_rev    |
+-----+-----+-----+-----+-----+
| 10001   | "Corrine Mecham" | "FEMALE"    | "15-20"    | "va"        | 197.0      |
+-----+-----+-----+-----+-----+
1 row selected
```



Using a View

```
select  
custview.membership,  
custview.cust_id  
from custview
```

Browse SQL

View Definition SQL:

```
select custview.membership, custview.cust_id  
from custview
```

Total Number of Records:: 12792

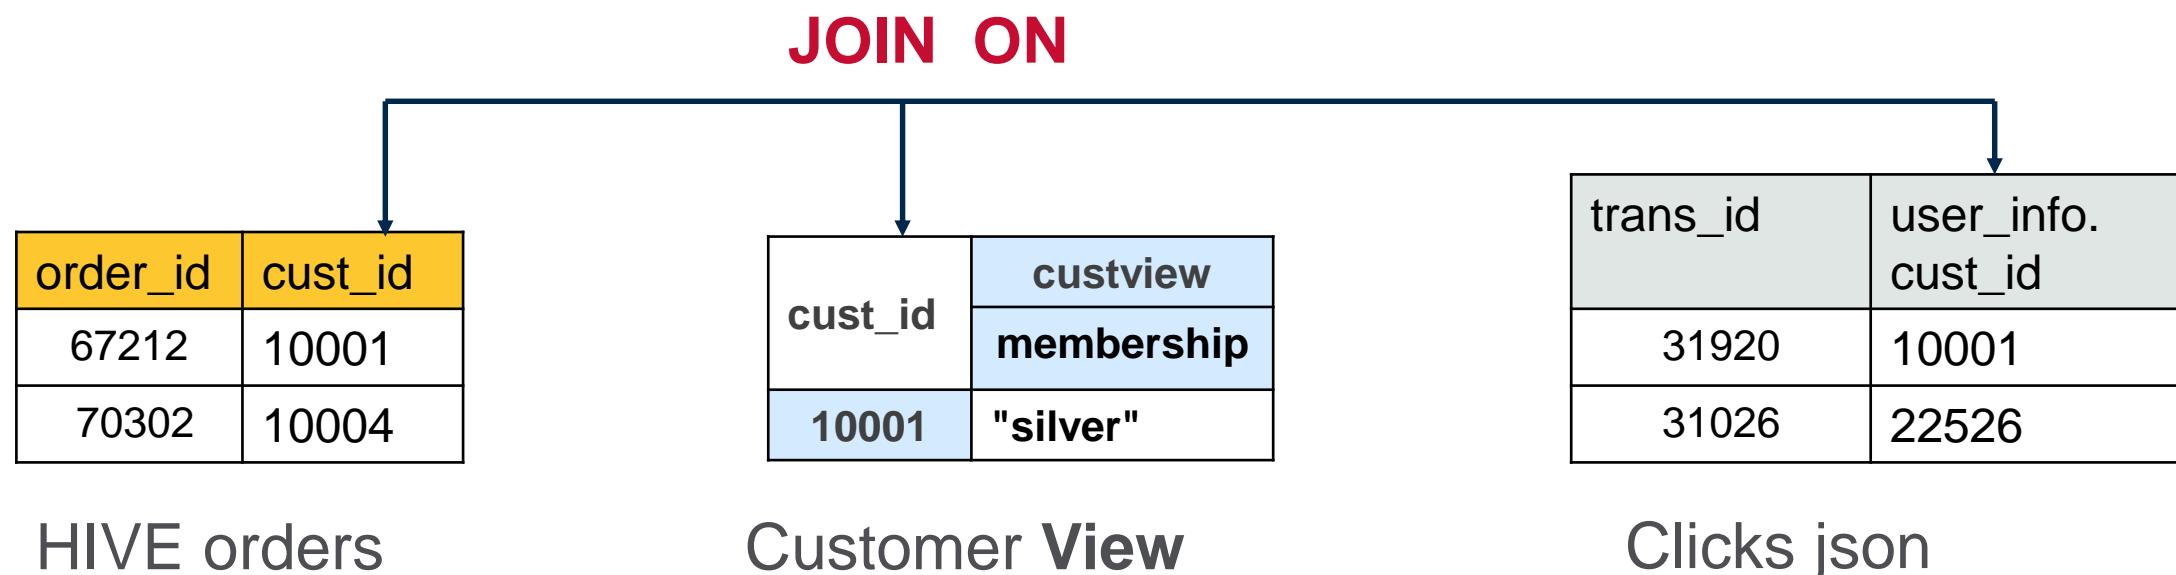
	membership	cust_id
1	"silver"	10001
2	"silver"	10005
3	"silver"	10006
4	"silver"	10007



How to Use a View with a JOIN

Business Scenario

- Show Customer Order information



 Joining with a View

Join the customers view and the orders table:

```
select custview.membership, sum(orders.order_total) as sales
from hive.orders, custview,
dfs.`/mapr/demo.mapr.com/data/nested/clicks/clicks.json` c
where orders.cust_id=custview.cust_id and
orders.cust_id=c.user_info.cust_id
group by custview.membership order by 2;
```

```
+-----+-----+
| membership |   sales   |
+-----+-----+
| "basic"    | 372866   |
| "silver"   | 728424   |
| "gold"     | 7050198  |
+-----+-----+
3 rows selected (11.374 seconds)
```



Combining and Aggregating Data from Files, Hive, and HBase using Drill and ANSI SQL

- GROUP BY
- UNION
- JOIN
- SUBQUERY
- CAST
- CREATING A VIEW
- ▶ CREATING A TABLE



Create a Table from multiple sources

Business Scenario

- Create a table with the Customer Order information from different data sources

CREATE TABLE

order_id	cust_id
67212	10001
70302	10004

HIVE orders

cust_id	custview
10001	"silver"

Customer View

trans_id	user_info.cust_id
31920	10001
31026	22526

Clicks json



Create a Table from the previous Join

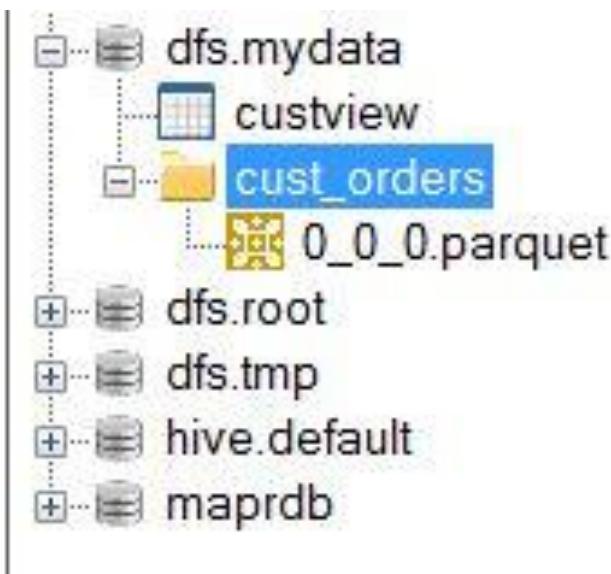
```
use dfs.mydata;

create table cust_orders as
select custview.membership, sum(orders.order_total) as sales
from hive.orders, custview,
dfs.`/mapr/demo.mapr.com/data/nested/clicks/clicks.json` c
where orders.cust_id=custview.cust_id
and orders.cust_id=c.user_info.cust_id
group by custview.membership order by 2;
```



Table created

- The “CREATE TABLE AS” command creates a directory that contains **parquet files** with the table data
 - Parquet is a columnar storage format



Data Preview:

	membership	sales
▶ 1	"basic"	372866
2	"silver"	728424
3	"gold"	7050198



Query the Table Created

Query the table created:

```
select * from cust_orders;
```

membership	sales
"basic"	372866
"silver"	728424
"gold"	7050198

3 rows selected (0.073 seconds)



Lab – Query HBase airline data with Drill

Import mapping to Row Key and Columns:

Row-key	delay			info			stats		timing	
Carrier-Flightnumber-Date-Origin-destination	Air Craft delay	Arr delay	Carrier delay	cncl	Cncl code	tailnum	distance	elaptime	arrtime	Dep time
AA-1-2014-01-01-JFK-LAX		13		0		N7704	2475	385.00	359	...

Drill Explorer

Browse SQL

Schemas:

- cp.default
- dfs.clicks
- dfs.default
- dfs.logs
- dfs.mydata
- dfs.root
- dfs.tmp
- hive.default
- maprdb
 - airline
 - delay
 - info
 - stats
 - timing
 - cust
 - customers
 - embeddedclicks
 - products
 - trades_flat
 - trades_tall

Metadata: `maprdb` `airline`

	Column_Name	Data_Type	Value_Width
1	aircraftdelay	Varchar(12)	12
2	arrdelay	Varchar(12)	12
3	carrierdelay	Varchar(12)	12
4	depdelay	Varchar(12)	12
5	nasdelay	Varchar(8)	8

Data Preview:

	Row_Key	aircraftdelay	arrdelay	carrierdelay	depdelay	nasdelay	securit
1	A-1-2014-01-01-JFK-LAX		13.00		14.00		
2	A-1-2014-01-02-JFK-LAX		1.00		-3.00		
3	A-1-2014-01-03-JFK-LAX						
4	A-1-2014-01-04-JFK-LAX	0.00	59.00	0.00	65.00	0.00	0.00
5	A-1-2014-01-05-JFK-LAX	0.00	110.00	0.00	110.00	0.00	0.00
6	A-1-2014-01-06-JFK-LAX		-8.00		17.00		
7	A-1-2014-01-07-JFK-LAX		-13.00		10.00		
8	A-1-2014-01-08-JFK-LAX		-10.00		23.00		
9	A-1-2014-01-09-JFK-LAX		-21.00		-1.00		
10	A-1-2014-01-10-JFK-LAX	0.00	20.00	0.00	29.00	0.00	0.00

Refresh Connect... Sample 100 Rows Close

v0.08.1.0618 (64 bit)



Drill Explorer

Browse SQL

View Definition SQL:

```
SELECT * FROM maprdb.airline LIMIT 10
```

Preview Create As...

Total Number of Records:: 10

	row_key	delay
1	0x41412D312D323031342D30312D30312D4A464B2D4C4158	{ "aircraftdelay" : "", "arrdelay" : "MTMuMDA=", "carrierdelay" : "", "depdelay" : "MTQuM
2	0x41412D312D323031342D30312D30322D4A464B2D4C4158	{ "aircraftdelay" : "", "arrdelay" : "MS4wMA==", "carrierdelay" : "", "depdelay" : "LTMuN
3	0x41412D312D323031342D30312D30332D4A464B2D4C4158	{ "aircraftdelay" : "", "arrdelay" : "", "carrierdelay" : "", "depdelay" : "", "nasdelay" : ""},
4	0x41412D312D323031342D30312D30342D4A464B2D4C4158	{ "aircraftdelay" : "MC4wMA==", "arrdelay" : "NTkuMDA=", "carrierdelay" : "MC4wMA=
5	0x41412D312D323031342D30312D30352D4A464B2D4C4158	{ "aircraftdelay" : "MC4wMA==", "arrdelay" : "MTEwLjAw", "carrierdelay" : "MC4wMA=
6	0x41412D312D323031342D30312D30362D4A464B2D4C4158	{ "aircraftdelay" : "", "arrdelay" : "LTguMDA=", "carrierdelay" : "", "depdelay" : "MTcuM
7	0x41412D312D323031342D30312D30372D4A464B2D4C4158	{ "aircraftdelay" : "", "arrdelay" : "LTEzLjAw", "carrierdelay" : "", "depdelay" : "MTAuM
8	0x41412D312D323031342D30312D30382D4A464B2D4C4158	{ "aircraftdelay" : "", "arrdelay" : "LTEwLjAw", "carrierdelay" : "", "depdelay" : "MjMuMD
9	0x41412D312D323031342D30312D30392D4A464B2D4C4158	{ "aircraftdelay" : "", "arrdelay" : "LTlxLjAw", "carrierdelay" : "", "depdelay" : "LTEuMD
10	0x41412D312D323031342D30312D31302D4A464B2D4C4158	{ "aircraftdelay" : "MC4wMA==", "arrdelay" : "MjAuMDA=", "carrierdelay" : "MC4wMA=

v0.08.1.0618 (64 bit)



Drill Explorer

Browse SQL

View Definition SQL:

```
SELECT CAST(row_key as VARCHAR(255)) Row_Key, CAST(`delay`['aircraftdelay'] AS Varchar(12))  
`aircraftdelay`, CAST(`delay`['arrdelay'] AS Varchar(12)) `arrdelay`, CAST(`delay`  
['carrierdelay'] AS Varchar(12)) `carrierdelay`, CAST(`delay`['depdelay'] AS Varchar(12))  
`depdelay`, CAST(`delay`['nasdelay'] AS Varchar(8)) `nasdelay`, CAST(`delay`['securitydelay']
```

Preview

Create As...

Total Number of Records: 100

	Row_Key	aircraftdelay	arrdelay	carrierdelay	depdelay	nasdelay	securitydelay	weatherdelay
▶ 1	AA-1-2014-01-01-JFK-LAX		13.00		14.00			
2	AA-1-2014-01-02-JFK-LAX		1.00		-3.00			
3	AA-1-2014-01-03-JFK-LAX							
4	AA-1-2014-01-04-JFK-LAX	0.00	59.00	0.00	65.00	0.00	0.00	59.00
5	AA-1-2014-01-05-JFK-LAX	0.00	110.00	0.00	110.00	0.00	0.00	110.00
6	AA-1-2014-01-06-JFK-LAX		-8.00		17.00			
7	AA-1-2014-01-07-JFK-LAX		-13.00		10.00			
8	AA-1-2014-01-08-JFK-LAX		-10.00		23.00			
9	AA-1-2014-01-09-JFK-LAX		-21.00		-1.00			
10	AA-1-2014-01-10-JFK-LAX	0.00	20.00	0.00	29.00	0.00	0.00	20.00
11	AA-1-2014-01-11-JFK-LAX	0.00	79.00	0.00	15.00	79.00	0.00	0.00
12	AA-1-2014-01-12-JFK-LAX		-17.00		-6.00			
13	AA-1-2014-01-13-JFK-LAX		-33.00		-5.00			
14	AA-1-2014-01-14-JFK-LAX		-41.00		-7.00			

v0.08.1.0618 (64 bit)



Create a View

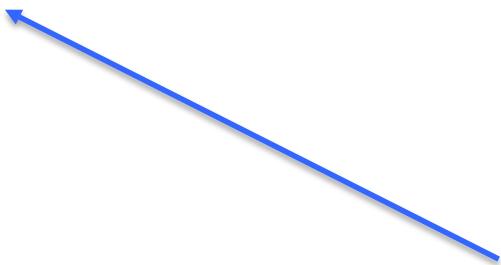
```
use dfs.mydata;
```

```
create or replace view airdelayview as  
SELECT convert_from(row_key , 'UTF8') as rowkey,  
convert_from(t.delay.aircraftdelay,'UTF8') as aircraftdelay,  
convert_from(t.delay.arrdelay,'UTF8') AS arrdelay,  
convert_from(t.delay.carrierdelay,'UTF8') AS carrierdelay,  
convert_from(t.delay.depdelay,'UTF8') AS depdelay,  
convert_from(t.delay.nasdelay,'UTF8') AS nasdelay,  
convert_from(t.delay.securitydelay,'UTF8') AS securitydelay,  
convert_from(t.delay.weatherdelay,'UTF8') AS weatherdelay  
FROM maprdb.airline t;
```



Mutable Workspace for Views and/or Tables

use **dfs.mydata**



Mutable workspace

```
"type": "file",
"enabled": true,
"connection": "maprfs:///",
"workspaces": {
  "mydata": {
    "location": "/mapr/data/views",
    "writable": true,
    ...
  },
}
```



Drill Explorer

Browse SQL

Schemas:

- cp.default
- dfs.clicks
- dfs.default
- dfs.logs
- dfs.mydata
 - airdelayview
 - airinfoview
 - airstatsview
 - airtimingview
 - custview
 - tradesview
- dfs.root
- dfs.tmp
- hive.default
- maprdb
 - airline
 - delay
 - info
 - stats
 - timing
 - cust
 - customers
 - embeddedclicks
 - products
 - trades.flat

Metadata: dfs . mydata . airdelayview

	COLUMN_NAME	DATA_TYPE	IS_NULLABLE
1	rowkey	ANY	YES
2	aircraftdelay	ANY	YES
3	arrdelay	ANY	YES
4	carrierdelay	ANY	YES
5	depdelay	ANY	YES

Data Preview:

	rowkey	aircraftdelay	arrdelay	carrierdelay	depdelay	nasde
1	AA-1-2014-01-01-JFK-LAX		13.00		14.00	
2	AA-1-2014-01-02-JFK-LAX		1.00		-3.00	
3	AA-1-2014-01-03-JFK-LAX					
4	AA-1-2014-01-04-JFK-LAX	0.00	59.00	0.00	65.00	0.00
5	AA-1-2014-01-05-JFK-LAX	0.00	110.00	0.00	110.00	0.00
6	AA-1-2014-01-06-JFK-LAX		-8.00		17.00	
7	AA-1-2014-01-07-JFK-LAX		-13.00		10.00	
8	AA-1-2014-01-08-JFK-LAX		-10.00		23.00	
9	AA-1-2014-01-09-JFK-LAX		-21.00		-1.00	
10	AA-1-2014-01-10-JFK-LAX	0.00	20.00	0.00	29.00	0.00

Refresh Connect... Sample 100 Rows Close

v0.08.1.0618 (64 bit)



Count number of cancellations by reason (code)

```
select count(*) as cancellations, cnclcode from  
dfs.mydata.airinfoview where cncl=1 group by cnclcode order  
by cancellations asc limit 100;
```

The screenshot shows the Drill Explorer interface with the SQL tab selected. The query window contains the provided SQL code. To the right, there are 'Preview' and 'Create As...' buttons. Below the query window, the text 'Total Number of Records:: 3' is displayed. A table below shows the results of the query:

	cancellations	cnclcode
1	4598	C
2	7146	A
3	19108	B



Find the longest airline delays

```
select rowkey, arrdelay, carrierdelay from dfs.mydata.airdelayview  
order by arrdelay desc limit 20;
```

Browse SQL

View Definition SQL:

```
select rowkey, arrdelay, carrierdelay from dfs.mydata.airdelayview  
order by arrdelay desc limit 20
```

Total Number of Records:: 20

	rowkey	arrdelay	carrierdelay
1	DL-2139-2014-01-02-LGA-MIA	996.00	0.00
2	F9-565-2014-01-20-MOT-DEN	992.00	992.00
3	AS-21-2014-01-04-ORD-SEA	99.00	0.00
4	AA-1199-2014-01-25-TUL-MIA	99.00	99.00
5	AA-2418-2014-01-14-DFW-DTW	99.00	45.00

↓ row



- ▶ Complex Data Types
 - ▶ Map, Array
- Query JSON file with Map Data type
- Query JSON file with Array Data type
- Join Hive and JSON file with complex data types
- Create a table from Join Hive and JSON file with complex data types
- Query Directories



Complex Data Types

- MAP
 - set of name/value pairs
- ARRAY
 - a repeated list of values

"user_info": {"cust_id":22526,"device":"IOS5","state":"il"}

"prod_id": [174,2]



	trans_id	date	time	user_info	trans_info
► 1	31920	2014-04-26	12:17:12	{ "cust_id": 22526, "device": "IOS5", "state": "il"}	{ "prod_id": [174, 2], "purch_flag": "false" }
2	31026	2014-04-20	13:50:29	{ "cust_id": 16368, "device": "AOS4.2", "state": "nc"}	{ "prod_id": [], "purch_flag": "false" }
3	33848	2014-04-10	04:44:42	{ "cust_id": 21449, "device": "IOS6", "state": "oh"}	{ "prod_id": [582], "purch_flag": "false" }
4	32383	2014-04-18	06:27:47	{ "cust_id": 20323, "device": "IOS5", "state": "oh"}	{ "prod_id": [710, 47], "purch_flag": "false" }
5	32359	2014-04-19	23:13:25	{ "cust_id": 15360, "device": "IOS5", "state": "ca"}	{ "prod_id": [0, 8, 170, 173, 1, 124, 46, 764, 30, 76] }



Selecting a column from a JSON Map Data Type

Business Scenario

- Select JSON Map “columns” cust_id, device... in a query

MAP

	trans_id	date	time	user_info	trans_info
► 1	31920	2014-04-26	12:17:12	{ "cust_id": 22526, "device": "IOS5", "state": "il"}	{ "prod_id": [174.2], "purch_flag": "false" }
2	31026	2014-04-20	13:50:29	{ "cust_id": 16368, "device": "AOS4.2", "state": "nc"}	{ "prod_id": [], "purch_flag": "false" }
3	33848	2014-04-10	04:44:42	{ "cust_id": 21449, "device": "IOS6", "state": "oh"}	{ "prod_id": [582], "purch_flag": "false" }
4	32383	2014-04-18	06:27:47	{ "cust_id": 20323, "device": "IOS5", "state": "oh"}	{ "prod_id": [710.47], "purch_flag": "false" }
5	32359	2014-04-19	23:13:25	{ "cust_id": 15360, "device": "IOS5", "state": "ca"}	{ "prod_id": [0.8, 170, 173, 1, 124, 46, 764, 30,] }



Results of Query JSON with Complex Data Types

table.column.column notation :

```
select t.user_info.cust_id  
as custid,  
t.user_info.device  
as device,  
t.user_info.state  
as state  
from `clicks/clicks.json` t  
limit 5;
```

Browse SQL

View Definition SQL:

```
select t.user_info.cust_id as custid,  
t.user_info.device as device,  
t.user_info.state as state  
from `clicks/clicks.json` t limit 5;
```

Total Number of Records:: 5

	custid	device	state
▶ 1	22526	IOS5	il
2	16368	AOS4.2	nc
3	21449	IOS6	oh
4	20323	IOS5	oh
5	15360	IOS5	ca



Result of Query with Map Data Type

table.column.column notation :

```
select
t.trans_info.prod_id
as prodid,
t.trans_info.purch_flag
as purchased
from
`clicks/clicks.json` t
limit 5;
```

Browse SQL

View Definition SQL:

```
select t.trans_info.prod_id as prodid,
t.trans_info.purch_flag as purchased
from `clicks/clicks.json` t limit 5;
```

Total Number of Records: 5

ARRAY

	prodid	purchased
1	[174, 2]	false
2	[]	false
3	[582]	false
4	[710, 47]	false
5	[0, 8, 170, 173, 1, 124, 46, 764, 30, 711, 0, 3, 25]	true



Selecting a column from a JSON Map Data Type

Business Scenario

- Select a Value in a JSON Array in a query
- **t.trans_info.prod_id[0]**

ARRAY

trans_info
{ "prod_id": [174, 2], "purch_flag": "false" }
{ "prod_id": [], "purch_flag": "false" }
{ "prod_id": [582], "purch_flag": "false" }
{ "prod_id": [710, 47], "purch_flag": "false" }
{ "prod_id": [0, 8, 170, 173, 1, 124, 46, 764, 30] }



Query JSON with Arrays

Get first product Id in prod_id array

```
select t.user_info.cust_id as cust_id,  
t.trans_info.prod_id[0] as prod_id  
from `clicks/clicks.json` t;
```

cust_id	prod_id
2587	174
1772	
1429	582



Query JSON with Arrays

Business Scenario

- **Which customers searched on at least 10 products?**

trans_id	date	user_info. cust_id	user_info. device	user_info. state	trans_info. prod_id	trans_info. purch_flag
31920	2014-04-26	10001	“iOS5”	“ca”	[174,2]	false
31026	2014-04-26	22526	AOS4.2	“nc”	[420]	true

`t.trans_info.prod_id[10]`
is not null

clicks.json file



Query JSON with Arrays

Which customers searched on at least 10 products?

```
select t.user_info.cust_id as cust_id,  
t.trans_info.prod_id[10] as prod_id  
from `clicks/clicks.json` t  
where t.trans_info.prod_id[10] is not null order by trans_id;
```

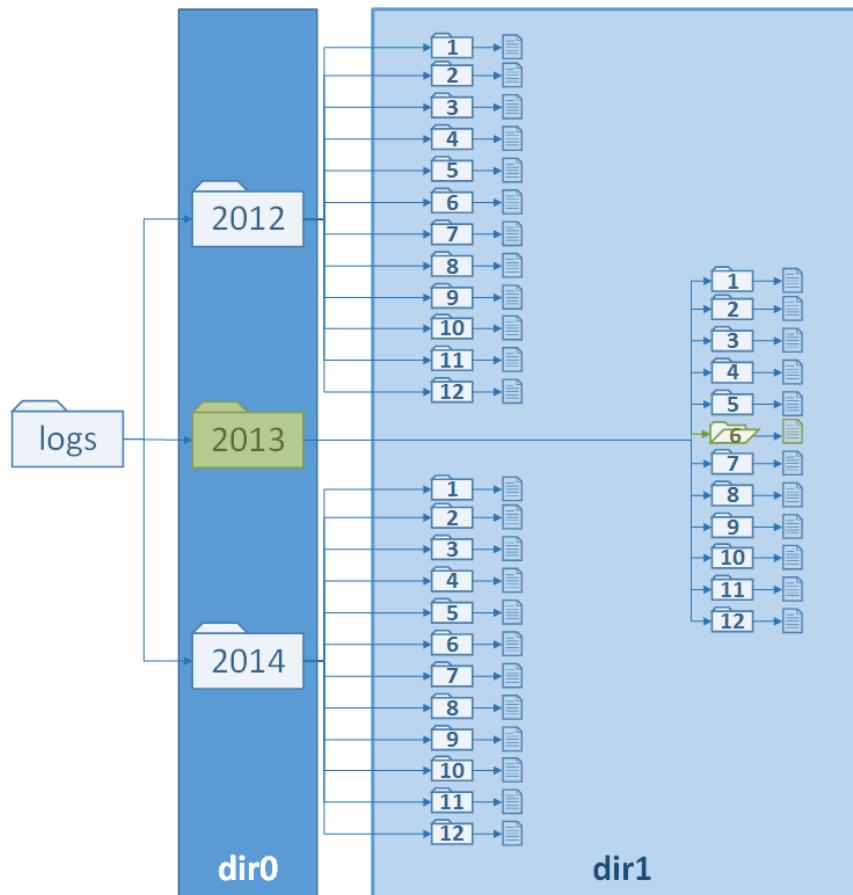
cust_id	prod_id
2587	156
1772	62
1429	18
4344	29
1818	0



Directories Partitions

- When Querying Files

- partitions are simply nested subdirectories of the table directory.
`drill> select * from logs where dir0=2013 and dir1=6;`





Query Partitioned directories

Business Scenario

- query files in directories and subdirectories in a single SELECT statement

Schemas:

- + cp.default
- + dfs.clicks
- + dfs.data
- + dfa.default
- + dfs.logs
 - + logs
 - + 2012
 - + 2013
 - + 2014
 - + 1
 - + log.json
 - + 2
 - + 3
 - + 4

Metadata:

	COLUMN_NAME	DATA_TYPE	IS_NULLABLE
▶ 1	Metadata does not exist		

Year subdirectory
Month subdirectory

Data Preview:

	trans_id	date	time	cust_id	device	state	camp_id	keyword
▶ 1	24101	01/13/2014	03:04:11	100	iOS5	ca	1	knew
2	24110	01/05/2014	22:25:04	101	iOS5	mn	6	look
3	24120	01/06/2014	00:32:23	1495	AOS4.4	mi	17	explores



Query directory logs

Business Scenario

- Query file data from logs top directory

Browse SQL

View Definition SQL:

```
select * from logs limit 10;
```

Total Number of Records:: 10

Year subdirectory
Month subdirectory

	dir0	dir1	trans_id	date	time	cust_id	device	state	camp_id	keywords	prod_id
► 1	2014	8	24181	08/02/2014	09:23:52	0	IOS5	il	2	wait	128
2	2014	8	24195	08/02/2014	07:58:19	243	IOS5	mo	6	hmm	107
3	2014	8	24204	08/01/2014	12:10:27	12048	IOS6	il	1	marge	324
4	2014	8	24222	08/02/2014	16:28:37	2488	IOS6	pa	2	to	391
5	2014	8	24227	08/02/2014	07:14:00	154687	IOS5	wa	2	on	376
6	2014	8	24242	08/02/2014	08:10:50	1	IOS5	il	1	in	348



Query subdirectory

Business Scenario

- Query log file data for the year 2014

Year subdirectory

	dir0	dir1	trans_id	date	time	cust_id	device	state	camp_id	keywords	prod_id
► 1	2014	8	24181	08/02/2014	09:23:52	0	IOS5	il	2	wait	128
2	2014	8	24195	08/02/2014	07:58:19	243	IOS5	mo	6	hmm	107
3	2014	8	24204	08/01/2014	12:10:27	12048	IOS6	il	1	marge	324
4	2014	8	24222	08/02/2014	16:28:37	2488	IOS6	pa	2	to	391
5	2014	8	24227	08/02/2014	07:14:00	154687	IOS5	wa	2	on	376
6	2014	8	24249	08/03/2014	23:18:59	1	IOS5	nv	1	hey	348
7	2014	8	24261	08/04/2014	23:32:16	7389	IOS5	mn	7	kapur	65
8	2014	8	24287	08/02/2014	18:21:26	2	IOS5	tx	13	church	283
9	2014	8	24323	08/04/2014	03:22:13	447	AOS4.4	in	9	here	496
10	2014	8	24378	08/01/2014	14:46:42	14006	IOS6	il	8	cecil	162



Query Selecting a subdirectory

Business Scenario

- Query log file data for the year 2014

```
> select * from logs where dir0='2014' ;
```

dir0	dir1	trans_id	date	time	cust_id	
2014	8	24181	08/02/2014	09:23:52	0	
2014	8	24195	08/02/2014	07:58:19	243	

▪ ▪ ▪



Query subdirectory

Business Scenario

- Query log file data for the month of August

Month subdirectory

	dir0	dir1	trans_id	date	time	cust_id	device	state	camp_id	keywords	prod_id
► 1	2014	8	24181	08/02/2014	09:23:52	0	IOS5	il	2	wait	128
2	2014	8	24195	08/02/2014	07:58:19	243	IOS5	mo	6	hmm	107
3	2014	8	24204	08/01/2014	12:10:27	12048	IOS6	il	1	marge	324
4	2014	8	24222	08/02/2014	16:28:37	2488	IOS6	pa	2	to	391
5	2014	8	24227	08/02/2014	07:14:00	154687	IOS5	wa	2	on	376
6	2014	8	24249	08/03/2014	23:18:59	1	IOS5	nv	1	hey	348
7	2014	8	24261	08/04/2014	23:32:16	7389	IOS5	mn	7	kapur	65
8	2014	8	24287	08/02/2014	18:21:26	2	IOS5	tx	13	church	283
9	2014	8	24323	08/04/2014	03:22:13	447	AOS4.4	in	9	here	496
10	2014	8	24378	08/01/2014	14:46:42	14006	IOS6	il	8	cecil	162



Query Selecting subdirectories

Business Scenario

- Query log file data for the month of August

```
> select * from logs where dir1='8' ;
```

dir0	dir1	trans_id	date	time	cust_id	
2014	8	24181	08/02/2014	09:23:52	0	
2014	8	24195	08/02/2014	07:58:19	243	
2014	8	24204	08/01/2014	12:10:27	12048	



Aggregate counts with subdirectories

Business Scenario

- **What is the count of customer clicks for August 2014 ?**

Year= 2014 Month=8

Count distinct

	dir0	dir1	trans_id	date	time	cust_id	device	state	camp_id	keywords	prod_id
► 1	2014	8	24181	08/02/2014	09:23:52	0	IOS5	il	2	wait	128
2	2014	8	24195	08/02/2014	07:58:19	243	IOS5	mo	6	hmm	107
3	2014	8	24204	08/01/2014	12:10:27	12048	IOS6	il	1	marge	324
4	2014	8	24222	08/02/2014	16:28:37	2488	IOS6	pa	2	to	391
5	2014	8	24227	08/02/2014	07:14:00	154687	IOS5	wa	2	on	376
6	2014	8	24249	08/03/2014	23:18:59	1	IOS5	nv	1	hey	348
7	2014	8	24261	08/04/2014	23:32:16	7389	IOS5	mn	7	kapur	65
8	2014	8	24287	08/02/2014	18:21:26	2	IOS5	tx	13	church	283
9	2014	8	24323	08/04/2014	03:22:13	447	AOS4.4	in	9	here	496
10	2014	8	24378	08/01/2014	14:46:42	14006	IOS6	il	8	cecil	162



Aggregate counts with subdirectories

Business Scenario

- What is the count of customer clicks for August 2014 ?

```
select count(distinct cust_id)
from logs
where dir0=2014 and dir1=8;
```

```
+-----+
|   EXPR$0   |
+-----+
| 167       |
+-----+
1 row selected
```



Free Hadoop On Demand Training

- <https://www.mapr.com/services/mapr-academy/big-data-hadoop-online-training>

DEV 320 - HBase Data Model and Architecture

[Register](#)

This course is intended for data analysts, data architects and application developers. DEV 320 provides you with a thorough understanding of the HBase data model and architecture, which is required before going on to designing HBase schemas and developing HBase applications.

[Learn More](#)

DEV 325 - HBase Schema Design

[Register](#)

NEW!
Targeted towards data analysts, data architects and application developers, the goal of this course is to enable you to design HBase schemas based on design guidelines. You will learn about the various elements of schema design and how to design for data access patterns. The course offers an in-depth look at designing row keys, avoiding hot-spotting and designing column families. It discusses how to transition from a relational model to an HBase model. You will learn the differences between tall tables and wide tables. Concepts are conveyed through lectures, hands-on labs and analysis of scenarios.

[Learn More](#)

DA 410 - Drill Essentials

[Register](#)

NEW!
This introductory Drill course, targeted at Data Analysts, Scientists and SQL programmers, covers how to use Drill to explore known or unknown data without writing code. You will write SQL queries on a variety of data types including structured data in a Hive table, semi-structured data in HBase or MapR-DB, and complex data file types, such as Parquet and JSON.



More Info

- <https://www.mapr.com/blog/how-use-sql-hadoop-drill-rest-json-nosql-and-hbase-simple-rest-client#.VT0AUCFViko>
- <https://cwiki.apache.org/confluence/display/DRILL/Apache+Drill+Tutorial>
- <https://www.dropbox.com/home/bigdatatechcon>



Next Steps



Q&A

Engage with us!

@mapr



maprtech

mapr-technologies



MapR

yourname@mapr.com



maprtech