



Real Time Data Analytics @ Uber

Ankur Bansal

Apache Big Data Europe

November 14, 2016

About Me

- Sr. Software Engineer, Streaming Team @ Uber
 - Streaming team supports platform for real time data analytics: Kafka, Samza, Flink, Pinot.. and plenty more
 - Focused on scaling Kafka at Uber's pace
- Staff software Engineer @ Ebay
 - Build & scale Ebay's cloud using openstack
- Apache Kylin: Committer, Emeritus PMC

Agenda

- Real time Use Cases
- Kafka Infrastructure Deep Dive
- Our own Development:
 - Rest Proxy & Clients
 - Local Agent
 - uReplicator (Mirrormaker)
 - Chaperone (Auditing)
- Operations/Tooling

Important Use Cases

Real-time Price Surging



Rider eyeballs



Open car information

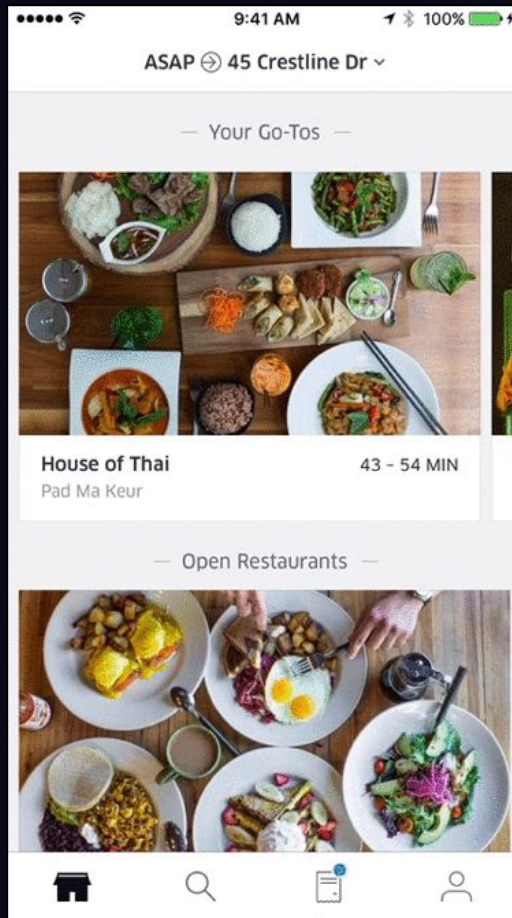
KAFKA

**Stream
Processing**



**SURGE
MULTIPLIERS**

Real-time Machine Learning - UberEats ETD

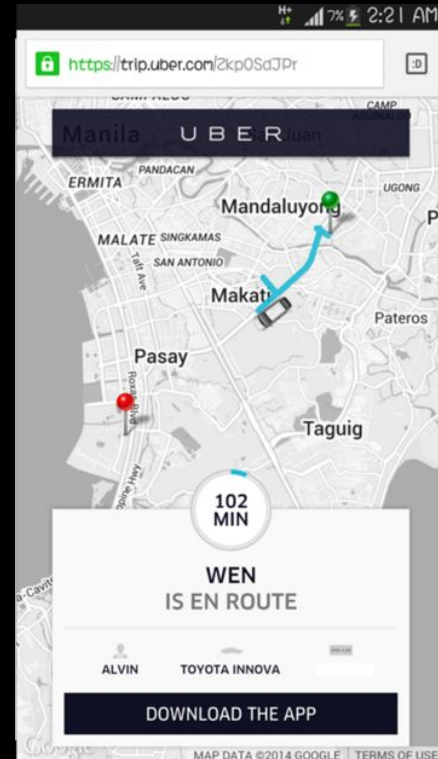


Experimentation Platform



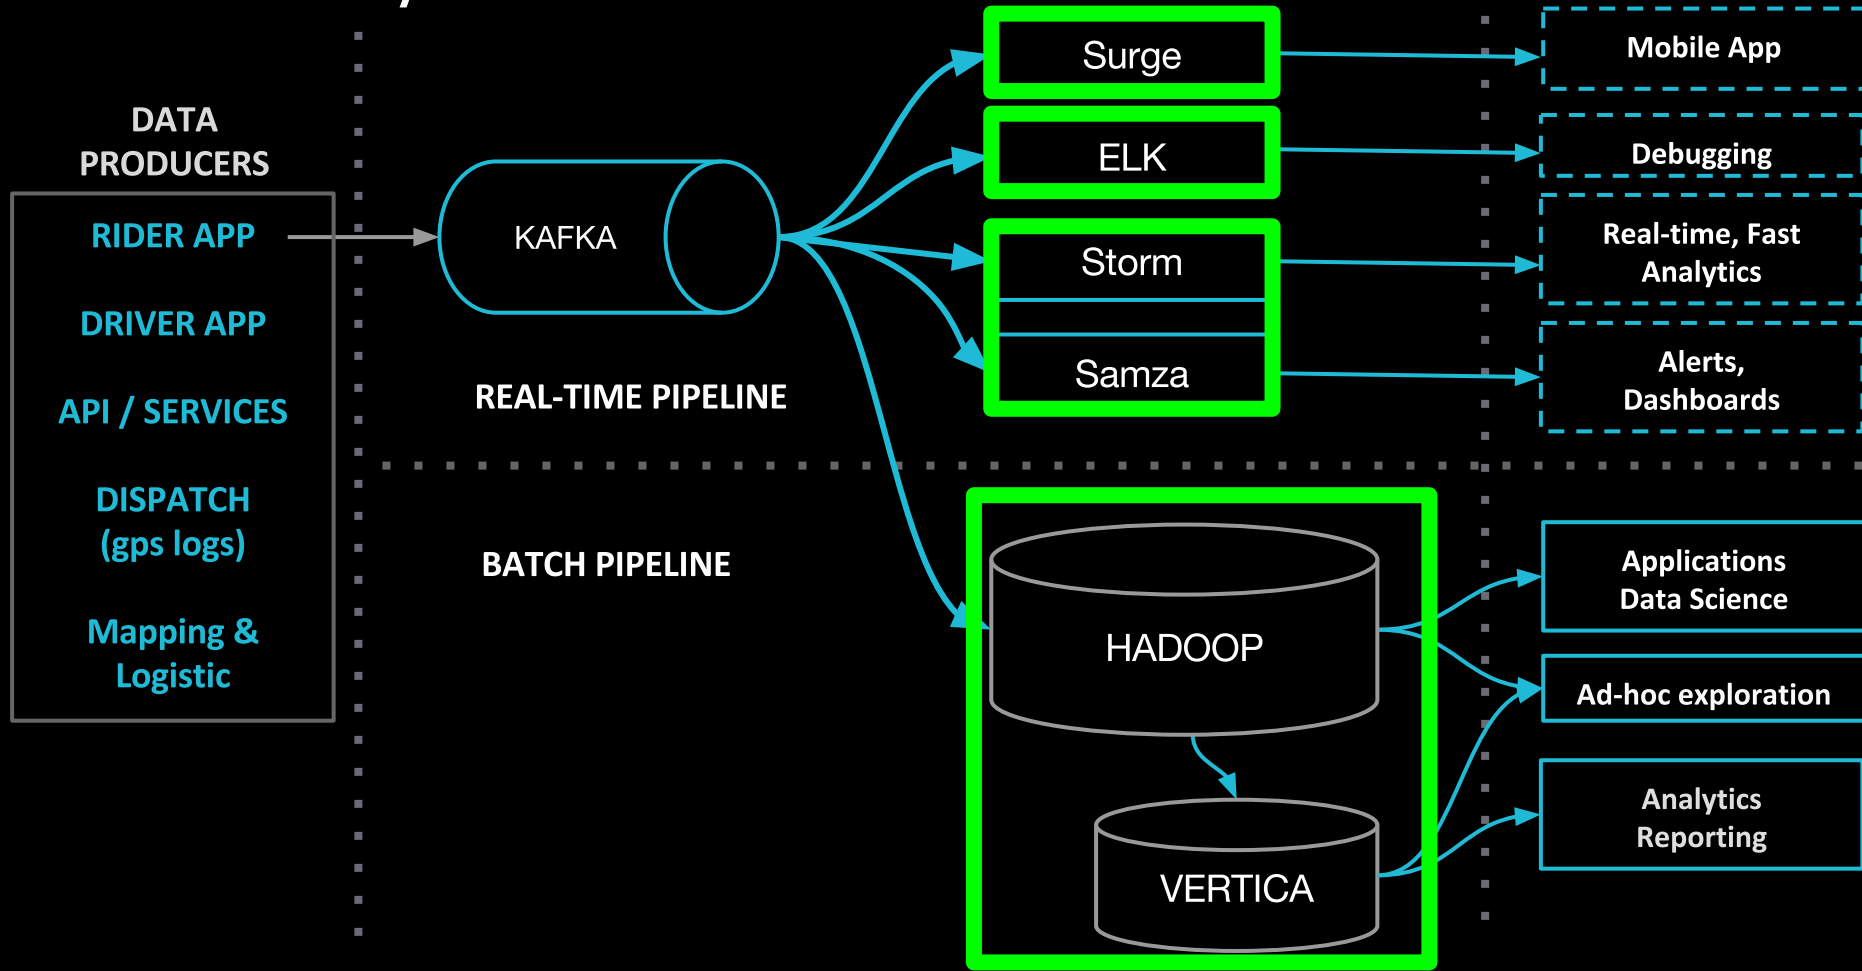
- Fraud detection
- Share my ETA

And many more ...



Apache Kafka is Uber's Lifeline

Kafka ecosystem @ Uber



Kafka cluster stats

100s of billion

Messages/day

100s TB

bytes/day

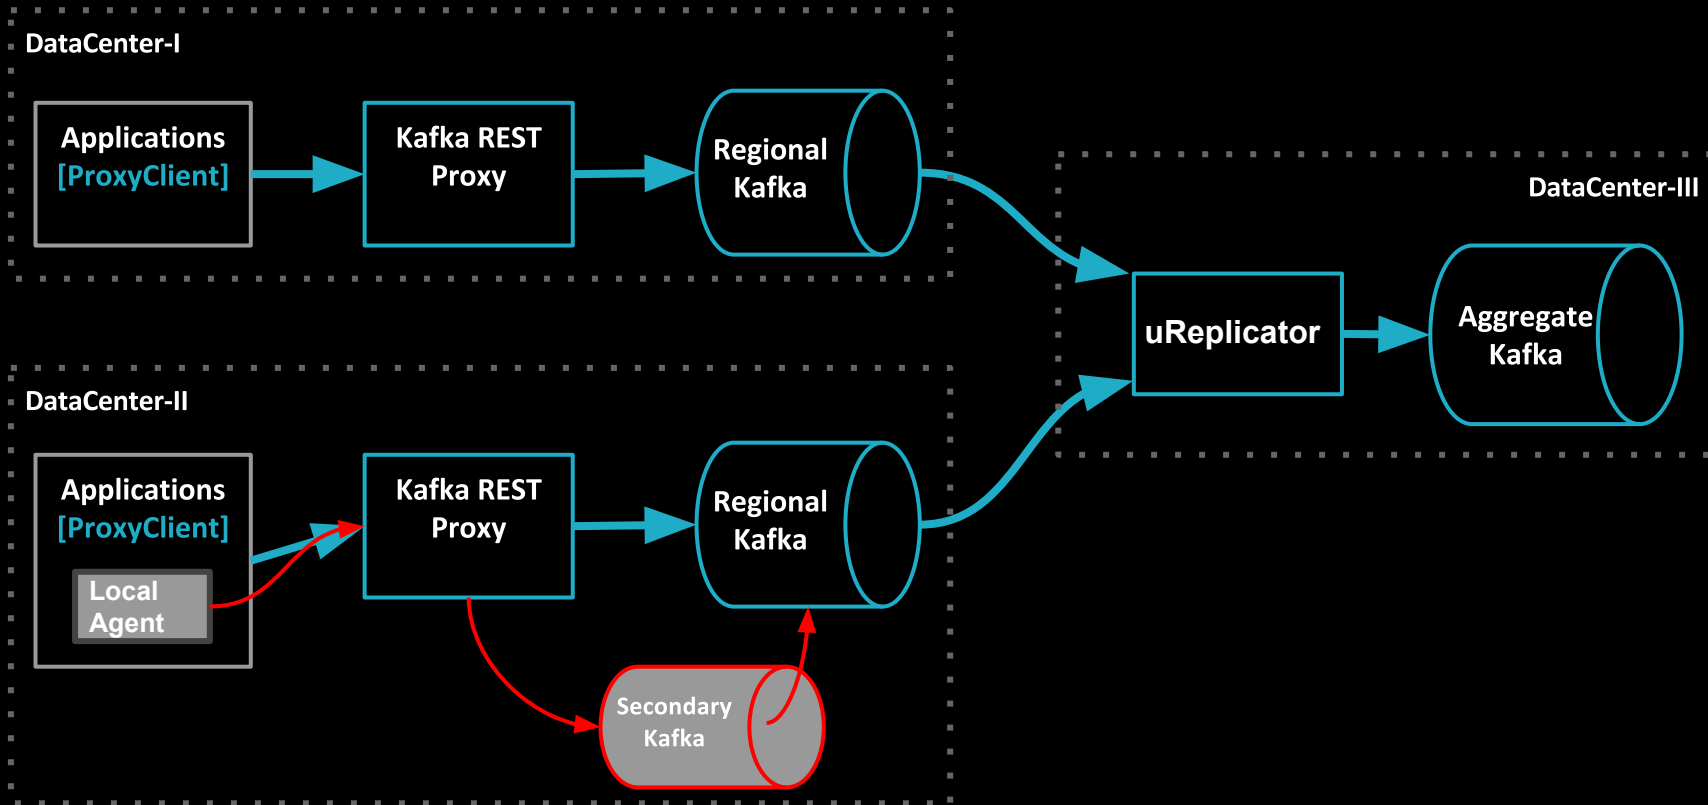
Multiple data centers

Kafka Infrastructure Deep Dive

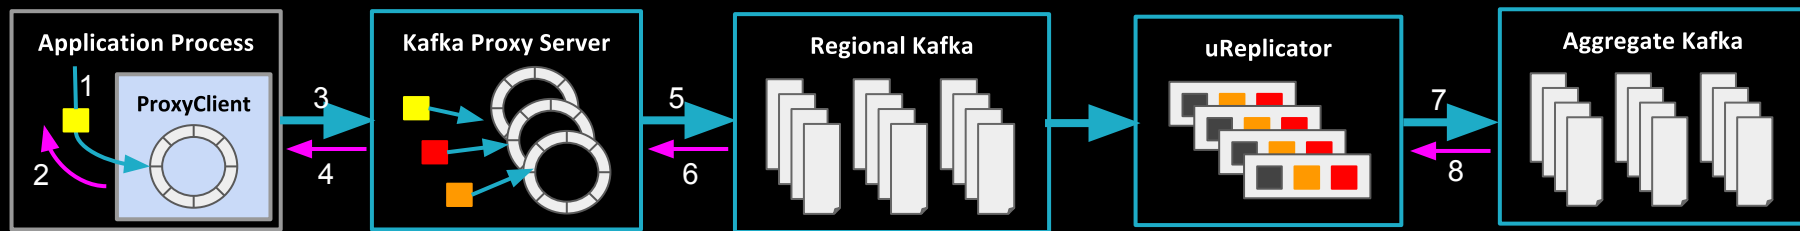
Requirements

- Scale to 100s Billions/day → 1 Trillion/day
- High Throughput (Scale: 100s TB → PB)
- Low Latency for most use cases(<5ms)
- Reliability - 99.99% (#Msgs Available /#Msgs Produced)
- Multi-Language Support
- Tens of thousands of simultaneous clients.
- Reliable data replication across DC

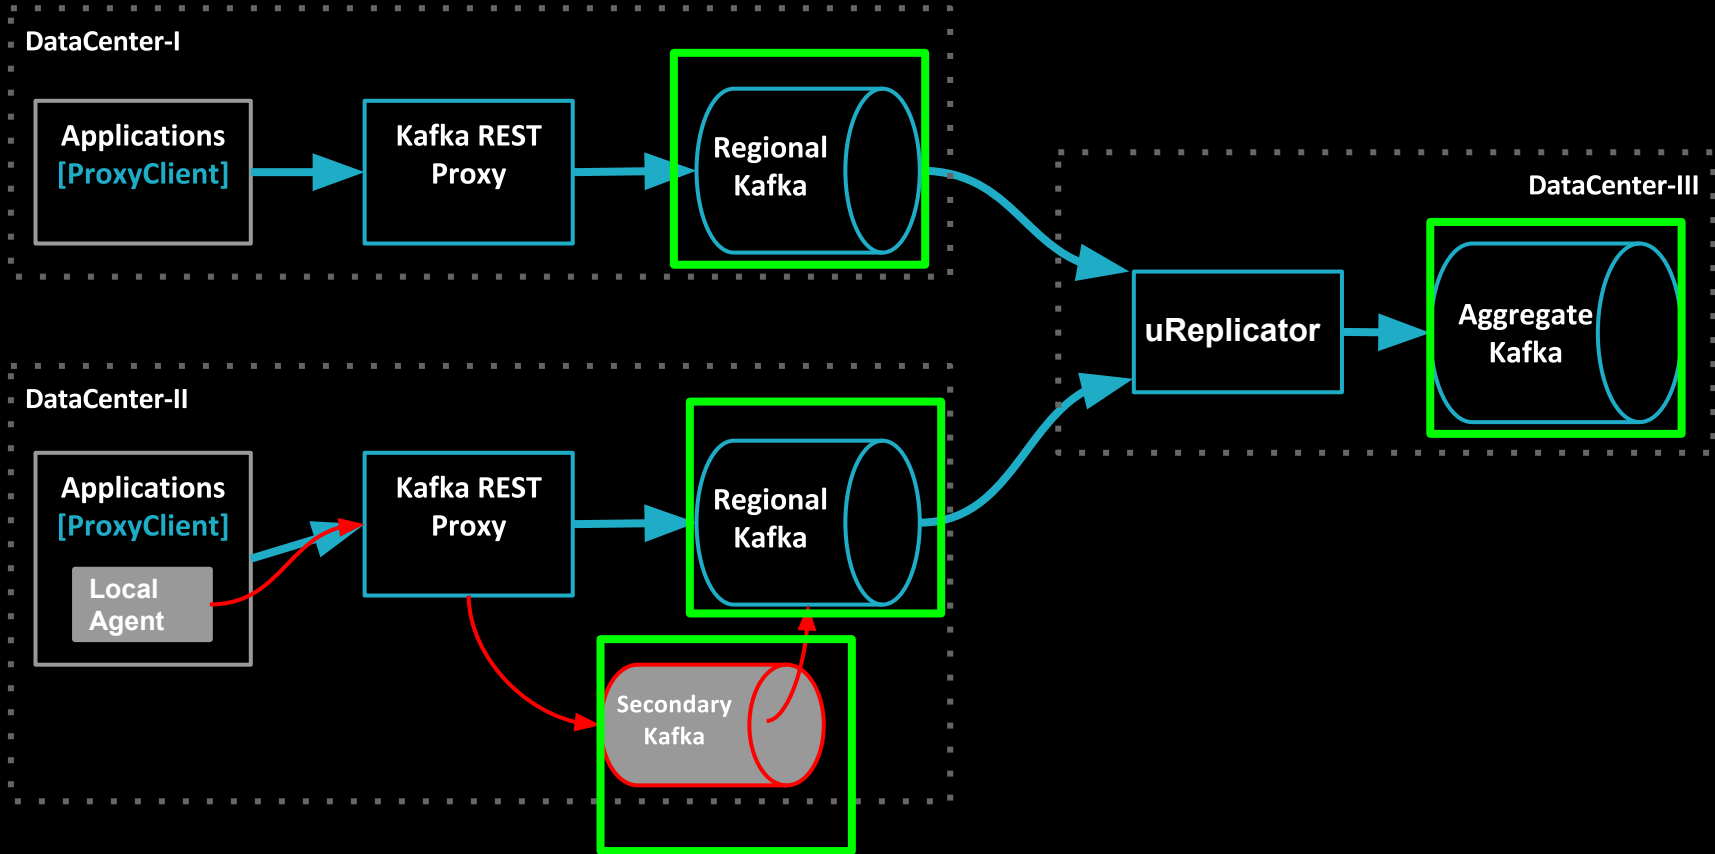
Kafka Pipeline



Kafka Pipeline: Data Flow



Kafka Clusters



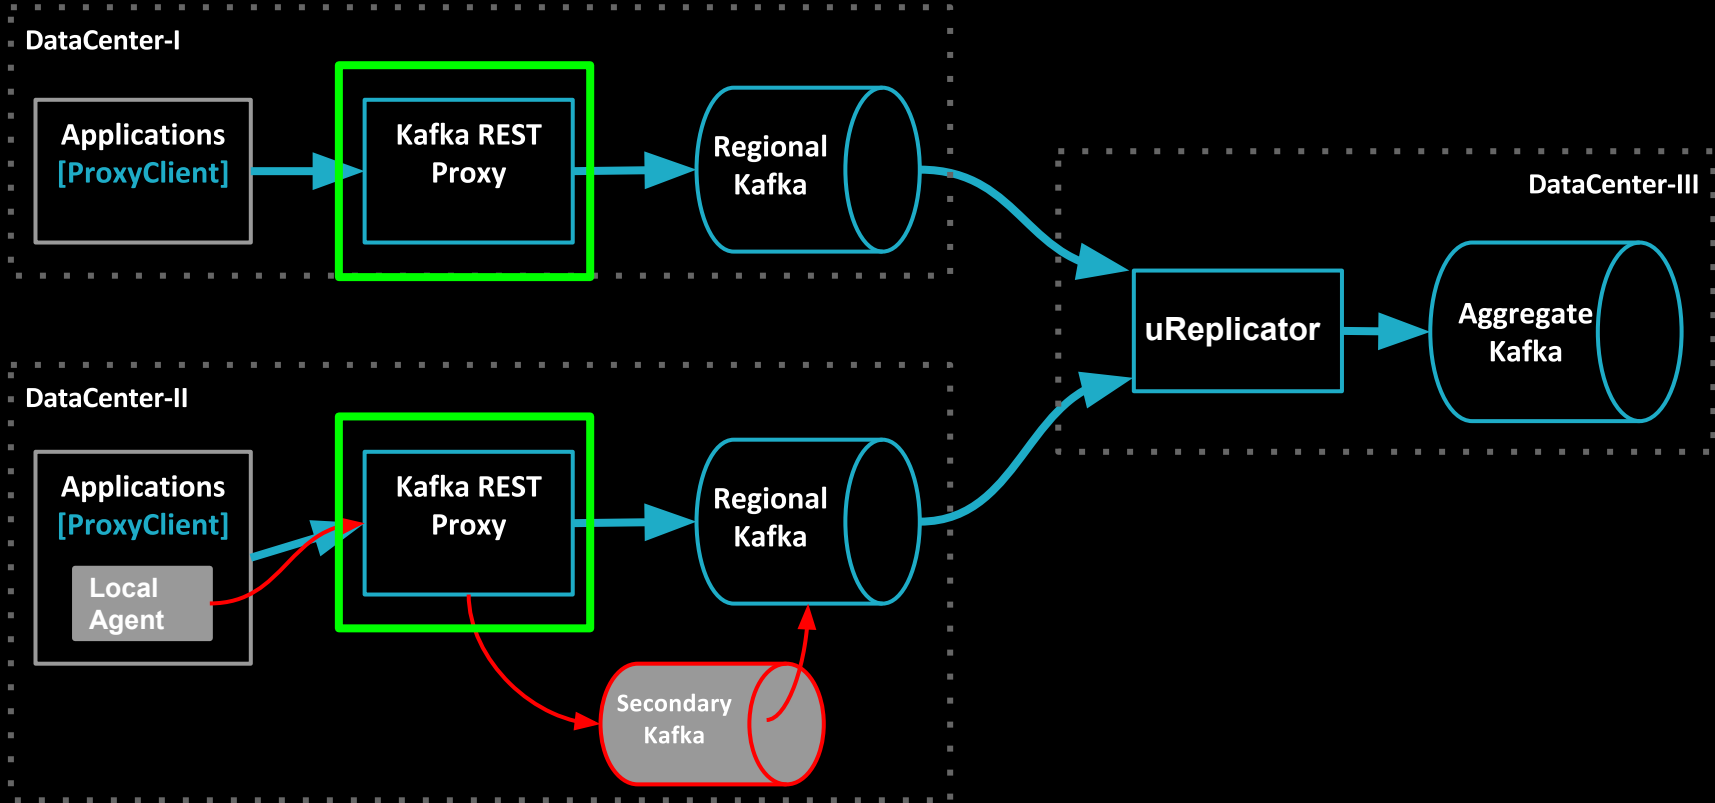
Kafka Clusters

- Use case based clusters
 - Data (async, reliable)
 - Logging (High throughput)
 - Time Sensitive (Low Latency e.g. Surge, Push notifications)
 - High Value Data (At-least once, Sync e.g. Payments)
- Secondary cluster as fallback
- Aggregate clusters for all data topics.

Kafka Clusters

- Scale to 100s Billions/day → 1 Trillion/day
- High Throughput (Scale: 100s TB → PB)
- Low Latency for most use cases(<5ms)
- Reliability - 99.99% (#Msgs Available /#Msgs Produced)
- Multi-Language Support
- Tens of thousands of simultaneous clients.
- Reliable data replication across DC

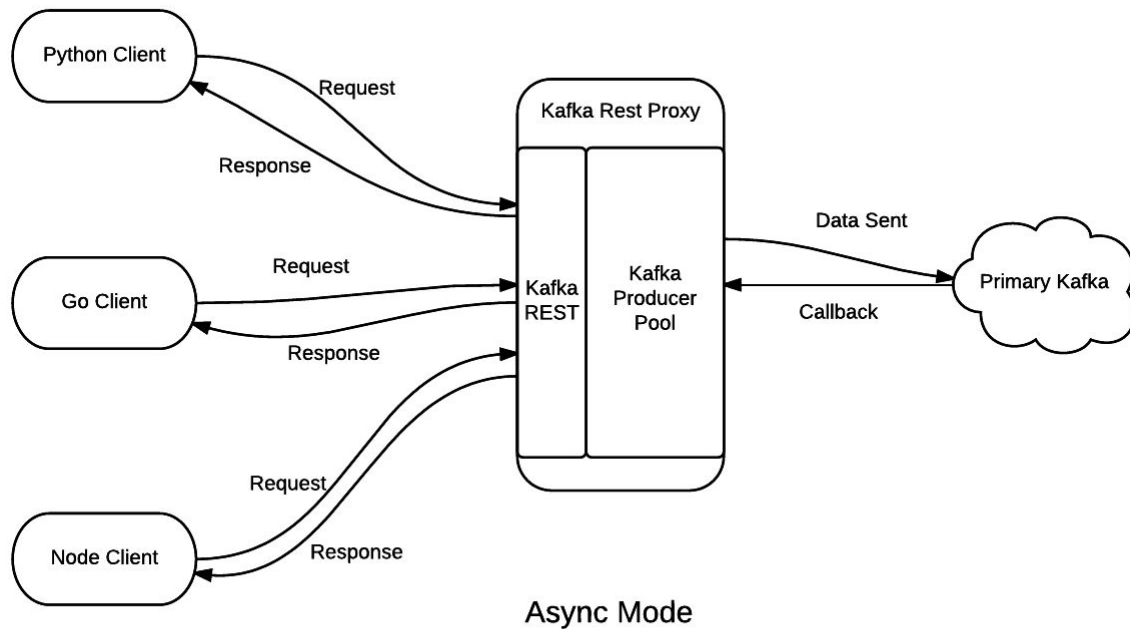
Kafka Rest Proxy



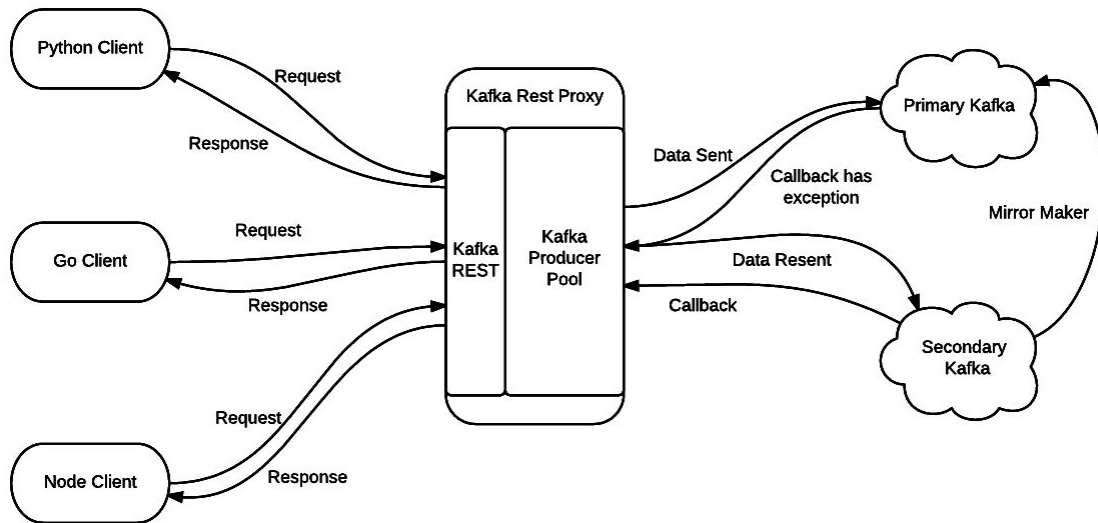
Why Kafka Rest Proxy ?

- Simplified Client API
- Multi-lang support (Java, NodeJs, Python, Golang)
- Decouple client from Kafka broker
 - Thin clients = operational ease
 - Less connections to Kafka brokers
 - Future kafka upgrade
- Enhanced Reliability
 - Primary & Secondary Kafka Clusters

Kafka Rest Proxy: Internals



Kafka Rest Proxy: Internals



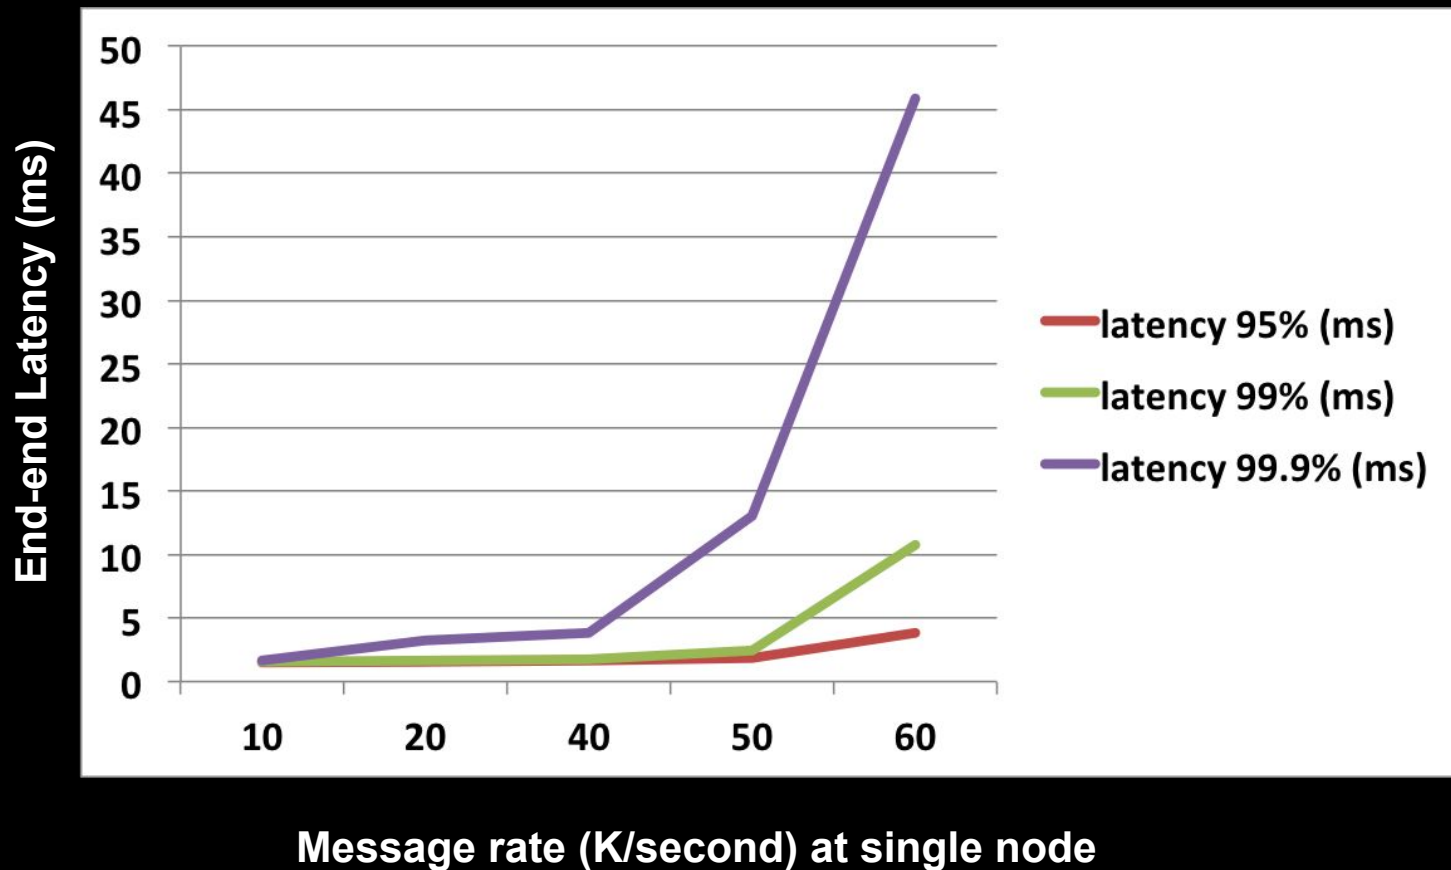
Async Mode

Kafka Rest Proxy: Internals

- Based on Confluent's open sourced Rest Proxy
- Performance enhancements
 - Simple http servlets on jetty instead of Jersey
 - Optimized for binary payloads.
 - Performance increase from 7K* to 45-50K QPS/box
- Caching of topic metadata.
- Reliability improvements*
 - Support for Fallback cluster
 - Support for multiple Producers (SLA based segregation)
- Plan to contribute back to community

*Based on benchmarking & analysis done in Jun '2015

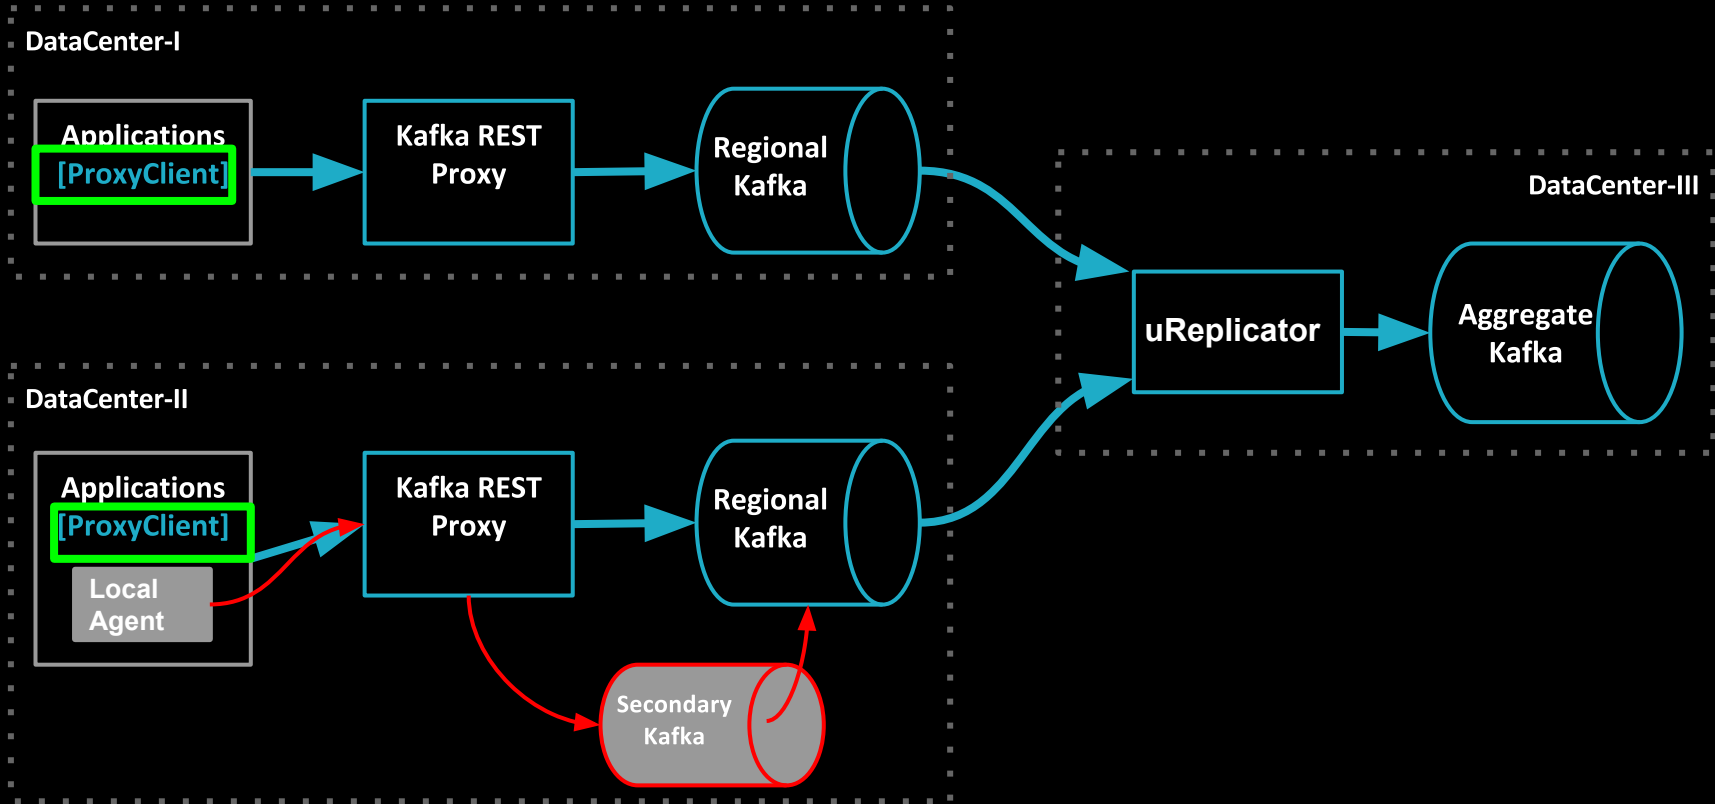
Rest Proxy: performance (1 box)



Kafka Clusters + Rest Proxy

- Scale to 100s Billions/day → 1 Trillion/day
- High Throughput (Scale: 100s TB → PB)
- Low Latency for most use cases(<5ms)
- Reliability - 99.99% (#Msgs Available /#Msgs Produced)
- Multi-Language Support
- Tens of thousands of simultaneous clients.
- Reliable data replication across DC

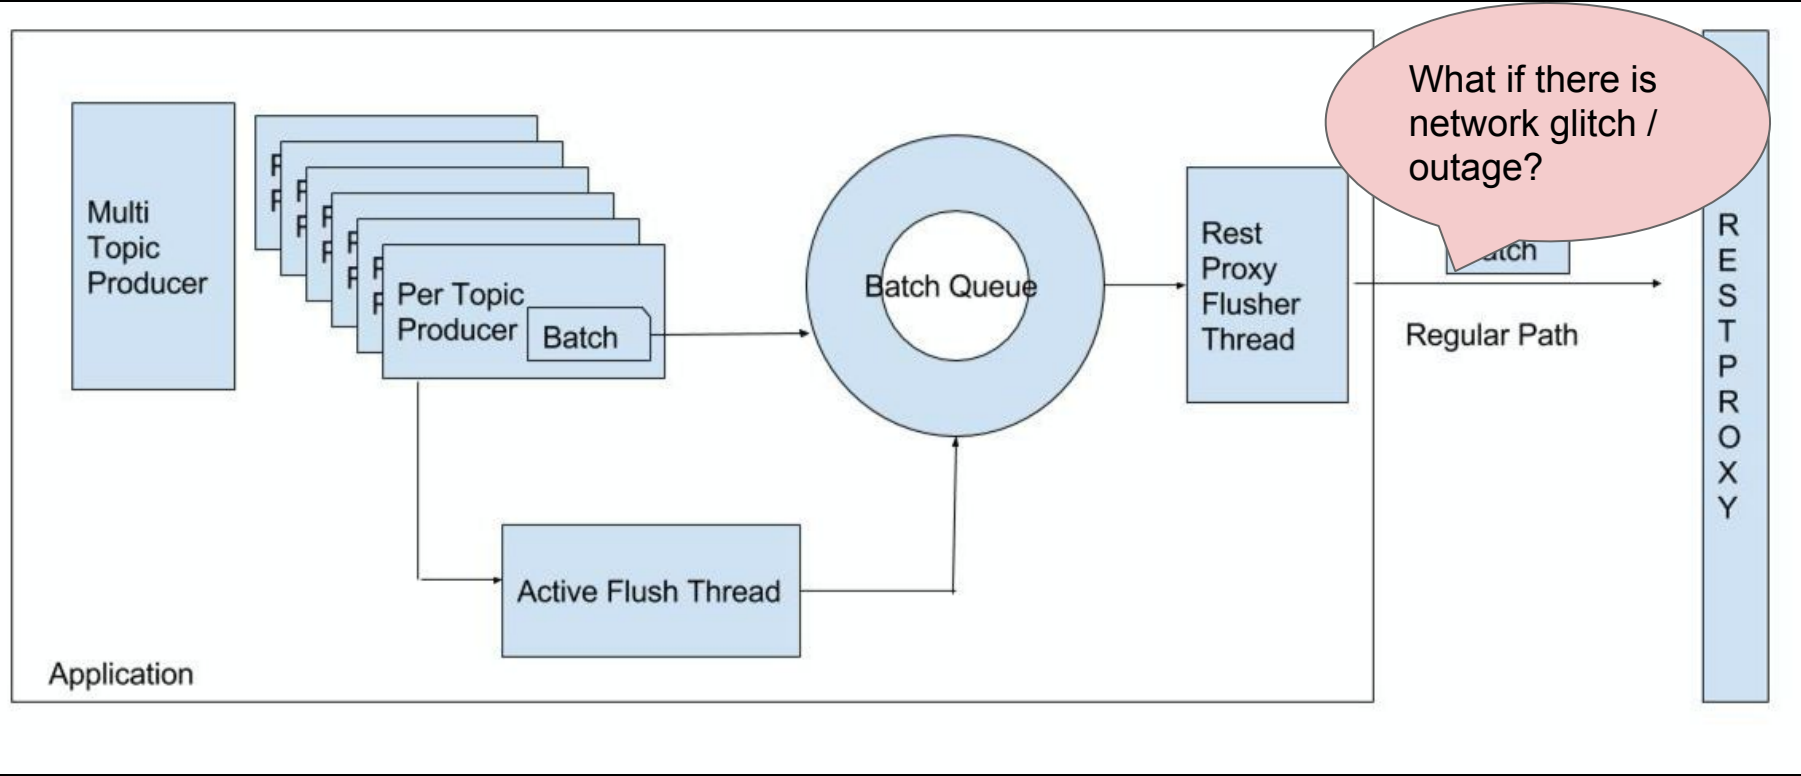
Kafka Clients



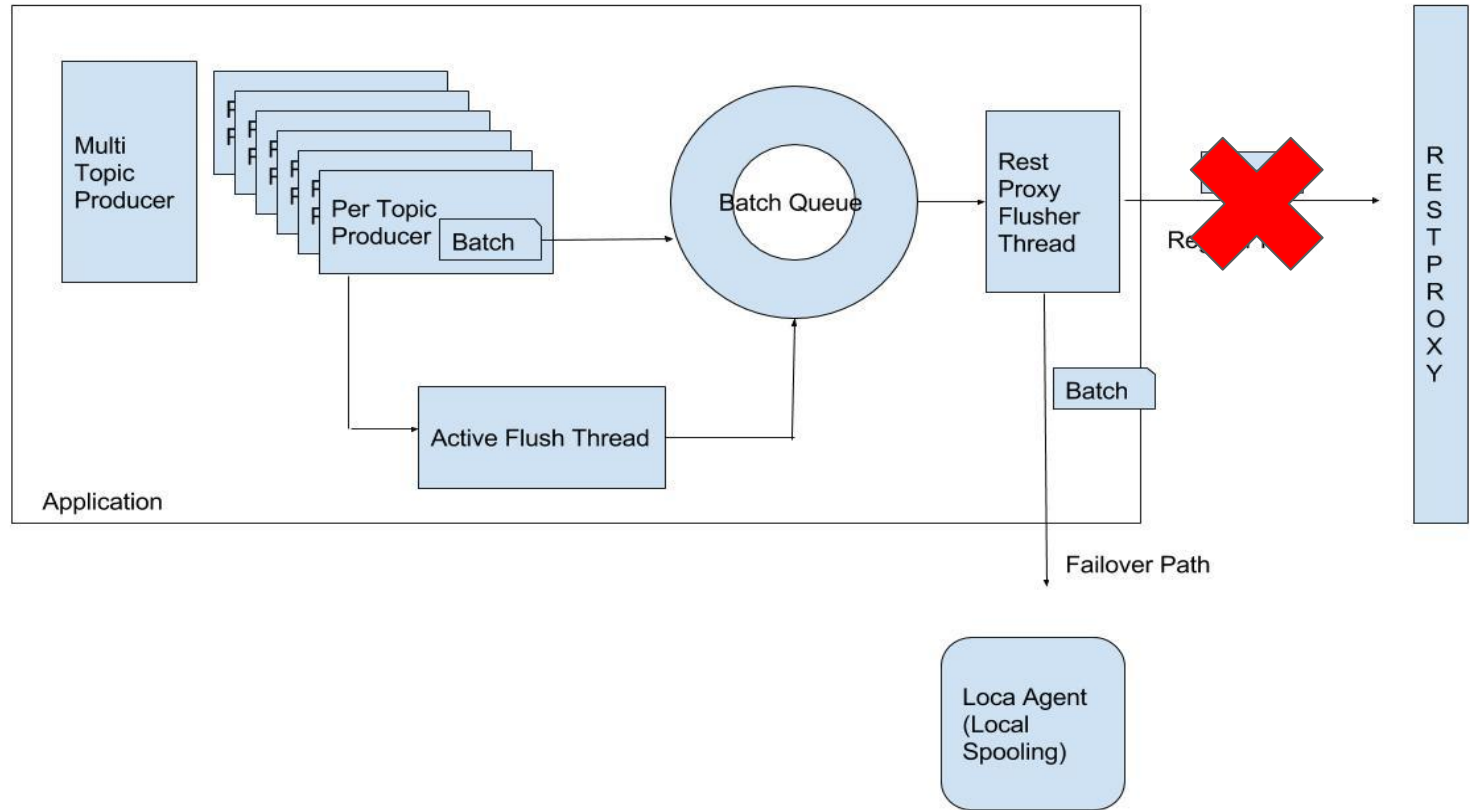
Client Libraries

- Support for multiple clusters.
- High Throughput
 - Non-blocking, async, batching
 - <1ms produce latency for clients
 - Handles Throttling/BackOff signals from Rest Proxy
- Topic Discovery
 - Discovers the kafka cluster a topic belongs
 - Able to multiplex to different kafka clusters
- Integration with Local Agent for critical data

Client Libraries



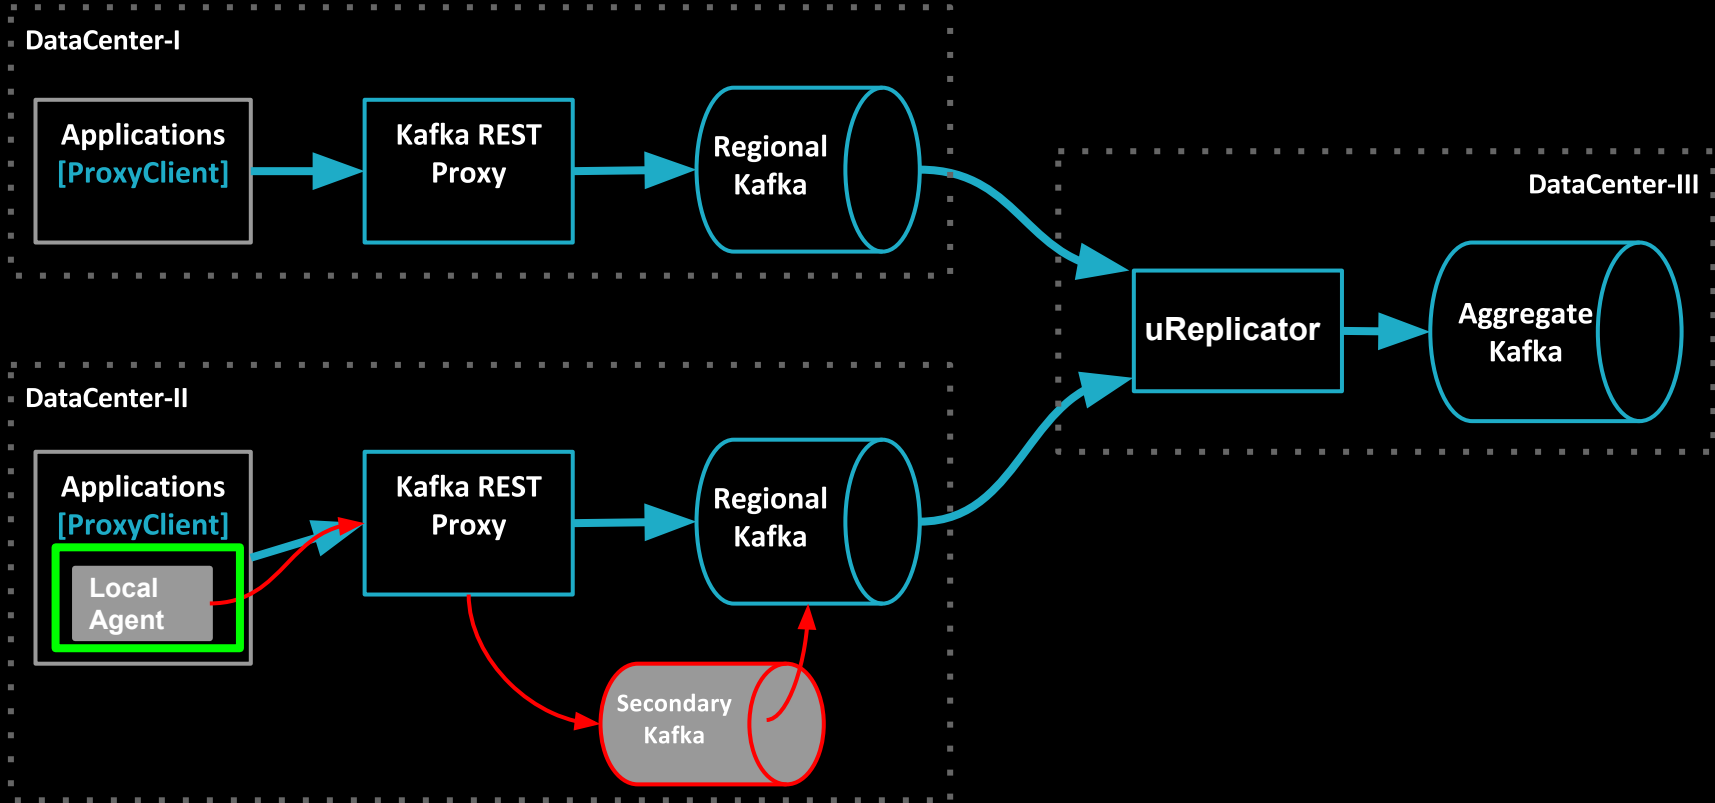
Client Libraries



Kafka Clusters + Rest Proxy + Clients

- Scale to 100s Billions/day → 1 Trillion/day
- High Throughput (Scale: 100s TB → PB)
- Low Latency for most use cases(<5ms)
- Reliability - 99.99% (#Msgs Available /#Msgs Produced)
- Multi-Language Support
- Tens of thousands of simultaneous clients.
- Reliable data replication across DC

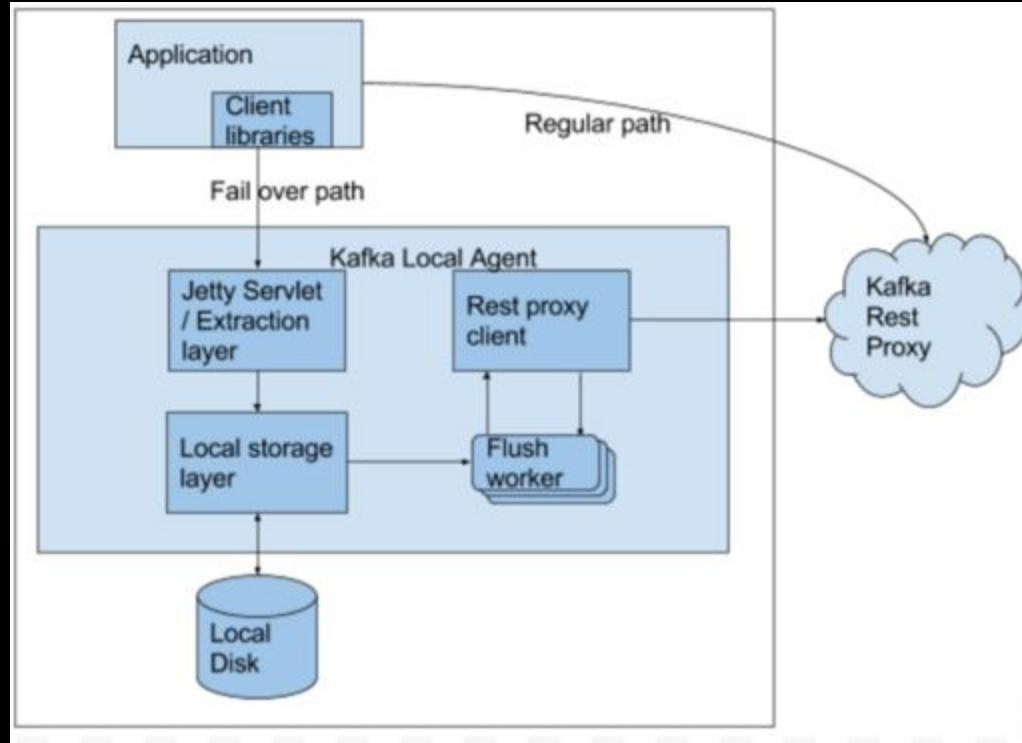
Local Agent



Local Agent

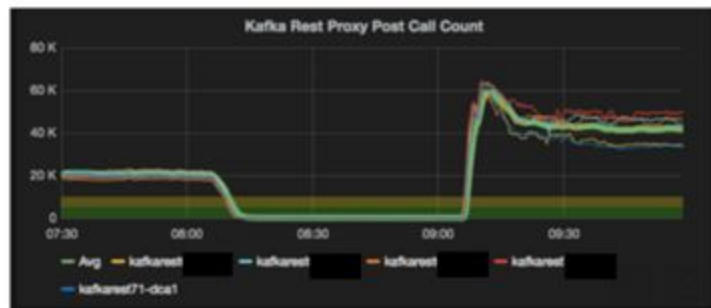
- Local spooling in case of downstream outage/backpressure
- Backfills at the controlled rate to avoid hammering infrastructure recovering from outage
- Implementation:
 - Reuses code from rest-proxy and kafka's log module.
 - Appends all topics to same file for high throughput.

Local Agent Architecture

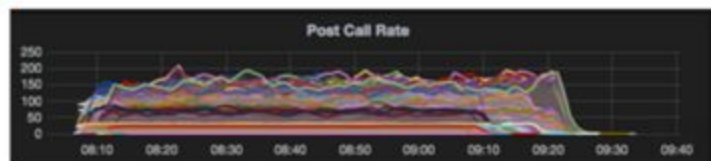


Local Agent in Action

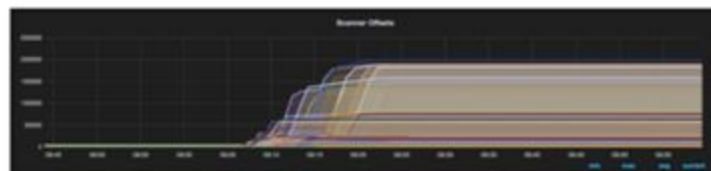
#1: Rest proxy data incoming traffic coming down to 0 around 8:10. Recovery around 9:10 with a spike from client libraries buffers



#2: Traffic starts moving to local agent in outage window



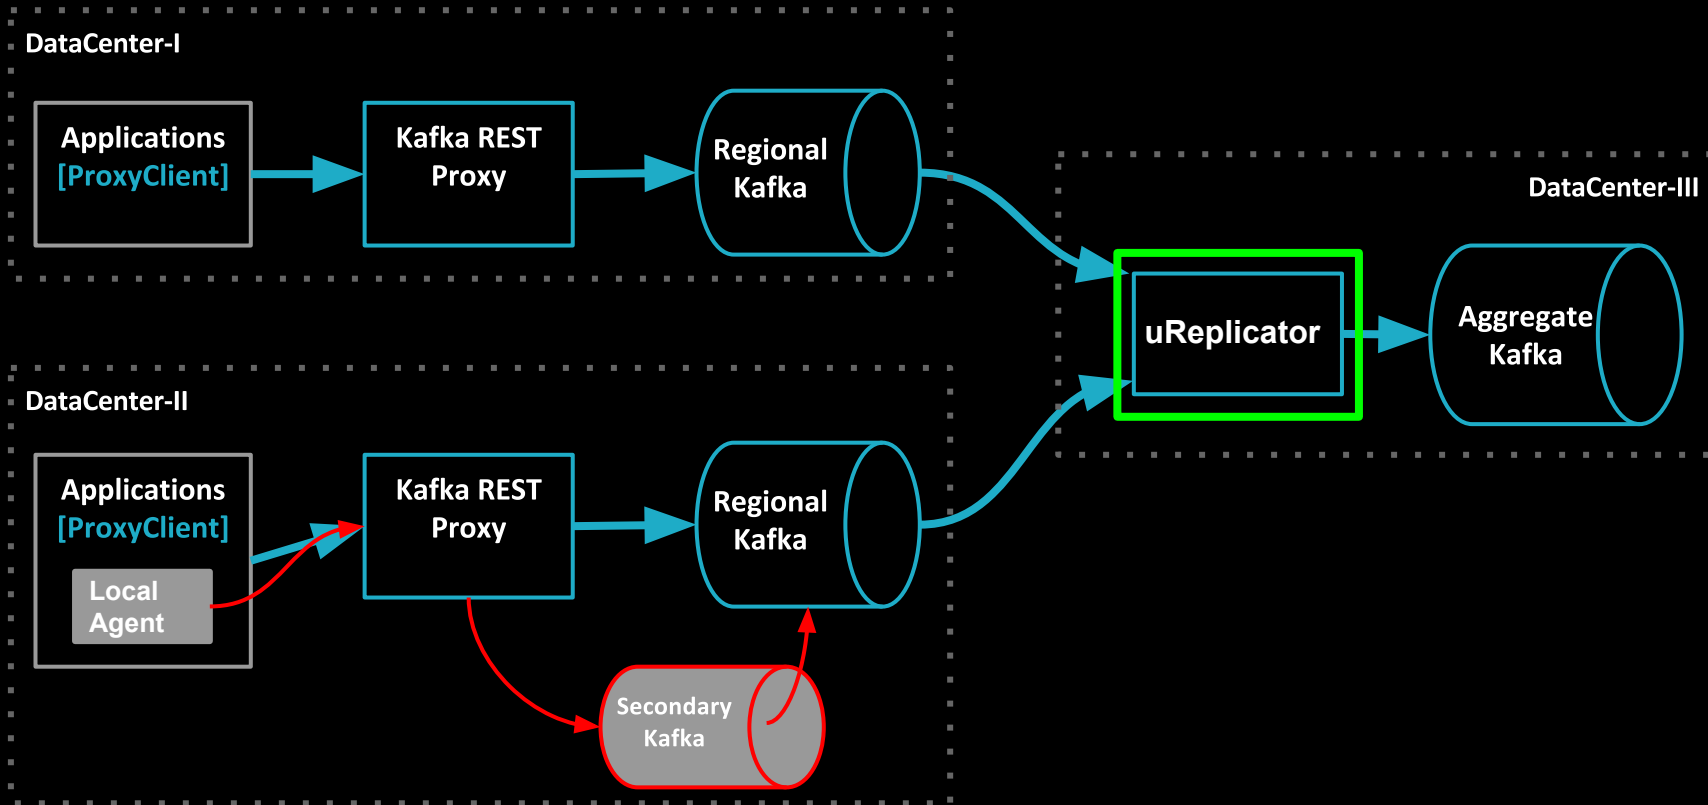
#3: Local agent producing to Rest proxy after recovery



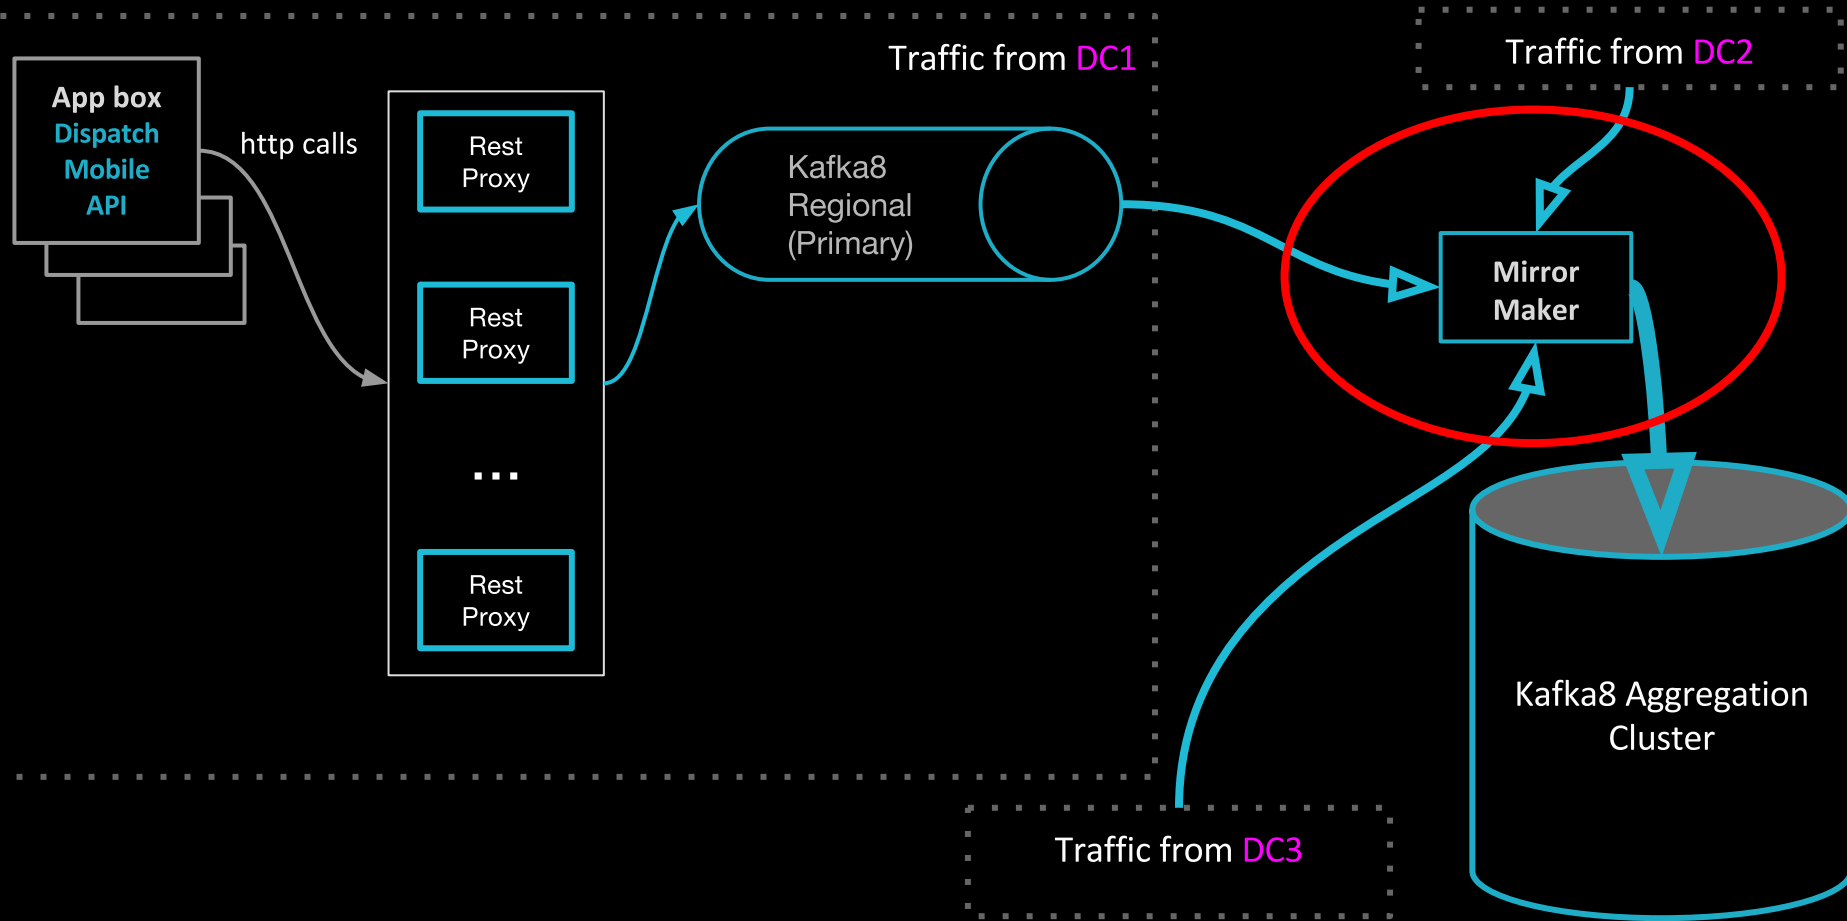
Kafka Clusters + Rest Proxy + Clients + Local Agent

- Scale to 100s Billions/day → 1 Trillion/day
- High Throughput (Scale: 100s TB → PB)
- Low Latency for most use cases(<5ms)
- Reliability - 99.99% (#Msgs Available /#Msgs Produced)
- Multi-Language Support
- Tens of thousands of simultaneous clients.
- Reliable data replication across DC

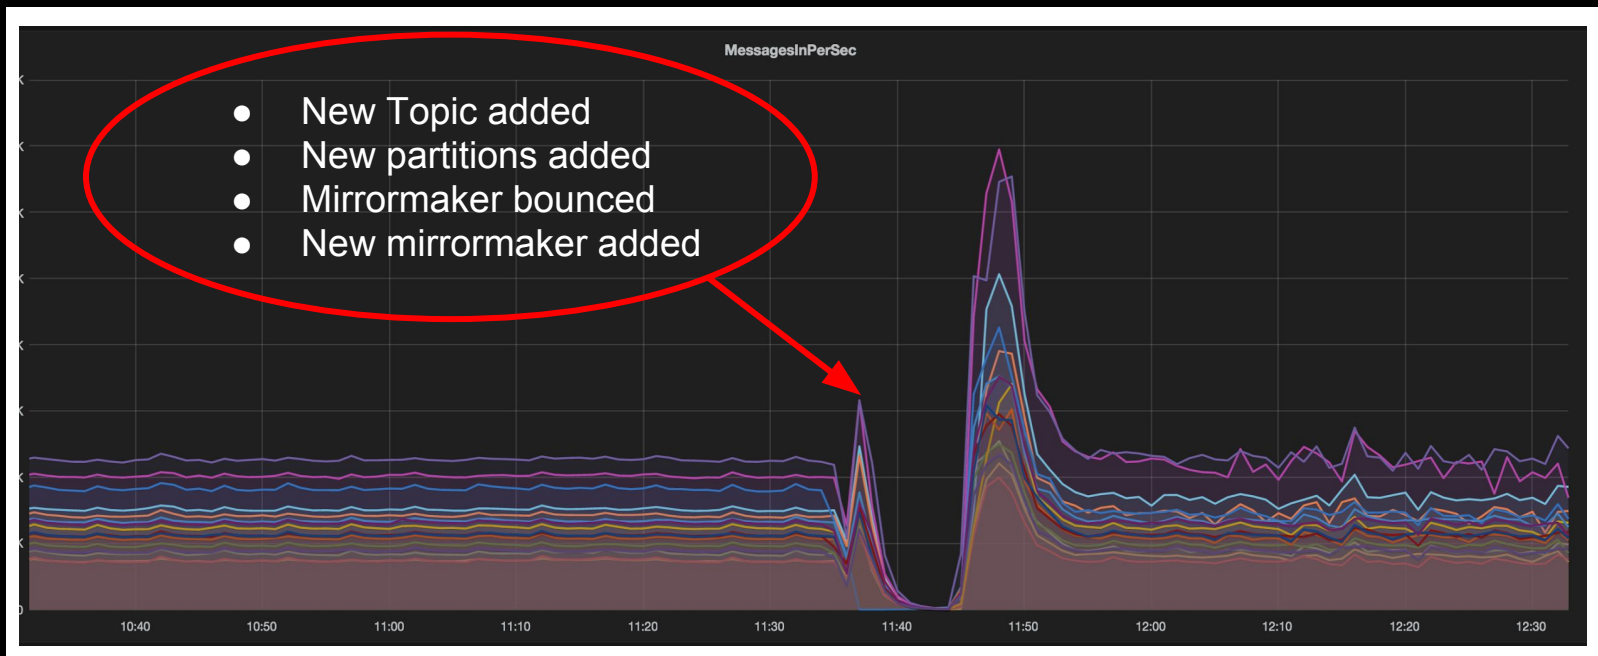
uReplicator



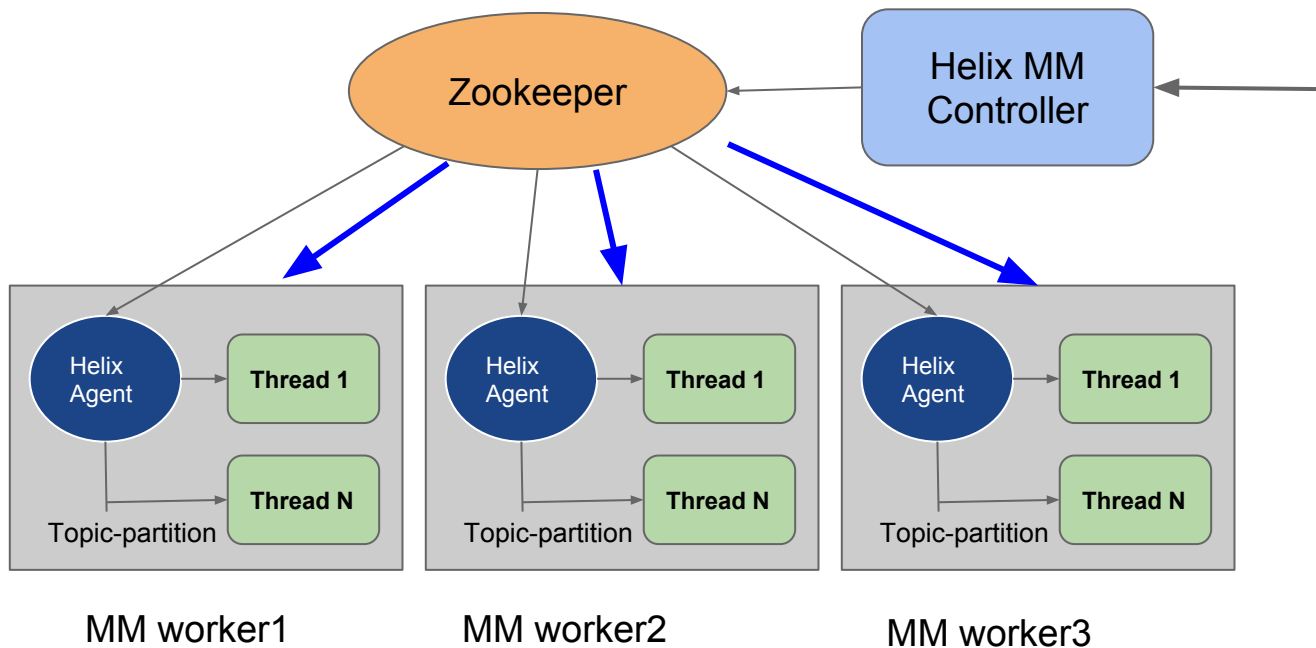
Multi-DC data flow



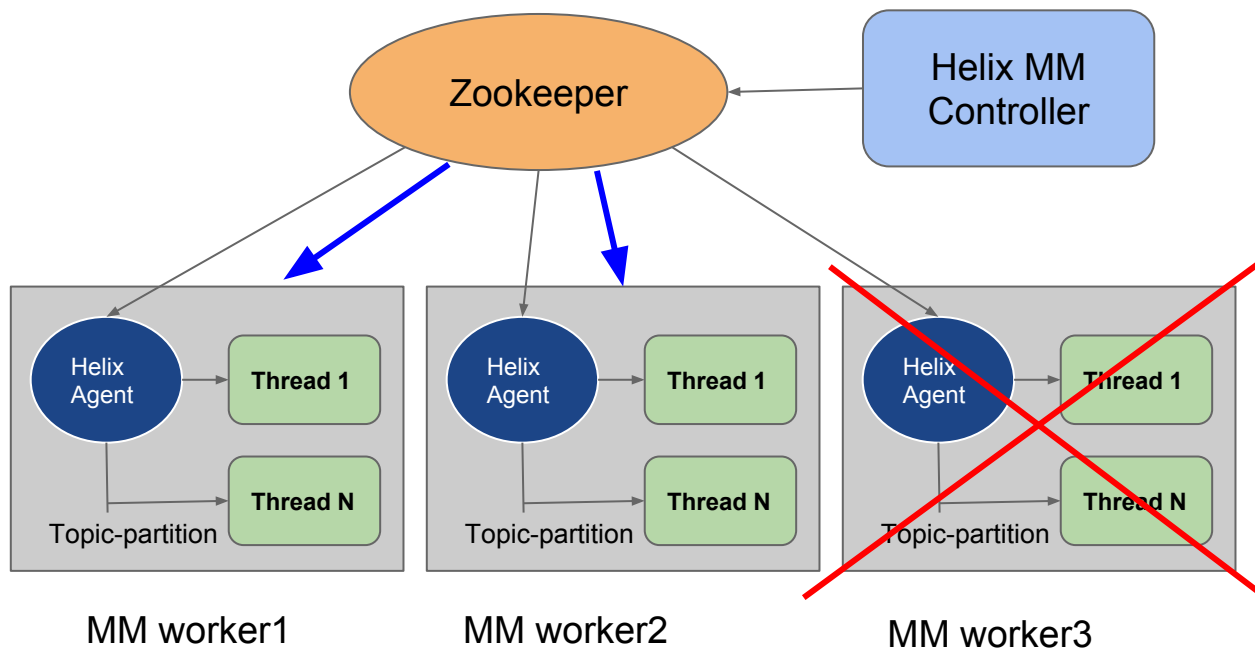
Mirrormaker : existing problems



uReplicator: In-house solution



uReplicator



Kafka Clusters + Rest Proxy + Clients + Local Agent

- Scale to 100s Billions/day → 1 Trillion/day
- High Throughput (Scale: 100s TB → PB)
- Low Latency for most use cases(<5ms)
- Reliability - 99.99% (#Msgs Available /#Msgs Produced)
- Multi-Language Support
- Tens of thousands of simultaneous clients.
- Reliable data replication across DC

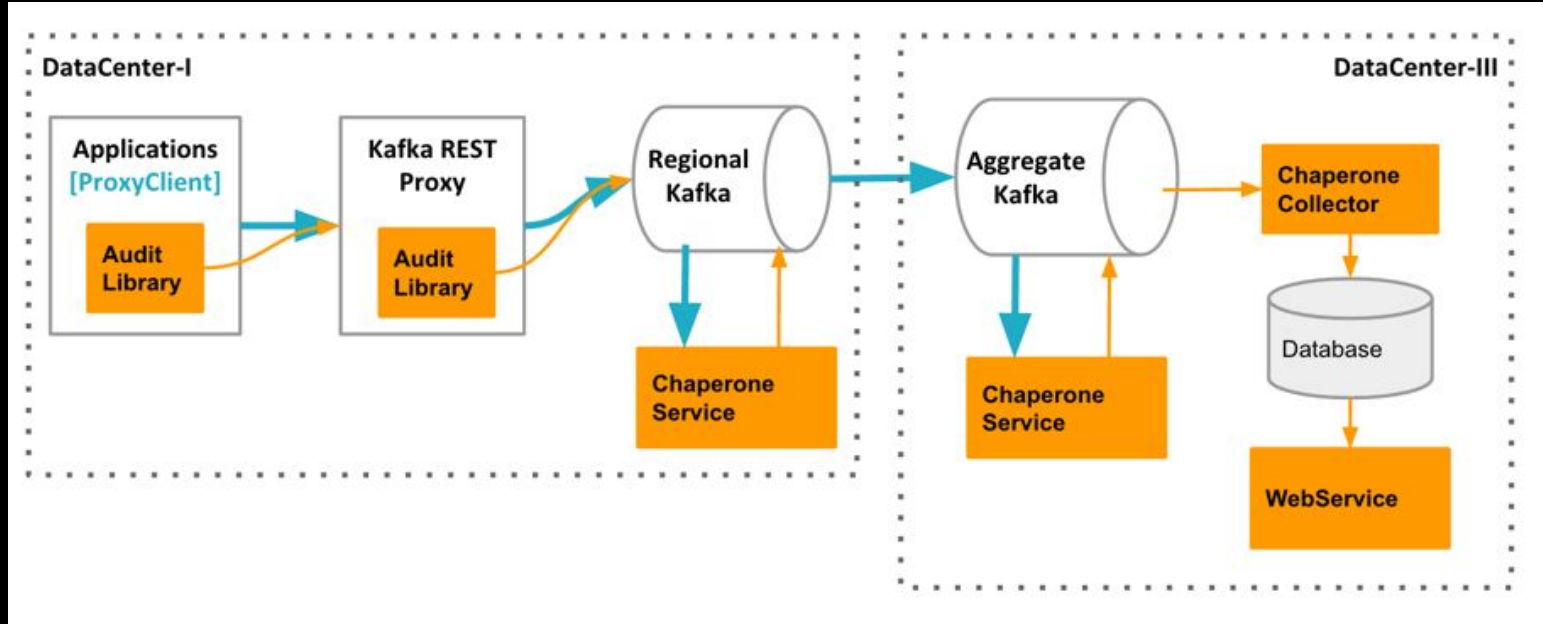
uReplicator

- Running in production for 1+ year
- Open sourced: <https://github.com/uber/uReplicator>
- Blog: <https://eng.uber.com/ureplicator/>

Chaperone - E2E Auditing



Chaperone Architecture

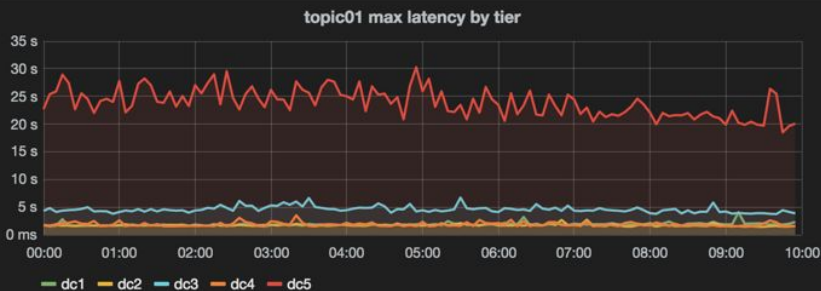
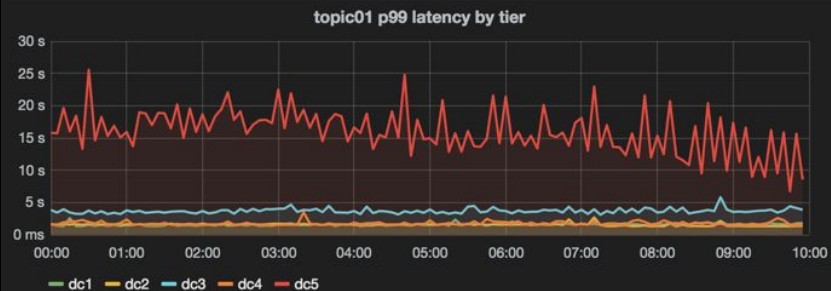
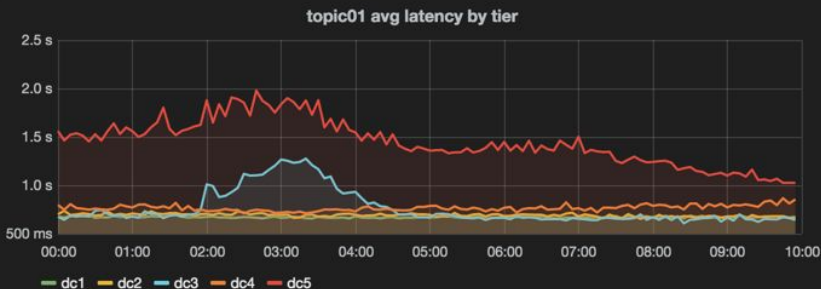
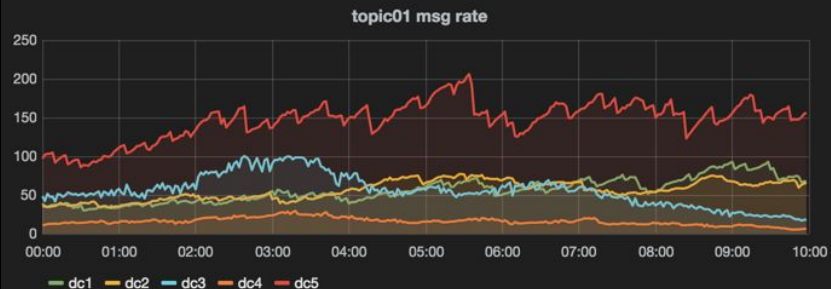


Chaperone : Track counts



Chaperone : Track Latency

MSGRATE AND LATENCY



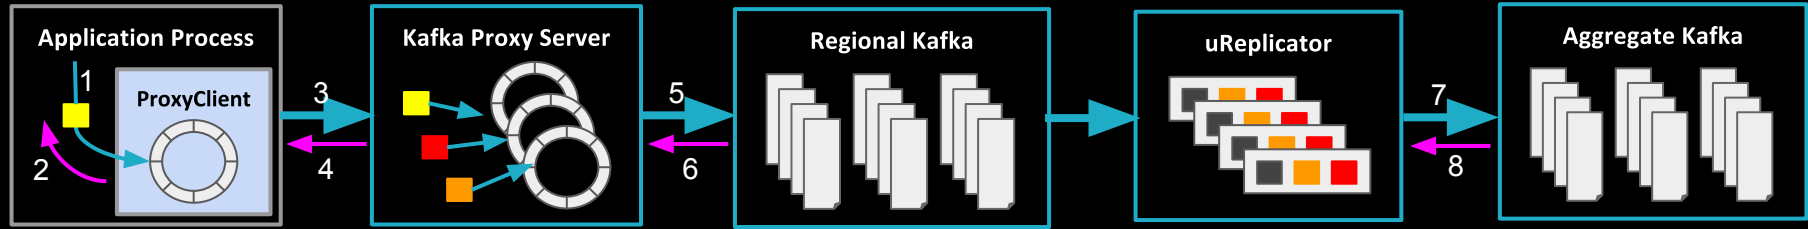
Chaperone

- Running in production for 1+ year
- Planning to open source in ~2 Weeks

At-least Once Kafka



Why do we need it?



- Most of infrastructure tuned for high throughput
 - Batching at each stage
 - Ack before produce (ack'ed != committed)
- Single node failure in any stage leads to data loss
- Need a reliable pipeline for High Value Data e.g. Payments

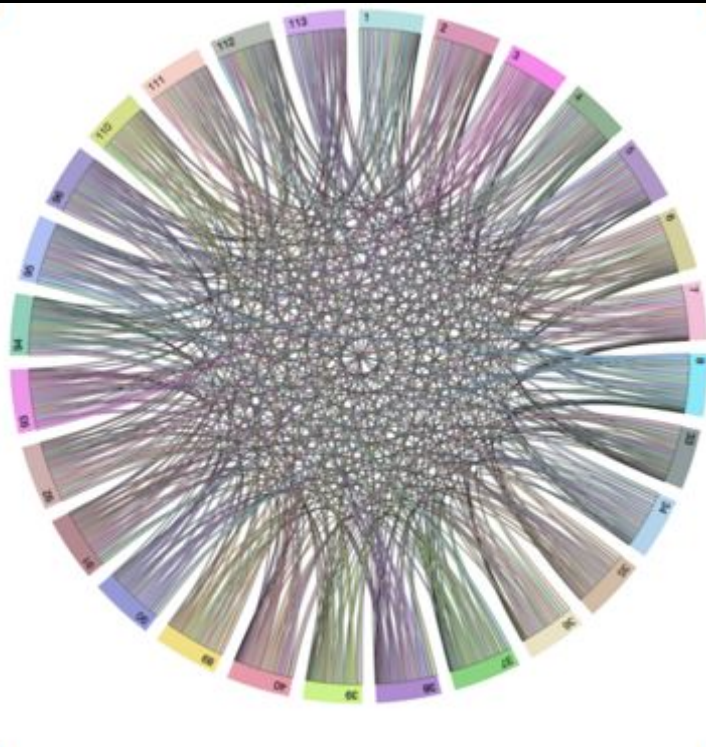
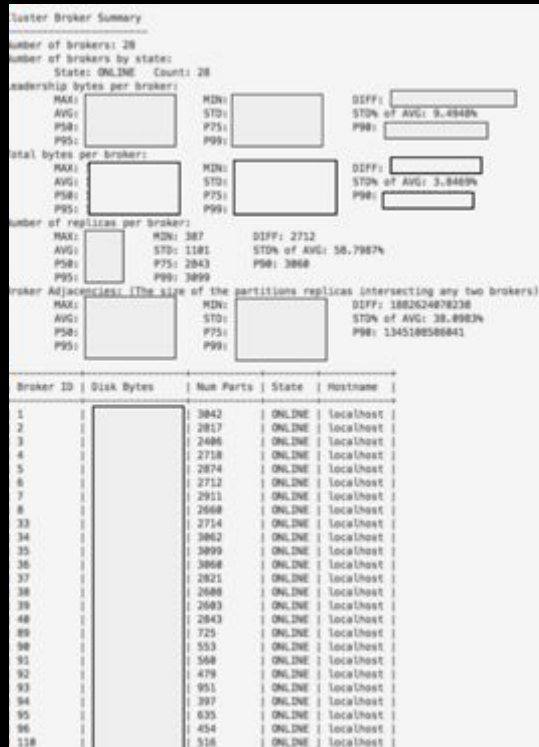
How did we achieve it?

- Brokers:
 - `min.insync.replicas=2`, can only tolerate one node failure
 - `unclean.leader.election=false`, need to wait until the old leader comes back
- Rest Proxy:
 - Partition Failover
- Improved Operations:
 - Replication throttling, to reduce impact of node bootstrap
 - Prevent catching up nodes to become ISR

Operations/Tooling

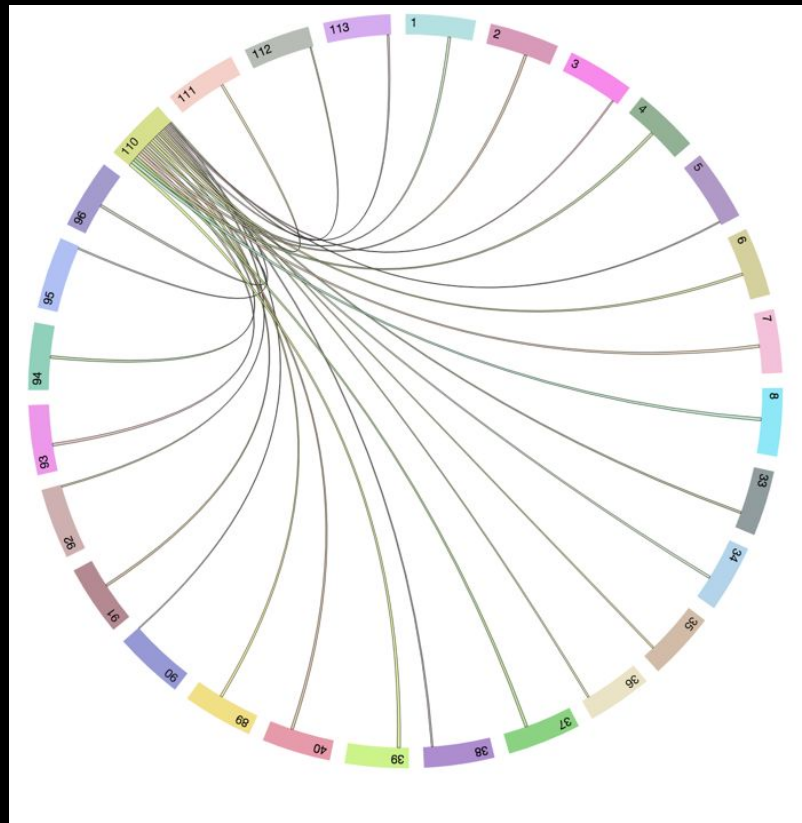


Partition Rebalancing



Partition Rebalancing

- Calculates partition imbalance and inter-broker dependency.
- Generates & Executes Rebalance Plan.
- Rebalance plans are incremental, can be stopped and resumed.
- Currently on-demand, Automated in the future.



XFS vs EXT4



Summary: Scale

- Kafka Brokers:
 - Multiple Clusters per DC
 - Use case based tuning
- Rest Proxy to reduce connections and better batching
- Rest Proxy & Clients
 - Batch everywhere, Async produce
 - Replace Jersey with Jetty
- XFS

Summary: Reliability

- Local Agent
- Secondary Clusters
- Multi Producer support in Rest Proxy
- uReplicator
- Auditing via Chaperone

Future Work

- Open source contribution
 - Chaperone
 - Toolkit
- Data Lineage
- Active Active Kafka
- Chargeback
- Exactly once mirroring via uReplicator

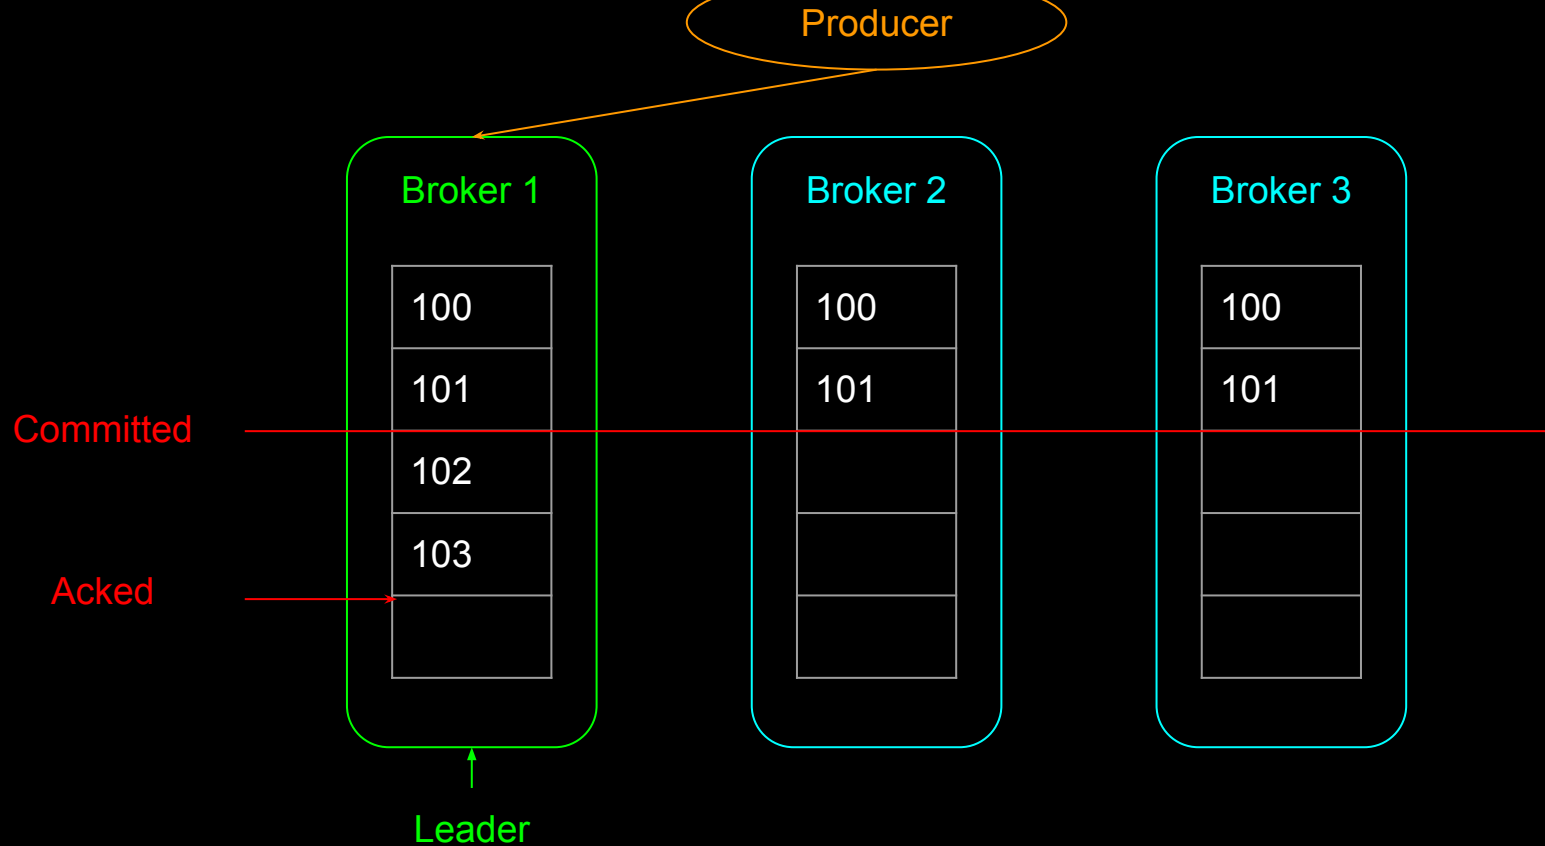
Questions ?

ankur@uber.com

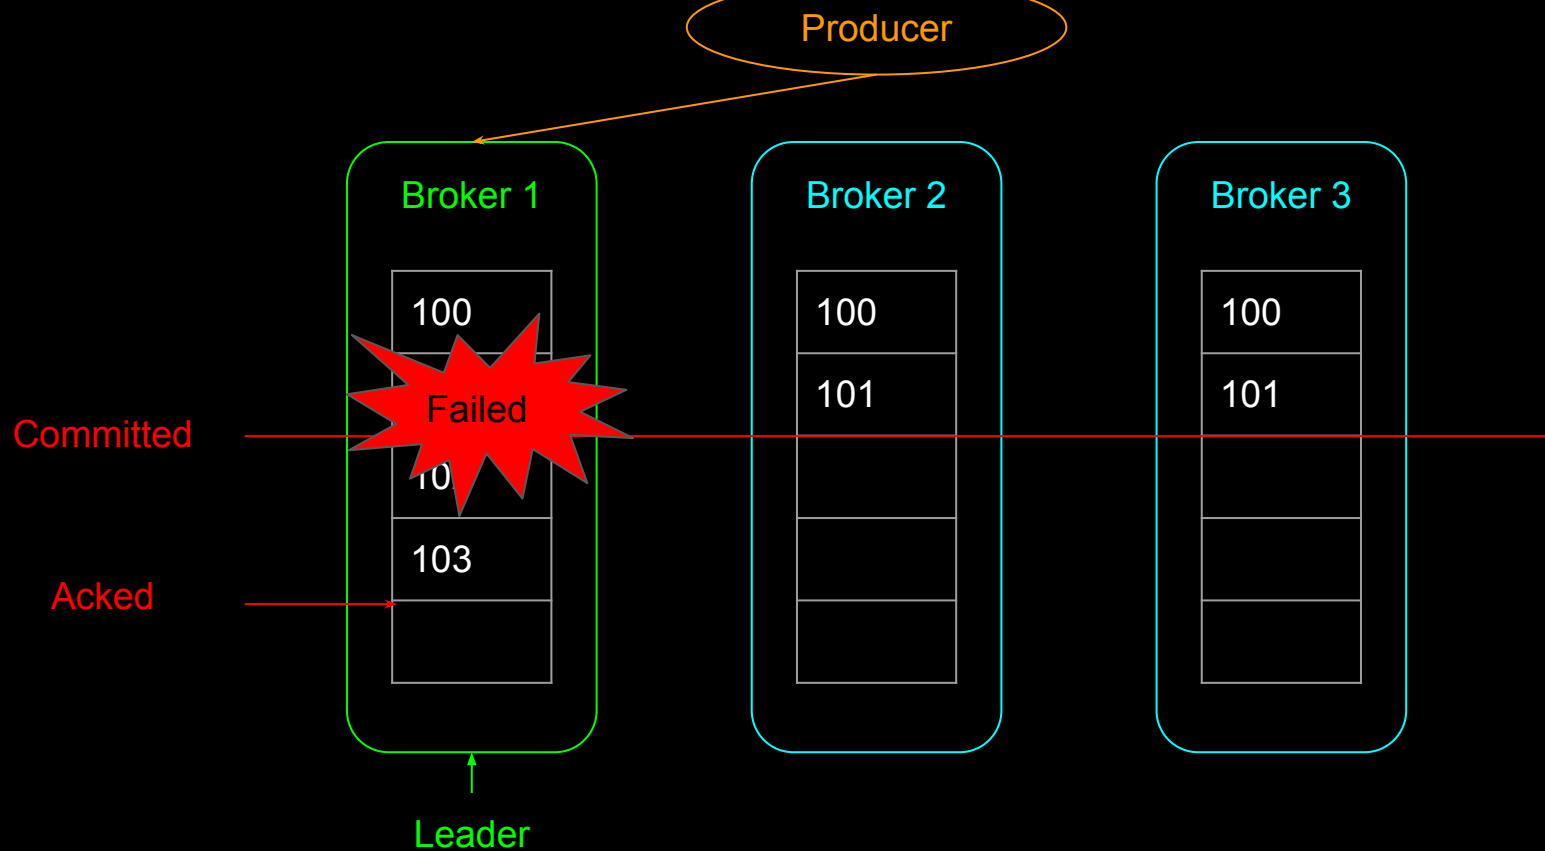
Extra Slides



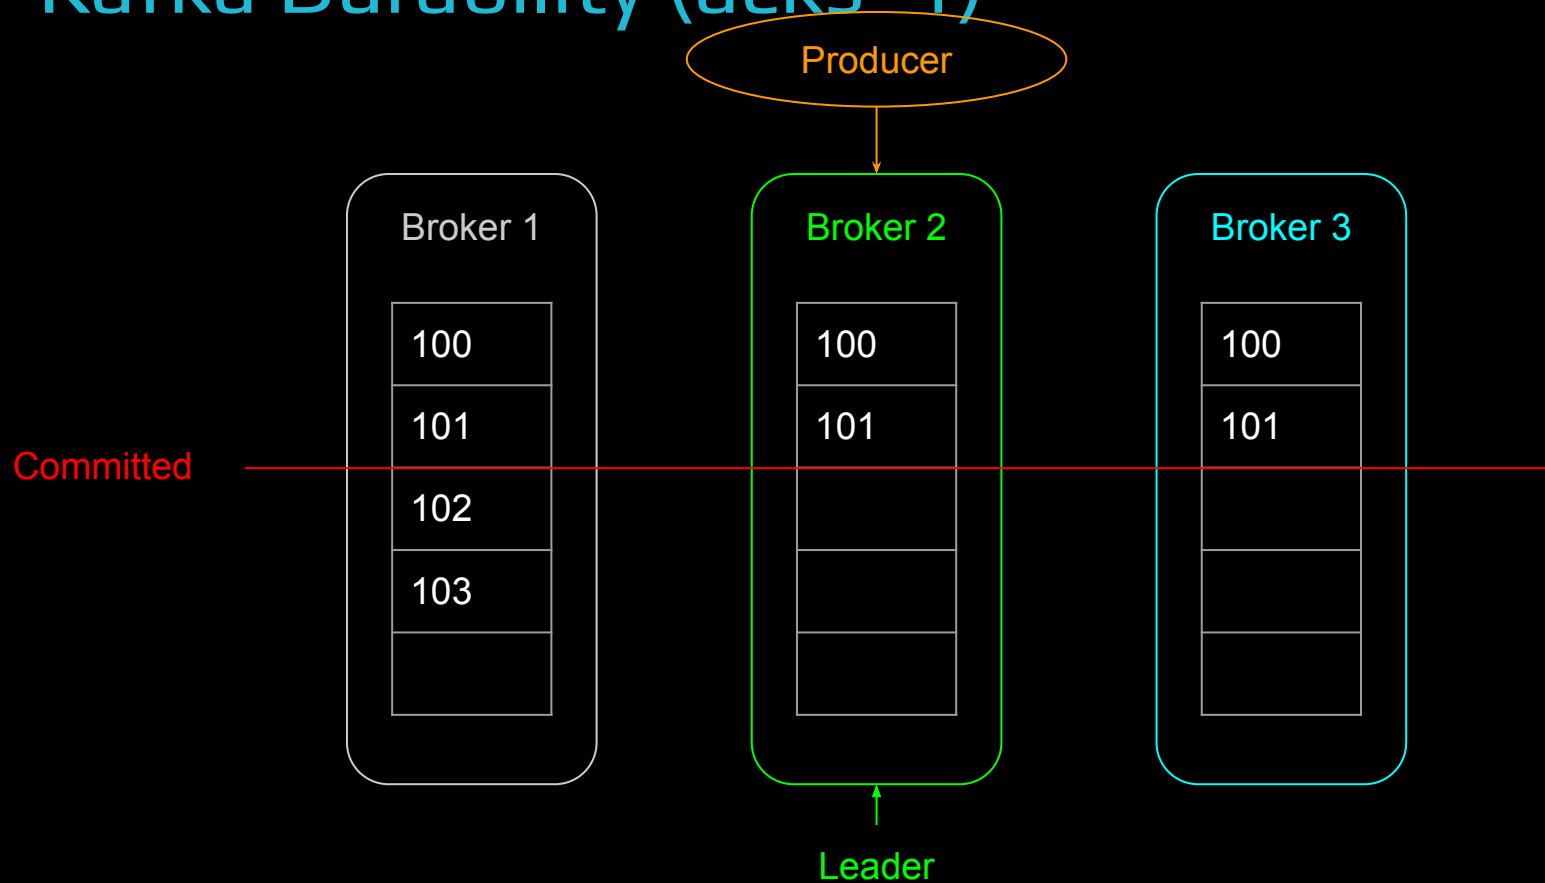
Kafka Durability (acks=1)



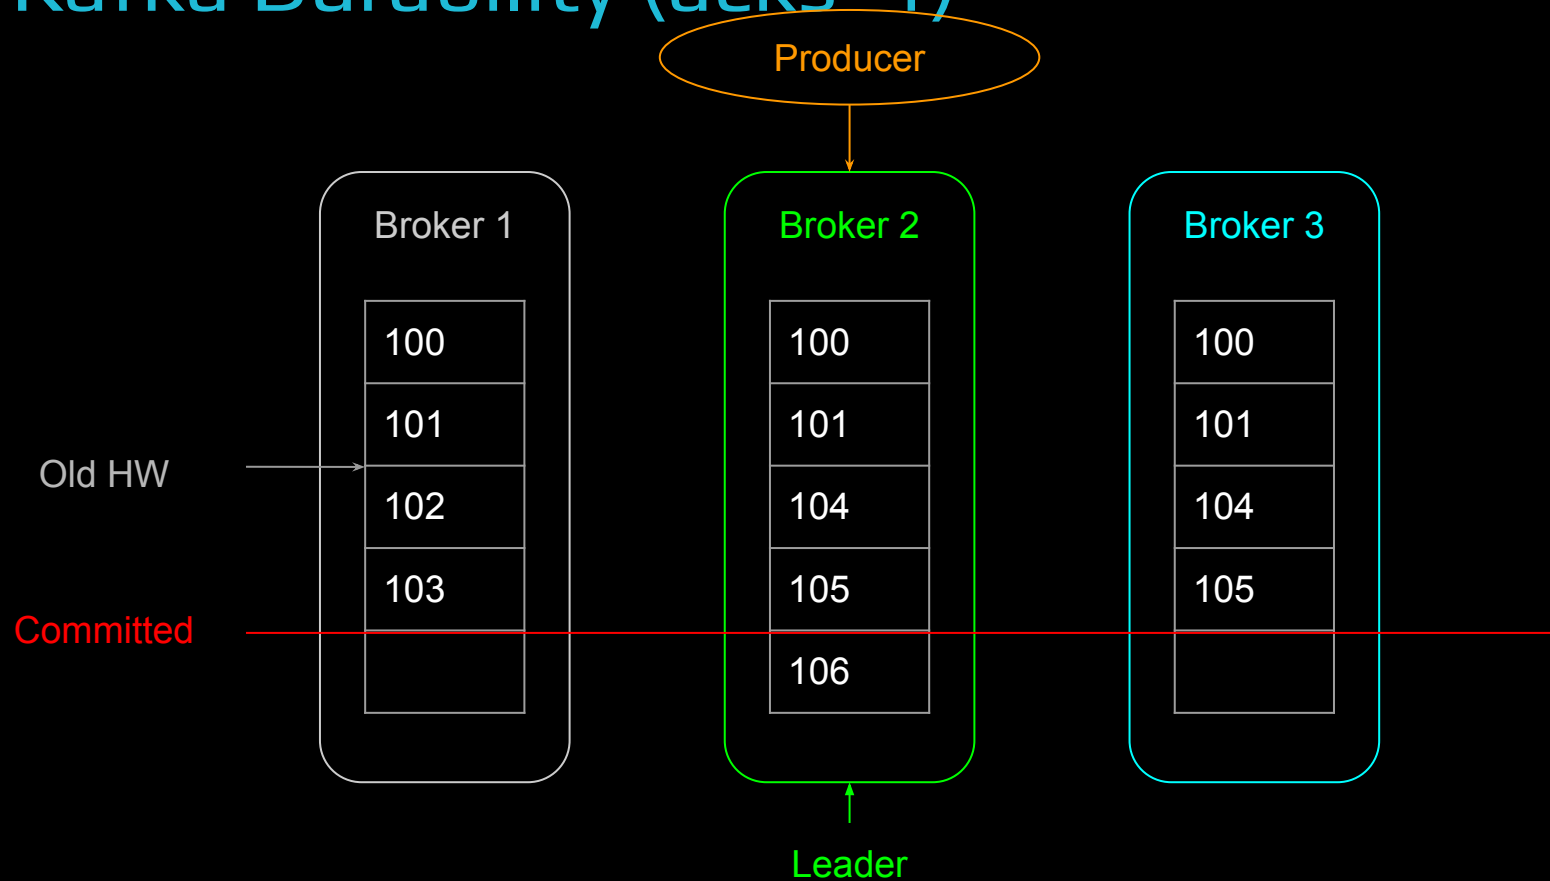
Kafka Durability (acks=1)



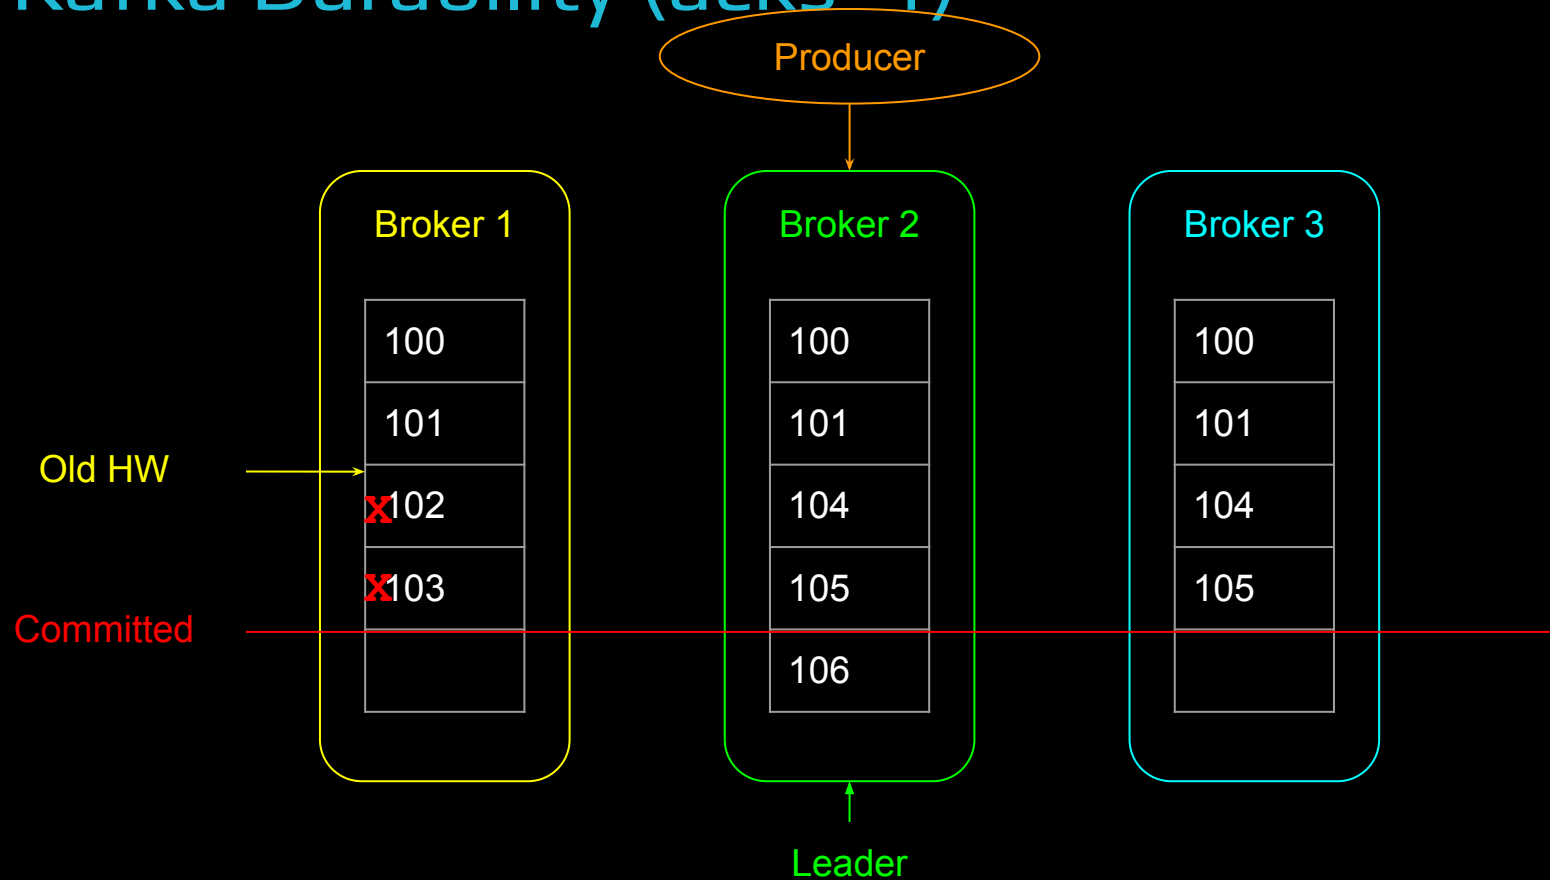
Kafka Durability (acks=1)



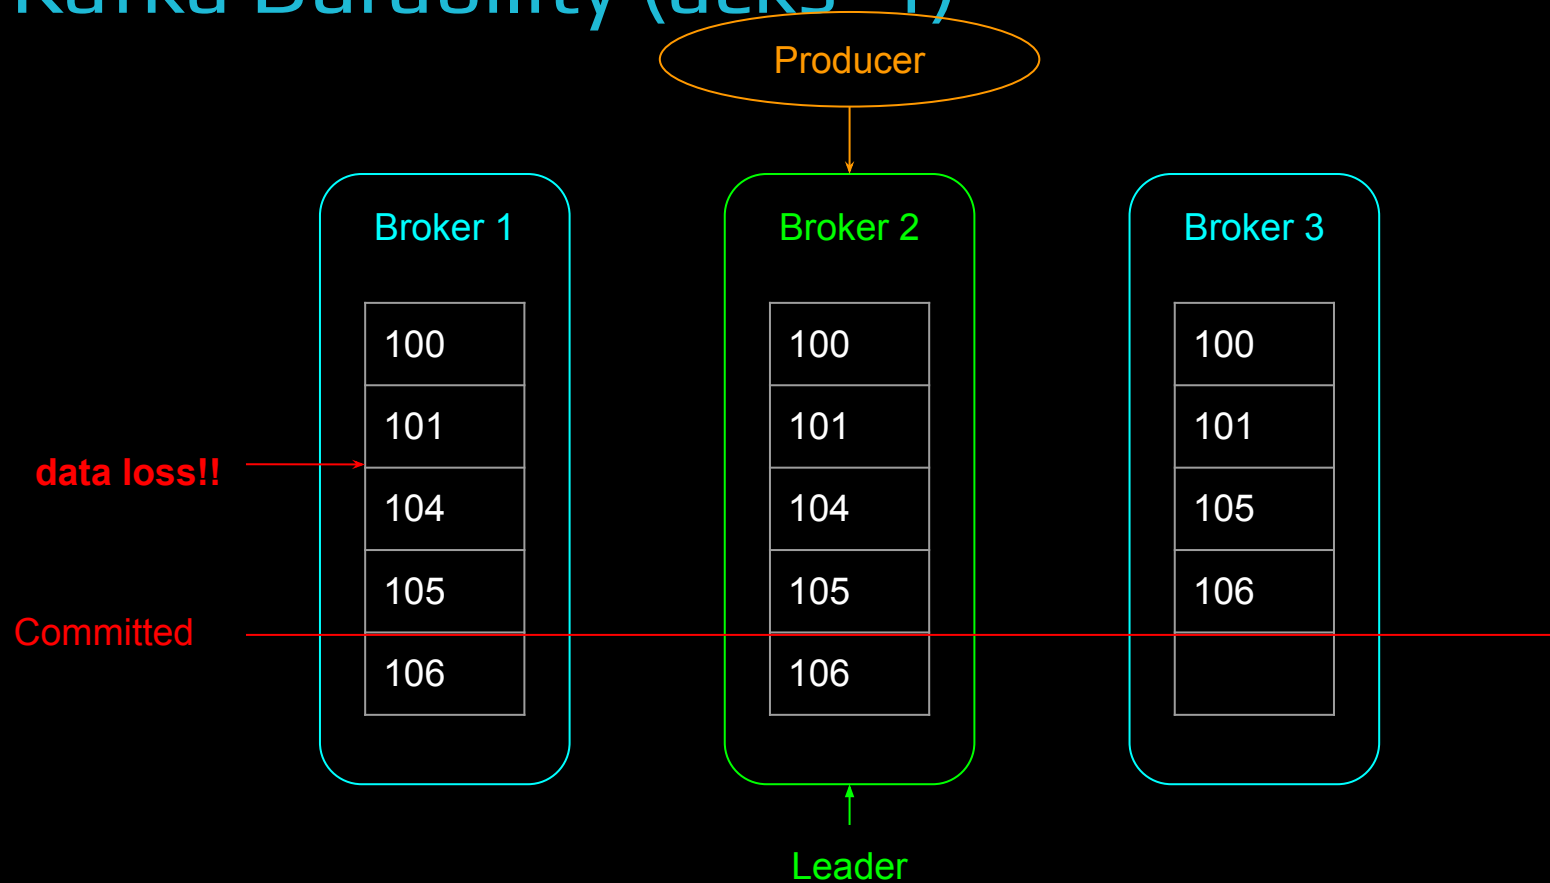
Kafka Durability (acks=1)



Kafka Durability (acks=1)

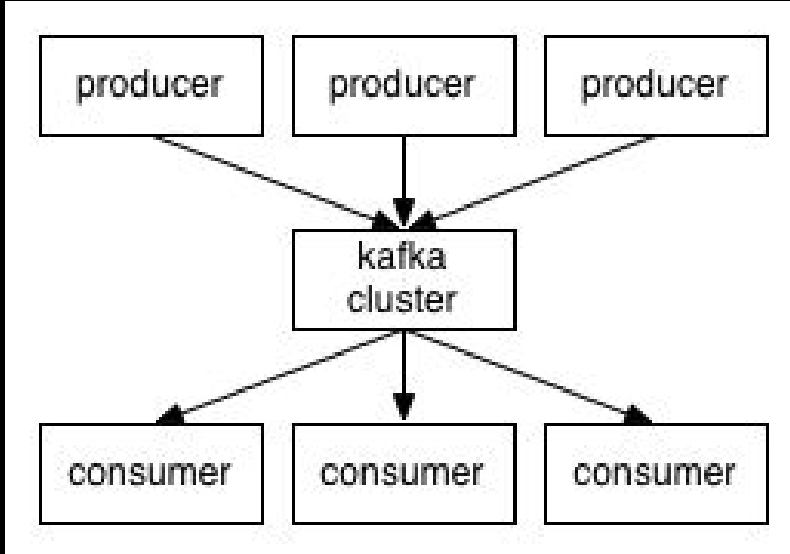


Kafka Durability (acks=1)





Distributed Messaging system

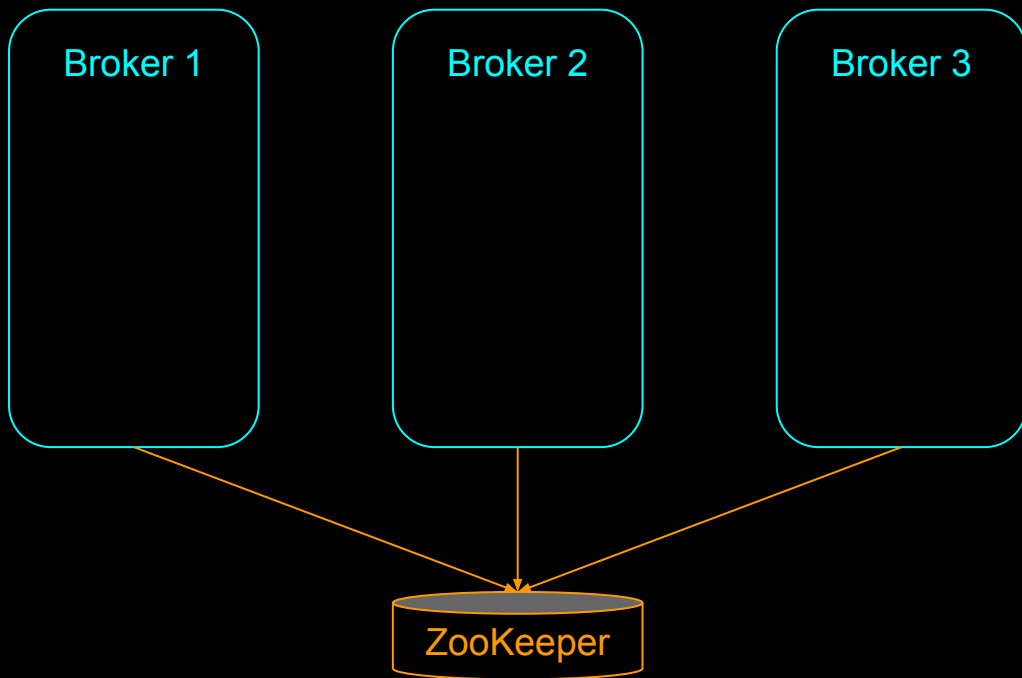


- High throughput
- Low latency
- Scalable
- Centralized
- Real-time

* Supported in Kafka 0.8+

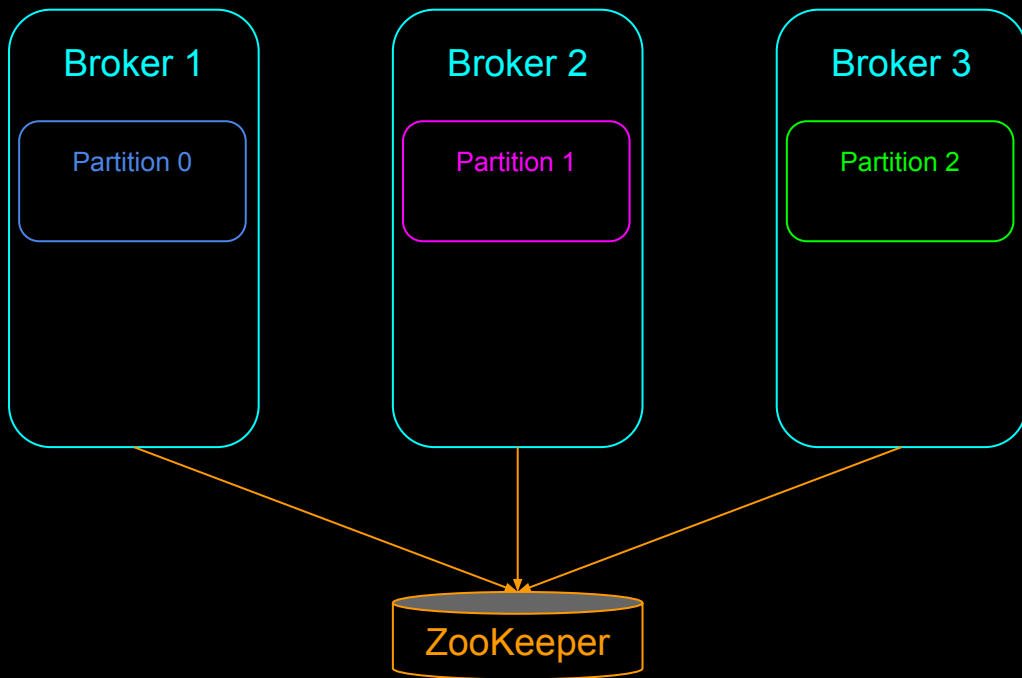
What is Kafka?

- **Distributed**
- Partitioned
- Replicated
- Commit Log



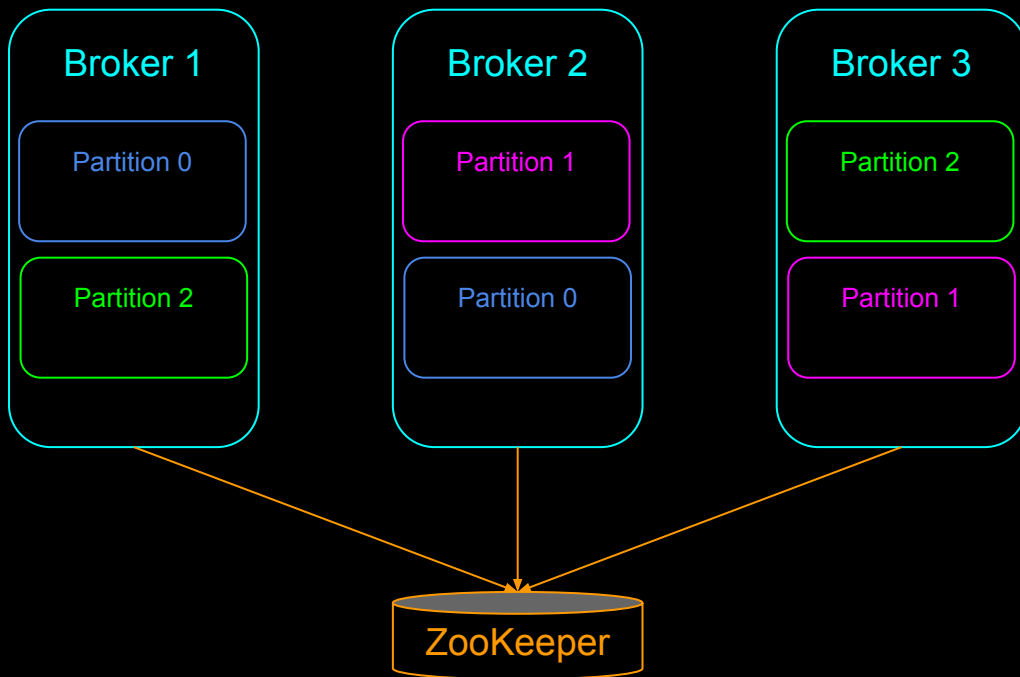
What is Kafka?

- Distributed
- **Partitioned**
- Replicated
- Commit Log



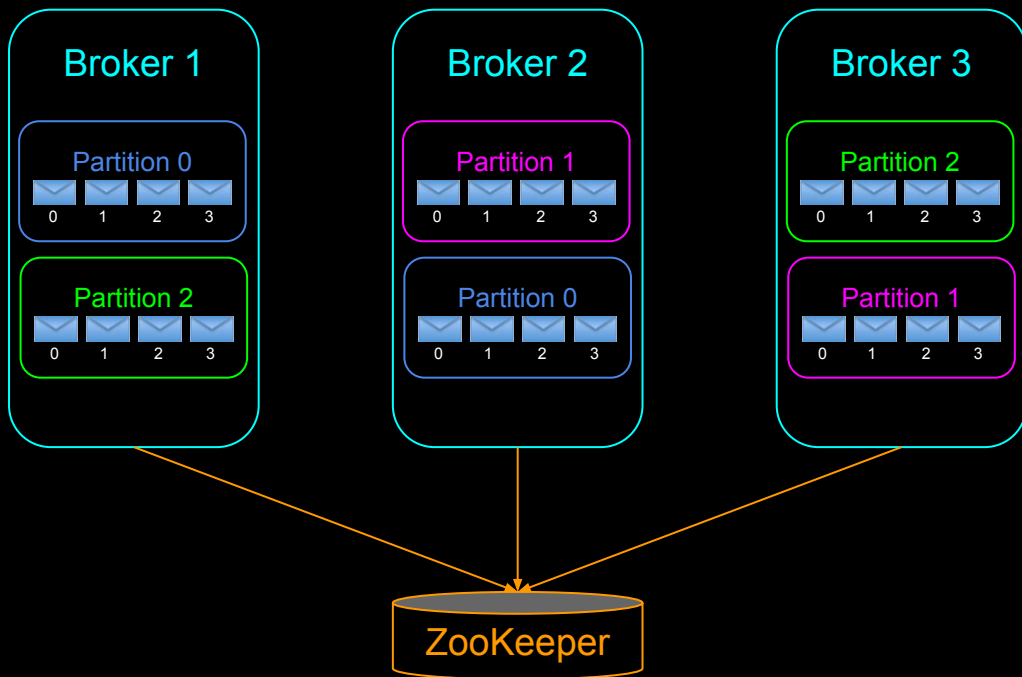
What is Kafka?

- Distributed
- Partitioned
- **Replicated**
- Commit Log



What is Kafka?

- Distributed
- Partitioned
- Replicated
- **Commit Log**



Kafka Concepts

